

is compiled into assembly code. The body of the code is as follows:

```

      xp at %ebp+8, yp at %ebp+12, zp at %ebp+16
1      movl    8(%ebp), %edi
2      movl    12(%ebp), %edx
3      movl    16(%ebp), %ecx
4      movl    (%edx), %ebx
5      movl    (%ecx), %esi
6      movl    (%edi), %eax
7      movl    %eax, (%edx)
8      movl    %ebx, (%ecx)
9      movl    %esi, (%edi)

```

Parameters *xp*, *yp*, and *zp* are stored at memory locations with offsets 8, 12, and 16, respectively, relative to the address in register *%ebp*.

Write C code for `decode1` that will have an effect equivalent to the assembly code above.

3.5 Arithmetic and Logical Operations

Figure 3.7 lists some of the integer and logic operations. Most of the operations are given as instruction classes, as they can have different variants with different operand sizes. (Only `leal` has no other size variants.) For example, the instruction class `ADD` consists of three addition instructions: `addb`, `addw`, and `addl`, adding bytes, words, and double words, respectively. Indeed, each of the instruction classes shown has instructions for operating on byte, word, and double-word data. The operations are divided into four groups: load effective address, unary, binary, and shifts. *Binary* operations have two operands, while *unary* operations have one operand. These operands are specified using the same notation as described in Section 3.4.

3.5.1 Load Effective Address

The *load effective address* instruction `leal` is actually a variant of the `movl` instruction. It has the form of an instruction that reads from memory to a register, but it does not reference memory at all. Its first operand appears to be a memory reference, but instead of reading from the designated location, the instruction copies the effective address to the destination. We indicate this computation in Figure 3.7 using the C address operator `&S`. This instruction can be used to generate pointers for later memory references. In addition, it can be used to compactly describe common arithmetic operations. For example, if register *%edx* contains value *x*, then the instruction `leal 7(%edx,%edx,4), %eax` will set register *%eax* to $5x + 7$. Compilers often find clever uses of `leal` that have nothing to do with effective address computations. The destination operand must be a register.

Instruction		Effect	Description
leal	S, D	$D \leftarrow \&S$	Load effective address
INC	D	$D \leftarrow D + 1$	Increment
DEC	D	$D \leftarrow D - 1$	Decrement
NEG	D	$D \leftarrow -D$	Negate
NOT	D	$D \leftarrow \sim D$	Complement
ADD	S, D	$D \leftarrow D + S$	Add
SUB	S, D	$D \leftarrow D - S$	Subtract
IMUL	S, D	$D \leftarrow D * S$	Multiply
XOR	S, D	$D \leftarrow D \wedge S$	Exclusive-or
OR	S, D	$D \leftarrow D \vee S$	Or
AND	S, D	$D \leftarrow D \& S$	And
SAL	k, D	$D \leftarrow D \ll k$	Left shift
SHL	k, D	$D \leftarrow D \ll k$	Left shift (same as SAL)
SAR	k, D	$D \leftarrow D \gg_A k$	Arithmetic right shift
SHR	k, D	$D \leftarrow D \gg_L k$	Logical right shift

Figure 3.7 **Integer arithmetic operations.** The load effective address (leal) instruction is commonly used to perform simple arithmetic. The remaining ones are more standard unary or binary operations. We use the notation \gg_A and \gg_L to denote arithmetic and logical right shift, respectively. Note the nonintuitive ordering of the operands with ATT-format assembly code.

Practice Problem 3.6

Suppose register `%eax` holds value x and `%ecx` holds value y . Fill in the table below with formulas indicating the value that will be stored in register `%edx` for each of the given assembly code instructions:

Instruction	Result
leal 6(%eax), %edx	_____
leal (%eax,%ecx), %edx	_____
leal (%eax,%ecx,4), %edx	_____
leal 7(%eax,%eax,8), %edx	_____
leal 0xA(,%ecx,4), %edx	_____
leal 9(%eax,%ecx,2), %edx	_____

3.5.2 Unary and Binary Operations

Operations in the second group are unary operations, with the single operand serving as both source and destination. This operand can be either a register or

```

8      } else if (_____)
9          val = ____;
10     return val;
11 }

```

gcc, with the command-line setting '-march=i686', generates the following assembly code:

```

      x at %ebp+8, y at %ebp+12
1     movl    8(%ebp), %ebx
2     movl    12(%ebp), %ecx
3     testl   %ecx, %ecx
4     jle     .L2
5     movl    %ebx, %edx
6     subl    %ecx, %edx
7     movl    %ecx, %eax
8     xorl    %ebx, %eax
9     cmpl    %ecx, %ebx
10    cmovl   %edx, %eax
11    jmp     .L4
12    .L2:
13    leal    0(,%ebx,4), %edx
14    leal    (%ecx,%ebx), %eax
15    cmpl    $-2, %ecx
16    cmovge  %edx, %eax
17    .L4:

```

Fill in the missing expressions in the C code.

3.6.7 Switch Statements

A switch statement provides a multi-way branching capability based on the value of an integer index. They are particularly useful when dealing with tests where there can be a large number of possible outcomes. Not only do they make the C code more readable, they also allow an efficient implementation using a data structure called a *jump table*. A jump table is an array where entry *i* is the address of a code segment implementing the action the program should take when the switch index equals *i*. The code performs an array reference into the jump table using the switch index to determine the target for a jump instruction. The advantage of using a jump table over a long sequence of if-else statements is that the time taken to perform the switch is independent of the number of switch cases. gcc selects the method of translating a switch statement based on the number of cases and the sparsity of the case values. Jump tables are used when there are a number of cases (e.g., four or more) and they span a small range of values.

Figure 3.18(a) shows an example of a C switch statement. This example has a number of interesting features, including case labels that do not span a contiguous

(a) Switch statement

```

1  int switch_eg(int x, int n) {
2      int result = x;
3
4      switch (n) {
5
6          case 100:
7              result += 13;
8              break;
9
10         case 102:
11             result += 10;
12             /* Fall through */
13
14         case 103:
15             result += 11;
16             break;
17
18         case 104:
19         case 106:
20             result *= result;
21             break;
22
23         default:
24             result = 0;
25     }
26
27     return result;
28 }

```

(b) Translation into extended C

```

1  int switch_eg_impl(int x, int n) {
2      /* Table of code pointers */
3      static void *jt[7] = {
4          &loc_A, &loc_def, &loc_B,
5          &loc_C, &loc_D, &loc_def,
6          &loc_D
7      };
8
9      unsigned index = n - 100;
10     int result;
11
12     if (index > 6)
13         goto loc_def;
14
15     /* Multiway branch */
16     goto *jt[index];
17
18 loc_def: /* Default case*/
19     result = 0;
20     goto done;
21
22 loc_C: /* Case 103 */
23     result = x;
24     goto rest;
25
26 loc_A: /* Case 100 */
27     result = x * 13;
28     goto done;
29
30 loc_B: /* Case 102 */
31     result = x + 10;
32     /* Fall through */
33
34 rest: /* Finish case 103 */
35     result += 11;
36     goto done;
37
38 loc_D: /* Cases 104, 106 */
39     result = x * x;
40     /* Fall through */
41
42 done:
43     return result;
44 }

```

Figure 3.18 Switch statement example with translation into extended C. The translation shows the structure of jump table `jt` and how it is accessed. Such tables are supported by GCC as an extension to the C language.

```

x at %ebp+8, n at %ebp+12
1  movl    8(%ebp), %edx          Get x
2  movl    12(%ebp), %eax        Get n
   Set up jump table access
3  subl    $100, %eax            Compute index = n-100
4  cmpl    $6, %eax             Compare index:6
5  ja      .L2                  If >, goto loc_def
6  jmp     *.L7(,%eax,4)         Goto *jt[index]
   Default case
7  .L2:                                loc_def:
8  movl    $0, %eax              result = 0;
9  jmp     .L8                  Goto done
   Case 103
10 .L5:                                loc_C:
11 movl    %edx, %eax            result = x;
12 jmp     .L9                  Goto rest
   Case 100
13 .L3:                                loc_A:
14 leal    (%edx,%edx,2), %eax    result = x*3;
15 leal    (%edx,%eax,4), %eax    result = x+4*result
16 jmp     .L8                  Goto done
   Case 102
17 .L4:                                loc_B:
18 leal    10(%edx), %eax        result = x+10
   Fall through
19 .L9:                                rest:
20 addl    $11, %eax             result += 11;
21 jmp     .L8                  Goto done
   Cases 104, 106
22 .L6:                                loc_D
23 movl    %edx, %eax            result = x
24 imull   %edx, %eax            result *= x
   Fall through
25 .L8:                                done:
   Return result

```

Figure 3.19 Assembly code for switch statement example in Figure 3.18.

range (there are no labels for cases 101 and 105), cases with multiple labels (cases 104 and 106), and cases that *fall through* to other cases (case 102) because the code for the case does not end with a break statement.

Figure 3.19 shows the assembly code generated when compiling `switch_eg`. The behavior of this code is shown in C as the procedure `switch_eg_impl` in Figure 3.18(b). This code makes use of support provided by gcc for jump tables,

as an extension to the C language. The array `jt` contains seven entries, each of which is the address of a block of code. These locations are defined by labels in the code, and indicated in the entries in `jt` by code pointers, consisting of the labels prefixed by '&&.' (Recall that the operator `&` creates a pointer for a data value. In making this extension, the authors of gcc created a new operator `&&` to create a pointer for a code location.) We recommend that you study the C procedure `switch_eg_impl` and how it relates assembly code version.

Our original C code has cases for values 100, 102–104, and 106, but the switch variable `n` can be an arbitrary int. The compiler first shifts the range to between 0 and 6 by subtracting 100 from `n`, creating a new program variable that we call `index` in our C version. It further simplifies the branching possibilities by treating `index` as an *unsigned* value, making use of the fact that negative numbers in a two's-complement representation map to large positive numbers in an unsigned representation. It can therefore test whether `index` is outside of the range 0–6 by testing whether it is greater than 6. In the C and assembly code, there are five distinct locations to jump to, based on the value of `index`. These are: `loc_A` (identified in the assembly code as `.L3`), `loc_B` (`.L4`), `loc_C` (`.L5`), `loc_D` (`.L6`), and `loc_def` (`.L2`), where the latter is the destination for the default case. Each of these labels identifies a block of code implementing one of the case branches. In both the C and the assembly code, the program compares `index` to 6 and jumps to the code for the default case if it is greater.

The key step in executing a switch statement is to access a code location through the jump table. This occurs in line 16 in the C code, with a `goto` statement that references the jump table `jt`. This *computed goto* is supported by gcc as an extension to the C language. In our assembly-code version, a similar operation occurs on line 6, where the `jmp` instruction's operand is prefixed with '*', indicating an indirect jump, and the operand specifies a memory location indexed by register `%eax`, which holds the value of `index`. (We will see in Section 3.8 how array references are translated into machine code.)

Our C code declares the jump table as an array of seven elements, each of which is a pointer to a code location. These elements span values 0–6 of `index`, corresponding to values 100–106 of `n`. Observe the jump table handles duplicate cases by simply having the same code label (`loc_D`) for entries 4 and 6, and it handles missing cases by using the label for the default case (`loc_def`) as entries 1 and 5.

In the assembly code, the jump table is indicated by the following declarations, to which we have added comments:

```

1      .section      .rodata
2      .align 4      Align address to multiple of 4
3      .L7:
4      .long  .L3      Case 100: loc_A
5      .long  .L2      Case 101: loc_def
6      .long  .L4      Case 102: loc_B
7      .long  .L5      Case 103: loc_C

```

```

8      .long      .L6      Case 104: loc_D
9      .long      .L2      Case 105: loc_def
10     .long      .L6      Case 106: loc_D

```

These declarations state that within the segment of the object-code file called “.rodata” (for “Read-Only Data”), there should be a sequence of seven “long” (4-byte) words, where the value of each word is given by the instruction address associated with the indicated assembly code labels (e.g., .L3). Label .L7 marks the start of this allocation. The address associated with this label serves as the base for the indirect jump (line 6).

The different code blocks (C labels `loc_A` through `loc_D` and `loc_def`) implement the different branches of the switch statement. Most of them simply compute a value for `result` and then go to the end of the function. Similarly, the assembly-code blocks compute a value for register `%eax` and jump to the position indicated by label .L8 at the end of the function. Only the code for case labels 102 and 103 do not follow this pattern, to account for the way that case 102 falls through to 103 in the original C code. This is handled in the assembly code and `switch_eg_impl` by having separate destinations for the two cases (`loc_C` and `loc_B` in C, .L5 and .L4 in assembly), where both of these blocks then converge on code that increments `result` by 11 (labeled `rest` in C and .L9 in assembly).

Examining all of this code requires careful study, but the key point is to see that the use of a jump table allows a very efficient way to implement a multiway branch. In our case, the program could branch to five distinct locations with a single jump table reference. Even if we had a switch statement with hundreds of cases, they could be handled by a single jump table access.

Practice Problem 3.28

In the C function that follows, we have omitted the body of the switch statement. In the C code, the case labels did not span a contiguous range, and some cases had multiple labels.

```

int switch2(int x) {
    int result = 0;
    switch (x) {
        /* Body of switch statement omitted */
    }
    return result;
}

```

In compiling the function, gcc generates the assembly code that follows for the initial part of the procedure and for the jump table. Variable `x` is initially at offset 8 relative to register `%ebp`.

<pre> 1 movl 8(%ebp), %eax <i>Set up jump table access</i> 2 addl \$2, %eax 3 cmpl \$6, %eax 4 ja .L2 5 jmp *.L8(,%eax,4) </pre>	<pre> <i>Jump table for switch2</i> 1 .L8: 2 .long .L3 3 .long .L2 4 .long .L4 5 .long .L5 6 .long .L6 7 .long .L6 8 .long .L7 </pre>
-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------	---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------

Based on this information, answer the following questions:

- A. What were the values of the case labels in the switch statement body?
- B. What cases had multiple labels in the C code?

Practice Problem 3.29

For a C function switcher with the general structure

```

1  int switcher(int a, int b, int c)
2  {
3      int answer;
4      switch(a) {
5          case _____: /* Case A */
6              c = _____;
7              /* Fall through */
8          case _____: /* Case B */
9              answer = _____;
10             break;
11         case _____: /* Case C */
12         case _____: /* Case D */
13             answer = _____;
14             break;
15         case _____: /* Case E */
16             answer = _____;
17             break;
18         default:
19             answer = _____;
20     }
21     return answer;
22 }

```

gcc generates the assembly code and jump table shown in Figure 3.20.

Fill in the missing parts of the C code. Except for the ordering of case labels C and D, there is only one way to fit the different cases into the template.

a at %ebp+8, b at %ebp+12, c at %ebp+16

1	movl	8(%ebp), %eax	1	.L7:
2	cmpl	\$7, %eax	2	.long .L3
3	ja	.L2	3	.long .L2
4	jmp	*.L7(,%eax,4)	4	.long .L4
5	.L2:		5	.long .L2
6	movl	12(%ebp), %eax	6	.long .L5
7	jmp	.L8	7	.long .L6
8	.L5:		8	.long .L2
9	movl	\$4, %eax	9	.long .L4
10	jmp	.L8		
11	.L6:			
12	movl	12(%ebp), %eax		
13	xorl	\$15, %eax		
14	movl	%eax, 16(%ebp)		
15	.L3:			
16	movl	16(%ebp), %eax		
17	addl	\$112, %eax		
18	jmp	.L8		
19	.L4:			
20	movl	16(%ebp), %eax		
21	addl	12(%ebp), %eax		
22	sall	\$2, %eax		
23	.L8:			

Figure 3.20 Assembly code and jump table for Problem 3.29.

3.7 Procedures

A procedure call involves passing both data (in the form of procedure parameters and return values) and control from one part of a program to another. In addition, it must allocate space for the local variables of the procedure on entry and deallocate them on exit. Most machines, including IA32, provide only simple instructions for transferring control to and from procedures. The passing of data and the allocation and deallocation of local variables is handled by manipulating the program stack.

3.7.1 Stack Frame Structure

IA32 programs make use of the program stack to support procedure calls. The machine uses the stack to pass procedure arguments, to store return information, to save registers for later restoration, and for local storage. The portion of the stack allocated for a single procedure call is called a *stack frame*. Figure 3.21 diagrams the general structure of a stack frame. The topmost stack frame is delimited by two pointers, with register %ebp serving as the *frame pointer*, and register %esp

Arguments: p1 at %ebp+8, p2 at %ebp+12, action at %ebp+16
 Registers: result in %edx (initialized to -1)
 The jump targets:

```

1  .L17:                                MODE_E
2      movl    $17, %edx
3      jmp     .L19
4  .L13:                                MODE_A
5      movl    8(%ebp), %eax
6      movl    (%eax), %edx
7      movl    12(%ebp), %ecx
8      movl    (%ecx), %eax
9      movl    8(%ebp), %ecx
10     movl    %eax, (%ecx)
11     jmp     .L19
12  .L14:                                MODE_B
13     movl    12(%ebp), %edx
14     movl    (%edx), %eax
15     movl    %eax, %edx
16     movl    8(%ebp), %ecx
17     addl    (%ecx), %edx
18     movl    12(%ebp), %eax
19     movl    %edx, (%eax)
20     jmp     .L19
21  .L15:                                MODE_C
22     movl    12(%ebp), %edx
23     movl    $15, (%edx)
24     movl    8(%ebp), %ecx
25     movl    (%ecx), %edx
26     jmp     .L19
27  .L16:                                MODE_D
28     movl    8(%ebp), %edx
29     movl    (%edx), %eax
30     movl    12(%ebp), %ecx
31     movl    %eax, (%ecx)
32     movl    $17, %edx
33  .L19:                                default
34     movl    %edx, %eax                Set return value
  
```

Figure 3.43 Assembly code for Problem 3.58. This code implements the different branches of a switch statement.

3.59 ♦♦

This problem will give you a chance to reverse engineer a switch statement from machine code. In the following procedure, the body of the switch statement has been removed:

```

1  int switch_prob(int x, int n)
2  {
3      int result = x;
4
5      switch(n) {
6
7          /* Fill in code here */
8      }
9
10     return result;
11 }

```

Figure 3.44 shows the disassembled machine code for the procedure. We can see in lines 4 and 5 that parameters *x* and *n* are loaded into registers *%eax* and *%edx*, respectively.

The jump table resides in a different area of memory. We can see from the indirect jump on line 9 that the jump table begins at address 0x80485d0. Using the GDB debugger, we can examine the six 4-byte words of memory comprising the jump table with the command *x/6w 0x80485d0*. GDB prints the following:

```

(gdb) x/6w 0x80485d0
0x80485d0: 0x08048438 0x08048448 0x08048438 0x0804843d
0x80485e0: 0x08048442 0x08048445

```

Fill in the body of the switch statement with C code that will have the same behavior as the machine code.

1	08048420 <switch_prob>:	
2	8048420: 55	push %ebp
3	8048421: 89 e5	mov %esp,%ebp
4	8048423: 8b 45 08	mov 0x8(%ebp),%eax
5	8048426: 8b 55 0c	mov 0xc(%ebp),%edx
6	8048429: 83 ea 32	sub \$0x32,%edx
7	804842c: 83 fa 05	cmp \$0x5,%edx
8	804842f: 77 17	ja 8048448 <switch_prob+0x28>
9	8048431: ff 24 95 d0 85 04 08	jmp *0x80485d0(,%edx,4)
10	8048438: c1 e0 02	shl \$0x2,%eax
11	804843b: eb 0e	jmp 804844b <switch_prob+0x2b>
12	804843d: c1 f8 02	sar \$0x2,%eax
13	8048440: eb 09	jmp 804844b <switch_prob+0x2b>
14	8048442: 8d 04 40	lea (%eax,%eax,2),%eax
15	8048445: 0f af c0	imul %eax,%eax
16	8048448: 83 c0 0a	add \$0xa,%eax
17	804844b: 5d	pop %ebp
18	804844c: c3	ret

Figure 3.44 Disassembled code for Problem 3.59.