

# How To Debug Programs, Part 2

## Debugging Tools

CS24 – Spring 2012

# CS24 and Debugging (2)

---

- ▶ Last time, covered a more general question:
  - ▶ Part 1: What are the basic principles and approaches of debugging?
  - ▶ [Mostly] independent of specific language, platform or toolset
- ▶ This time:
  - ▶ Part 2: What tools and approaches can I use to debug my C and IA32 assembly language programs in CS24?
- ▶ Tonight: Cover GDB, DDD, and Valgrind tools



# Last Time: Reproducing Failures

---

- ▶ Last time, discussed importance of reproducing failures:
  1. **So you can watch it fail.** You can see exactly what the program was doing as it crashed and burned.
  2. **So you can zero in on the cause.** If the program fails in some circumstances but not in others, this will give you *hints* as to what part of the program actually contains the defect.
  3. **So you can test if you actually fixed it.** If you can reliably cause the failure, and your fix makes it go away, you win!
- ▶ Debugging tools help you watch your program execute in a less intrusive way
  - ▶ Alternative is adding debug-printout statements...
  - ▶ Not always possible – e.g. if you can't change the binary, or if program is in assembly language, etc.



# GDB: The GNU Debugger

---

- ▶ gdb is the GNU Debugger
  - ▶ Designed to play very well with gcc, g++, as, etc.
- ▶ Understands different source languages you might use
  - ▶ Properly displays C code, C++ code, IA32 assembly code...
  - ▶ Understands many expressions written in these languages too
- ▶ Idea: run your program in the debugger
  - ▶ Invoke gdb with your binary file...
    - ▶ gdb starts up and loads your program into memory
  - ▶ Possibly perform initial setup for debugging your program
    - ▶ e.g. set breakpoints at known trouble-spots
  - ▶ Tell the debugger to run your program!
    - ▶ Interact with the debugger to execute instructions step by step, watch how state is manipulated by your program, etc.



# Setting Up For Debugging

---

- ▶ Normally, the compiler doesn't record variable names, function names, source code, etc. in the binary file
  - ▶ Names are translated into numeric addresses
  - ▶ C statements are translated into (possibly many) IA32 machine code instructions
  - ▶ Don't want to store information about the sources, since it takes up space, makes the program slower to load, etc.
- ▶ But, gdb needs this info to facilitate easier debugging
- ▶ Must tell compiler to leave this info in the resulting binary
  - ▶ Can pass `-g` flag to gcc (and to as)
  - ▶ Causes compiler to include debugging symbols in the binary file
  - ▶ Then, gdb can load this info when it loads your program



# Setting Up For Debugging (2)

---

- ▶ Should also tell the compiler to turn off optimizations
  - ▶ Compiler may reorder machine-code operations to optimize execution of the program
    - ▶ (applies rules to generate an equivalent program – produces same results, but may not be in exactly the same order you specified)
  - ▶ Can easily make debugging very confusing!
  - ▶ Specify `-O0` to gcc (as doesn't optimize anything)
- ▶ Can set this up in the Makefile:
  - ▶ At top of Makefile, put these variable definitions:  
`CFLAGS = -g -O0`  
`ASFLAGS = -g`
  - ▶ (Read about implicit build rules in the Make Primer to understand how these variables are used in compilation.)



# Setting Up For Debugging (3)

---

- ▶ Don't forget to recompile everything when setting up to debug!
  - ▶ Every object-file needs to be regenerated to include debug info
  - ▶ Sometimes students forget to recompile after changing the makefile to include debugging symbols
    - ▶ make clean
    - ▶ make
  - ▶ Should see “-g -O0” in the commands issued by make
  - ▶ If not, your edits may not be correct
- ▶ Most of the CS24 makefiles include the -g flag
  - ▶ (We kinda expect you'll be needing to debug your stuff...)
- ▶ Once you have included debug info, ready to use gdb!



# Example: “My program crashes!”

---

- ▶ A common experience:

```
[user@host]> ./nummain
```

Segmentation fault

[On MacOSX: “Bus error”]

```
[user@host]>
```

- ▶ What is going on?!? Where is my program crashing?!
- ▶ gdb makes it easy to find where your program is crashing

```
[user@host]> gdb nummain
```

*...gdb startup information...*

```
(gdb)
```

- ▶ Next, start your program running within the debugger

```
(gdb) run
```

*[starts the program from main()]*

Starting program: nummain

- ▶ When your program crashes, gdb will come back and tell you





# Example: “My program crashes!” (2)

---

- ▶ When your program crashes:

(gdb) run

Starting program: nummain

Program received signal SIGSEGV, Segmentation fault.

0x08048546 in **add\_num (lst=0xbffff3c4, num=3) at numlist.c:16**

16    lst->tail->next = n;

- ▶ Notice: gdb tells you *exactly* where the crash occurred!
    - ▶ What function, what the arguments to the function were, and what source file the function is in
    - ▶ Additionally, shows the exact line that had the problem!
- 



# Example: “My program crashes!” (3)

---

- ▶ Can look at the surrounding code in context:

(gdb) **list** *[shows 10 lines of code]*

```
11 void add_num(list *lst, int num) {  
12     node *n = (node *) malloc(sizeof(node));  
13     n->value = num;  
14     n->next = NULL;  
15  
16     lst->tail->next = n;  
17  
18     lst->tail = n;  
19     lst->size++;  
20 }
```

- ▶ Can see the code in context
- ▶ “Oh, I forgot to check if the list was initially empty...”



# Example: “My program crashes!” (4)

---

- ▶ Can also look at variables in the program:

(gdb) **print lst** *[Hmm, lst looks okay...]*

\$1 = (list \*) 0xbffff3c4

(gdb) **print \*lst**

\$2 = {head = 0x0, tail = 0x0, size = 0}

(gdb) **print lst->tail** *[Ah ha! lst->tail is NULL.]*

\$3 = (node \*) 0x0

(gdb) **print n**

\$4 = (node \*) 0x804b008

(gdb) **print \*n** *[Well, at least n is okay...]*

\$5 = {value = 3, next = 0x0}

- ▶ Note that gdb understands C syntax!

- ▶ Can even do things like: `print *(lst->head->next->next)`



# Example: “My program crashes!” (4)

---

- ▶ Can see the exact call-sequence for your program:

(gdb) **where** *[shows the entire call-stack]*

#0 0x08048546 in add\_num (lst=0xbffff3c4, num=3) at numlist.c:16

#1 0x08048651 in main () at nummain.c:10

- ▶ (this program isn't very exciting)



# GDB Breakpoints

---

- ▶ Can also set breakpoints in your program
  - ▶ When gdb reaches the breakpoint, it will stop automatically
  - ▶ Can only set breakpoints while the program is stopped!
  - ▶ (Frequently must set these up before running the program.)
- ▶ Fixed previous bug, but still have odd behavior:
  - ▶ `[user@host]> ./nummain`
  - ▶ Original list: 3 | 4 | 5
  - ▶ Reversed list: 5 | 4 | 3
  - ▶ Original list: 5 9 2 6 5 4
  - ▶ Reversed list: 4 5 6 2 9 5
  - ▶ Cleaning up list.
- ▶ Hmm, losing some of our digits in the last part of the program
- ▶ Maybe `reverse()` has a bug in it...



# GDB Breakpoints (2)

---

- ▶ Set breakpoint in reverse():

gdb nummain

*...gdb startup information...*

(gdb) **break reverse**

Breakpoint 1 at 0x80485eb: file numlist.c, line 50.

(gdb)

- ▶ Every time gdb reaches the breakpoint, it will stop executing and let you poke around

(gdb) **run**

Starting program: nummain

Original list: 3 | 4 | 5

Breakpoint 1, reverse (lst=0xbffff3c4) at numlist.c:50

50 prev = NULL;

(gdb)

---



# GDB Breakpoints (3)

---

- ▶ As before, can **list** the code, or **print** various values
- ▶ Can also single-step through the code:
  - ▶ **step** Steps to next line of code (steps into function calls)
  - ▶ **next** Steps to next line of code (steps over function calls)
  - ▶ **continue** Resumes running program, until next breakpoint, or whenever the program terminates.

- ▶ Example, continued:

Breakpoint 1, reverse (lst=0xbffff3c4) at numlist.c:50

```
50    prev = NULL;
```

```
(gdb) next
```

```
51    curr = lst->head;
```

```
(gdb) next
```

```
53    while (curr != NULL) {
```

```
(gdb) next
```

```
54        next = curr->next;
```



# GDB Breakpoints (4)

---

- ▶ Can set breakpoints in several ways:
  - ▶ `break label`
    - ▶ Can specify a function name, or a label in your IA32 assembly code
  - ▶ `break filename:line`
    - ▶ Can specify a filename and line number as well
- ▶ Every breakpoint is given a numeric ID  
(gdb) `break reverse`  
Breakpoint **1** at 0x80485eb: file numlist.c, line 50.  
(gdb)
- ▶ Can enable and disable breakpoints as well:  
(gdb) `disable 1` *[disabled reverse() breakpoint]*  
(gdb) `enable 1` *[re-enables reverse() breakpoint]*





# GDB Commands

---

- ▶ All commands can be abbreviated with one character:

▶ <u>Cmd</u>	<u>Abbrev</u>	<u>Description</u>
▶ run	r	Starts program from main()
▶ break	b	Set a breakpoint
▶ continue	cont, c	Continues execution of program
▶ step	s	Steps to next line (steps into functions)
▶ next	n	Steps to next line (steps over functions)
▶ print	p	Print out various values
▶ quit	q	Exits gdb

- ▶ Many other commands in GDB as well!

▶ gdb has an extensive help facility	
▶ help	h
▶ help <i>command</i>	

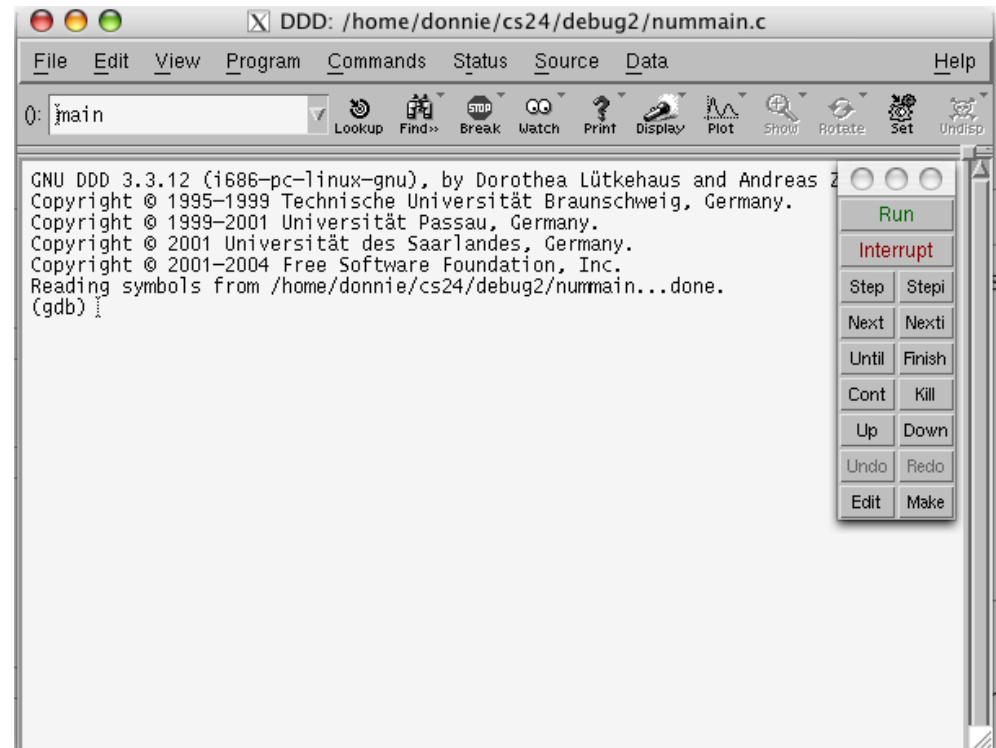


# DDD: Graphical Frontend to GDB

- ▶ ddd is a graphical frontend to gdb

- ▶ (if you like such things)
- ▶ Start with: `ddd program`

- ▶ Provides some basic hand-holding with gdb
- ▶ (I don't think it helps much, but whatever.)



# GDB and IA32 Nuances

---

- ▶ GDB can also step through IA32 code
  - ▶ Everything is basically the same, except for print expressions!
- ▶ To print individual registers:
  - ▶ **print \$eax** [not %eax!]
    - ▶ Shows contents of %eax register
  - ▶ **print \*(int \*) \$esp**
    - ▶ Shows value stored at (%esp)
    - ▶ Note: must tell gdb what type is at the pointer!
  - ▶ **print \*(int \*) (\$esp + 4)**
    - ▶ Shows value stored at 4(%esp)
- ▶ To display all registers at once:
  - ▶ **info registers** [or: i reg]
    - ▶ Lists all registers (eax, ebx, ..., eip, eflags), in base-10 and base-16



# GDB Summary

---

- ▶ gdb is very good for watching your program's state and behavior as it executes
- ▶ Extensive set of commands
  - ▶ Will post a gdb cheat-sheet on the Moodle for you to use
- ▶ Isn't the best for finding dynamic memory mgmt bugs:
  - ▶ Accessing memory that you freed
  - ▶ Freeing a memory block multiple times
  - ▶ Accessing memory through an uninitialized pointer



# Valgrind

---

- ▶ Valgrind is a suite of dynamic analysis tools built on a single unified platform
  - ▶ Simulates execution of Intel IA32 programs in a virtual machine
  - ▶ Allows sophisticated debugging of runtime behavior
  - ▶ Also allows sophisticated performance analysis of behavior
- ▶ Valgrind was designed to run on Linux
  - ▶ A port to MacOSX exists, but doesn't always work very well ☹️
  - ▶ No valgrind port to Cygwin ☹️
- ▶ Most widely used tool is the memcheck tool
  - ▶ Invoked by default when you run valgrind at command-line
- ▶ Valgrind also benefits from debugging symbols!
  - ▶ Always want to compile with `-g -O0` when using valgrind



# Valgrind (2)

---

- ▶ Usually use valgrind after you have identified and resolved other issues with gdb

- ▶ gdb and valgrind are really good at different things

- ▶ Simple to use:

```
[user@host]> valgrind ./nummain
```

```
==30421== Memcheck, a memory error detector
```

```
==30421== Copyright (C) 2002-2009, and GNU GPL'd, by Julian Seward et al.
```

```
==30421== Using Valgrind-3.5.0 and LibVEX; rerun with -h for copyright info
```

```
==30421== Command: ./nummain
```

```
==30421==
```

- ▶ Valgrind begins executing your program!

- ▶ Prints out memory management issues as it runs



# Memory Management Errors

---

- ▶ If your program does anything bad, valgrind immediately prints out an error message
- ▶ Example: nummain accesses previously-freed memory

Cleaning up list.

```
==30421== Invalid read of size 4
```

```
==30421==    at 0x8048595: empty (numlist.c:31)
```

```
==30421==    by 0x80487DC: main (nummain.c:41)
```

```
==30421== Address 0x41bb224 is 4 bytes inside a block of size 8 free'd
```

```
==30421==    at 0x4026996: free (in /usr/lib/valgrind/...memcheck-x86-linux.so)
```

```
==30421==    by 0x8048591: empty (numlist.c:30)
```

```
==30421==    by 0x80487DC: main (nummain.c:41)
```

- ▶ As with gdb, valgrind tells you exactly what happened, and where it happened
- 



# Memory Management Errors (2)

---

- ▶ At end of program execution, valgrind prints out a summary of the program's heap usage
  - ▶ This output identifies leaked memory blocks!
- ▶ Example: nummain doesn't call free() on list nodes

Cleaning up list.

==30484==

==30484== HEAP SUMMARY:

==30484== in use at exit: 80 bytes in 10 blocks

==30484== total heap usage: 10 allocs, 0 frees, 80 bytes allocated

==30484==

==30484== LEAK SUMMARY:

==30484== definitely lost: 8 bytes in 1 blocks

==30484== indirectly lost: 72 bytes in 9 blocks

==30484== possibly lost: 0 bytes in 0 blocks

==30484== still reachable: 0 bytes in 0 blocks

==30484== suppressed: 0 bytes in 0 blocks

==30484== Rerun with --leak-check=full to see details of leaked memory





# Memory Management Errors (3)

---

- ▶ Rerunning with `--leak-check=full` causes valgrind to record where leaked blocks were allocated

```
[user@host]> valgrind --leak-check=full ./nummain
```

```
...
```

```
Cleaning up list.
```

```
==30485==
```

```
==30485== HEAP SUMMARY:
```

```
==30485==    in use at exit: 80 bytes in 10 blocks
```

```
==30485== total heap usage: 10 allocs, 0 frees, 80 bytes allocated
```

```
==30485==
```

```
==30485== 80 (8 direct, 72 indirect) bytes in 1 blocks are definitely lost ...
```

```
==30485==    at 0x402760A: malloc (in /usr/lib/valgrind/...-x86-linux.so)
```

```
==30485==    by 0x80484E7: add_num (numlist.c:12)
```

```
==30485==    by 0x8048722: main (nummain.c:29)
```

```
==30485==
```

```
==30485== LEAK SUMMARY:
```

```
...
```



# GDB and Valgrind

---

- ▶ GDB and Valgrind are very powerful debugging tools!
- ▶ GDB is best for tracking down crashing bugs and other subtle defects
  - ▶ Primarily useful when need to watch what the program is doing as the bug occurs
  - ▶ Doesn't necessarily identify more subtle memory management issues
- ▶ Valgrind is best for finding subtle memory mgmt issues
  - ▶ Forgetting to free memory, freeing memory multiple times, accessing memory after freeing it, etc.
- ▶ Always use the right tool for the job!

