



CS 24: INTRODUCTION TO COMPUTING SYSTEMS

Spring 2014

Lecture 2

LAST TIME

- Began exploring the concepts behind a simple programmable computer
- Construct the computer using Boolean values (a.k.a. “bits”) and logic gates to process them
- Represent unsigned and signed integers as vectors of bits

$$\text{B2U}_w(\mathbf{x}) = \sum_{i=0}^{w-1} x_i 2^i$$

$$\text{B2T}_w(\mathbf{x}) = -x_{w-1}2^{w-1} + \sum_{i=0}^{w-2} x_i 2^i$$

- Briefly explored how to construct more complex computations using gates
 - e.g. unsigned and signed arithmetic

SIGNED INTEGER REPRESENTATION (2)


- Easy trick for converting an integer to its two's complement representation:
 - Invert the bits, then add one
- Example:
 - Find two's complement representation for -42
 - Unsigned representation for 42 is 00101010_2
 - Invert the bits: 11010101_2
 - Add one: 11010110_2
- Converting back, following $B2T_{w=8}$ function:
 - $= -1 \times 2^7 + 1 \times 2^6 + 1 \times 2^4 + 1 \times 2^2 + 1 \times 2^1$
 - $= -128 + 64 + 16 + 4 + 2$
 - $= -42$

FUNCTIONAL COMPONENTS OF A SIMPLE PROCESSOR

- Can use our logic gates to construct various components to use in a processor
 - Already saw how to implement addition with logic
- Minimal components for a simple processor:
- Signal Buses
 - Ability to route signals within our processor
- Arithmetic/Logic Unit (ALU)
 - Performs various arithmetic and logical operations on data inputs, based on control inputs
- Memory
 - Addressable locations to store and retrieve values

BUSES



- A bus is a set of wires that transfer signals from one component to another
 - Transmits values of a fixed bit-width, e.g. 32 bits
- Common uses for buses in a computer:
 - Transfer data between CPU and memory
 - Transfer data between CPU and peripherals
- Buses often drawn as a single line with a slash across it



- Individual signals drawn as a line with no slash


ROUTING BUSES

- Multiplexers and demultiplexers (decoders) are used to route buses between multiple components
- Example: a 4-input multiplexer (MUX)
 - Has two address inputs
 - Address selects one of 4 data inputs
 - Corresponding data input is fed to the data output
- A 4-input demultiplexer (DEMUX)
 - Again, two address inputs
 - Address selects one of 4 data outputs
 - Single data input fed to corresponding data output



ARITHMETIC/LOGIC UNIT

- A component that can perform various arithmetic and logic functions
- Symbol:



- Given two w -bit inputs and a set of control inputs
 - Control inputs specify the operation to perform
- Produces a w -bit result, and status outputs
 - Example status outputs:
 - sign flag (topmost bit of R)
 - carry-out flag (unsigned overflow)
 - zero flag (is $R == 0$?)
 - overflow flag (signed overflow)

EXAMPLE ALU OPERATIONS


- Control signals specify what operation to perform
- Example: for our contrived ALU

Control	Operation
0001	ADD A B
0011	SUB A B
0100	NEG A
1000	AND A B
1001	OR A B
1010	XOR A B
1011	INV A
1100	SHL A
1110	SHR A

- By feeding appropriate control and data signals to ALU in sequence, can perform computations
- Some operations require only one argument
 - Second argument ignored

MEMORY

- Need a component to store both instructions and data to feed to the ALU
- Memory:
 - An array of linearly addressable locations
 - Each location has its own address
 - Each location can hold a single w -bit value
- Inputs and outputs:
 - Address of location to read or write
 - Read/Write control signal
 - Data-input bus
 - Data-output bus




ASIDE: CPU COMPONENTS AND GATES

- It is a big claim that we can construct all of these components entirely from logic gates...
- Unfortunately, beyond the scope of CS24 to explore all the ways such components can be constructed, different approaches, etc.
- If you are curious how these things work, see the primer on the CS24 Moodle
 - Shows some *basic* ways these components can be constructed
- **Don't need to know this material in depth!**
 - For CS24, really only need to understand the basics of how to implement logic equations with gates

ASSEMBLING THE COMPUTER

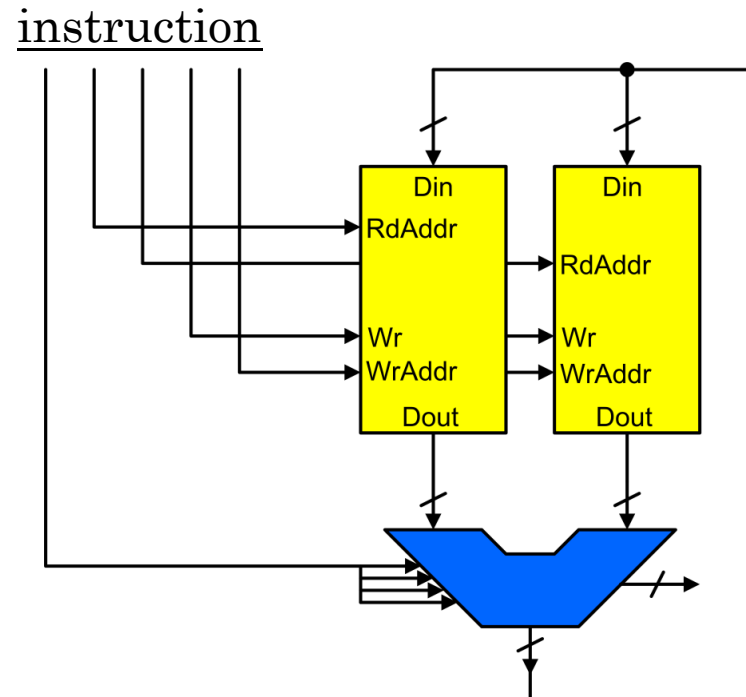
- Hook these components together, like this:



- Simplifications in our computer:
 - Two memory banks; identical copies of each other
 - Don't care about ALU status outputs

INSTRUCTING THE COMPUTER

- This set of inputs forms an instruction
- Consists of:
 - The operation the ALU should perform
 - Addresses of two input values
 - Whether result should be stored
 - If so, what address to store the result at
- To program our computer:
 - Devise a set of instructions to implement our desired computation



WRITING A PROGRAM

- Instructions for the processor are very limited
 - Can only compute one value, from one or two values
- Usually can't implement a program in only one instruction
- Instead:
 - String together a sequence of instructions to implement the computation
 - Instructions will communicate via memory locations
- Computation we will implement:
 - $C = (A - 2B) \& 00001111_2$
 - Given inputs A and B
 - Multiply B by 2, subtract result from A, then bitwise-AND with a mask

IMPLEMENTING OUR COMPUTATION (1)

- Computation: $C = (A - 2B) \& 00001111_2$
- Step 1: Assign locations for inputs and outputs
- Inputs:
 - A and B (obvious)
 - Also, our mask: 00001111_2
 - Program needs to include our constants, too
- Output:
 - C
- Givens:
 - Our memory has 8 locations
 - Memory addresses are 3 bits wide
 - Data values are 8 bits wide ($w = 8$)

ASSIGNING DATA LOCATIONS

- Locations for our initial and final data values:

Address	Value
0	A
1	B
2	
3	
4	00001111 ₂
5	
6	
7	C

IMPLEMENTING OUR COMPUTATION (2)

- Step 2: decompose our program into instructions the processor can actually handle
- Program:
 - $C = (A - 2B) \& 00001111_2$
- Need to know processor's operations for this step.
- Steps:
 - Perform $2B$ first, as $B + B$
 - Then, subtract previous result from A
 - Finally, bitwise-AND this with mask to produce C

Control	Operation	
0001	ADD	A B
0011	SUB	A B
0100	NEG	A
1000	AND	A B
1001	OR	A B
1010	XOR	A B
1011	INV	A
1100	SHL	A
1110	SHR	A

IMPLEMENTING OUR COMPUTATION (3)

- Step 3: need to assign locations to these intermediate values!

- Result of $B + B =$ location 2
- Result of $A - 2B =$ location 3
- Result of bitwise-AND stored in location 7
 - This is our result

Address	Value
0	A
1	B
2	2B
3	$A - 2B$
4	00001111 ₂
5	
6	
7	C

IMPLEMENTING OUR COMPUTATION (4)

- Step 4: Translate our program into instructions!
- Need to know form of instructions:
 - Operation Rd1Addr Rd2Addr Wr WrAddr
- Also need our memory layout and operation codes

Control	Operation	
0001	ADD	A B
0011	SUB	A B
0100	NEG	A
1000	AND	A B
1001	OR	A B
1010	XOR	A B
1011	INV	A
1100	SHL	A
1110	SHR	A

Address	Value
0	A
1	B
2	2B
3	A - 2B
4	00001111 ₂
5	
6	
7	C

IMPLEMENTING OUR COMPUTATION (5)

Control	Operation		Address	Value
0001	ADD	A B	0	A
0011	SUB	A B	1	B
0100	NEG	A	2	$2B$
1000	AND	A B	3	$A - 2B$
1001	OR	A B	4	00001111_2
1010	XOR	A B	5	
1011	INV	A	6	
1100	SHL	A	7	C
1110	SHR	A		

○ Operation Rd1Addr Rd2Addr Wr WrAddr

○ Writing our program:

- Slot 2 = $2B$ 000: 0001 001 001 1 010
- Slot 3 = $A - 2B$ 001: 0011 000 010 1 011
- Slot 7 = (...) & mask 010: 1000 011 100 1 111

RUNNING OUR PROGRAM

- To run our program:
 - Load instructions into instruction memory
 - Load initial data into data memory
 - Start program counter at 0
- Each instruction executes in sequence, updating memory locations
 - Uses results of previous instructions
- State of our computer:
 - Instruction memory
 - Data memory
 - Program counter

RUNNING OUR PROGRAM: INITIAL STATE

- Instruction memory:

000: 0001 001 001 1 010

001: 0011 000 010 1 011

010: 1000 011 100 1 111

- Data memory:

0: A

1: B

2: ???

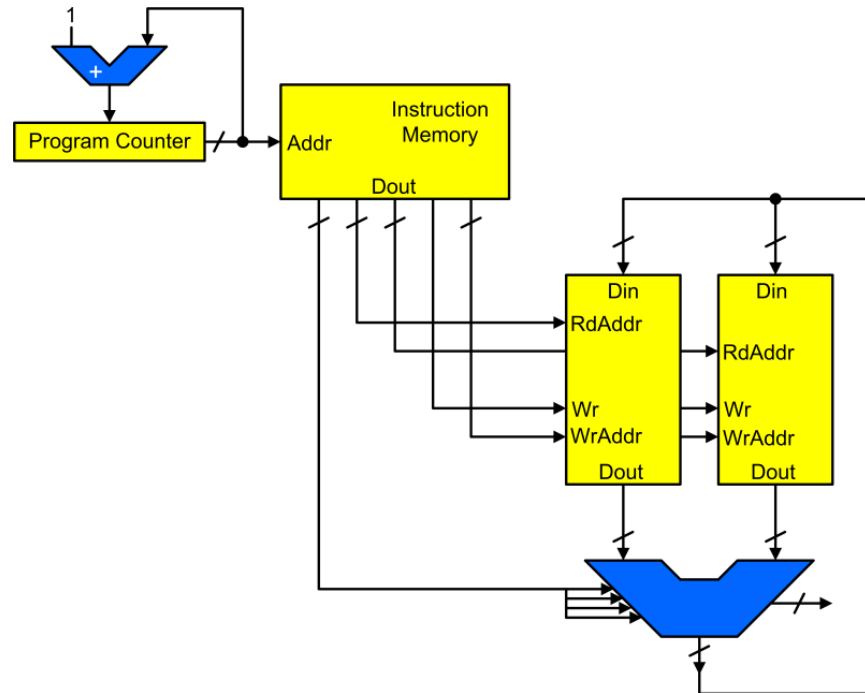
3: ???

4: 00001111₂

5: ???

6: ???

7: ???



Program Counter: 000

STEP 1: SLOT 2 = B + B

- Instruction memory:

000: 0001 001 001 1 010

001: 0011 000 010 1 011

010: 1000 011 100 1 111

- Data memory:

0: A

1: B

2: $2B = B + B$

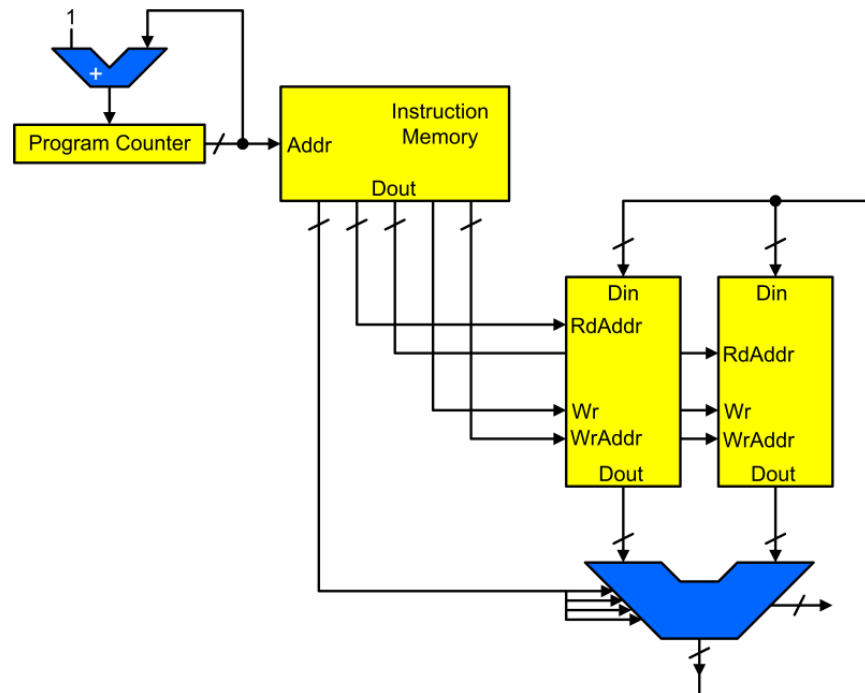
3: ???

4: 00001111_2

5: ???

6: ???

7: ???



Program Counter: 000

STEP 1: UPDATE PROGRAM COUNTER

- Instruction memory:

000: 0001 001 001 1 010

001: 0011 000 010 1 011

010: 1000 011 100 1 111

- Data memory:

0: A

1: B

2: 2B

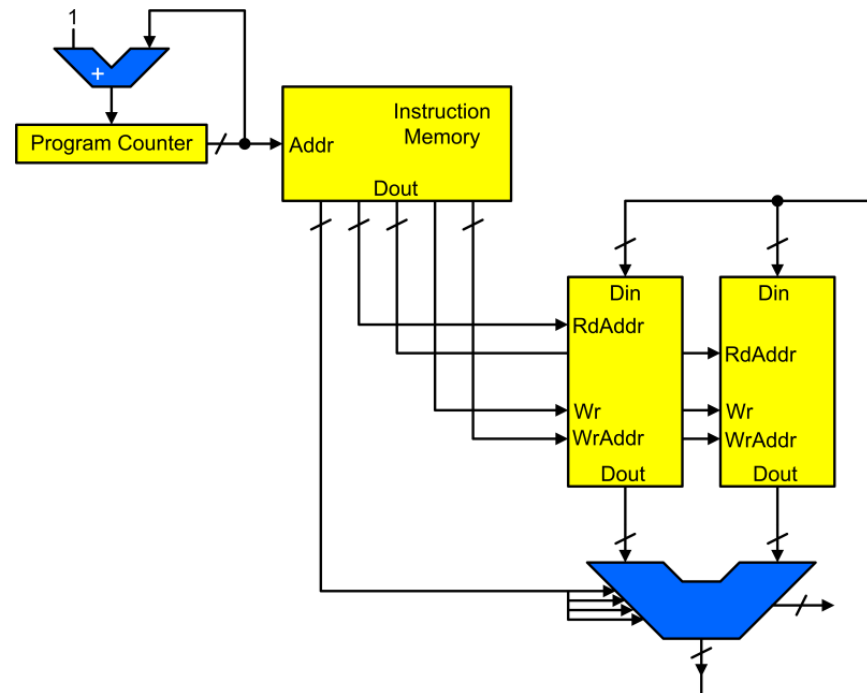
3: ???

4: 00001111₂

5: ???

6: ???

7: ???



Program Counter: 001

STEP 2: SUBTRACT 2B FROM A

- Instruction memory:

000: 0001 001 001 1 010

001: 0011 000 010 1 011

010: 1000 011 100 1 111

- Data memory:

0: A

1: B

2: 2B


3: **A - 2B**

4: 00001111₂

5: ???

6: ???

7: ???



Program Counter: 001

STEP 2: UPDATE PROGRAM COUNTER

- Instruction memory:

000: 0001 001 001 1 010

001: 0011 000 010 1 011

010: 1000 011 100 1 111

- Data memory:

0: A

1: B

2: 2B

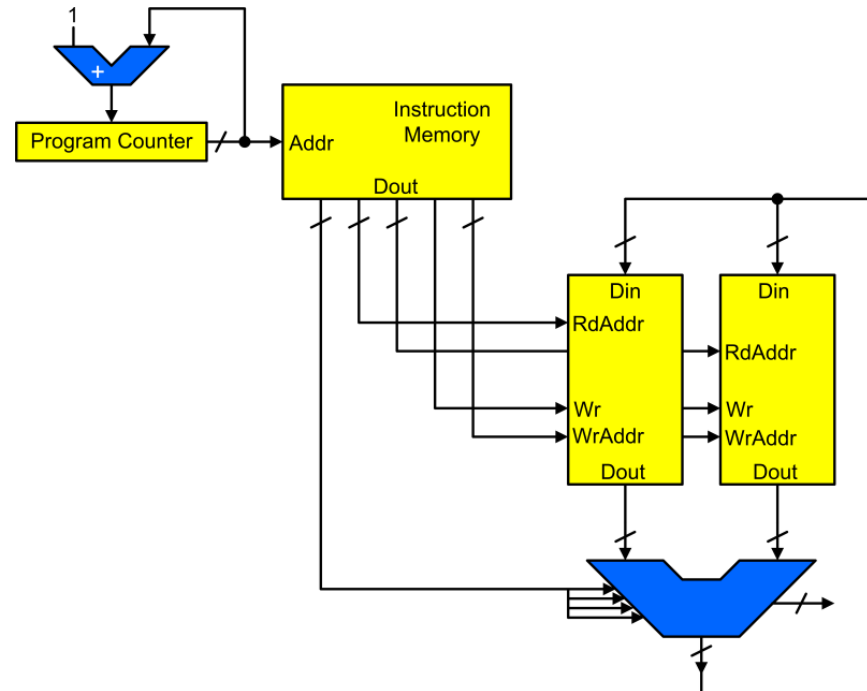
3: A – 2B

4: 00001111₂

5: ???

6: ???

7: ???



Program Counter: 010

STEP 3: $C = (A - 2B) \& 00001111_2$

○ Instruction memory:

000: 0001 001 001 1 010

001: 0011 000 010 1 011

010: 1000 011 100 1 111

○ Data memory:

0: A

1: B

2: 2B


3: $A - 2B$

4: 00001111_2

5: ???

6: ???

7: $(A - 2B) \& 00001111_2$



Program Counter: 010

RUNNING OUR PROGRAM: FINAL RESULT

- Instruction memory:

000: 0001 001 001 1 010

001: 0011 000 010 1 011

010: 1000 011 100 1 111

- Data memory:

0: A

1: B

2: 2B

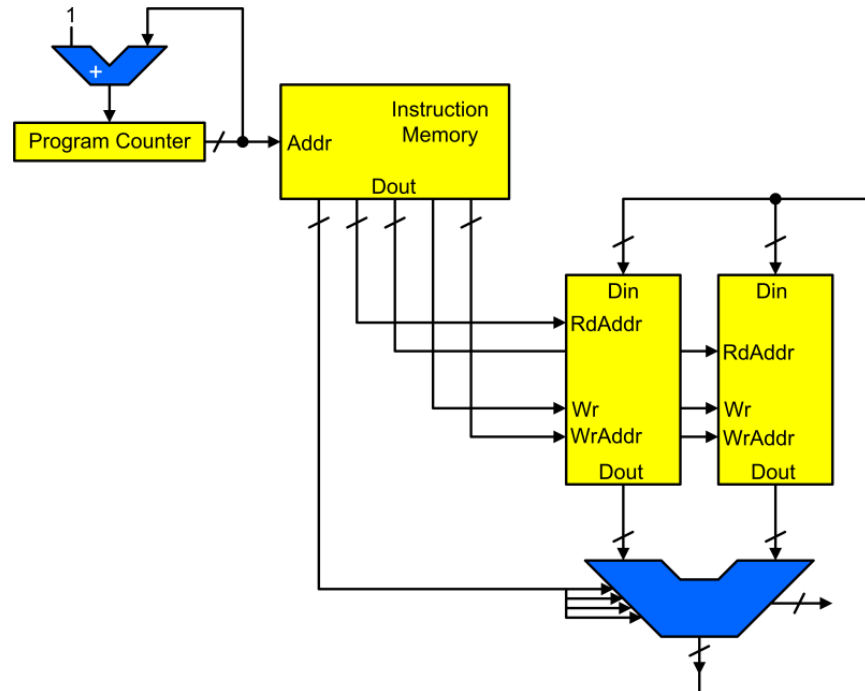
3: $A - 2B$

4: 00001111_2

5: ???

6: ???

7: $(A - 2B) \& 00001111_2$



A PROGRAMMABLE COMPUTER!

- Using our basic functional components, we are able to build a simple programmable computer!
- Implemented a computation using our processor's instruction set:
 1. Assigned memory locations to inputs and output
 2. Decomposed computation into processor instructions
 3. Assigned memory locations for intermediate values
 4. Encoded sequence of instructions for our program
- By feeding instructions to computer in sequence, we can perform our computation
 - Individual instructions communicate by reading and writing various memory locations

MACHINE CODE, ASSEMBLY LANGUAGE

- Our program:

000: 0001 001 001 1 010

001: 0011 000 010 1 011

010: 1000 011 100 1 111

- This is called machine code

- The actual data values that comprise the program
- Hard to read and write!

- Humans normally use assembly language

- A more human-readable language that is translated into machine code using an assembler

ADD R1, R1, R2 # R2 = 2B

SUB R0, R2, R3 # R3 = A - 2B

AND R3, R4, R7 # C = R3 & 00001111

- Allows human-readable names, operations, comments

C LOGICAL AND BITWISE OPERATIONS

- Before going forward, need to review what C offers for logical and bitwise operations
- C uses integers to represent Boolean values
 - 0 = false; any nonzero value = true
- Logical Boolean operators:
 - Logical AND: **a && b**
 - Logical OR: **a || b**
 - Logical NOT: **!a**
 - Result is 1 if true, 0 if false
- **&&** and **||** are short-circuit operators
 - Evaluated left-to-right
 - For **&&**, if LHS is false then RHS is not evaluated
 - For **||**, if LHS is true then RHS is not evaluated

C LOGICAL AND BITWISE OPERATIONS (2)

- C also has many bit-manipulation operations
- Given $a = 00010100_2$ (20_{10}), $b = 00110010_2$ (50_{10})
 - $a \ \& \ b = 00010000$ Bitwise AND
 - $a \ | \ b = 00110110$ Bitwise OR
 - $\sim a = 11101011$ Bitwise negation (invert)
 - $a \ ^ \ b = 00100110$ Bitwise XOR
- Note:
 - C has no way of specifying base-2 literals
 - Normal approach: use *hexadecimal* literals instead
- Hexadecimal: base-16 numbers
 - Digits are 0..9, A..F (or a..f, makes no difference)
 - A = 10, B = 11, ..., F = 15

HEXADECIMAL VALUES AND BIT-MASKS

- Example: **0x0F** is a hexadecimal literal in C
 - Each digit of a hexadecimal value represents 4 bits – a compact, simple way to write bit-fields
 - **0x0F** = 0000 1111 (also **0x0f**)
 - **0x03C7** = 0000 0011 1100 0111 (also **0x03c7**)
- Use bitwise AND to mask out or clear specific bits
 - **a & 0x0F**
 - Clears high nibble of **a**; retains low nibble of **a**
- Use bitwise OR to set specific bits
 - **a = a | 0x28**
 - Sets bits 3 and 5 of value in **a** (**0x28** = 0010 1000)
 - Other bits in **a** remain unchanged
- Use bitwise XOR to toggle specific bits
 - **a = a ^ 0x28**
 - Toggles bits 3 and 5 of value in **a**; other bits are left unchanged

C BIT-SHIFTING OPERATIONS

- C also includes bit-shifting operations
- Shift bits in **a** left by **n** bits: **a << n**
 - New bits on right are 0
 - Shifting left by n bits is identical to multiplying by 2^n

```
a = 42;          /* a = 00101010 = 42 */  
a = a << 1;      /* a = 01010100 = 84 */
```
- Shift bits in **a** right by **n** bits: **a >> n**
 - Shifting right by n bits is identical to dividing by 2^n
- Question: What should new bits on left be?
 - Depends on whether **a** is signed or unsigned!

```
a = -24;          /* Two's complement: 11101000 */  
a = a >> 1;      /* Should be -12 (11110100) now */
```
 - Leftmost bit represents sign
 - Preserve sign by using same value as original sign-bit

ARITHMETIC VS. LOGICAL SHIFT-RIGHT

- Distinguish between arithmetic shift-right and logical (i.e. bitwise) shift-right
 - Arithmetic shift-right preserves the value's sign
 - Logical shift-right always adds 0-bits to left of value
- Some languages make this distinction
 - Java: `>>` is arithmetic, `>>>` is logical
 - IA32 assembly: **SAR** is arithmetic, **SHR** is logical
- In C:
 - If argument is signed, shift-right is arithmetic

```
char a = -24;          /*      -24 = 11101000 */
printf("%d", a >> 1);  /* Prints -12 = 11110100 */
```
 - If argument is unsigned, shift-right is logical

```
/* Prints 116 = 01110100; topmost bit is 0 */
printf("%d", (unsigned char) a >> 1);
```

BIT-SHIFTS AND MASKS

- Can use bit-shifts with masks to extract sub-byte values
- `(a >> 4) & 0x0F`
 - Retrieves high nibble of **a**
- Does it matter if **a** is signed or unsigned?
 - Nope. We chop off the sign bit after we shift.

MULTIPLICATION?

- Our processor's instruction set:
- Hmm, no multiply instruction.
- No problem; implement multiply with addition and shifting

$$\text{mul}_w(a, b) = \sum_{i=0}^{w-1} a_i b 2^i$$

```
int mul(int a, int b) {  
    int p = 0;  
    while (a != 0) {  
        if (a & 1 == 1)  
            p = p + b;  
        a = a >> 1;  
        b = b << 1;  
    }  
    return p;  
}
```

Control	Operation	
0001	ADD	A B
0011	SUB	A B
0100	NEG	A
1000	AND	A B
1001	OR	A B
1010	XOR	A B
1011	INV	A
1100	SHL	A
1110	SHR	A

MULTIPLICATION ???

- Our multiply program:

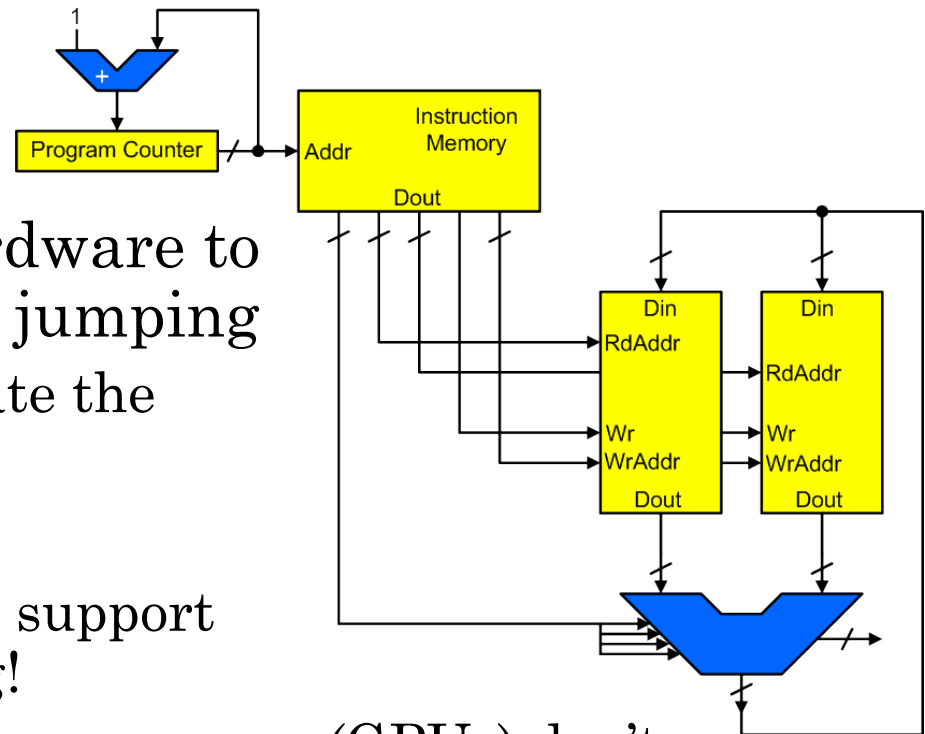
```
int mul(int a, int b) {  
    int res = 0;  
    while (a != 0) {  
        if (a & 1 == 1)  
            p = p + b;  
        a = a >> 1;  
        b = b << 1;  
    }  
    return p;  
}
```

Control	Operation	
0001	ADD	A B
0011	SUB	A B
0100	NEG	A
1000	AND	A B
1001	OR	A B
1010	XOR	A B
1011	INV	A
1100	SHL	A
1110	SHR	A

- Can we write a program to execute this code?
 - NO! ☹
 - Our processor doesn't support any branching or jumping operations

BRANCHING SUPPORT

- Our current processor architecture *can't* support branching or jumping!
 - Can only execute code in sequential order
- Need to extend the hardware to support branching and jumping
 - Need to be able to update the Program Counter field
- Note:
 - Not always essential to support branching and jumping!
 - Most dedicated graphics processors (GPUs) don't support looping or branching at all
 - However, is essential for a general-purpose processor



SUMMARY

- We designed a simple programmable computer!
 - Assembled functional components so we can perform a variety of simple computations
 - Feed instructions into our processor in sequence, from instruction memory
 - Instructions communicate by reading and writing various memory locations
- But, our computer has substantial limitations...
 - Can't even implement a simple loop yet. ☹
 - Need to extend our processor to support branching
- Also, our computer only has 8 bytes of memory
 - Need to examine impact of increasing memory size