



CS 24: INTRODUCTION TO COMPUTING SYSTEMS

Spring 2015

Lecture 3

LAST TIME

- Basic components of processors:
 - Buses, multiplexers, demultiplexers
 - Arithmetic/Logic Unit (ALU)
 - Addressable memory
- Assembled components into a simple processor
 - Can perform a wide range of operations, but all very simple
 - Need to string together a sequence of instructions, which communicate via memory locations
- Implemented a computation on this processor:
 - Assigned memory locations for inputs, outputs, constants
 - Decomposed the computation into steps the processor could actually handle
 - Assigned locations to intermediate values
 - Translated the program into machine-code instructions

MULTIPLICATION...

- Then, we wanted to implement multiplication:

```
int mul(int a, int b) {  
    int res = 0;  
    while (a != 0) {  
        if (a & 1 == 1)  
            p = p + b;  
        a = a >> 1;  
        b = b << 1;  
    }  
    return p;  
}
```

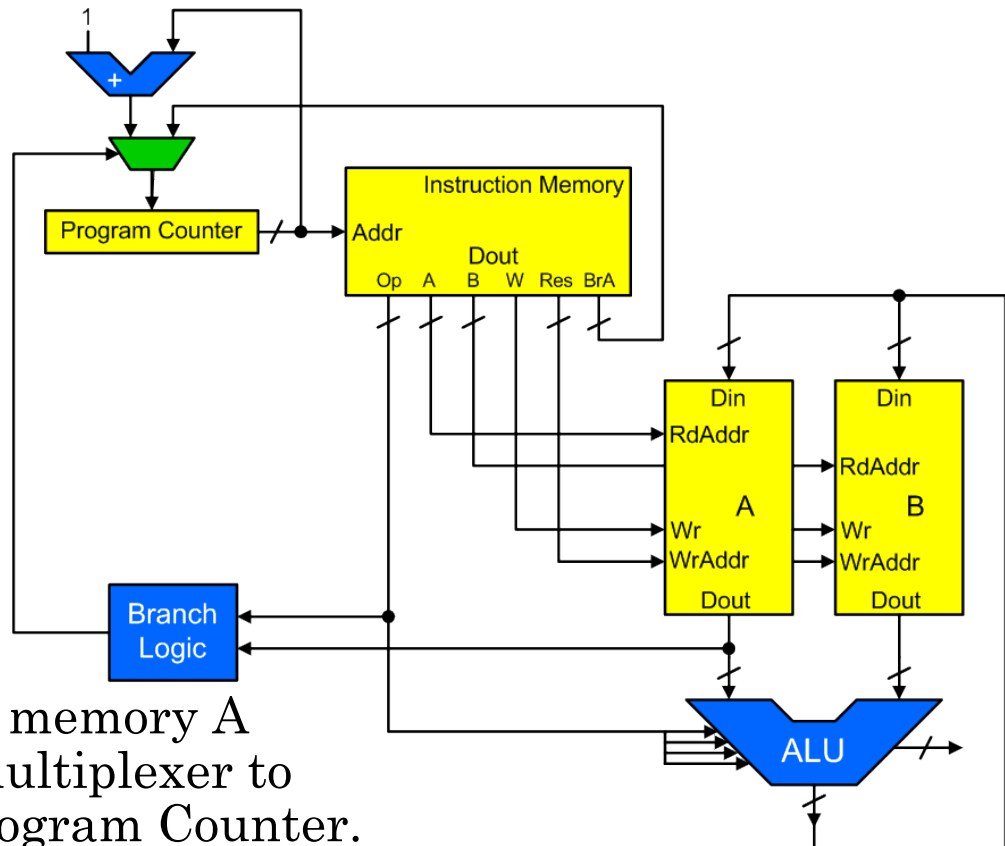
Control	Operation	
0001	ADD	A B
0011	SUB	A B
0100	NEG	A
1000	AND	A B
1001	OR	A B
1010	XOR	A B
1011	INV	A
1100	SHL	A
1110	SHR	A

- But, we couldn't write this program:
 - Our processor doesn't support branching operations!

UPDATE OUR ISA AND PROCESSOR

- Add a new instruction: BRZ A, Addr (Branch if Zero)
 - If value in slot A is 0, change Prog Ctr to address Addr

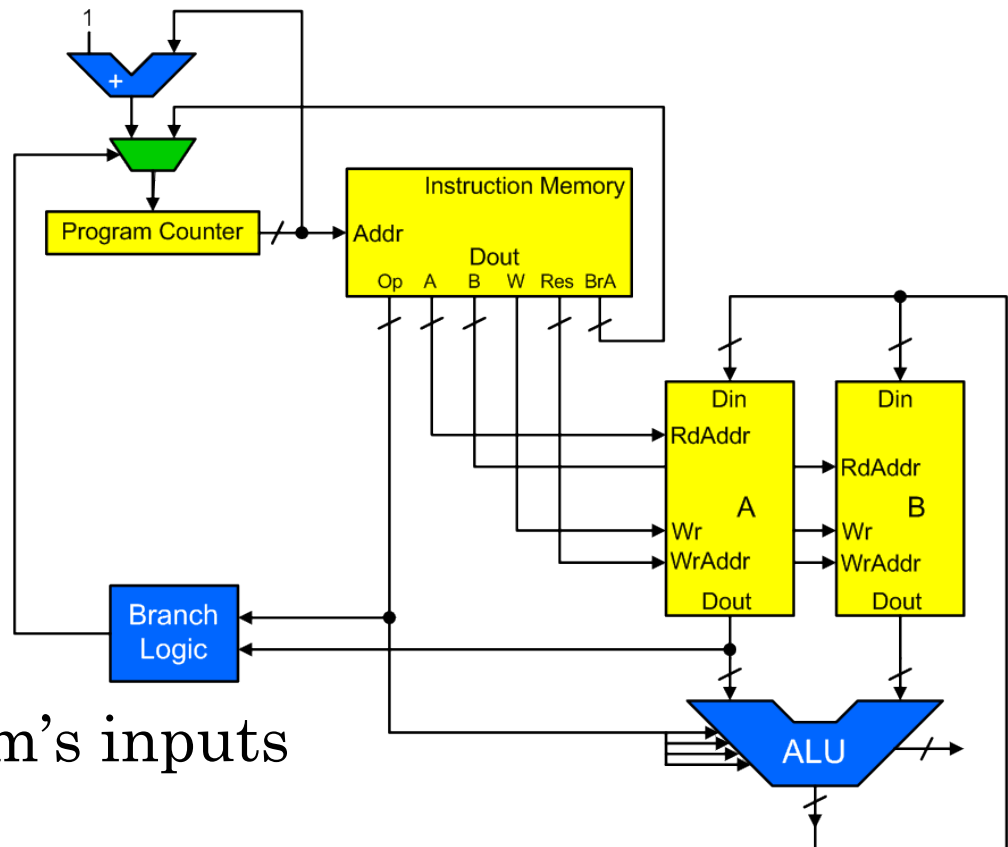
- New logic to support this instruction:



- Branch Logic:
 - If opcode is BRZ, and memory A outputs 0, then tell multiplexer to load Addr into the Program Counter.

UPDATED PROCESSOR

- With updated processor, can reuse instructions by creating loops in our programs
 - Perform many computations with just a few instructions
- Can now implement computations where number of steps is dependent on program's inputs



BACK TO MULTIPLICATION

- Should have enough capability now to encode our program

```
int mul(int a, int b) {  
    int p = 0;  
    while (a != 0) {  
        if (a & 1 == 1)  
            p = p + b;  
        a = a >> 1;  
        b = b << 1;  
    }  
    return p;  
}
```

- Coding is more complex now!
 - Need to plan out our loops...
 - Need to know the *addresses* to jump to!

Control	Operation	
0001	ADD	A B
0011	SUB	A B
0100	NEG	A
0111	BRZ	A Addr
1000	AND	A B
1001	OR	A B
1010	XOR	A B
1011	INV	A
1100	SHL	A
1110	SHR	A

WRITING OUR PROGRAM

- Use same general process as before
- Step 1: identify inputs and outputs

```
int mul(int a, int b) {  
    int p = 0;  
    while (a != 0) {  
        if (a & 1 == 1)  
            p = p + b;  
        a = a >> 1;  
        b = b << 1;  
    }  
    return p;  
}
```

- Inputs: A, B, some constant(s)
- Outputs: P

WRITING OUR PROGRAM (2)

- Step 2:
 - Decompose program into processor instructions.
- This step will be much more involved now:
 - May need quite a few temporary values
 - (Don't know how many though...)
 - Need to keep track of various addresses for branching
 - (Also don't know how many...)
- Strategy: use names for unknown values
 - Put in placeholders for temp-variables, places to jump
 - Write the code you *do* understand...
 - Once code is written, replace names with addresses

WRITING OUR PROGRAM (3)

○ Code:

```
int p = 0;
while (a != 0) {
    if (a & 1 == 1)
        p = p + b;
    a = a >> 1;
    b = b << 1;
}
```

Control	Operation	
0001	ADD	A B
0011	SUB	A B
0100	NEG	A
0111	BRZ	A Addr
1000	AND	A B
1001	OR	A B
1010	XOR	A B
1011	INV	A
1100	SHL	A
1110	SHR	A

○ First step: Set P = 0

○ Options?

- Subtract P from itself: **SUB P, P, P**
- XOR P with itself: **XOR P, P, P**

○ XOR option appears frequently in assembly code

WRITING OUR PROGRAM (4)

- Code:

```
int p = 0;
while (a != 0) {
    if (a & 1 == 1)
        p = p + b;
    a = a >> 1;
    b = b << 1;
}
```

- Steps:

```
XOR P, P, P
WHILE:
...
```

- Will certainly need to go back to start of loop
- Add a *label* to the code
 - Use the label in branching instructions
 - At end of translation, replace label with an actual address
- We already do this with our variables...

WRITING OUR PROGRAM (5)

- Code:

```
int p = 0;
while (a != 0) {

    if (a & 1 == 1)
        p = p + b;
    a = a >> 1;
    b = b << 1;
}
```

- Steps:

```
XOR P, P, P
WHILE:
    BRZ A, DONE
    ...

DONE:
```

- Don't know what will go inside loop yet, but we know we will need to exit it when finished!
- Add a **DONE** label, and use **BRZ** instruction to test A.

WRITING OUR PROGRAM (6)

- Code:

```
int p = 0;
while (a != 0) {
    if (a & 1 == 1)
        p = p + b;
    a = a >> 1;
    b = b << 1;
}
```

- Steps:

```
XOR P, P, P
WHILE:
    BRZ A, DONE
    AND A, 1, Tmp
    BRZ Tmp, SKIP
    ...
SKIP:
    ...
DONE:
```

- Compute $A \& 1$, then store in a temporary location
- Then, use branching instruction to skip body of if statement
 - (Can use **BRZ** since Tmp will be 0 or 1 here...)

WRITING OUR PROGRAM (7)

- Code:

```
int p = 0;
while (a != 0) {

    if (a & 1 == 1)

        p = p + b;

    a = a >> 1;
    b = b << 1;
}
```

- Steps:

```
XOR P, P, P
WHILE:
    BRZ A, DONE
    AND A, 1, Tmp
    BRZ Tmp, SKIP
    ADD P, B, P
SKIP:
    SHR A, A
    SHL B, B
    BRZ 0, WHILE
DONE:
```

- Remaining operations are mostly easy to write...
- Need an *unconditional* branching operation
 - Just use **BRZ** with a value of 0

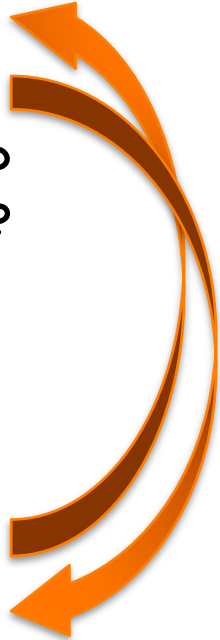
CONTROL-FLOW IN OUR PROGRAM

- Code:

```
int p = 0;
while (a != 0) {
    if (a & 1 == 1)
        p = p + b;
    a = a >> 1;
    b = b << 1;
}
```

- Steps:

```
XOR P, P, P
WHILE:
    BRZ A, DONE
    AND A, 1, Tmp
    BRZ Tmp, SKIP
    ADD P, B, P
SKIP:
    SHR A, A
    SHL B, B
    BRZ 0, WHILE
DONE:
```



- Manually implemented our **while** loop:
 - Test condition at start; exit if result is false
 - At end of loop, unconditionally return to start

CONTROL-FLOW IN OUR PROGRAM (2)

- Code:

```
int p = 0;
while (a != 0) {


    if (a & 1 == 1)

        p = p + b;

    a = a >> 1;
    b = b << 1;
}
```

- Steps:

```
XOR P, P, P
WHILE:
    BRZ A, DONE
    AND A, 1, Tmp
    BRZ Tmp, SKIP
    ADD P, B, P
SKIP:
    SHR A, A
    SHL B, B
    BRZ 0, WHILE
DONE:
```



- Also manually implemented the **if** statement:
 - Tested inverse of the condition, and skip over if-body if result is false

FINISHING OUR PROGRAM

- Our assembly code so far:
- Assign locations to values:

Address	Value
0	A
1	B
2	Tmp
3	1
4	0
5	
6	
7	P

```
XOR P, P, P
WHILE:
    BRZ A, DONE
    AND A, 1, Tmp
    BRZ Tmp, SKIP
    ADD P, B, P
SKIP:
    SHR A, A
    SHL B, B
    BRZ 0, WHILE
DONE:
```

- Note: two constants to include as well...

FINISHING OUR PROGRAM (2)

- Update our code with the memory addresses:

Address	Value
0	A
1	B
2	Tmp
3	1
4	0
5	
6	
7	P

```
        XOR 111, 111, 111
WHILE:
        BRZ 000, DONE
        AND 000, 011, 010
        BRZ 010, SKIP
        ADD 111, 001, 111
SKIP:
        SHR 000, 000
        SHL 001, 001
        BRZ 100, WHILE
DONE:
```

FINISHING OUR PROGRAM (3)

- Now to figure out the instruction addresses

WHILE = address 1

SKIP = address 5

DONE = address 8

```
0: XOR 111, 111, 111
WHILE:
1: BRZ 000, DONE
2: AND 000, 011, 010
3: BRZ 010, SKIP
4: ADD 111, 001, 111
SKIP:
5: SHR 000, 000
6: SHL 001, 001
7: BRZ 100, WHILE
DONE:
8:
```

FINISHING OUR PROGRAM (3)

- Update code with the instruction addresses

WHILE = address 1

SKIP = address 5

DONE = address 8

```
0: XOR 111, 111, 111
1: BRZ 000, 1000
2: AND 000, 011, 010
3: BRZ 010, 0101
4: ADD 111, 001, 111
5: SHR 000, 000
6: SHL 001, 001
7: BRZ 100, 0001
8:
```

FINISHING OUR PROGRAM (4)

- Finally, encode each instruction into machine code

Control	Operation	
0001	ADD	A B
0011	SUB	A B
0100	NEG	A
0111	BRZ	A Addr
1000	AND	A B
1001	OR	A B
1010	XOR	A B
1011	INV	A
1100	SHL	A
1110	SHR	A

XOR 111, 111, 111	1010 111 111 111 0000
BRZ 000, 1000	0111 000 000 000 1000
AND 000, 011, 010	1000 000 011 010 0000
BRZ 010, 0101	0111 010 000 000 0101
ADD 111, 001, 111	0001 111 001 111 0000
SHR 000, 000	1110 000 000 000 0000
SHL 001, 001	1100 001 000 001 0000
BRZ 100, 0001	0111 100 000 000 0001

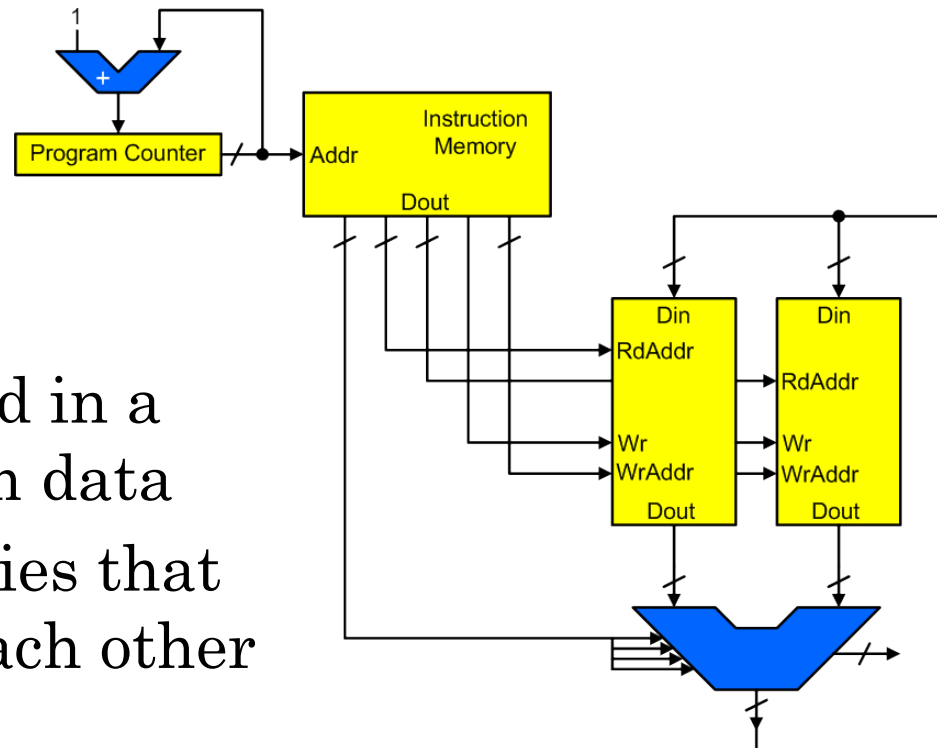
THE FINAL PROGRAM

- Our final machine code:
 - Unused/ignored bits are grayed out
- 136 bits; 44 bits unused
 - 30% of our program's space is unused!
- This instruction encoding is not very efficient
- Many processors employ variable-length instruction encodings to reduce program size
 - Particularly common in CISC processors
 - Increases complexity in instruction processing logic

1010	111	111	111	0000
0111	000	000	000	1000
1000	000	011	010	0000
0111	010	000	000	0101
0001	111	001	111	0000
1110	000	000	000	0000
1100	001	000	001	0000
0111	100	000	000	0001

PROCESSOR MEMORIES

- Our processor has a curious memory layout
 - Doesn't seem much like our modern computers...



- Instructions are stored in a separate memory from data
- Uses two data memories that are mirror copies of each other

MEMORY ARCHITECTURES

- Several different memory architectures have been employed in computer systems
- Harvard Architecture:
 - Instruction memory and data memory are separate
 - May even have different word sizes from each other
 - Instruction memory is read-only
 - Data memory is read/write
 - *Cannot* treat instructions as data.
 - Constants must be explicitly loaded into data memory
- Named after Harvard Mark I computer
 - A relay-based computer with separate instruction and data memories
- Our example processor uses a Harvard architecture

MEMORY ARCHITECTURES (2)

- Von Neumann Architecture:
 - Instructions and data are stored in a single read/write memory
- A big benefit over Harvard architecture:
 - Much easier to load and manage programs on the computer
- Some big problems too:
 - Programs can easily corrupt their own code!
 - (Many exploits take advantage of this characteristic...)
 - CPU only has one bus to read both instructions and data
 - Harvard architecture has two buses; can use them in parallel
- Modern processors often blend these approaches
 - Overall system design uses von Neumann architecture
 - Internally, CPU has separate instruction, data memories
 - Used in very different ways, so separating them allows for much greater hardware optimization

INDIRECT MEMORY ACCESS

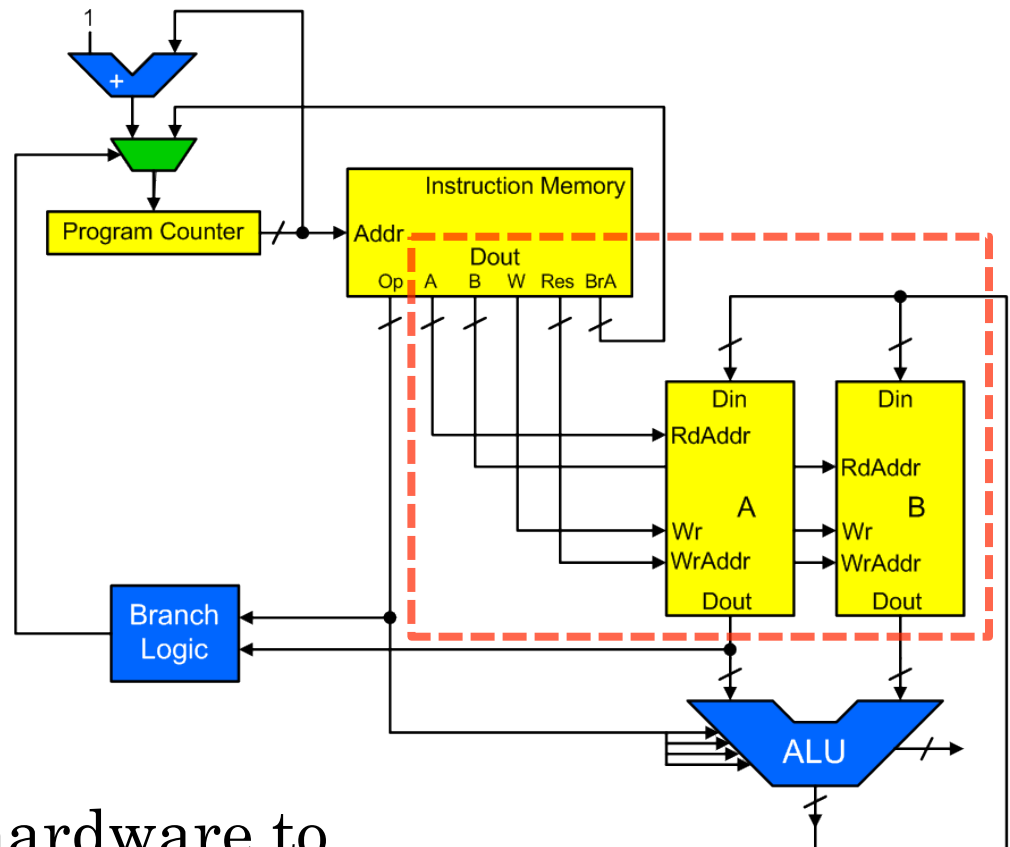
- Still plenty of problems our simple processor design can't handle
- Example function:
 - Add two vectors together, component by component, storing the result into a third vector
 - Vectors to use are arguments to the function
- C code:

```
void vector_add(int *a, int *b, int *r, int length) {  
    int i;  
    for (i = 0; i < length; i++)  
        r[i] = a[i] + b[i];  
}
```
- This code is *independent* of where **a**, **b**, **r** are stored
 - Code can be reused on *any* set of three vectors
 - Addresses of the vectors are specified in the arguments

INDIRECT MEMORY ACCESS (2)

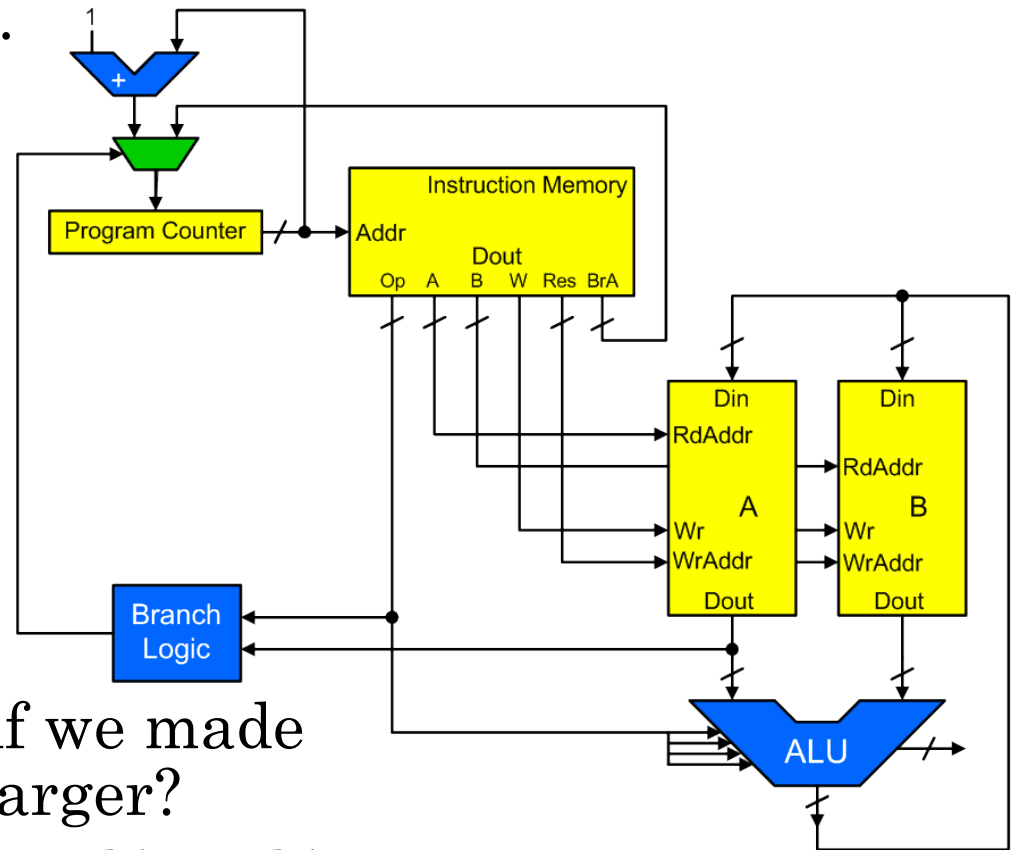
○ Problem:

- This processor can only specify literal address-values in the instructions
- Need to introduce a way to get memory addresses from variables instead...
- Need to update the hardware to support *indirect* memory access



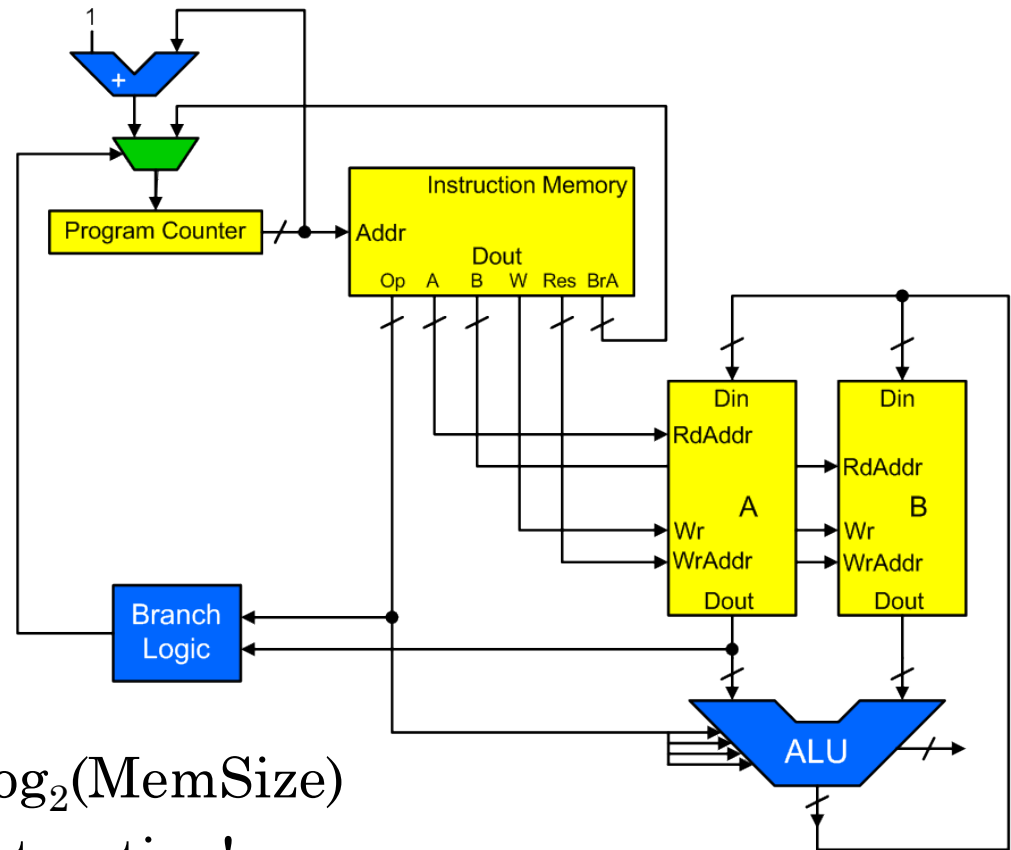
LARGE MEMORIES

- But, before we talk about new functionality, this design has a problem.
- Our current memory only has 8 locations
 - Not very useful...
- Really need *much* larger memories than this!
- What would happen if we made our memories *much* larger?
 - e.g. 4GB, like modern 32-bit architectures



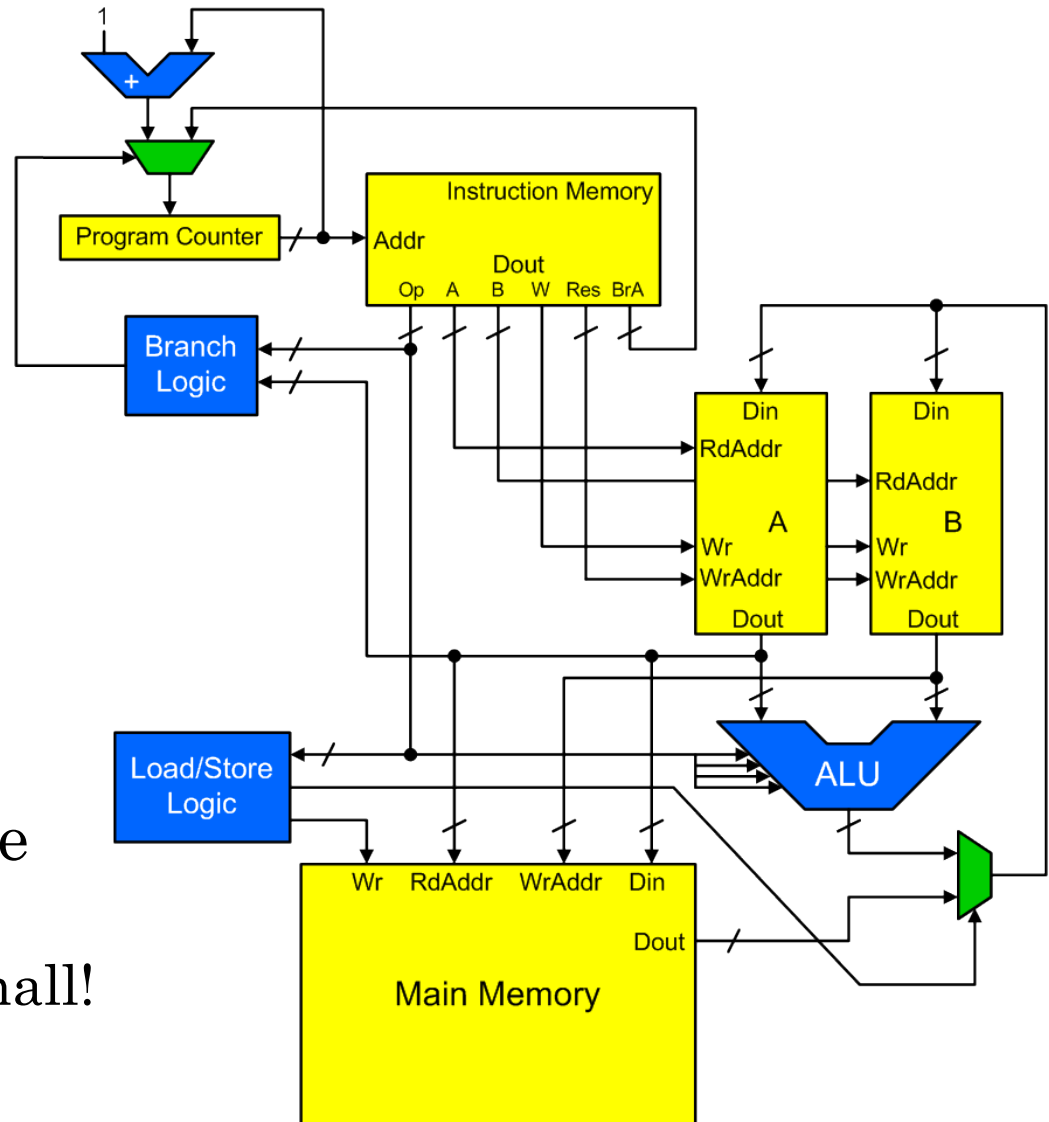
LARGE MEMORIES AND INSTRUCTION SIZE

- We *directly encode* memory addresses into our instructions.
 - Memory has 8 slots
 - $\log_2(8) = 3$ bits for memory addresses
- If memory is 4GB in size, we need 32 bits to specify addresses!
- Instruction size \approx
 - $\log_2(\text{NumOps}) + 3 \times \log_2(\text{MemSize})$
 - ≈ 100 bits for each instruction!



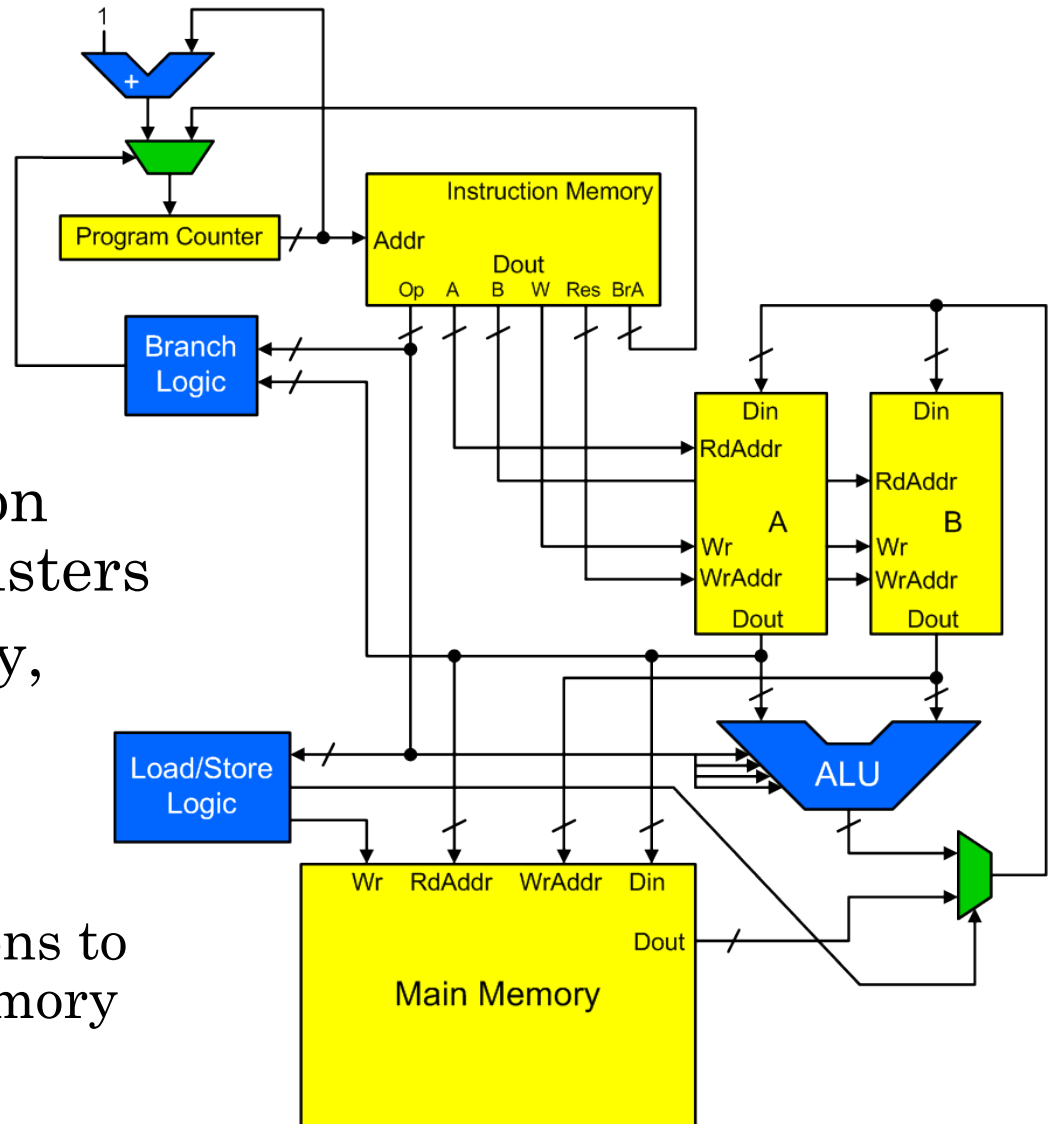
SUPPORTING LARGE MEMORIES

- A better design for large memories:
- Small memory is called a *register file*
 - Individual locations are called *registers*
 - Small number of registers (e.g. 8, 64)
- Instructions only use register addresses
 - Instructions stay small!



SUPPORTING LARGE MEMORIES (2)

- A better design for large memories:
- ALU only operates on values stored in registers
- To use main memory, must explicitly load or store values into registers
 - Need new instructions to work with main memory



LOADING VALUES FROM MAIN MEMORY

- To load values from main memory, introduce:

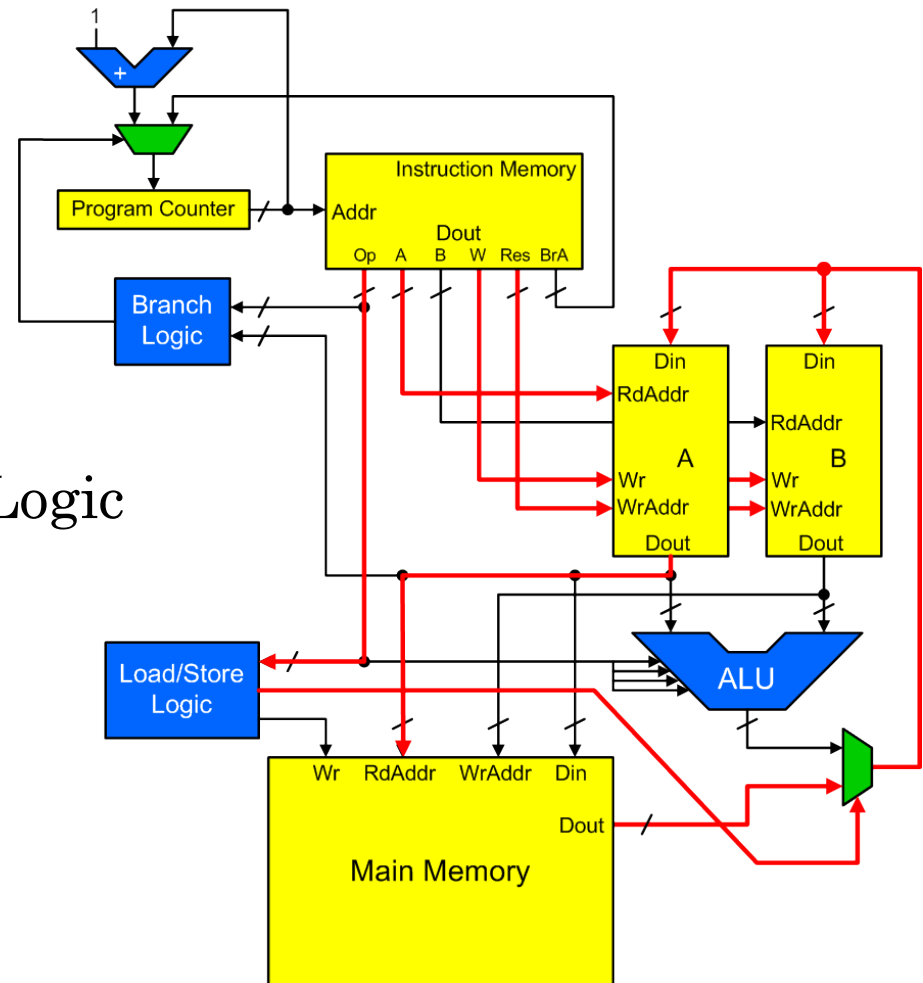
- **LD SRC, DST**
- SRC, DST are registers

- Meaning:

- $DST = Memory[Src]$

- **LD** activates Load/Store Logic

- SRC value fed to RdAddr of main memory
- Output of main memory fed to DST register
- ALU is not utilized at all



STORING VALUES TO MAIN MEMORY

○ To store values to main memory, introduce:

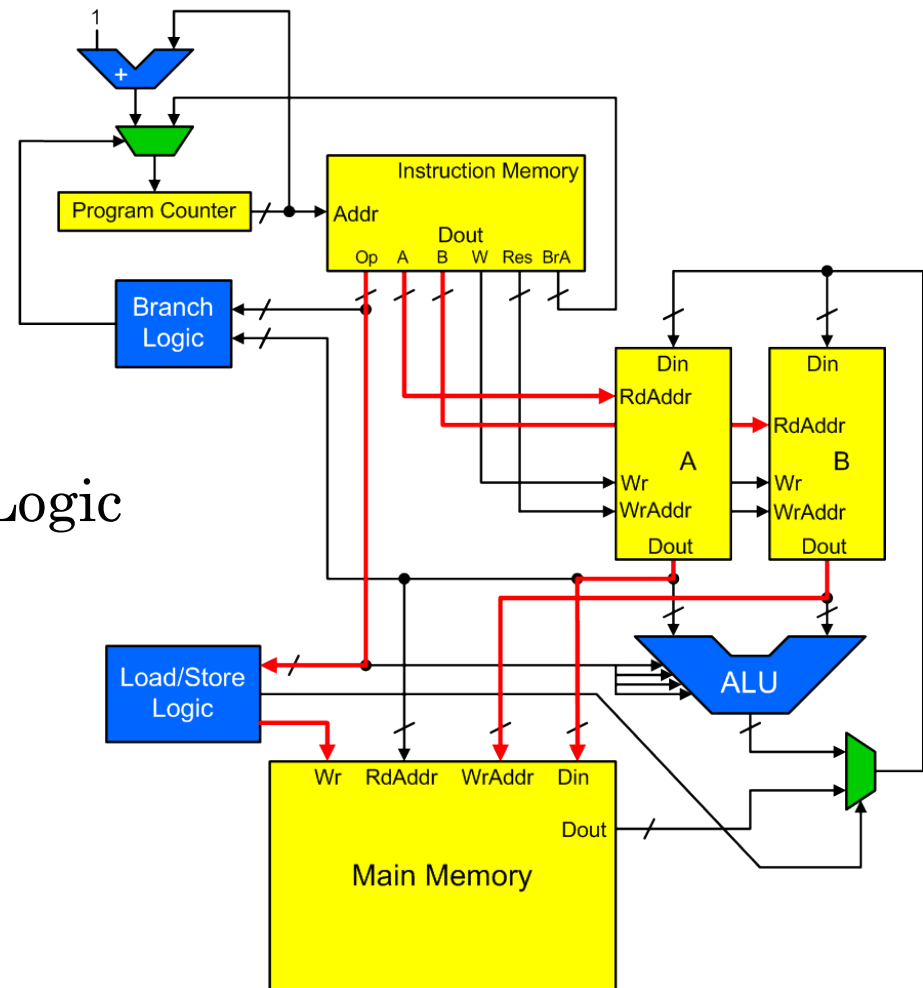
- **ST SRC, DST**
- Again, SRC, DST are registers

○ Meaning:

- $\text{Memory}[\text{DST}] = \text{SRC}$

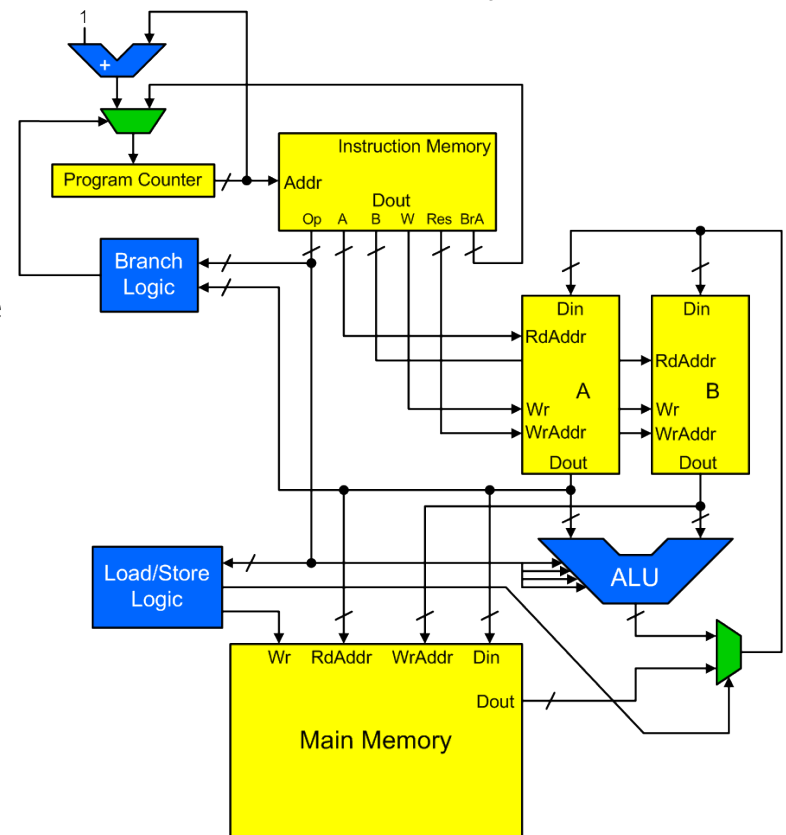
○ **ST** activates Load/Store Logic

- DST value fed to WrAddr of main memory
- SRC register fed to input of main memory
- Again, ALU not utilized



LOAD/STORE ARCHITECTURE

- This is called a Load/Store Architecture
 - ALU only operates on registers
 - Must explicitly load/store values to main memory
 - (Often seen in RISC ISAs)
- Benefits:
 - Very simple architecture
 - All instructions are same size
 - Instructions execute quickly!
- Drawbacks:
 - Memory-intensive programs require many operations to implement!



EXAMPLE: VECTOR-ADD FUNCTION

- Consider the body of our vector-add function:

```
int i;  
for (i = 0; i < length; i++)  
    r[i] = a[i] + b[i];
```

- Arguments **a**, **b**, **r** are addresses of the *start* of their respective arrays

- One implementation of $r[i] = a[i] + b[i]$:

```
ADD A, I, ADDR      # Compute address of a[i]  
LD  ADDR, AI        #      and retrieve it.  
ADD B, I, ADDR      # Compute address of b[i]  
LD  ADDR, BI        #      and retrieve it.  
ADD AI, BI, RI      # Compute a[i] + b[i].  
ADD R, I, ADDR      # Compute address of r[i]  
ST  RI, ADDR        #      and store sum.
```

- Hope the processor can make this fast, somehow...*

ALTERNATIVE TO LOAD/STORE

- Instead of Load/Store architecture, we can also support a rich set of operand-types

- For example, an operand could be:

- A constant (i.e. an immediate value)
- A register

ADD 14, R1, R2 # R2 = 14 + R1

- A direct address

ADD [150], 14, R7 # R7 = Memory[150] + 14

- An indirect address

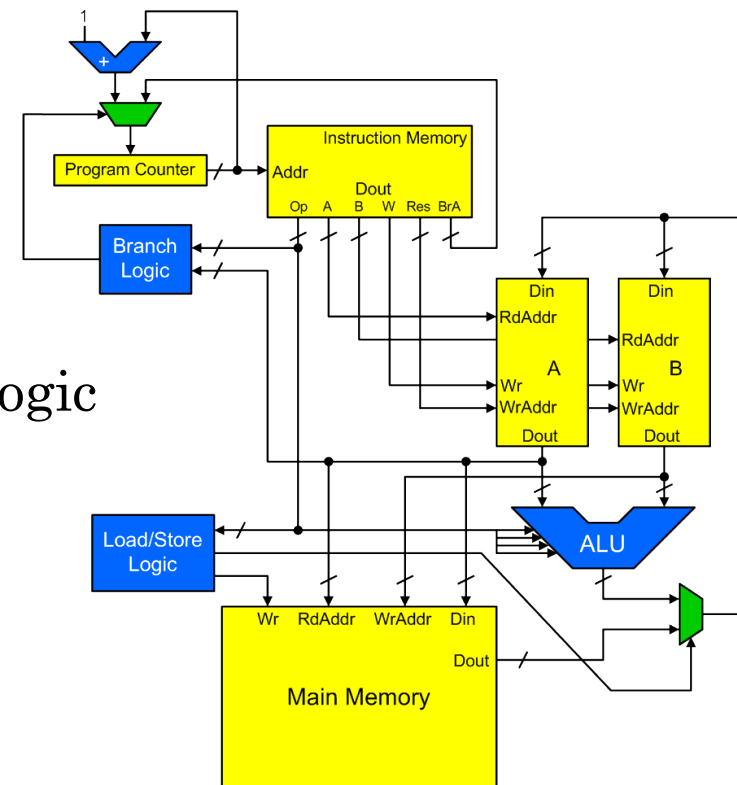
```
ADD [R5 + 150], 14, R6 # R6 = Memory[R5+150] + 14
```

- Instruction encodings may need to include:

- Immediate value, Register index, Address, Address + Register index
- Instruction encodings are variable-length, and must also include operand-mode bits to indicate types of operands

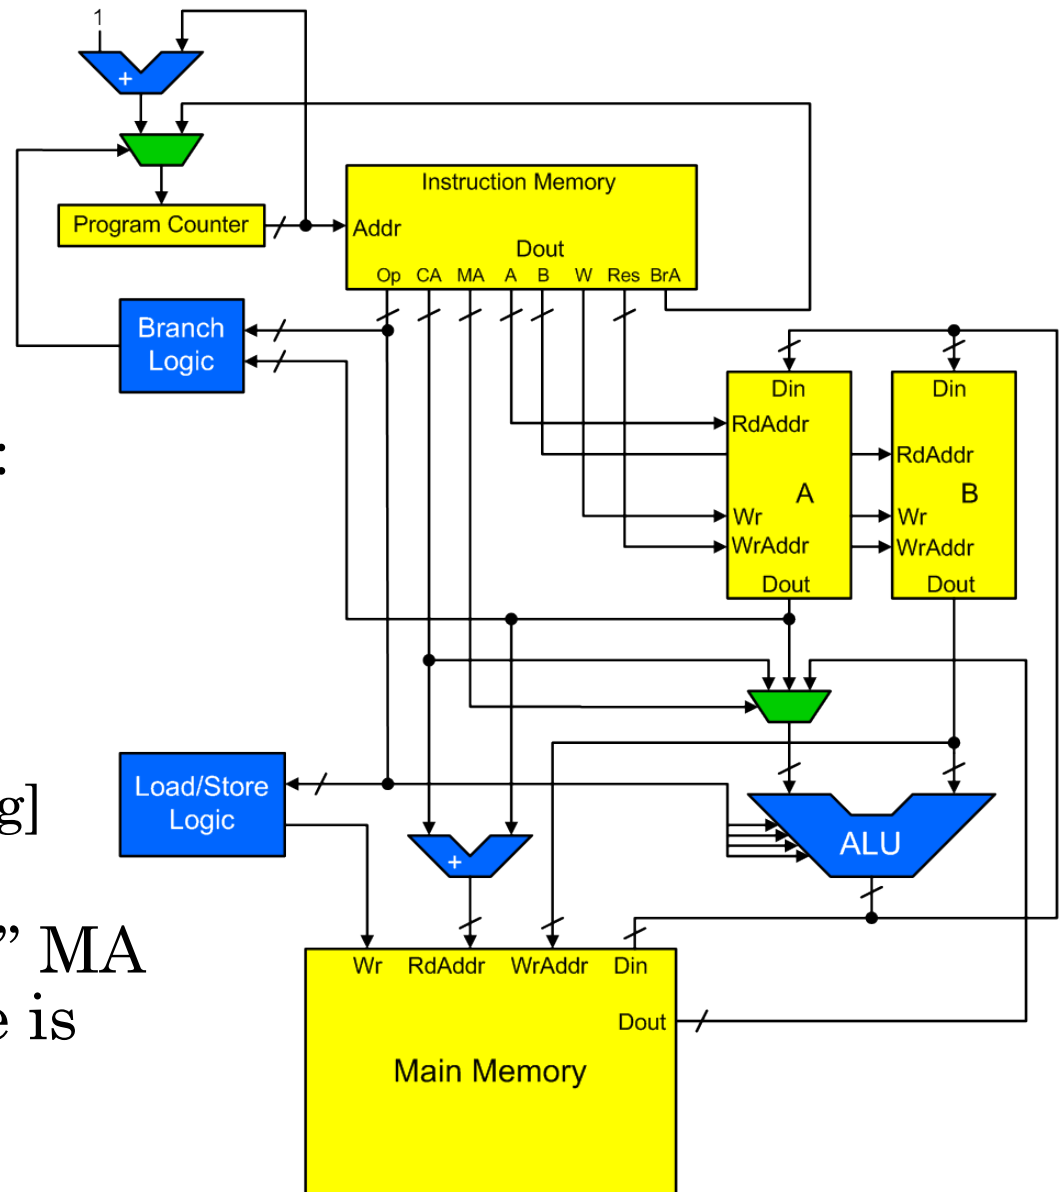
SUPPORTING VARIOUS OPERAND TYPES

- Instruction encodings can include:
 - Immediate value, Register index, Address, Address + Register index
- Clearly, we will need new logic to feed the ALU!
- Also, instructions become much more complicated
 - (CISC processors employ multiple operand types)
 - Requires complex, dedicated instruction-fetch and decode logic
- Finally, may need to do arithmetic on addresses fed to the main memory



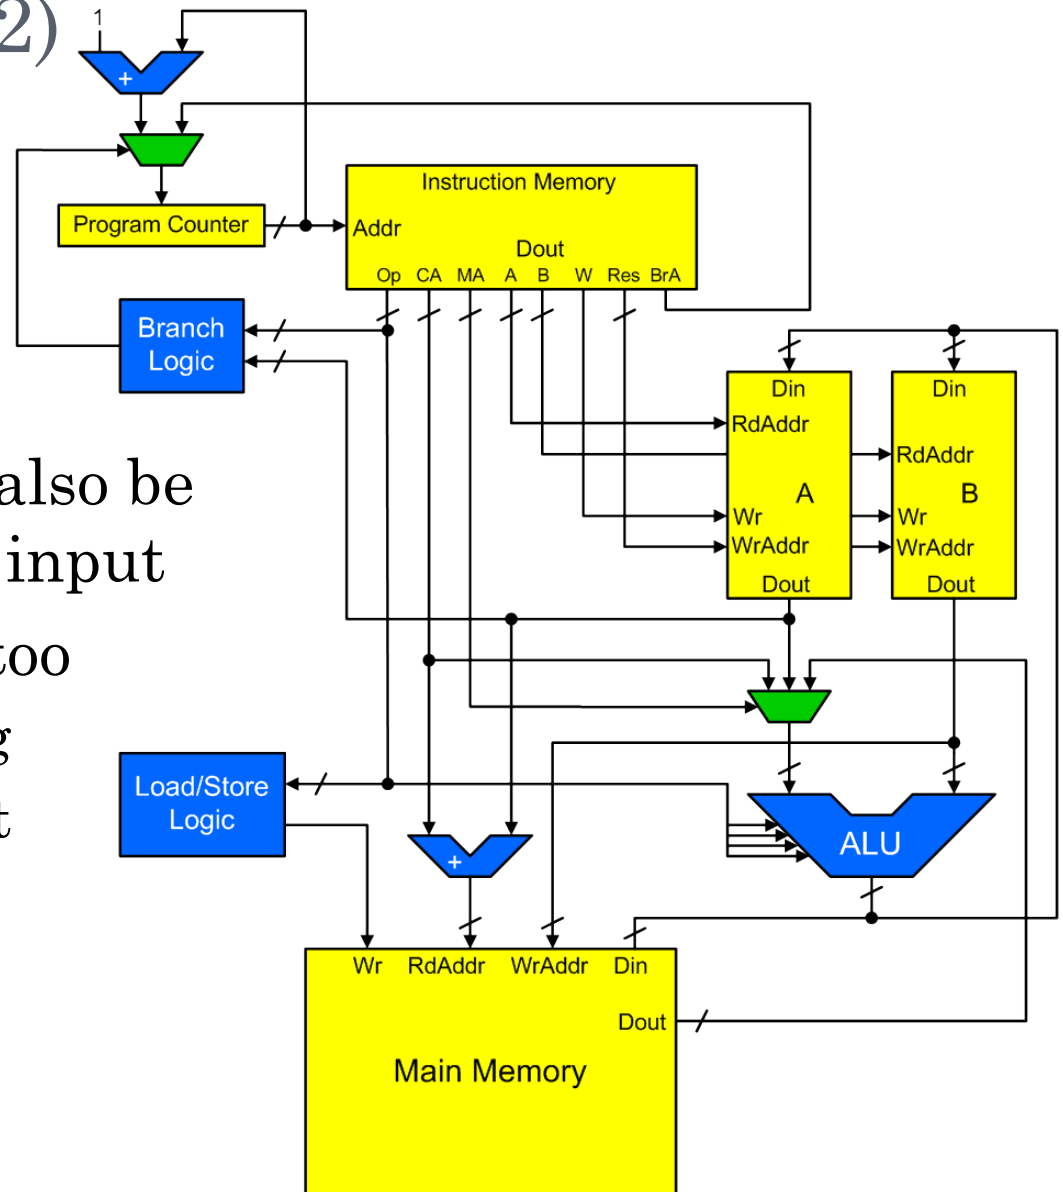
OPERAND TYPES

- Example logic for supporting multiple operand types:
- Operand A supports three operand types:
 - Constant
 - Encoded in CA value
 - Register
 - Encoded in A value
 - Memory[Const + Reg]
 - Uses both A and CA
- New “operand mode” MA controls which value is fed into the ALU



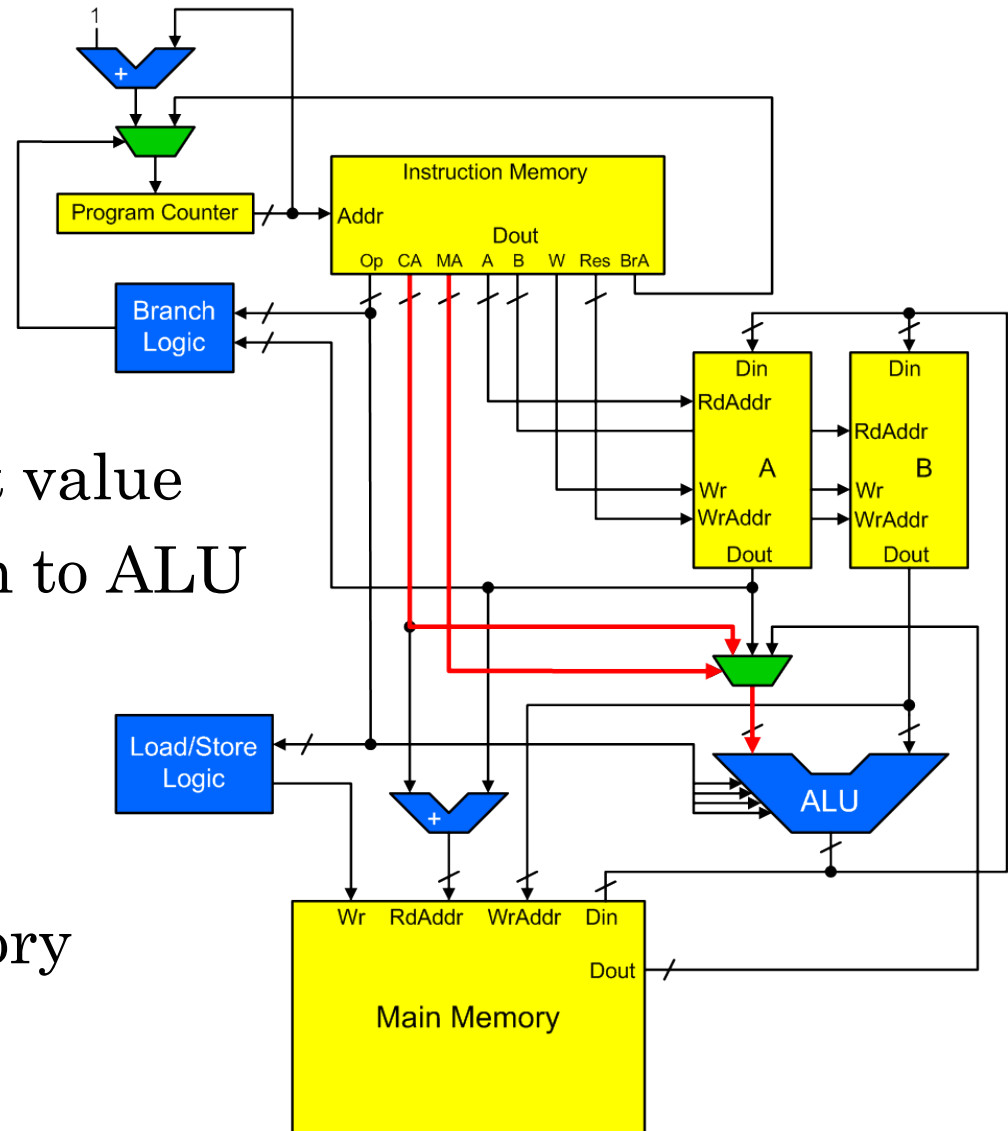
OPERAND TYPES (2)

- Example logic for supporting multiple operand types:
- Similar logic would also be added to other ALU input
- Many other details too
 - Instruction decoding
 - Routing ALU output to registers/memory
 - etc.



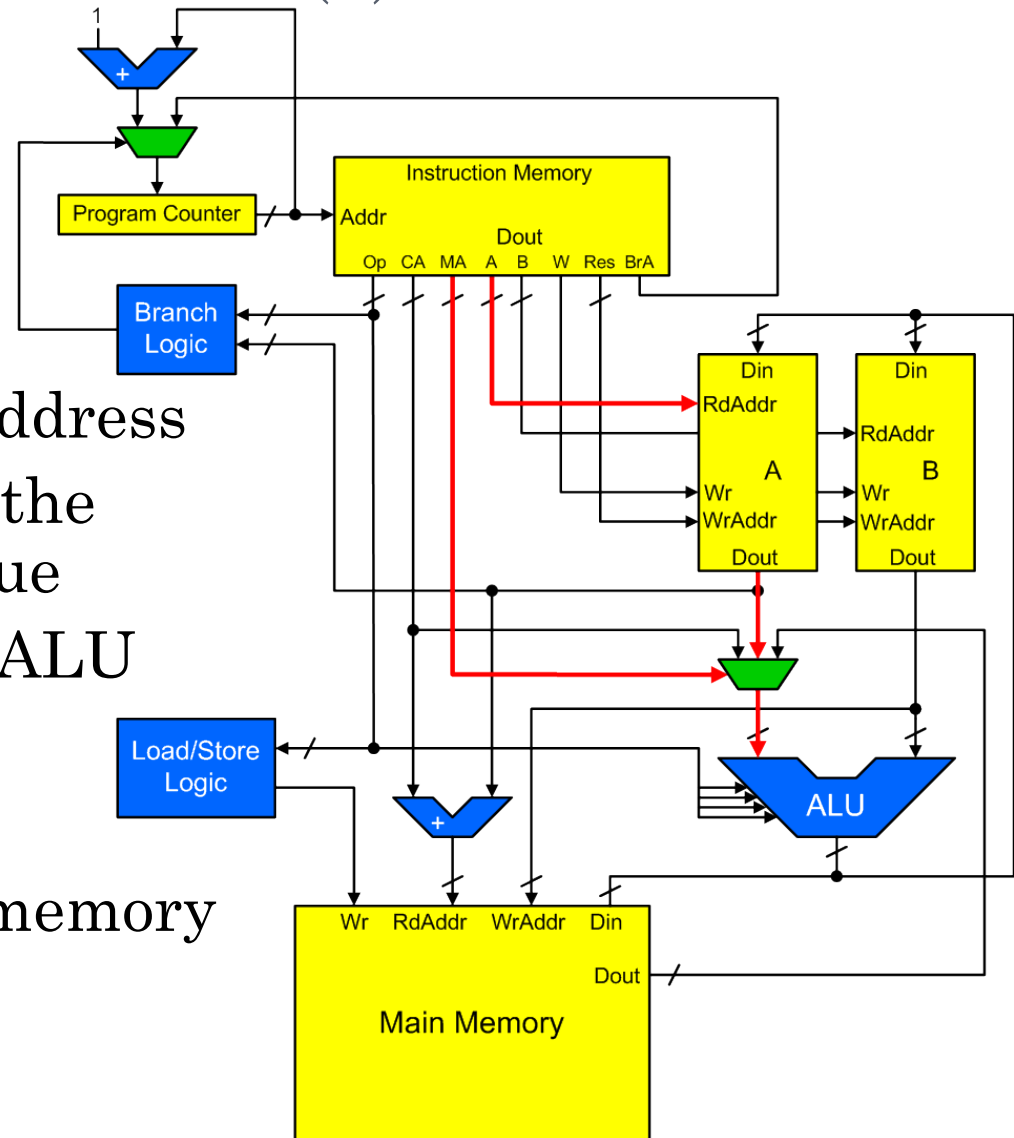
EXAMPLE OPERAND TYPES

- Constant operand:
- CA specifies constant value
- MA feeds CA through to ALU
- Register A and memory are both unused



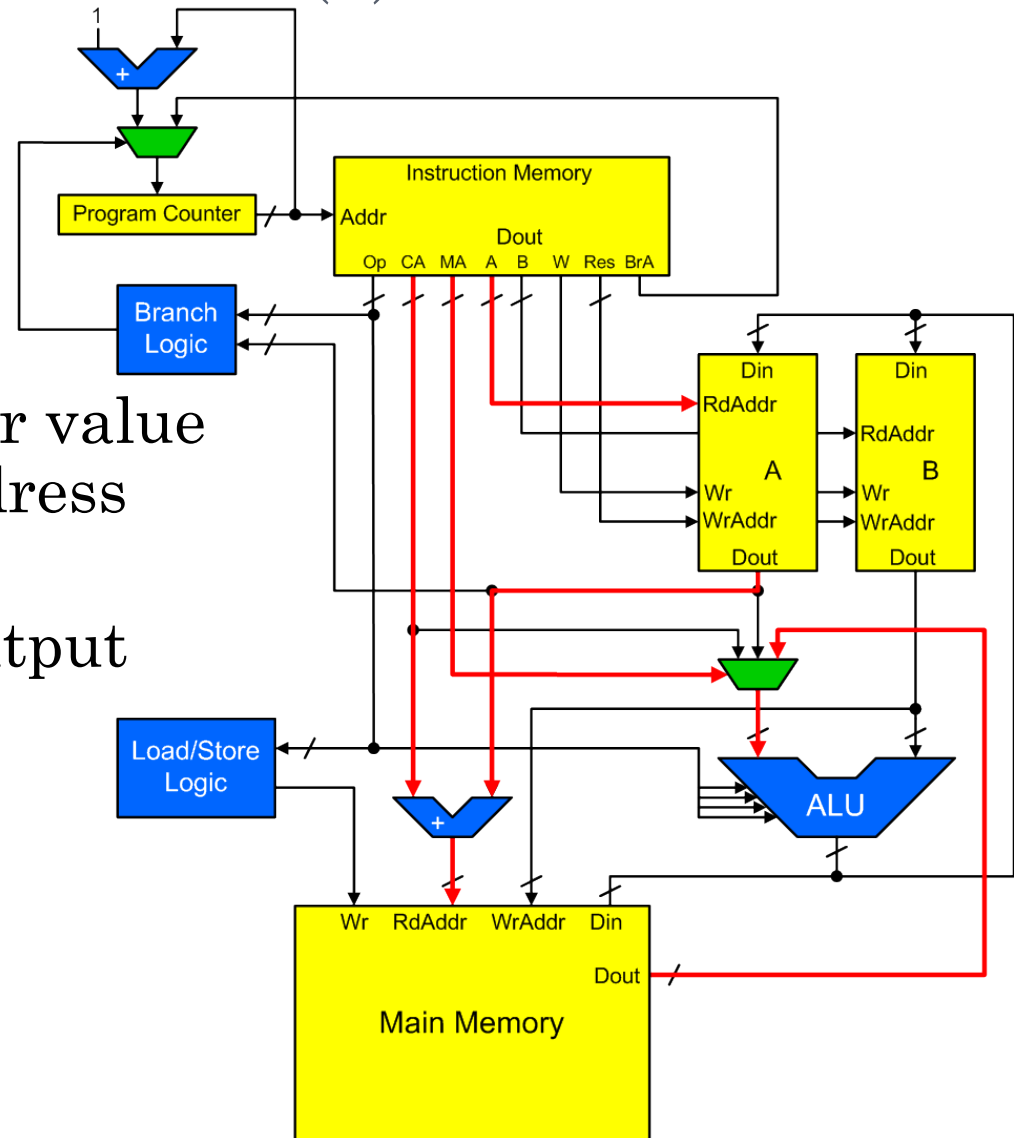
EXAMPLE OPERAND TYPES (2)

- Register operand:
- A specifies register address
- Register file outputs the specified register value
- MA feeds register to ALU
- Constant value and memory are both unused



EXAMPLE OPERAND TYPES (3)

- Indirect memory access operand:
- Constant and register value added to produce address for main memory
- MA feeds memory output to ALU



SUMMARY

- Added branching to our processor
 - Can implement larger computations with fewer instructions
 - Reuse instructions on different data by looping
 - Tailor computations based on the specific data values
 - Allows us to perform computations where the work performed is proportional to the input values
- Explored various memory architecture details
 - Harvard vs. Von Neumann architectures
 - Separating instruction and data processing paths
 - Load/Store architecture vs. multiple operand types
 - Allows us to implement computations that are independent of the specific memory location
 - Allows us to access *much* larger memories, using a relatively small instruction set

NEXT TIME

- Start looking at the Intel IA32 instruction set
- Start writing programs you can run on *real* computers! 😊
- Begin to develop abstractions to facilitate construction of larger programs
 - Subroutines, stacks, recursion