



CS24: INTRODUCTION TO COMPUTING SYSTEMS

Spring 2015

Lecture 6

LAST TIME: CDECL

- How to implement basic C abstractions in IA32?
 - C subroutines with arguments, local/global variables
- Began discussing the *cdecl* calling convention
 - Widely used on *NIX systems running on x86
- Both the procedure caller and the callee have to coordinate the operation!
 - Shared resources: the stack, the register file
- Calling convention specifies:
 - Who sets up which parts of the call
 - What needs to be saved, and by whom
 - How to return values back to the caller
 - Who cleans up which parts of the call

CDECL CHEAT SHEET (1)

- Caller pushes arguments in reverse order
- Caller uses **call** to invoke subroutine
- Callee pushes caller's **%ebp** onto stack, then sets **%ebp = %esp**

```
pushl    %ebp  
movl     %esp, %ebp
```

- Arguments are at positive offsets from **%ebp**

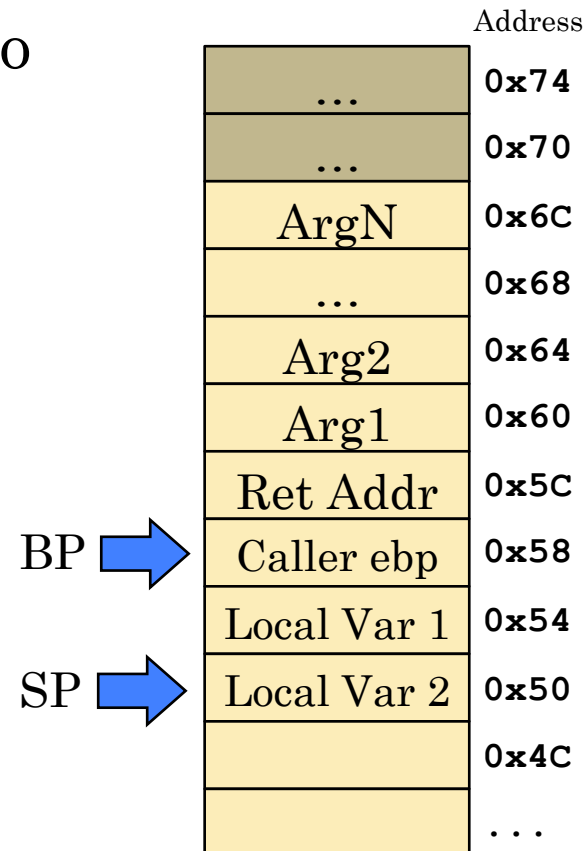
8(%ebp) = Arg1

12(%ebp) = Arg2

- Local variables at negative offsets from **%ebp**

-4(%ebp) = Local Var 1

-8(%ebp) = Local Var 2

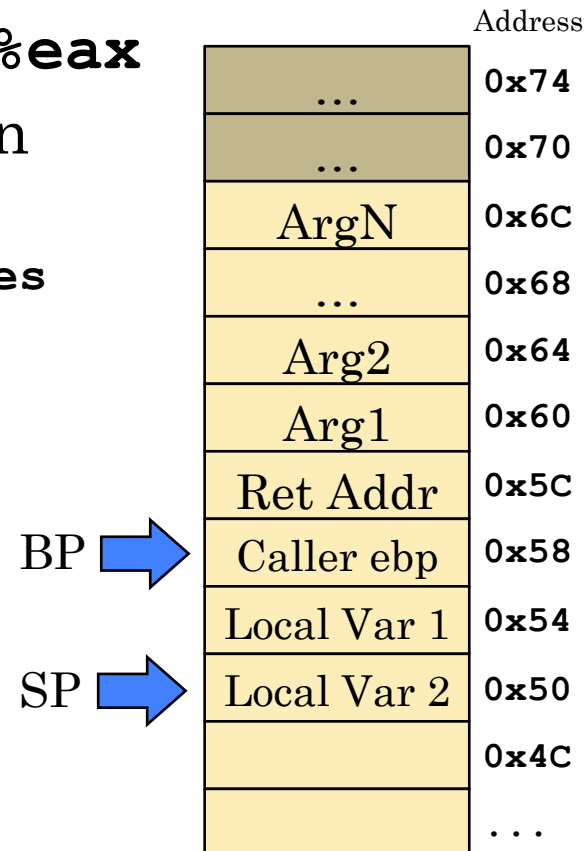


CDECL CHEAT SHEET (2)

- Caller-save registers: **%eax, %ecx, %edx**
- Callee-save registers: **%ebp, %ebx, %esi, %edi**
- Callee saves return-value into **%eax**
- Callee restores stack state, then uses **ret** to return

```
# Also discards local variables
movl %ebp, %esp
popl %ebp
ret
```

- Caller removes arguments from the stack
 - Either using **pop** instructions, or by adding a constant to **%esp**



BACK TO OUR EXAMPLE C PROGRAM

- A simple accumulator:
- Uses a global variable to store current value
- Functions to update accumulator, or reset it
- Main function to exercise the accumulator

- *How is this program implemented in IA32?*

```
int value;

int accum(int n) {
    value += n;
    return value;
}

int reset() {
    int old = value;
    value = 0;
    return old;
}

int main() {
    int i, n;

    reset();
    for (i = 0; i < 10; i++) {
        n = rand() % 1000;
        printf("n = %d\taccum = %d\n",
               n, accum(n));
    }
    return 0;
}
```

OUR EXAMPLE PROGRAM

- Can look at **gcc** assembly language output for our accumulator example
 - **gcc -O2 -S main.c**
 - **-S** generates assembly output, not a binary file
 - Result is in **main.s**
 - **-O2** applies some optimizations to generated code
 - Otherwise, assembly output includes some pretty silly code
- Results vary *widely* based on target platform!
 - We will look at Linux **gcc** output
 - (MacOS X output is *very* different... Ask if you want an explanation of what's going on. It's very cool!)

GENERATED ASSEMBLY CODE

- Some of the output:
- Lines starting with `.` are assembler directives
 - e.g. `.text` tells assembler to generate machine code for instructions that follow
- Lines with a colon are labels
 - e.g. `accum`, `reset` are labels specifying start of our functions
- `.size` directive specifies the size of various symbols, in bytes
 - `accum` = address of function's start
 - `.` = current address
 - `.-accum` is size of fn. body

```
.file    "amain.c"
        .text
        .p2align 4,,15
.globl accum
        .type    accum, @function
accum:
        pushl    %ebp
        movl     %esp, %ebp
        movl     8(%ebp), %eax
        addl     value, %eax
        popl     %ebp
        movl     %eax, value
        ret
        .size    accum, .-accum
        .p2align 4,,15
.globl reset
        .type    reset, @function
reset:
        pushl    %ebp
        movl     value, %eax
        ...
```

GLOBAL VARIABLES

- End of our output:

```
...  
    popl    %esi  
    popl    %ebp  
    leal    -4(%ecx), %esp  
    ret  
    .size   main, .-main  
    .comm   value,4,4  
    .ident  "GCC: (GNU) ...  
    .section                               ...
```

- **.size main, .-main**
is end of **main()** function

- Global variable **value** specified with **.comm** directive
 - A “common symbol,” possibly shared across multiple files
 - Specifies name, size, optional alignment of variable
 - Address is assigned when assembling the code
 - The actual memory is uninitialized

ACCUMULATOR CODE

- Accumulator function:

```
int accum(int n) {  
    value += n;  
    return value;  
}
```

- Translated into:

accum:

```
pushl    %ebp  
movl     %esp, %ebp  
movl     8(%ebp), %eax  
addl     value, %eax  
popl     %ebp  
movl     %eax, value  
ret
```


Set up stack frame

Move n into eax

eax += value

Restore caller ebp

Store updated value



RESET CODE

- Reset function:

```
int reset() {  
    int old = value;  
    value = 0;  
    return old;  
}
```

- Translated into:

```
reset:  
    pushl    %ebp  
    movl     value, %eax    # Result goes into eax  
    movl     %esp, %ebp  
    popl     %ebp  
    movl     $0, value  
    ret
```

- Clearly involves some unnecessary steps...
 - ebp** isn't used at all! Could reduce down to 3 instructions.

MAIN FUNCTION

- Main function code:

```
int main() {  
    int i, n;  
  
    reset();  
  
    for (i = 0; i < 10; i++) {  
        n = rand() % 1000;  
        printf("n = %d\taccum = %d\n",  
              n, accum(n));  
    }  
  
    return 0;  
}
```

MAIN FUNCTION (2)

- Main function code:

```
int main() {  
    int i, n;  
  
    reset();  
    ...
```

- Assembly code:

```
main:  
    leal    4(%esp), %ecx    # Stack init:  aligns stack with  
    andl    $-16, %esp      #    16-byte boundary, then  
    pushl   -4(%ecx)        #    copies return-addr to TOS.  
    pushl   %ebp            # Set up stack frame  
    movl    %esp, %ebp  
    pushl   %esi            # Callee-save registers  
    xorl    %esi, %esi      # %esi is i; sets i = 0.  
    pushl   %ebx  
    pushl   %ecx  
    subl    $12, %esp       # Alloc space for fn. args  
    call    reset          # Clear accumulator value  
    ...
```

MAIN FUNCTION (3)

○ Main function code, cont.

```
for (i = 0; i < 10; i++) {  
    n = rand() % 1000;  
    printf("n = %d\taccum = %d\n",  
          n, accum(n));  
}
```

○ Assembly code:

- **esi** is loop variable **i**
- **ebx** is **n**
- **.L6** is start of loop

○ **rand() % 1000** implemented in a *very* unintuitive way...

- Integer division/modulus with a constant can be implemented as multiplication
- (See the book Hacker's Delight)

.L6:

```
call    rand  
movl    $274877907, %edx  
addl    $1, %esi  
movl    %eax, %ecx  
imull    %edx  
movl    %ecx, %eax  
sarl    $31, %eax  
movl    %ecx, %ebx  
sarl    $6, %edx  
subl    %eax, %edx  
imull    $1000, %edx, %edx  
subl    %edx, %ebx  
movl    %ebx, (%esp)  
call    accum  
movl    %ebx, 4(%esp)  
movl    $.LC0, (%esp)  
movl    %eax, 8(%esp)  
call    printf  
cmpl    $10, %esi  
jne     .L6  
...
```

MAIN FUNCTION (4)

- Main function code, cont.

```
for (i = 0; i < 10; i++) {  
    n = rand() % 1000;  
    printf("n = %d\taccum = %d\n",  
          n, accum(n));  
}
```

- Calls to **accum(n)**, **printf(...)**

- Note that **gcc** doesn't explicitly push arguments onto stack!

- Also doesn't pop off stack when done

- This is a compiler optimization

- Why do the pushes and pops, when it can be faked? ☺
- **gcc** allocates extra space on the stack to speed up these calls

.L6:

```
call    rand  
movl    $274877907, %edx  
addl    $1, %esi  
movl    %eax, %ecx  
imull    %edx  
movl    %ecx, %eax  
sarl    $31, %eax  
movl    %ecx, %ebx  
sarl    $6, %edx  
subl    %eax, %edx  
imull    $1000, %edx, %edx  
subl    %edx, %ebx  
movl    %ebx, (%esp)  
call    accum  
movl    %ebx, 4(%esp)  
movl    $.LC0, (%esp)  
movl    %eax, 8(%esp)  
call    printf  
cmpl    $10, %esi  
jne     .L6  
...
```

MAIN FUNCTION (5)

- Main function code, cont.

```
for (i = 0; i < 10; i++) {  
    n = rand() % 1000;  
    printf("n = %d\taccum = %d\n",  
        n, accum(n));  
}
```

- Also need a string constant to pass to **printf()**

- Before **main()** code:

.LC0:

.string "n = %d\taccum = %d\n"

- **as** copies this data to the output binary file
- Address of data is **.LC0**

.L6:

```
call    rand  
movl    $274877907, %edx  
addl    $1, %esi  
movl    %eax, %ecx  
imull    %edx  
movl    %ecx, %eax  
sarl    $31, %eax  
movl    %ecx, %ebx  
sarl    $6, %edx  
subl    %eax, %edx  
imull    $1000, %edx, %edx  
subl    %edx, %ebx  
movl    %ebx, (%esp)  
call    accum  
movl    %ebx, 4(%esp)  
movl    $.LC0, (%esp)  
movl    %eax, 8(%esp)  
call    printf  
cmpl    $10, %esi  
jne     .L6  
...
```

CALLING CONVENTION AND RECURSION

- The cdecl calling convention:
 - Each function call has its own region of the stack
 - Caller pushes arguments onto stack, then calls the callee
 - Callee saves caller's frame pointer, then sets up its own frame pointer
 - Callee stores its local variables after the frame pointer
 - When callee returns to caller, stack is restored to prev state
- This calling convention easily supports recursion
- A procedure can call itself:
 - Each recursive invocation of the procedure will have its own stack space, as long as the conventions are followed!
- You get to explore this more on Assignment 2! ☺

C FLOW-CONTROL STATEMENTS

- C provides a variety of flow-control statements
 - **if** statements

```
if (test-expr)
    then-statement
else
    else-statement
```
 - **while** loops, **for** loops, **do** loops

```
while (test-expr)
    body-statement
```
- Conceptually straightforward to use in your C programs
- How are these normally translated to IA32 assembly language?
 - Helps us better understand what the compiler generates
 - Also helps us know how to write them in IA32 ourselves!

C FLOW-CONTROL STATEMENTS (2)

- C flow-control statements implemented in IA32 using a combination of conditional and unconditional jumps
- Example: **if** statements

```
if (test-expr)  
    then-statement ;  
else  
    else-statement ;
```

- A common translation:

```
t = test-expr ;  
if (t)  
    goto true_branch;  
    else-statement ;  
    goto done;  
true_branch:  
    then-statement ;  
done:
```

- Compiler frequently optimizes/rearranges this flow

DO-WHILE LOOPS

- **do**-loops not used as frequently in programs, but very easy to implement in assembly language
 - Requires minimum number of branching operations

```
do  
    body-statement  
while (test-expr);
```

- A simple translation:

```
loop:  
    body-statement ;  
    t = test-expr ;  
    if (t)  
        goto loop;
```

WHILE LOOPS

- **while**-loops are much more common, but more involved at the assembly-language level

```
while (test-expr)  
    body-statement ;
```

- One translation:

```
loop:  
    t = test-expr ;  
    if (!t)  
        goto done;  
    body-statement ;  
    goto loop;  
done:
```

- Problem: This code is slow to execute.
- Branching has a *big* performance impact
 - Affects instruction caching and pipelining

WHILE LOOPS (2)

- A faster implementation “peels off” the first test

```
while (test-expr)  
    body-statement ;
```

- Equivalent to:

```
if (!test-expr)  
    goto done ;  
do  
    body-statement ;
```

```
while (test-expr)
```

```
done :
```

- Translating this to assembly code yields a loop body containing only one branching operation

WHILE LOOPS (3)

- Our original while-loop:

```
while (test-expr)  
    body-statement ;
```

- Completing the translation:

```
    t = test-expr ;  
    if (!t)  
        goto done ;  
loop:  
    body-statement ;  
    t = test-expr ;  
    if (t)  
        goto loop ;  
done:
```

FOR LOOPS

- **for**-loops are more sophisticated **while** loops:

```
for (init-expr ; test-expr ; update-expr)  
    body-statement ;
```

- Equivalent to:

```
init-expr ;  
while (test-expr) {  
    body-statement ;  
    update-expr ;  
}
```

- We know how to translate these components into assembly language
 - Transform **while** loop into **do-while** loop
 - Insert additional **for**-loop operations into appropriate places

FOR LOOPS (2)

- Implementing **for** loops:

```
for (init-expr ; test-expr ; update-expr)  
    body-statement ;
```

- Translate into:

```
    init-expr ;  
    t = test-expr ;  
    if (!t)  
        goto done;  
loop:  
    body-statement ;  
    update-expr ;  
    t = test-expr ;  
    if (t)  
        goto loop;  
done:
```


MORE ADVANCED PROGRAMS...

- We can now map basic C programs into IA32 assembly language
 - ...including programs that use recursion
- What about this problem?
 - Write a function that takes an argument n
 - Return a collection of all prime numbers $\leq n$
 - e.g. `int * find_primes(int n)`
- Don't yet have the tools to implement this!
 - Requires a variable amount of memory
 - Memory lifetime must extend beyond a single function call

DYNAMIC ALLOCATION AND HEAP

- When programs need a variable amount of memory, they can allocate it from the heap
 - A large, resizable pool of memory for programs to use
 - Programs can request blocks of memory from the heap
 - When finished, programs release blocks back to the heap
 - This is a *run-time facility* for programs to utilize
- For C programs, standard functions to support heap:
 - Allocate a block of memory using **malloc()**
 - **void * malloc(size_t size)**
 - Returns pointer to block of memory with specified size, or **NULL** if **size** bytes are not available
 - **size_t** is an unsigned integer data type
 - Release a block of memory using **free()**
 - **void free(void *ptr)**
 - Specified memory block is returned to heap

HEAP AND STACK

- Programs and stack usage:
 - Stack space automatically reclaimed when function returns
 - Stack values can last for up to the lifetime of a procedure call
 - Procedures are specifically encoded to use the stack
 - Explicit accesses relative to base pointer **ebp**
 - Adjustments to stack pointer to allocate/release space
 - Stack space used by a procedure doesn't vary substantially during its execution
 - Set of local variables within each code block is fixed
 - Set of arguments passed to a procedure is also fixed
- Programs and heap usage:
 - Memory required often depends on the input values
 - Can vary quite dramatically!
 - Programs must *explicitly* allocate and release blocks of memory on the heap

SIMPLE EXAMPLE: VECTOR ADDITION

- Variation on an example from lecture 3:

```
int * vector_add(int a[], int b[], int length) {  
    int *result;  
    int i;  
  
    result = malloc(length * sizeof(int));  
    for (i = 0; i < length; i++)  
        result[i] = a[i] + b[i];  
  
    return result;  
}
```

- Now the procedure dynamically allocates a result vector from heap, then sums inputs into this memory
- Now that we understand IA32 more deeply, let's explore how to implement this function

POINTERS AND ARRAYS

- Our vector-add function:

```
int * vector_add(int a[], int b[], int length) {  
    int *result;  
    int i;  
  
    result = (int *) malloc(length * sizeof(int));  
    for (i = 0; i < length; i++)  
        result[i] = a[i] + b[i];  
  
    return result;  
}
```

- Clear that pointers and arrays are closely related
 - Declare **result** as **int***
 - Access it as an array, just like **a** and **b**

ARRAYS IN C

- C arrays are collections of elements, all of same data type
 - Array elements are contiguous in memory
 - Elements are accessed by indexing into the array
- For an array declaration: ***T A[N]***
 - ***T*** is the data type
 - ***A*** is the array's variable name
 - ***N*** is the number of elements
- C allocates a contiguous region of memory for the array
 - Allocates ***N × sizeof(T)*** bytes for the array
 - ***sizeof(type)*** is a standard C operator that returns the size of the specified data type, in bytes
 - e.g. ***sizeof(int)*** = 4 for ***gcc*** on IA32
 - ***sizeof(type)*** is resolved to a value at compile-time
- ***A*** holds a pointer to the start of the array
 - Can use ***A*** to access various elements of the array

ARRAYS IN C (2)

- For an array declaration: **$T \ A[N]$**
 - **T** is the data type
 - **A** is the array's variable name
 - **N** is the number of elements
- **A** holds a pointer to the start of the array
 - Can use **A** to access various elements of the array
- What address does array-element i reside at?
 - **A** points to first element in array
 - Each element is **`sizeof(T)`** bytes in size
 - Element i resides at address **$A + \text{sizeof}(T) * i$**
- This is what the C array-index operator **`[]`** does
 - **$A[i]$** computes index of element i , then reads/writes the element

POINTERS AND ARRAYS IN C

- C also supports pointer arithmetic
 - Similar idea to array indexing
- For a C pointer variable: **T *p**
- Adding 1 to **p** advances it one *element*, not one byte
 - e.g. for **int *p**, **p + 1** actually advances **p** by 4 bytes
- Adding/subtracting an offset *N* from a pointer to **T** will move forward or backward *N* elements
- Implication:
 - **A[i]** is identical to saying ***(A + i)**
 - **A** is a pointer to first element in array
 - **A + i** moves forward *i* elements (*not i bytes!*)
 - **A + i** is a *pointer* to element *i*, so dereference to get to the actual element **A[i]**

VARIATION ON THEME

- Our original function:

```
int * vector_add(int a[], int b[], int length) {
    int *result;
    int i;

    result = (int *) malloc(length * sizeof(int));
    for (i = 0; i < length; i++)
        result[i] = a[i] + b[i];

    return result;
}
```

- Could also write something crazy like this:

```
int * vector_add(int *a, int *b, int length) {
    int *result, *elem;

    result = (int *) malloc(length * sizeof(int));
    for (elem = result; length != 0; length--) {
        *elem = *a + *b;
        elem++; a++; b++;
    }
    return result;
}
```

- Optimizing compilers may generate code like this
- Take advantage of C's equivalence between arrays and pointers

VECTOR-ADD AND IA32

- How do we implement this function in IA32?

```
int * vector_add(int *a, int *b, int length) {  
    int *result;  
    int i;  
  
    result = (int *) malloc(length * sizeof(int));  
    for (i = 0; i < length; i++)  
        result[i] = a[i] + b[i];  
  
    return result;  
}
```

- Next time, will go through the entire process of implementing this, from scratch. 😊