



CS24: INTRODUCTION TO COMPUTING SYSTEMS

Spring 2015

Lecture 7

LAST TIME

- Dynamic memory allocation and the heap:
 - A *run-time facility* that satisfies multiple needs:
 - Programs can use widely varying, possibly large amounts of memory for computations
 - Allocated memory is available beyond a single function call
 - Allocate with: `void * malloc(size_t size)`
 - Release with: `void free(void *ptr)`
- C arrays: `T A[N]`
 - Allocates $N \times \text{sizeof}(T)$ bytes for the array
 - Element i resides at address `A + sizeof(T) * i`
 - `A[i]` is equivalent to `*(A + i)`
 - Adding i to a pointer moves forward i elements, not i bytes
 - C infers the element size from the pointer's type

TODAY: IMPLEMENTING VECTOR-ADD

- How to implement this function in IA32?

```
int * vector_add(int *a, int *b, int length) {
    int *result;
    int i;

    result = (int *) malloc(length * sizeof(int));
    for (i = 0; i < length; i++)
        result[i] = a[i] + b[i];

    return result;
}
```

- Note: This example works on Linux.
 - e.g. MacOS X introduces many additional complexities
 - Prepends an underscore _ on public symbol names
 - Must use relative/indirect references to external symbols
- If on a confusing platform, use **gcc -S** on test code to figure out what's needed

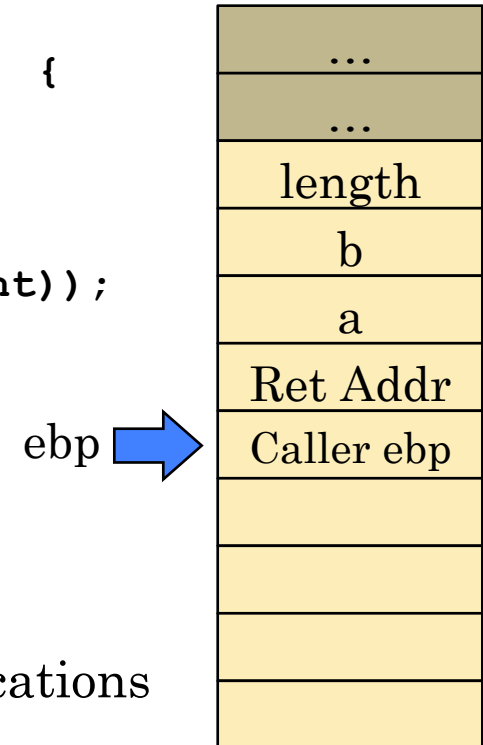
VECTOR-ADD ARGUMENTS

- How to implement this function in IA32?

```
int * vector_add(int *a, int *b, int length) {  
    int *result;  
    int i;  
  
    result = (int *) malloc(length * sizeof(int));  
    for (i = 0; i < length; i++)  
        result[i] = a[i] + b[i];  
  
    return result;  
}
```

- Using cdecl calling convention:

- a**, **b**, **length** pushed by caller; *ebp + offset* locations
- Arguments are pushed last-to-first...
- a** stored at: **8 (%ebp)**
- b** stored at: **12 (%ebp)**
- length** stored at: **16 (%ebp)**



IMPLEMENTING VECTOR-ADD (1)

○ Initial code:

- Put in placeholders where implementation is unknown

```
vector_add:
    # a      = 8(%ebp)
    # b      = 12(%ebp)
    # length = 16(%ebp)
    pushl    %ebp          # Save caller ebp
    movl     %esp, %ebp    # Set up stack frame ptr

    # TODO:   Save callee-save registers we alter

    # TODO:   Allocate space for result with malloc

    # TODO:   Implement vector-sum loop

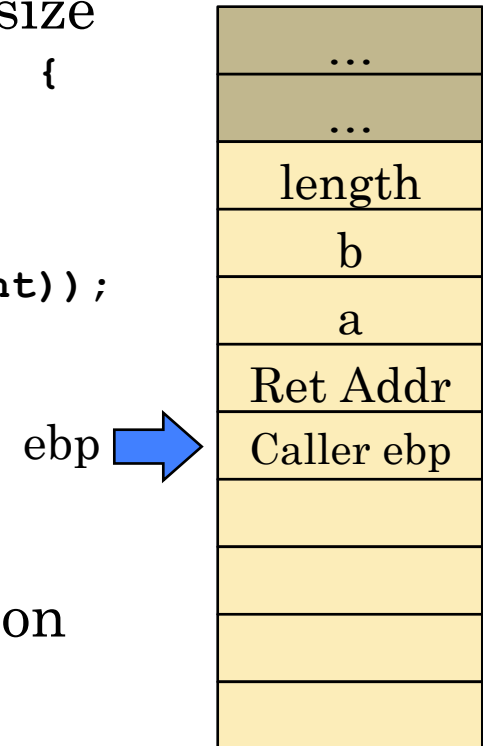
    # TODO:   Restore callee-save registers we alter

    # Clean up stack before returning to caller
    movl     %ebp, %esp    # Remove local vars
    popl     %ebp          # Restore caller ebp
    ret
```

CALLING MALLOC

- Next task is to call **malloc()** with array size

```
int * vector_add(int *a, int *b, int length) {  
    int *result;  
    int i;  
  
    result = (int *) malloc(length * sizeof(int));  
    for (i = 0; i < length; i++)  
        result[i] = a[i] + b[i];  
  
    return result;  
}
```



- malloc()** also uses cdecl calling convention
 - Push **length** * 4 onto stack
 - Call **malloc()** subroutine
 - Memory address will be returned in **eax**
- We want to return this memory, so leave **malloc()**'s result in **eax** and return it to our caller as well

IMPLEMENTING VECTOR-ADD (2)

- Continuing our code:

```
vector_add:
```

```
# a      = 8(%ebp)
```

```
# b      = 12(%ebp)
```

```
# length = 16(%ebp)
```

```
...
```

```
# Allocate space for result with malloc
```

```
movl    16(%ebp), %ecx
```

```
shll    $2, %ecx    # ecx = length * 4
```

```
pushl   %ecx
```

```
call    malloc      # allocate space for result
```

```
popl    %ecx        # cdecl: caller removes args
```

```
# TODO: Implement vector-sum loop
```

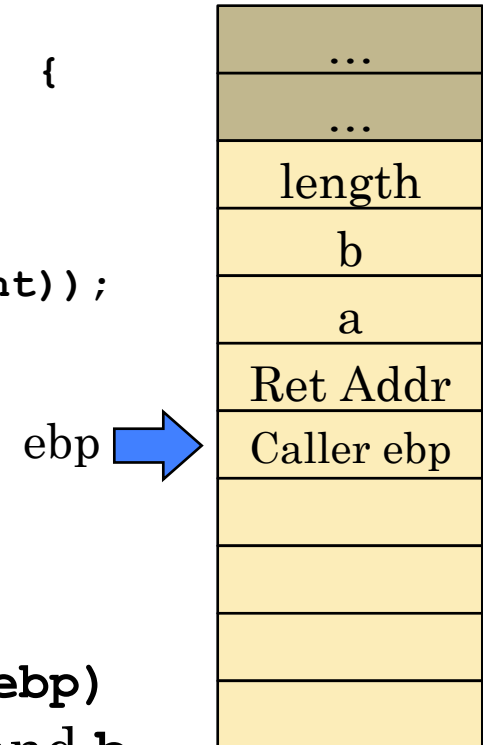
- We return this memory to caller; leave pointer in **eax**

IMPLEMENTING OUR LOOP

- Next, implement our loop!

```
int * vector_add(int *a, int *b, int length) {  
    int *result;  
    int i;  
  
    result = (int *) malloc(length * sizeof(int));  
    for (i = 0; i < length; i++)  
        result[i] = a[i] + b[i];  
  
    return result;  
}
```

- Need a loop variable
 - Use **esi**. Compare directly to **length**, 16(%ebp)
- Also need to access element i of arrays **a** and **b**
 - Use scaled, indexed memory addressing
 - (R_b, R_i, s) syntax accesses $M[R_b + R_i \times s]$
 - Need to store start-addresses of **a**, **b** into registers too



IMPLEMENTING VECTOR-ADD (3)

- Continuing our code:

```
vector_add:
```

```
    # a      = 8(%ebp)
```

```
    # b      = 12(%ebp)
```

```
    # length = 16(%ebp)
```

```
    ...
```

```
    # Implement vector-sum loop
```

```
    movl     8(%ebp), %ebx      # Start of a
```

```
    movl     12(%ebp), %ecx     # Start of b
```

```
    movl     $0, %esi          # Loop variable i
```

- Now we need to implement our **for**-loop

- for**-loop = **while**-loop + initialization and update
- Turn **while**-loop into **do**-loop by factoring out first test

IMPLEMENTING VECTOR-ADD (4)

- Continuing our code:

```
...
# Implement vector-sum loop
movl    8(%ebp), %ebx    # Start of a
movl    12(%ebp), %ecx   # Start of b
movl    $0, %esi        # Loop variable i

cmpl    16(%ebp), %esi   # First for-test
jge     vadd_done        # Is i >= length ?

vadd_loop:
movl    (%ebx, %esi, 4), %edx # edx = a[i]
addl    (%ecx, %esi, 4), %edx # edx += b[i]
movl    %edx, (%eax, %esi, 4) # r[i] = edx

incl    %esi            # i++
cmpl    16(%ebp), %esi   # Subsequent for-tests
jl      vadd_loop        # Is i < length ?

vadd_done:
```

LOOP CONDITIONS

- AT&T syntax is really confusing for conditions
- C code: **for (i = 0; i < length; i++)**
- Assembly code:

```
...  
    cmpl    16(%ebp), %esi    # First for-test  
    jge     vadd_done        # Is i >= length ?  
  
vadd_loop:  
    movl    (%ebx, %esi, 4), %edx    # edx = a[i]  
    addl    (%ecx, %esi, 4), %edx    # edx += b[i]  
    movl    %edx, (%eax, %esi, 4)    # r[i] = edx  
  
    incl    %esi              # i++  
    cmpl    16(%ebp), %esi    # Subsequent for-tests  
    jl      vadd_loop        # Is i < length ?  
vadd_done:
```

- **cmpl Src2, Src1** sets flags as for Src1 – Src2
 - For A op B, generally want to specify **cmp B, A**
 - This way, the conditional jump's type mirrors the C code

FINISHING UP THE IMPLEMENTATION

- Just need to save/restore our callee-save registers

- Code:

```
vector_add:
    # a      = 8(%ebp)
    # b      = 12(%ebp)
    # length = 16(%ebp)
    pushl    %ebp          # Save caller ebp
    movl     %esp, %ebp    # Set up stack frame ptr

    # Save callee-save registers.
    pushl    %ebx
    pushl    %esi

    ... (rest of implementation)

    # Restore callee-save registers.
    popl     %esi
    popl     %ebx

    # Clean up stack before returning to caller
    movl     %ebp, %esp    # Remove local vars
    popl     %ebp          # Restore caller ebp
    ret
```

USING OUR VECTOR-ADD CODE

- Save assembly code into **vector_add.s**
- To make function callable, need to put this at top:

```
# Function to add two vectors together.  
# (bla bla bla)  
.globl vector_add  
vector_add:  
...
```

- **.globl** makes the symbol visible to other functions
- Many other directives you can add, but this is minimal to make the function callable.
- Assemble our function into an object file:
 - **as -o vector_add.o vector_add.s**
 - Can also include debug info by adding **-g** to args

USING OUR VECTOR-ADD CODE (2)

- A simple test program, **va_main.c**:

```
#include <stdio.h>
#include <stdlib.h>

/* Normally would declare this in file vector_add.h */
int * vector_add(int *a, int *b, int length);

int main() {
    int a[] = {1, 2, 3, 4, 5};
    int b[] = {6, 7, 8, 9, 10};
    int *res;
    int i;

    res = vector_add(a, b, 5);
    for (i = 0; i < 5; i++)
        printf("res[%d] = %d\n", i, res[i]);

    free(res);
    return 0;
}
```

USING OUR VECTOR-ADD CODE (2)

- Compile **va_main.c**, link with **vector_add.o**
`gcc -c va_main.c`
`gcc -o vadd va_main.o vector_add.o`
- Run the program:

```
[user@host:~]> ./vadd  
res[0] = 7  
res[1] = 9  
res[2] = 11  
res[3] = 13  
res[4] = 15
```
- Success! 😊
 - And if not, **gdb** is always available for debugging...

VECTOR-ADD FUNCTION

- Using scaled, indexed memory access mode, very easy to implement array indexing in IA32
 - Once array's base address and index are in registers, very easy to retrieve/store a particular element
- Can also call C runtime functions from assembly
 - Called **malloc()** function by adhering to calling conventions
- Returned a heap-allocated chunk of memory to our caller
 - Now our functions can return results larger than **eax**
 - Results also last longer than a single procedure call since they aren't allocated on the stack
 - The program has to free the memory when it's done...

FINDING PRIMES

- Our other program to return primes:
 - Write a function that takes an argument n
 - Returns a collection of all prime numbers $\leq n$
 - e.g. `find_primes(int n)`
- Size of result depends on inputs
- Need to build up result as we go
- Might be better to use a linked-list structure:

```
typedef struct IntList {  
    int value;           /* Value for this node */  
    struct IntList *next; /* Pointer to next node */  
} IntList;
```

```
IntList * find_primes(int n);
```

FINDING PRIMES (2)

- A simple algorithm to generate result:
 for $i = 2$ **to** n **do**
 if *is_prime*(i) **then**
 Append new IntList node for i onto result.
 done
 return result
- Simple to implement, except possibly list part...
- To keep track of list pointers:
 IntList *result = NULL;
 IntList *last = NULL;
- Use **malloc()** to allocate new nodes for list

FINDING PRIMES (3)

- To create and append new list nodes:

```
IntList *new;
...
if (is_prime(i)) {
    /* Allocate new node. */
    new = (IntList *) malloc(sizeof(IntList));
    new->value = i;
    new->next = NULL;

    /* Append new node onto existing list. */
    if (last != NULL)
        last->next = new;
    else
        result = new;

    /* Get ready to append next new node. */
    last = new;
}
```

PRIME NUMBER LINKED-LIST


- There is no way to construct such a data structure on the stack!
 - Simply doesn't match the usage pattern for the stack
- Introduces another reason to use memory heap!
- Previous reasons:
 - Size of result depends on value of input(s)
 - Need result to live beyond lifetime of the procedure
- New reason:
 - Data structure simply cannot be represented on stack

PROCESS MEMORY LAYOUT

- Where do the stack and heap actually live?
- Where does the program itself live?
- Every program is laid out in memory following a specific pattern
 - Memory regions devoted to stack, heap, instructions, constants, global variables, etc.
- Some of these regions vary in size from program to program
 - Instruction data depends on program size
 - Program may not have global variables
 - Compiler determines these details at compile-time
- Some regions depend on program behavior
 - e.g. heap usage or stack usage is a run-time behavior


PROCESS MEMORY LAYOUT (2)

- A common memory layout:
- Program text (instructions), globals, and constants are placed by compiler
 - Reside toward bottom of address space
- Stack resides near top of address space
 - Grows downward as space is consumed
- Heap resides above fixed-size program data
 - Grows upward as space is consumed
- Arrangement gives stack and heap maximum room to grow



PROCESS MEMORY LAYOUT (3)

- Variable-size memory areas have soft and hard limits
 - When soft limit is hit, the region is extended, if possible
 - When hard limit is hit, the program is aborted by the operating system
- Book mentions **brk** value (§9.9)
 - “Break” address, where the program’s memory heap ends
 - **sbrk()** (“set break”) requests more heap space
- Also see **getrlimit()/setrlimit()** functions
 - Gets/sets limits on **brk**
 - Also limits on stack size, etc.



STACK AND HEAP MANAGEMENT


- Stack is relatively easy to manage
 - State consists of memory region + stack pointer
- Heap is *much* more complex!
 - Programs allocate and free blocks of memory, depending on their own needs
 - Heap must keep track of what memory is currently used, and what memory is available for use
 - Data structures for managing heap memory are necessarily complex
- Heap memory is managed by a heap allocator
 - Receives and attempts to satisfy allocation requests from a program
 - Manages these data structures internal to the heap

HEAP ALLOCATORS

- Two kinds of heap allocators
 - Explicit allocators require applications to manually release memory when finished
 - e.g. C `malloc()/free()`, C++ `new/delete`
 - Implicit allocators detect when a memory region is no longer used
 - Employ garbage collectors for finding unused regions
 - e.g. Python, Java, Scheme, ...
- Today: begin looking at explicit heap allocators
- You will get to write one in HW3! ☺


ALLOCATION AND FRAGMENTATION

- Allocators must handle any sequence of allocation requests
 - May affect ability to satisfy requests
- Example: heap with 6KB total space
 - A = allocate 1KB
 - B = allocate 1KB
 - C = allocate 2KB
 - D = allocate 1KB
 - Free B
- Try to allocate 2KB?
 - 2KB of memory is *available*, but it's not contiguous!
- Heap memory can become fragmented from use
 - Allocator must minimize occurrence of fragmentation...
 - ...but this also depends heavily on program's behavior!



ALLOCATION AND FRAGMENTATION (2)


- Allocators cannot modify a memory block once in use
 - e.g. to move it to another part of memory
- Our example:
 - A = allocate 1KB
 - B = allocate 1KB
 - C = allocate 2KB
 - D = allocate 1KB
 - Free B
 - Try to allocate 2KB?
- Cannot modify allocated blocks to compact free space
 - Program has pointers into the allocated region
 - Allocator simply cannot find and update those pointers!
- Allocators can only manipulate free blocks
 - Program isn't using the free blocks, so it doesn't care...



COALESCING FREE BLOCKS

○ Another sequence of allocations:

- A = allocate 1KB
- B = allocate 1KB
- C = allocate 1KB
- D = allocate 2KB
- Free C
- Free B
- Try to allocate 2KB?



- Allocator needs to coalesce adjacent free blocks
 - Failure to do so leads to false fragmentation
 - Suitable free memory is available, but allocator can't find a region of the proper size to satisfy request
 - Ideally, coalescing free blocks will be *fast*

DATA ALIGNMENT AND ALLOCATORS

- Allocators often must care about data alignment
- So far, we have presented memory as an array of individually addressable bytes

00	01	02	03	04	05	06	07	08	09	0A	0B	0C	0D	0E	0F
10	11	12	13	14	15	16	17	18	19	1A	1B	1C	1D	1E	1F
20	21	22	23	24	25	26	27	28	29	2A	2B	2C	2D	2E	2F

- The processor may have a much larger data bus, e.g. $w = 32$ bits
 - Processor can read or write 32 bits of data at once
- To take full advantage of this, memory must also be able to read or write the same size data value

DATA ALIGNMENT AND ALLOCATORS (2)

- For e.g. 32-bit word size, memory looks like this:

00	01	02	03
04	05	06	07
08	09	0A	0B

- Each addressable cell actually holds 4 bytes
- Not possible to address individual bytes in each cell; must read or write the entire cell
- (Memories often allow you to specify a mask saying which bytes in the cell you actually want to use...)
- The processor is able to present an abstraction that individual bytes are addressable...
 - ...but, the CPU will still be interacting with memory in 32-bit words, not 8-bit bytes

DATA ALIGNMENT AND ALLOCATORS (3)

- For e.g. 32-bit word size, memory looks like this:

00:	00	01	02	03	01:	04	05	06	07	02:	08	09	0A	0B	03:	0C	0D	0E	0F
04:	10	11	12	13	05:	14	15	16	17	06:	18	19	1A	1B	07:	1C	1D	1E	1F
08:	20	21	22	23	09:	24	25	26	27	0A:	28	29	2A	2B	0B:	2C	2D	2E	2F

- (Boxes contain byte-addresses; cell address is on left)
- CPU presents an abstraction that each byte is individually addressable...
- CPU maps the program's memory accesses to their actual addresses in physical memory
- Example: Write value 0xFF to byte-address 0x06
 - CPU writes 0x0000FF00 to memory address 1, and tells memory that only the 3rd byte should be written

DATA ALIGNMENT AND ALLOCATORS (4)

- For e.g. 32-bit word size, memory looks like this:

00:	00	01	02	03	01:	04	05	06	07	02:	08	09	0A	0B	03:	0C	0D	0E	0F
04:	10	11	12	13	05:	14	15	16	17	06:	18	19	1A	1B	07:	1C	1D	1E	1F
08:	20	21	22	23	09:	24	25	26	27	0A:	28	29	2A	2B	0B:	2C	2D	2E	2F

- What happens when a program tries to access a word that spans multiple memory cells?
 - e.g. read word starting at byte-address 7
- The CPU must perform two reads, then assemble the result into a single word
 - Read memory cell at address 1, keep only the 4th byte
 - Read memory cell at address 2, keep bottom 3 bytes
 - Assemble these into a single word

DATA ALIGNMENT AND ALLOCATORS (5)

- For e.g. 32-bit word size, memory looks like this:

00:	00	01	02	03	01:	04	05	06	07	02:	08	09	0A	0B	03:	0C	0D	0E	0F
04:	10	11	12	13	05:	14	15	16	17	06:	18	19	1A	1B	07:	1C	1D	1E	1F
08:	20	21	22	23	09:	24	25	26	27	0A:	28	29	2A	2B	0B:	2C	2D	2E	2F

- This obviously complicates the CPU logic...
- Some CPUs simply disallow non-word-aligned memory accesses
 - If a program tries to perform such an access, the CPU reports an error
 - (IA32 supports non-word-aligned memory accesses...)
- The compiler and the memory allocator must be aware of these data-alignment issues.*

DATA ALIGNMENT AND ALLOCATORS (6)

- For e.g. 32-bit word size, memory looks like this:

00:	00	01	02	03	01:	04	05	06	07	02:	08	09	0A	0B	03:	0C	0D	0E	0F
04:	10	11	12	13	05:	14	15	16	17	06:	18	19	1A	1B	07:	1C	1D	1E	1F
08:	20	21	22	23	09:	24	25	26	27	0A:	28	29	2A	2B	0B:	2C	2D	2E	2F

- Solution: the compiler positions all values so that they always start at word boundaries
 - e.g. on a 32-bit system, only store values at addresses that are evenly divisible by 4 bytes
- Programmers generally don't have to think about these word-alignment issues...
 - ...but sometimes you can make your program *much* faster by knowing about them...

DATA ALIGNMENT AND ALLOCATORS (7)

- For e.g. 32-bit word size, memory looks like this:

00:	00	01	02	03	01:	04	05	06	07	02:	08	09	0A	0B	03:	0C	0D	0E	0F
04:	10	11	12	13	05:	14	15	16	17	06:	18	19	1A	1B	07:	1C	1D	1E	1F
08:	20	21	22	23	09:	24	25	26	27	0A:	28	29	2A	2B	0B:	2C	2D	2E	2F

- Similarly, allocators generally only hand out chunks of memory that are word-aligned
 - Request size is also rounded up to next word-aligned boundary
- Example: an allocator that keeps regions aligned on 8-byte boundaries, and a request for 53 bytes
 - Allocator returns block with address 0x00417E58 (bottom 3 bits of the address are 0), size is 56 bytes

NEXT TIME

- More about how explicit allocators work
 - How do allocators track which memory regions are free and which are used?
 - How do allocators deal with fragmentation?
 - What approaches and data structures are used?