




CS24: INTRODUCTION TO COMPUTING SYSTEMS

Spring 2014

Lecture 8

LAST TIME

- Began examining explicit heap allocators
 - The program is responsible for releasing memory when it's no longer needed
- Allocator must deal with several challenges:
 - Avoiding or minimizing memory fragmentation
 - Coalescing adjacent blocks of free memory
 - Dealing with data alignment issues



HEAP ALLOCATOR INTERFACE

- Common interface exposed by explicit allocators:

void * malloc(size_t size)

- Allocates a block of memory of [at least] **size** bytes
- Returns a pointer to start of the block, or **NULL** if **size** bytes are not available

void free(void *ptr)

- Releases a block of memory back to the heap


- Sometimes other operations as well:

void * realloc(void *ptr, size_t size)

- Attempts to change the size of an existing allocation
- If reallocation succeeds, original contents are copied to the new region, and original allocation is freed
- If reallocation fails, returns **NULL** and original allocation is left unchanged

REPRESENTING MEMORY BLOCKS


- Common way to represent blocks of memory on heap:



- Header specifies:
 - Total block size in bytes, including all parts
 - “Allocated” / “free” flag
 - Since memory block is usually word-aligned, the block-size value won’t actually use all bits
 - e.g. for aligning blocks on 8 bytes, bottom 3 bits of size will be 0
 - Can use bottom-most bit(s) to store allocated/free flag
- Payload: area that the program gets to use
 - Payload’s size is what the program requested
- Padding: any space necessary to make the block word-aligned (if necessary or desirable for platform)

REPRESENTING MEMORY BLOCKS (2)

- Heap memory blocks:



- When a program requests memory:


- Allocator works with block structure, creates/updates a block header to satisfy request
- Allocator returns a pointer to start of the *payload*, not start of the *header*

- Abstraction:

- Caller *doesn't care* how allocator manages the heap
- Just wants some memory to use for their program!

REPRESENTING MEMORY BLOCKS (3)


- Heap memory blocks:



- What can happen if program accidentally writes outside of the payload region?
 - e.g. due to a bug, the program writes past the end of its payload, or perhaps before the start
- In these cases, the heap can become corrupted
 - Can no longer keep track of heap's allocation state
- These bugs usually manifest at the *next* allocation or deallocation operation
 - This is when the allocator tries to use its state...

REPRESENTING MEMORY BLOCKS (4)

- Heap memory blocks:




- When program frees memory:

- Program passes its pointer to start of the payload back to the allocator
- Allocator must adjust pointer to gain access to the header

IMPLICIT FREE LIST


- Heap memory blocks form an implicit free list
 - Can determine start of next memory block by looking at size in header of current block




- When allocation request is made, allocator searches thru list of blocks to find a free block that can satisfy the request
- Several strategies for finding a suitable free block
 - First fit – start at beginning of list, stop when first suitable free block is found
 - Next fit – similar to first-fit, but remember where the last suitable free block was found, and start next search there
 - Best fit – check *all* free blocks; choose the smallest free block that can satisfy the request

COALESCING BLOCKS


- When a block is freed, may need to coalesce it with adjacent free blocks



- Example: 4th block is freed by application




- Allocator checks if next block is also free
- If so, coalesce into a larger free block




COALESCING PREVIOUS BLOCKS

- This is only fast when coalescing with *next* block



- Example: 3rd block is freed by application



- How can we efficiently find the *previous* free block?!
- A simple solution: boundary tags (Knuth)
 - Give each block a footer as well as a header
 - (Footer is identical to header)
 - Can easily find prev. block size, and whether it's free

EXPLICIT FREE LISTS

- Implicit free list approach is simple, but not fast
 - Allocation requires a scan of *all* blocks, whether allocated or free
 - No point in looking at allocated blocks!
- Two observations:
 - Really only need to keep the *free* blocks in a list, since these are the ones we check during allocation
 - Since free blocks aren't used by the program, could even store free-list pointers within the block payloads
- For example:
 - Use a linked list to chain together free blocks
- Called an explicit free list approach
 - Explicitly arranging free blocks into a data structure

ORGANIZING FREE BLOCKS

- Many strategies for organizing/selecting free blocks
- Example 1: Maintain free-list in *LIFO order*
 - Newly freed blocks always go at front of free-list
 - Use first-fit policy for assigning new blocks
 - Even though free-list is in LIFO order, can use boundary tags on memory blocks to support constant-time coalescing with neighboring blocks
 - Very fast approach, but more susceptible to memory fragmentation
- Example 2: Maintain free-list in *address order*
 - Keep free blocks sorted by increasing address in free-list
 - Use first-fit policy for assigning new blocks
 - Slower due to linear-time insertion sort when freeing
 - Better memory utilization than LIFO-order free list

ORGANIZING FREE BLOCKS (2)

- Other strategies maintain multiple free-lists
- Example: Segregated fits strategy
 - Each free-list is assigned a size class
 - e.g. {0-1024}, {1025-2048}, {2049-3072}, etc.
 - Each list contains free blocks of that size class
 - When allocating memory:
 - First-fit strategy, starting with appropriate size class
 - If no block is found, go to next larger size class
 - If still no block found, request more memory from OS
 - Once available block is found, may optionally split block and put free part into free-list for corresponding size class
 - When freeing memory:
 - Coalesce with adjacent free blocks, then put result into free list of appropriate size class
- This approach frequently used by standard allocators
 - Fast: allocation searches target blocks of appropriate size
 - Efficient: approaches memory usage of best-fit strategies

DYNAMIC MEMORY ALLOCATION

- Like branching instructions, a heap facility greatly expands the programs we can write
- Heap management is a hard problem to solve!
 - Don't force every program to solve it separately...
 - Provide a run-time facility that implements this capability
 - Programs can leverage this facility as needed
- Understanding how heaps are implemented will help you use them more effectively. For example:
 - Code that performs allocations and deallocations of varying sizes can suffer from memory fragmentation issues
 - Code that always allocates/deallocates blocks of the same size won't suffer from fragmentation
 - Allocating many small objects will tend to waste lots of time and memory, due to bookkeeping overhead

STACK VS. HEAP: COMPARISON

- Stacks are much faster and simpler than heaps
 - Bookkeeping complexity/overhead of heaps is why we don't allocate each local variable on the heap!
 - Much faster to keep local variables and other temporaries on stack
- Heaps provide different capabilities than stacks
 - Enables different usage patterns, but also has additional costs
- Important to understand these distinctions so you use the right tool for the job

IMPORTANT OPTIMIZATION PRINCIPLES

- Two major optimization principles in systems:
- **Make the common case fast.**
 - Determine the most common behaviors of programs, and optimize hardware to execute these cases fast.
 - (Will explore this principle in 2nd half of the term!)
- **Make the fast case common.**
 - If hardware is good at certain cases, arrange your computations to make them as frequent as possible!

DATA ALIGNMENT AND OPTIMIZATION

- See this very clearly with data alignment:
 - “Make the fast case common.”
 - Compiler and assembler adjust the layout of your code to align it with word boundaries for the CPU

- Accumulator code from lecture 5:

```
int value;

int accum(int n) {
    value += n;
    return value;
}

int reset() {
    int old = value;
    value = 0;
    return old;
}
...

.file    "amain.c"
.text
.p2align 4,,15
.globl accum
.type    accum, @function
accum:
... # code for accum()
ret
.size    accum, .-accum
.p2align 4,,15
.globl reset
.type    reset, @function
reset:
... # code for reset()
```

- Same thing also happens with data structures

HETEROGENEOUS DATA STRUCTURES IN C

- C represents heterogeneous data structures with **struct** declarations
 - **struct** members can be different data types, if desired
 - Each member has its own name
 - Members can be primitive types, pointers, arrays, other structs, etc.

- Example: a linked-list node

```
struct node {  
    int number;  
    char *string;  
    struct node *next;  
};
```

- Using our struct:

```
struct node n;  
n.number = 42;    /* Refer to members of struct. */  
n.string = "answer";  
n.next = NULL;
```

C STRUCTURES IN IA32

- Compiler computes important details for each struct:

- Relative offset of each member from start of struct
- Size of each member's data type

Member	Offset	Size
int number	0	4 bytes
char *string	4	4 bytes
node *next	8	4 bytes

- When C code accesses struct members:

- Compiler adds computed offsets to starting address of struct to access specific members

```
void init(struct node *n) {    init:
    n->number = 42;           pushl %ebp
    n->string = "answer";      movl  %esp, %ebp
    n->next = NULL;           movl  8(%ebp), %eax    # eax = n
                              movl  $42,    (%eax)    # n->number
                              movl  $.LC0,  4(%eax)    # n->string
                              movl  $0,     8(%eax)    # n->next
                              popl  %ebp
                              ret
```

STRUCTURES AND DATA ALIGNMENT

- Previous example had all 4-byte members
 - No data alignment issues
- Can use members that are less than a word


```
struct s1 {
```

```
    int i;
```


```
    char ch;
```

```
    int j;
```

```
};
```



- If compiler uses only one byte for **ch**, can't properly align **j** with word boundaries
- Compiler *pads* the struct to properly align all members




- **ch** is followed by 3 unused bytes, to properly align **j**

STRUCTURES AND DATA ALIGNMENT (2)

- In this example:

```
struct s1 {  
    int i;  
    char ch;  
    int j;  
};
```



- Accesses to **ch** use byte-width operations

```
void init(struct s1 *rec) {  
    rec->i = 1234;  
    rec->ch = 'a';  
    rec->j = 5678;  
}  
  
init:  
    pushl %ebp  
    movl %esp, %ebp  
    movl 8(%ebp), %eax # eax = rec  
    movl $1234, (%eax) # rec->i  
    movb $97, 4(%eax) # rec->ch  
    movl $5678, 8(%eax) # rec->j  
    popl %ebp  
    ret
```


- **sizeof(s1)** reports
12 bytes, not 9 bytes

STRUCTURES AND DATA ALIGNMENT (3)

- Some compilers support packing these structures into minimal space necessary (non-standard!)

- e.g. **gcc** supports a **__packed__** attribute

```
struct s1 {  
    int i;  
    char ch;  
    int j;  
} __attribute__((__packed__));
```




```
void init(struct s1 *rec) {  
    rec->i = 1234;  
    rec->ch = 'a';  
    rec->j = 5678;  
}  
  
init:  
    pushl %ebp  
    movl %esp, %ebp  
    movl 8(%ebp), %eax # eax = rec  
    movl $1234, (%eax) # rec->i  
    movb $97, 4(%eax) # rec->ch  
    movl $5678, 5(%eax) # rec->j  
    popl %ebp  
    ret
```

- sizeof(s1)** also reports 9 bytes now


STRUCTURES AND DATA ALIGNMENT (4)

- Packed version of structure will be significantly slower to work with in memory

- e.g. need multiple reads to access value of **j**



- If working with an array of **s1** values, *many* member accesses will not be word-aligned!




- Primarily useful for easily breaking apart packed values – e.g. network packets, disk/file structures
 - In memory, work with properly word-aligned struct
 - For IO, use packed version of struct to easily convert to/from a byte-sequence

STRUCTURES AND DATA ALIGNMENT (5)

- Even if small member is at end of structure, still need to pad to word boundary

```
struct s2 {  
    int i;  
    int j;  
    char ch;  
};
```



- With local variables, don't want to affect other variables' word alignment

```
void foo(int i, int j) {  
    struct s2;          /* -12(%ebp) */  
    int d;              /* -16(%ebp) */  
    ...  
}
```

- If using in an array, want every element to be word-aligned

```
struct s2 values[100];
```

- Note:** 25% of space is wasted on padding to word boundaries! Important to consider in data structure design.

OPTIMIZATION AND DATA ALIGNMENT

- Optimization: “Make the fast case common.”
 - Computers are fastest when accessing data aligned on word boundaries...
 - ...so the compiler lays out instructions and data to be aligned on word boundaries.
- Wastes a certain amount of space, but yields a substantial performance improvement

C UNIONS

- C also provides unions:

- Unlike structs, all union members occupy the same memory location

```
union value {  
    int int_val;  
    char *str_val;  
    float float_val;  
};
```

Member	Offset	Size
int int_val	0	4 bytes
char *str_val	0	4 bytes
float float_val	0	4 bytes

- All members are assigned same offset by the compiler
- Size of union is size of largest member
- Very useful for representing different, mutually-exclusive kinds of values in a single structure
- Also can be bug-prone, since C type-system can't keep you from misinterpreting a value!

C UNIONS (2)

- Unions normally used in concert with a tag field
 - Unlike structs, union members all occupy the same memory location

```
struct value {  
    enum ValType type;  
    union {  
        int int_val;  
        char *str_val;  
        float float_val;  
    };  
};
```

- Code can use **type** to determine what union-member to access

```
if (v->type == IntValue)  
    set_result(a * v->int_val);
```

C LANGUAGE RUN-TIME FACILITIES

- C is a relatively low-level language
 - Virtually all C abstractions translate easily to assembly language
 - Primitive data types, procedures, arrays and composite data types, flow-control statements
 - Not many run-time facilities for C programmers
- Example: array bounds checking
 - C does not stop you from indexing past an array's bounds!

```
int a;  
int r[4];  
int b;  
...  
r[4] = 12345;      /* Compiles!      */  
r[-1] = 67890;    /* Also compiles! */
```

- May affect **a** and/or **b**, depending on relative placement of the variables by the compiler

C PROGRAMS AND ARRAY BOUNDS

- Lack of array bounds-checking can cause problems
- Buffer overflows:

- Program includes a **char** buffer for receiving input data

```
/* Buffer for reading in lines of input data. */
```

```
char buf[100];
```

```
...
```

```
gets(buf);    /* Standard C function in stdio.h */
```

- An example implementation of **gets()**:

```
char * gets(char *s) {
```

```
    int i = 0;
```

```
    int ch = getchar(); /* Get char from console */
```

```
    while (ch != EOF && ch != '\n') {
```

```
        s[i] = ch; ch = getchar(); i++;
```

```
    }
```

```
    s[i] = 0;    /* Zero-terminate the string. */
```

```
    return (ch == EOF && !feof(stdin)) ? NULL : s;
```

```
}
```

C PROGRAMS AND ARRAY BOUNDS (2)

- **gets()** has no way of knowing how large its buffer is!

```
char * gets(char *s) {  
    int i = 0;  
    int ch = getchar();  
    while (ch != EOF && ch != '\n') {  
        s[i] = ch; ch = getchar(); i++;  
    }  
    s[i] = 0;    /* Zero-terminate the string. */  
    return (ch == EOF && !feof(stdin)) ? NULL : s;  
}
```

- **gets()** is a common source of buffer overflows!

- NEVER EVER use **gets()** in your C programs!
- Much better to use:

```
char * fgets(char *s, int size, FILE *stream)
```

- Takes the buffer's size as an argument
- Function makes sure to stay within the buffer

BUFFER OVERFLOW EXPLOITS

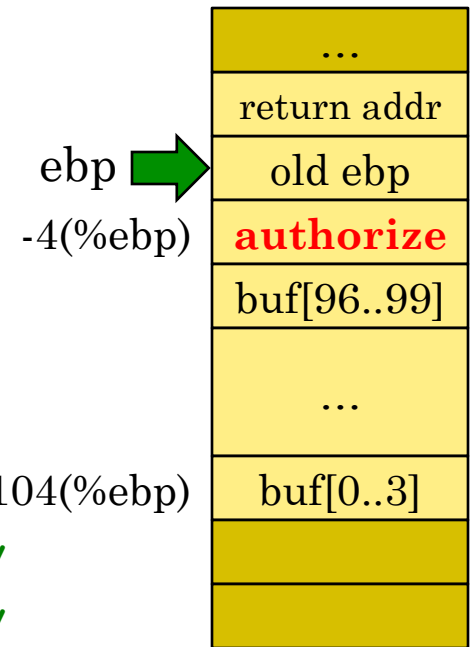
- Buffer overflows don't just cause your program to crash!
- They can frequently be leveraged to compromise system security
- Denial of service:
 - Cause the server program to crash unexpectedly when fed bad input
 - Example: Ping of Death
 - IP packets > 64KB in size are illegal
 - Many OSes had a packet buffer size of exactly 64KB
 - When sent a packet larger than 64KB, target machine would crash, hang, or otherwise freak out

BUFFER OVERFLOW EXPLOITS (2)

- Modify program state:

- Example server code:

```
void handle_request() {  
    int authorize = 0;  
    char buf[100];  
  
    ... /* Read request into buffer */  
    ... /* Verify user, etc.          */  
  
    if (authorize) {  
        ... /* Request is allowed!  Do it. */  
    }  
}
```



- By overflowing the input buffer, can modify value of other local variables.

BUFFER OVERFLOW EXPLOITS (3)

- Executing arbitrary code:

```
void handle_request() {  
    char buf[100];  
    int authorize = 0;
```


```
... /* Read request into buffer */  
... /* Verify user, etc. */
```

```
    if (authorize) {  
        ... /* Request is allowed! Do it. */  
    }  
}
```

-100(%ebp)

-104(%ebp)

ebp →



- Instead of modifying state, exploit aims to change the actual *return address* on stack

- Input includes malicious code loaded into the input buffer
- Set return-address to jump into the buffer

BUFFER OVERFLOW EXPLOITS (4)

- These examples store the buffer on the stack
- Heap-based buffer overflows are just as feasible
 - Overwrite jump-tables, function pointers, other code
- Details of these attacks depend very much on system details! For example:
 - Which direction the stack grows
 - Stack layout of function(s) being exploited
 - What local variables are present, and what effect they have
 - Where the return-address is stored
 - Heap layout of program being exploited
 - ...many other details of server being compromised...
- Nonetheless, a very large percentage of exploits use buffer overflows to compromise the system.

AVOIDING BUFFER OVERFLOWS

- A simple solution to buffer overflows?
 - Build array bounds-checking into your language
 - All array indexes are verified before access occurs
 - Invalid indexes are flagged with some kind of error
- To support this, need to store more information about arrays
 - Add metadata to the array representation
- Example:

```
struct array_t {  
    int length;  /* Number of elements */  
    struct value_t values[];  
};
```

- Arrays include length information in their run-time representation

ARRAY BOUNDS-CHECKING

○ New array representation:


```
struct array_t {  
    int length; /* Number of elements */  
    struct value_t values[];  
};
```

- Last member of a struct can be an array with no size
- Supports variable-size arrays in structs

```
array_t *a = (array_t *)  
    malloc(sizeof(array_t) + n * sizeof(value_t));  
a->length = n;
```

- **values** is a pointer to start of variable-size array

○ Memory layout:



ARRAY BOUNDS-CHECKING (2)

- Can expose array length as a member for programs to reference

```
for (int i = 0; i < a.length; i++) {  
    compute(a[i]);  
}
```

- In language implementation, all array-indexing operations are bounds-checked against length
- Still need to further constrain our language!
- Close off other potential holes:
 - No more pointer manipulation and pointer arithmetic
 - Present a simplified, opaque “reference” abstraction for programmers to use
 - Need a clean way to report errors when they occur

MEMORY LEAKS

- Another common issue: memory leaks
- Explicit allocators require program to inform allocator when memory is no longer used
- If the program fails to do this properly:
 - Program no longer has a pointer to allocated memory
 - But, allocator thinks memory block is still in use!
 - Cannot reclaim memory until program terminates
- *Implicit allocators* assume the responsibility for reclaiming unused memory
 - Garbage collection: reclaiming unused heap storage
 - Program no longer has to free memory itself
 - Memory-leak issues largely disappear

LOOKING FORWARD

- Languages like C don't make it very easy to write correct and bug-free programs!
 - Much easier than assembly language, but can easily have buffer overflows, memory leaks, exploits...
- Take another step up the abstraction hierarchy:
 - Programming languages that provide more powerful abstractions can mitigate many of these problems
- Will focus on three very common features:
 - Memory management: references, implicit allocators
 - Object-oriented programming and polymorphism
 - Another flow-control mechanism: exception handling