



CS24: INTRODUCTION TO COMPUTING SYSTEMS

Spring 2015

Lecture 9

LAST TIME

- Finished covering most of C's abstraction capabilities
 - Dynamic memory allocation on heap, structs, unions
- Began to see something very disturbing:
 - It's *very easy* to write incorrect or unsafe programs in C!
- Unchecked array accesses and buffer overflows:
 - Allow an attacker to crash a program, modify its data in unintended ways, or even execute arbitrary code!
- Other memory management problems as well:
 - Programs don't free memory when they are done with it
 - Programs allocate memory and then access beyond its end
 - Programs access memory after they have freed it
- Idea: *Most people don't actually need all this power!*
 - Provide a simplified memory management abstraction that makes it much easier to implement correct programs

HIGHER-LEVEL LANGUAGE FACILITIES

- Starting to enter realm of higher-level languages
- Much safer programming models:
 - Easier to write correct programs
 - Fewer potential security holes!
- Much greater abstracting capabilities
 - Greater modularity, encapsulation, code reuse
- Also requires much larger run-time support for various facilities
 - Usually slower than C/C++ programs, but much safer!

HIGHER-LEVEL LANGUAGE FEATURES (2)


- In next few lectures, will explore three language features:
 - Implicit allocators and garbage collection
 - The object-oriented programming model
 - Exception handling
- Specifically, how to map these features to C, IA32
 - Taking another step up the abstraction hierarchy...
- Examples will draw from Java language features
 - A higher level language than C, with C-like syntax
 - (Not a *huge* step up the abstraction hierarchy...)
 - Includes all of the above language features
 - Itself implemented in C/C++ and assembly language

IMPLICIT HEAP ALLOCATORS

- Explicit heap allocators rely on the program to release memory when no longer in use
 - *...but programs are notoriously bad at this...*
- A first step towards better programs:
 - Implicit allocators assume the responsibility for identifying when a program is finished with memory
 - Employ a process called garbage collection to identify when a memory block is no longer used by a program
- Use of an implicit allocator eliminates *many* memory management issues for programs
 - Additional overhead for performing garbage collection
 - (A few other issues as well, all relatively minor)

REACHABILITY GRAPH

- Garbage collectors are often built on the concept of a reachability graph
- Some allocated nodes are root nodes
 - Referenced from global environment, or stored on the stack
- All nodes reachable from the root nodes are live
- Unreachable nodes are garbage
 - May be reclaimed by allocator and reused for subsequent allocations
- How to determine this reachability graph?



IMPLICIT ALLOCATORS AND POINTERS

- What about performing garbage collection in a language like C or C++?
- With such languages, garbage collection can be *very* difficult
- Unfortunately, C/C++ allows:
 - Pointers into the middle of memory blocks
 - Pointers into the middle of structs and arrays
 - Recasting pointers into pointers of other types
 - Recasting pointers into integers and vice-versa
 - All kinds of crazy pointer arithmetic!
- Makes it *very* difficult to build an accurate reachability graph in a C/C++ run-time system

CONSERVATIVE GARBAGE COLLECTORS

- In languages like C and C++:
- Garbage collector attempts to identify any value that “looks like” a pointer
 - References a memory location that is currently valid
 - Assume this is a pointer to memory that is in use
- If the value isn’t actually a pointer?
 - Only real drawback is that the garbage-collector thinks memory is in use, when it really isn’t in use
 - (Hopefully) doesn’t happen often enough to be a problem
- Conservative garbage collection:
 - All reachable blocks are identified as reachable
 - Some unreachable blocks are also identified as reachable
 - **Not all garbage is reclaimed.** But we can live with that.

PRECISE GARBAGE COLLECTORS

- Another approach is to *strictly control* how programs can use and manipulate pointers
 - Specify rules on casting pointers, and disallow casting to/from non-pointer types
 - Completely forbid pointer arithmetic!
- Allows precise garbage collection
 - Garbage collector can determine the reachability graph *exactly*, for all memory blocks in the heap
 - All garbage can be reclaimed by the allocator
- Given other issues with pointer manipulation (e.g. buffer overflows), *makes tons of sense* to constrain pointers this way!!!

JAVA REFERENCES


- Like many higher-level languages, Java includes references
- A “reference” is simply a means for looking up and accessing a particular object
 - A pointer is a very primitive kind of reference: it contains the exact memory address of the object
- Introduce a higher-level reference abstraction:
 - The reference is opaque to the program!
 - Programs can no longer directly access or manipulate the memory address associated with a reference
 - (The reference’s type can also only be manipulated in very specific, controlled ways.)
 - Gives the run-time system much greater flexibility in managing memory (e.g. moving allocated blocks)

INDIRECTION

- **Indirection** is a very important technique used in computer system design
 - The ability to reference something using a name or other value that represents the *actual* value
- Have already seen this technique multiple times
 - e.g. “branch to register,” labels as jump-targets in code, using pointers to access values, jump tables, ...
- References allow programs to *indirectly* access objects, while the runtime directly accesses them
- “All problems in computer science can be solved by another level of indirection.”
 - David Wheeler, inventor of the subroutine


GARBAGE COLLECTION ALGORITHMS

- Variety of algorithms used for garbage collection
- Simplest algorithm is called “mark and sweep”
- Every object has a flag associated with it
- Initially: all flags are 0
- Two phases:
 - First phase involves traversing entire object graph, marking all reachable objects
 - Second phase involves removing all unreachable objects




MARK-AND-SWEEP (1)

- Start by marking root nodes as reachable
- (Note: would implement marking phase as depth-first traversal, not breadth-first...)




MARK-AND-SWEEP (2)

- Next, nodes reachable from root nodes



MARK-AND-SWEEP (3)


- Continue until all reachable nodes are marked.
- Now, any node with a 0 is unreachable, and may be reclaimed.



MARK-AND-SWEEP (4)


○ Second phase:

- If object is unreachable, reclaim it.
- (If object is reachable, reset flag to 0 for next time.)



MARK-AND-SWEEP (5)

- Final result:



MARK-AND-SWEEP CHARACTERISTICS

- This GC algorithm has several drawbacks
- Most important one:
 - To ensure that all garbage is identified, the program cannot run while the garbage collector is working!
 - Even on a multicore system, program and GC cannot run concurrently
- This is called a “stop-the-world” garbage collector
 - Entire program must be suspended while garbage collection is performed
- Clearly unacceptable for applications where response-time is critical
 - e.g. real-time applications, interactive applications

GARBAGE COLLECTOR CHARACTERISTICS


- Several kinds of garbage collector characteristics
- Stop-the-world GC vs. concurrent GC
 - Stop-the-world garbage collectors must suspend the program while performing garbage collection
 - Concurrent garbage collectors work in the background, while the program continues to run
 - Algorithm may also require “stop-the-world” phases, but they are kept as short as possible
 - Frequently, the “stop-the-world” phases are parallelized to minimize wait-time
- Serial GC vs. Parallel GC
 - Serial garbage collectors are not able to use multiple processors during GC phases
 - Parallel collectors are able to employ multiple processors to speed up collection phases

GC CHARACTERISTICS (2)

- Compacting GC vs. non-compacting GC
 - Non-compacting collectors do nothing with the remaining live objects
 - This is the only real option for languages that allow programs to use explicit pointers into memory
 - Memory fragmentation can become a big problem!
 - Compacting garbage collectors move remaining live objects together, maximizing size of free space
 - Only possible to compact memory if language doesn't expose explicit pointers to programs!
 - Another reason to only expose opaque references

MARK-AND-SWEEP CHARACTERISTICS (2)

- In what situations is mark-and-sweep garbage collection the least expensive?
 - (Compacting mark-and-sweep, in particular?)
- If most objects are not reclaimed, mark-and-sweep will be relatively inexpensive
 - Mark phase always touches *all* objects, regardless...
 - Sweep phase will have to do less work if most objects don't need reclaimed
 - Particularly true if GC must also compact memory!



OTHER GC STRATEGIES


- Mark-and-sweep garbage collection must do a lot of work if many objects are deleted during sweep:
 - Must deallocate each garbage object individually (this overhead definitely adds up)
 - Must compact live objects to avoid fragmentation
- Mark-and-sweep prefers long-living objects ☺
- Another strategy: Copying garbage collectors
 - Instead of compacting live objects together, live objects are all copied (“evacuated”) to another, contiguous region of memory
 - Compaction is performed automatically!
 - All dead objects can be deallocated in one operation!

STOP-AND-COPY GARBAGE COLLECTORS


- Stop-and-copy garbage collectors divide memory into two regions: “from-space” and “to-space”
- New objects are allocated in the to-space
- When to-space becomes full, it becomes the “from-space,” and vice versa:
 - Starting with root objects, all reachable objects are copied from from-space into to-space
 - At end, entire from-space can be reclaimed at once
- Program resumes execution, using new to-space
- Automatically compacts live objects, but also effectively halves memory available to programs

STOP-AND-COPY EXAMPLE


- Memory divided into from-space and to-space
- New objects are allocated in the to-space, until it becomes full



- Program is stopped; starting with root objects, reachable objects are copied to new to-space




- At end, entire from-space is reclaimed



STOP-AND-COPY GC

- Stop-and-copy garbage collection is fast when:
 - Not a lot of objects live through the GC phase!




- The fewer objects you have to copy, the better.
- Stop-and-copy prefers short-lived objects ☺
- Observation: Different garbage collection algorithms are best for different situations...
- Questions:
 - Do all heap-allocated objects behave “the same” with respect to garbage collection?
 - Are there differences that we can take advantage of?

GENERATIONAL GARBAGE COLLECTION

- Some empirical observations about programs:
- Most objects in a program are very short-lived
 - e.g. used for local variables, intermediate results, etc.
 - “Most objects die young.”
- Longer-lived objects generally do not reference shorter-lived ones
- These ideas called the *generational hypothesis*
 - (Also “infant mortality,” but that’s just macabre.)
- Idea:
 - Design a garbage collector that takes advantage of this behavioral characteristic of OO program state
 - Called generational garbage collection
 - Optimization: “Make the common case fast.”


SUN HOT-SPOT GENERATIONAL GC

- Sun Java VMs implement generational garbage collection
- Divides objects into “young objects,” “old objects”
- Young objects kept in a separate memory area
 - Uses stop-and-copy GC, since most will not live long
 - Three memory areas: Eden, and two survivor spaces
- Newest objects are allocated in Eden
 - Most never make it out alive...
- If new objects survive the first GC pass, moved into to-space




SUN HOT-SPOT GENERATIONAL GC (2)

- Survivor spaces for slightly older “young objects”
 - Give young objects “additional chances to die” before they are considered “old objects”
- As before, when to-space fills up, turn into from-space, and perform stop-and-copy GC
- An important difference:
 - Objects considered “old enough” are promoted from from-space into “old generation” memory



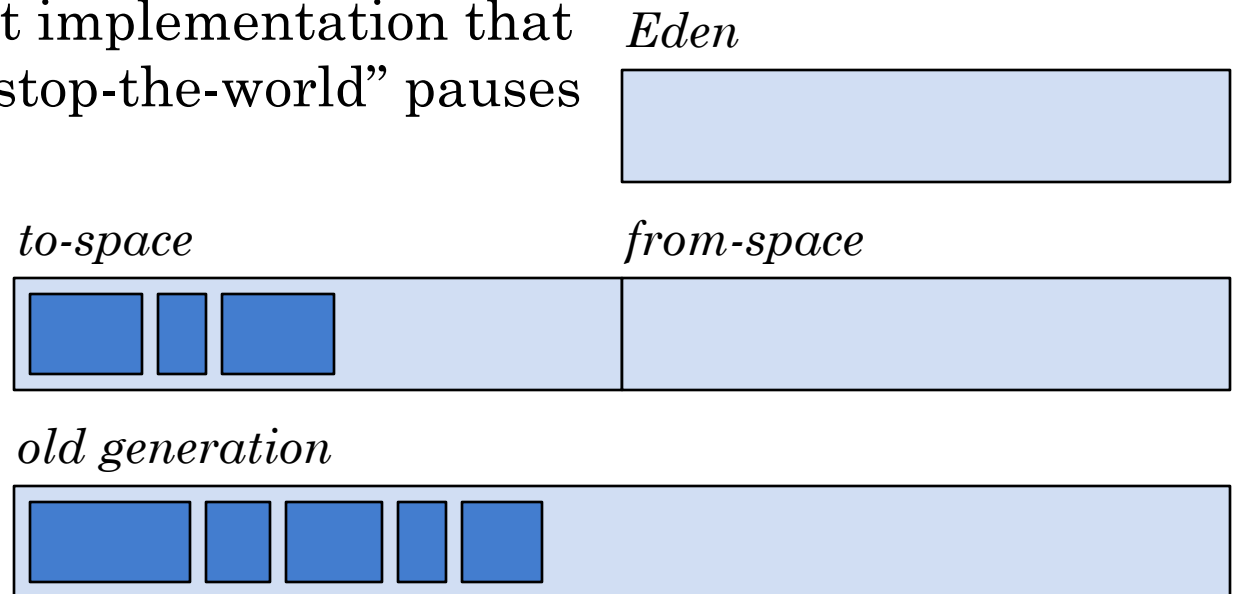
SUN HOT-SPOT GENERATIONAL GC (3)

- After young-object garbage collection, both Eden and the from-space are now empty
 - Objects have either been reclaimed, or copied to another memory area
- Two benefits of stop-and-copy GC here:
 - Most objects die, so GC is fast
 - Allocating new objects also fast



SUN HOT-SPOT GENERATIONAL GC (4)

- Finally, old generation also needs periodic GC
- Most of these objects are expected to survive...
 - Stop-and-copy garbage collection is not appropriate!
- Old generation is managed with a compacting mark-and-sweep algorithm
 - A concurrent implementation that minimizes “stop-the-world” pauses



GENERATIONAL GARBAGE COLLECTION

- Generational garbage collection is very complex!
- Takes advantage of two important details:
 - Different garbage collection algorithms are good in different situations
 - Program state tends to fall into two major categories: young, short-lived objects, and old, long-lived objects
- Provides a much more effective garbage collection system than the individual GC algorithms could possibly provide on their own!

GENERAL-PURPOSE SOLUTIONS

- Another very important system-design pattern: creating general solutions from specialized ones
- Can frequently solve a problem in multiple ways
 - e.g. mark-and-sweep GC vs. stop-and-copy GC
 - Each solution works well in different situations
- We want our computers to be general-purpose...
- We want our operating systems to support a wide range of program behaviors and usage scenarios
- The most powerful, generic solutions frequently blend multiple techniques in a very elegant way
 - Generational garbage collection is a great example!
 - This is a common theme in computer system design

REFERENCE COUNTING

- Some implicit allocators are based on reference counting, instead of on a reachability graph
- Each object keeps a reference count
 - A simple integer count of how many other objects reference the object
- When code first references an object, its reference-count is automatically incremented
- When code finishes working with the object, its reference count is automatically decremented
- When an object's reference count hits zero, it is automatically reclaimed
 - Instead of periodic, complicated garbage-collection sweeps, objects are immediately reclaimed


REFERENCE COUNTING AND CYCLES

- Major drawback of reference counting:
 - Cannot properly release cycles of objects!

- Our earlier example:


- All objects have a nonzero reference-count...
- ...but some of the objects are unreachable!

- Despite this limitation, many systems still use reference-counting for automatically freeing objects
 - e.g. the Python runtime




REFERENCE COUNTING AND CYCLES (2)

- Several techniques for dealing with cycles
- Cycle detection algorithms
 - Use more standard reachability-graph techniques to detect and reclaim cycles
 - Since it's only needed for a subset of objects, won't have a heavy impact
 - Approaches usually focus on identifying objects that *could be* part of a cycle, and starting from there
- More complex allocator, but keeps the coder's life simpler!



REFERENCE COUNTING AND CYCLES (3)

- Another approach: weak references
 - Simply don't allow cycles in reference graph
 - When objects need to refer to each other, one object uses a weak reference
 - Doesn't increment target object's reference count
 - Target of a weak reference may go away unexpectedly
- Properly breaks cycles...
- Programmer must design the program to use weak references properly
 - Can be *very* difficult to get this correct!



REF-COUNTING, GARBAGE COLLECTION

- Other drawbacks with reference counting
 - Cost of incrementing and decrementing reference-counts really adds up
 - Languages with garbage collection don't incur this overhead
 - Some ways to optimize away reference-count updates
- Nonetheless, still quite a common approach
 - Easy to implement!

SUMMARY

- Implicit allocators allow us to eliminate *many* memory management issues
 - Provide a simplified abstraction to programs
- Programs allocate memory, but implicit allocator uses garbage collection to determine when to free
 - e.g. mark and sweep, stop and copy
 - More advanced allocators blend these techniques to create very powerful, general-purpose approaches, e.g. generational garbage collection
- Many languages replace pointers with references
 - Program indirectly references objects, while runtime can still directly access and manipulate them
 - Eliminates many other kinds of serious security bugs!