



CS24: INTRODUCTION TO COMPUTING SYSTEMS

Spring 2015

Lecture 11

EXCEPTION HANDLING

- Many higher-level languages provide exception handling
- Concept:
 - One part of the program knows how to detect a problem, but not how to handle it in a general way
 - Another part of the program knows how to handle the problem, but can't detect it
- When a problem is detected, the code throws an exception
 - An “exception” is a value representing the error
 - Frequently, an object that contains the error's details
 - The exception's type indicates the category of error
- Code that knows how to handle the problem can catch the exception
 - Provides an exception handler that responds to the error

JAVA EXCEPTION HANDLING

- Java exception handling uses **try/catch** blocks

```
public static void main(String[] args) {  
    loadConfig(args);  
    try {  
        double x = readInput();  
        double result = computeValue(x);  
        System.out.println("Result = " + result);  
    }  
    catch (IllegalArgumentException e) {  
        printError(e.getMessage());  
    }  
}
```

- If input is invalid, **computeValue()** throws an exception
- Execution immediately transfers to the exception handler for **IllegalArgumentException**

JAVA EXCEPTION HANDLING (2)

- Only exceptions from within **try** block are handled!

```
public static void main(String[] args) {  
    loadConfig(args);  
    try {  
        double x = readInput();  
        double result = computeValue(x);  
        System.out.println("Result = " + result);  
    }  
    catch (IllegalArgumentException e) {  
        printError(e.getMessage());  
    }  
}
```

- If **loadConfig()** throws, the exception isn't handled here
- try**: “If *this* code throws, I want to handle the exceptions.”
 - (Assuming the exception matches one of the **catch** blocks...)

JAVA EXCEPTION HANDLING (3)

- Code can report an exception by **throw**-ing it:

```
double computeValue(double x) {  
    if (x < 3.0) {  
        throw new IllegalArgumentException(  
            "x must be at least 3");  
    }  
    return Math.sqrt(x - 3.0);  
}
```

- Now the function can complete in two ways:
 - Normal completion: returns the computed result
 - Abnormal termination:
 - Function stops executing immediately when **throw** occurs
 - Program execution jumps to the nearest enclosing **try/catch** block with a matching exception type

EXCEPTIONS WITHIN A FUNCTION

- Exceptions can be used within a single function

```
static void loadConfig(String[] args) {  
    try {  
        for (int i = 0; i < args.length; i++) {  
            if (args[i].equals("-n")) {  
                i++;  
                if (i == args.length)  
                    throw new Exception("-n requires a value");  
                ...  
            }  
            else if ...  
        }  
    }  
    catch (Exception e) {  
        System.err.println(e.getMessage());  
        showUsage(); System.exit(1);  
    }  
}
```

- Used to signal an error in argument-parsing code

EXCEPTIONS SPANNING FUNCTIONS

- Exceptions can also span multiple function calls
 - Doesn't have to be handled by immediate caller of function that throws!

- Example:

```
Webpage loadPage(URL url) {  
    try {  
        InputStream in = sendHttpRequest(url);  
        ...  
    }  
    catch (UnknownHostException e) ...  
}  
  
InputStream sendHttpRequest(URL url) {  
    Socket sock = new Socket(url.getHost(), url.getPort());  
    ...  
}
```

- **Socket** constructor could throw an exception
 - Propagates out of **sendHttpRequest()** function...
 - Exception is handled in **loadPage()** function

EXCEPTION HANDLING REQUIREMENTS

- A challenging feature to implement!
 - Can throw objects containing arbitrary information
 - Exception can stay within a single function, or propagate across multiple function calls
 - Actual **catch**-handler that receives the exception, depends on who called the function that threw
 - A function can be called from multiple places...
 - A thrown exception should be handled by the nearest *dynamically-enclosing* **try/catch** block
- Also want exception passing to be fast
 - Ideally, won't impose any overhead on the program until an exception is actually thrown
 - Assumption: exceptions aren't thrown very often
 - ...*hence the name "exception"*...
 - (Not always a great assumption these days, but oh well!)

IMPLEMENTING EXCEPTION HANDLING

- With exception handling, there are two important points in program execution
- When execution enters a **try** block:
 - Some exceptions might be handled by this **try/catch** block...
 - May need to do some kind of bookkeeping so we know where to jump back to in case an exception is thrown
- When an exception is actually thrown:
 - Need to jump to the appropriate **catch** block
 - Need to access information from previous **try**-point, so that we can examine the proper set of catch blocks
 - This will frequently span multiple stack frames

EXCEPTIONS WITHIN A FUNCTION

- When exception is thrown and caught within a single function:

```
void foo() {  
    try {  
        ...  
        if (failed)  
            throw new Exception();  
        ...  
    }  
    catch (Exception e) {  
        ... // Handle the exception  
    }  
}
```

- In this case, can translate **throw** into a simple jump to the appropriate exception handler
 - Types are available at compile time

EXCEPTIONS WITHIN A FUNCTION (2)

- Still need some way to pass exception object to the handler...

```
void foo() {  
    try {  
        ...  
        if (failed)  
            throw new Exception();  
        ...  
    }  
    catch (Exception e) {  
        ... // Handle the exception  
    }  
}
```

- Assume there will be at most one exception in flight at any given time
- Store [reference to] the exception in a global variable

EXCEPTIONS WITHIN A FUNCTION (3)

- One possible translation of our code:

```
void foo() {
    ...           // Code that sets up failed flag
    if (failed) {
        set_exception(new Exception()); // throw
        goto foo_catch_Exception;
    }
    ...           // Other code within try block
foo_end_try:     // End of try-block
    goto foo_end_trycatch;

foo_catch_Exception: {
    e = get_exception();
    ... // Handle the exception
    goto foo_end_trycatch;
}

foo_end_trycatch:
    return;
}
```

EXCEPTIONS SPANNING FUNCTIONS

- Not a good general solution! Normal case is to have exceptions spanning multiple function calls.

- Can't implement with **goto**, since **goto** can't span multiple functions!

- Really can't hard-code the jump now, anyway...

- Really want a way to *record* where to jump, dynamically

- i.e. when we enter **try** block

- Then, a way to jump back to that place, even across multiple function calls

- i.e. when exception is thrown

```
int f(int x) {
    try {
        return g(3 * x);
    } catch (Exception e) {
        return -1;
    }
}

int g(int x) {
    return h(15 - x);
}

int h(int x) {
    if (x < 5)
        throw new Exception();

    return Math.sqrt(x - 5);
}
```

setjmp() AND longjmp()

- C standard includes two very interesting functions:
- **int setjmp(jmp_buf env)**
 - Records current execution state into **env**, at exact time of **setjmp()** call
 - Information recorded includes callee-save registers, **esp**, and caller return-address
 - Always returns 0
- **void longjmp(jmp_buf env, int val)**
 - Restores execution state from **env**, back into all registers saved by **setjmp()**
 - **esp** is restored from **env**:
 - Any intervening stack frames are discarded
 - Stack is restored to the state as when **setjmp()** was called
 - Caller return-address on stack when **setjmp()** was called
 - Then **longjmp()** returns, with **val** in **%eax**
 - (or **%eax** = 1 if **val** is 0)
- To caller, it appears that **setjmp()** returned again!

setjmp() AND longjmp() (2)

- Previous example is simple enough to implement using `setjmp()` and `longjmp()`:

```
int f(int x) {  
    try {  
        return g(3 * x);  
    } catch (Exception e) {  
        return -1;  
    }  
}
```

```
int g(int x) {  
    return h(15 - x);  
}
```

```
int h(int x) {  
    if (x < 5)  
        throw new Exception();  
  
    return Math.sqrt(x - 5);  
}
```

```
static jmp_buf env;
```

```
int f(int x) {  
    if (setjmp(env) == 0)  
        return g(3 * x);  
    else  
        return -1;  
}
```

```
int g(int x) {  
    return h(15 - x);  
}
```

```
int h(int x) {  
    if (x < 5)  
        longjmp(env, 1);  
  
    return sqrtl(x - 5);  
}
```

setjmp() AND longjmp() (3)

- When we enter **try** block, record execution state in case an exception is thrown
- If an exception is thrown, use **longjmp()** to return to where **try** was entered
 - *Stack frames of intervening function calls are discarded!*
- Return value of **setjmp()** indicates whether an exception was thrown
 - Example has only one kind of exception, so any return-value will do

```
static jmp_buf env;

int f(int x) {
    if (setjmp(env) == 0)
        return g(3 * x);
    else
        return -1;
}

int g(int x) {
    return h(15 - x);
}

int h(int x) {
    if (x < 5)
        longjmp(env, 1);

    return sqrtl(x - 5);
}
```


setjmp()/longjmp() EXAMPLE

- What happens with `f(5)` ?
- Things will go badly... 😊

```
static jmp_buf env;

int f(int x) {
    if (setjmp(env) == 0)
        return g(3 * x);
    else
        return -1;
}

int g(int x) {
    return h(15 - x);
}

int h(int x) {
    if (x < 5)
        longjmp(env, 1);

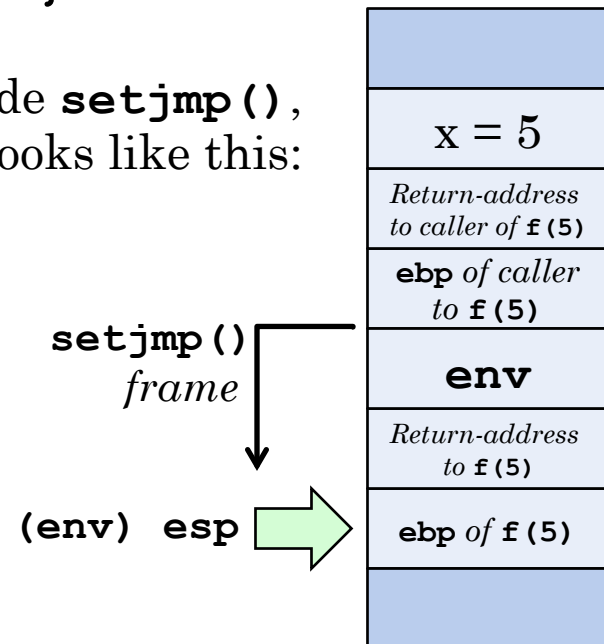
    return sqrtl(x - 5);
}
```

setjmp() / longjmp() EXAMPLE (2)

- **f(5)** calls **setjmp()** to prepare for an exception
 - Will return 0 since it's actually the **setjmp()** call
- **setjmp()** stores:
 - Callee-save registers, including current **esp**
 - **setjmp()**-caller's **ebp** and return address
 - (grab these from stack)
- **env** now holds everything necessary for **longjmp()** to act like it's **setjmp()** ...

```
static jmp_buf env;  
  
int f(int x) {  
    if (setjmp(env) == 0)  
        return g(3 * x);  
    else  
        return -1;  
}
```

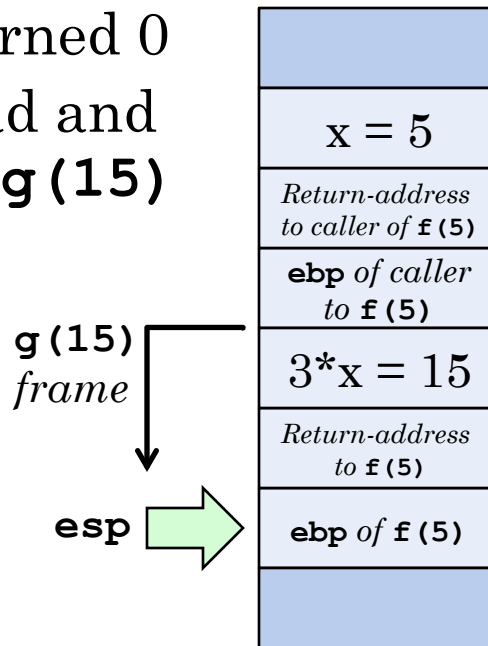
Inside **setjmp()**,
stack looks like this:



setjmp()/longjmp() EXAMPLE (3)

- `setjmp()` returned 0
- `f(5)` goes ahead and calls `g(3*x) = g(15)`

- Now the stack looks like this:



- Note that the stack frame from calling `setjmp()` is long gone...
 - (along with the return-address to where `setjmp()` was called from)
 - `env` still contains these values!

```
static jmp_buf env;

int f(int x) {
    if (setjmp(env) == 0)
        return g(3 * x);
    else
        return -1;
}

int g(int x) {
    return h(15 - x);
}

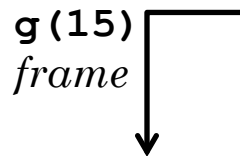
int h(int x) {
    if (x < 5)
        longjmp(env, 1);

    return sqrtl(x - 5);
}
```

setjmp()/longjmp() EXAMPLE (4)

- Now in **g(15)** call
- g** calls **h(15-x)**
= **h(0)**
- Stack looks like this:


g(15)
frame



h(0)
frame

- Problem:

 - h()** can't handle values less than 5
 - h()** needs to abort the computation



```
static jmp_buf env;

int f(int x) {
    if (setjmp(env) == 0)
        return g(3 * x);
    else
        return -1;
}

int g(int x) {
    return h(15 - x);
}

int h(int x) {
    if (x < 5)
        longjmp(env, 1);

    return sqrtl(x - 5);
}
```

setjmp()/longjmp() EXAMPLE (5)

- **h(0)** needs to abort!
 - Got a bad argument...
- **h()** “throws an exception”
 - Calls **longjmp()** to switch back to nearest enclosing **try**-block
 - **env** contains details of where nearest enclosing **try**-block is...
- **longjmp()** restores execution state back to execution in **f()**
- **f()** “catches the exception”
 - It now sees **setjmp()** return a nonzero result, indicating there was an exception...
 - **f()** returns -1 as final result

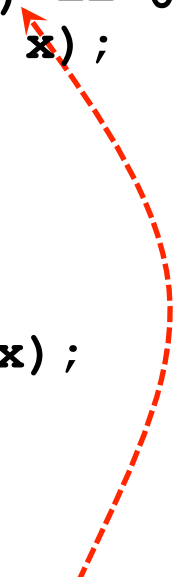
```
static jmp_buf env;

int f(int x) {
    if (setjmp(env) == 0)
        return g(3 * x);
    else
        return -1;
}

int g(int x) {
    return h(15 - x);
}

int h(int x) {
    if (x < 5)
        longjmp(env, 1);

    return sqrtl(x - 5);
}
```





setjmp()/longjmp() EXAMPLE (6)

```
int f(int x) {
    if (setjmp(env) == 0)
        ...
}
...
int h(int x) {
    if (x < 5)
        longjmp(env, 1);
    ...
}
```

- When **h()** calls **longjmp()**, **esp** and caller **eip/ebp** are restored from **env**
- When **longjmp()** returns, execution resumes in **f()**, “back at **setjmp()**”
 - Caller has no idea who returned back!
- Result of **setjmp()** indicates error
 - (but it's technically **longjmp()**'s result...)
- f()** handles the error appropriately

env -> esp



HOW DO THESE THINGS WORK?!

- **setjmp()** and **longjmp()** must be implemented in assembly language
 - No C mechanism for saving and restoring registers
 - No C mechanism for manipulating the stack this way
- Implementation is also very platform-specific!
 - Size of **jmp_buf** corresponds to how many registers need to be saved and restored
 - Linux on IA32 uses 8 dwords
 - MacOS X on IA32 uses 18 dwords (!)
 - MacOS X on PPC uses 192 dwords (!!!)
 - RISC processors tend to have a large number of registers, due to load/store architecture
 - Specification is ambiguous about exactly what needs to be saved...

HOW DO THESE THINGS WORK?! (2)

- Implementing **setjmp()**/**longjmp()** is surprisingly straightforward
 - Simply requires understanding of stack frames in cdecl
- In **setjmp()**, must know how to save the caller's execution state, to fake a return from **setjmp()**
 - Return-address where caller invoked **setjmp()** from
 - **esp** value inside **setjmp()**
 - (also the callee-save registers, since they will change before **longjmp()** is called...)
- In **longjmp()**, just need to manipulate the stack to restore this execution state, then **ret** !
 - Caller will see return-value in **eax** like usual
 - Returns to where caller invoked **setjmp()** from
 - *They'll never know the difference!*

MULTIPLE CATCH BLOCKS

- A **try** block can have multiple **catch** blocks

```
Webpage loadPage(String urlText) {  
    try {  
        Socket s = httpConnect(urlText);  
        ...  
    }  
    catch (MalformedURLException e) {  
        ...  
    }  
    catch (UnknownHostException e) {  
        ...  
    }  
}
```

- Easy to support this with **setjmp()** and **longjmp()**
 - **longjmp()** can simply pass a different integer value for each kind of exception
 - Compiler can assign integer values to all exception types

MULTIPLE CATCH BLOCKS (2)

- One possible translation:

```
jmp_buf env;
...
switch (setjmp(env)) {
case 0:      /* Normal execution */
    ...      // Translation of
    ...      //   Socket s = httpConnect(urlText);
    break;

case 1037:   /* Caught MalformedURLException */
    ...
    break;
case 1053:   /* Caught UnknownHostException */
    ...
    break;
}
```

- Code that calls **longjmp()** passes exception-type in call
- Many details left out, involving variable scoping, etc.*

NESTED EXCEPTION HANDLERS

- One major flaw in our implementation:
 - **try/catch** blocks can be nested within each other!
 - We only have one **jmp_buf** variable in our example
 - A nested **try/catch** block would overwrite the outer **try**-block's values stored in the **jmp_buf**
- Solution is straightforward:
 - Introduce a “try-stack” for managing the **jmp_buf** values of nested **try/catch** blocks
 - When we enter into a new **try**-block, push the new **try/catch** handler state (**jmp_buf**) onto try-stack
 - This is separate from the regular program stack
 - (It doesn't strictly *have* to be separate, but to keep things simple, we will keep it separate!)

NESTED EXCEPTION HANDLERS (2)

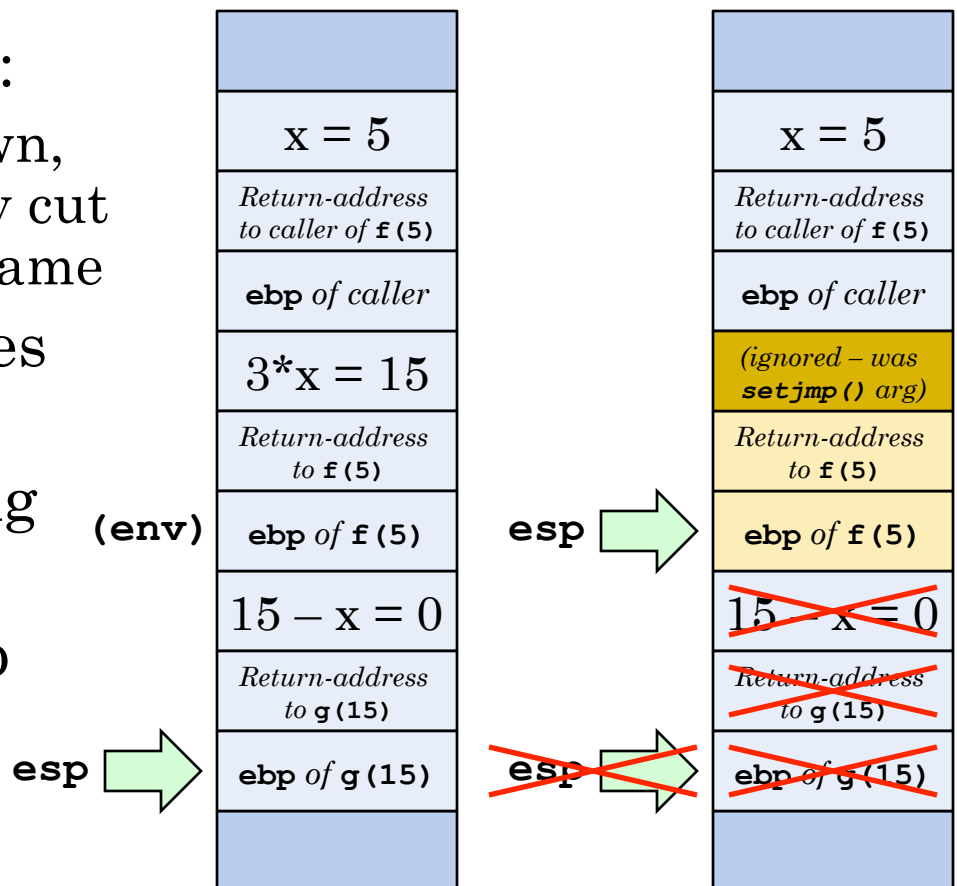
- Once we have a try-stack for nested handlers, need some basic exception handling operations
- When program enters a **try** block:
 - Call **setjmp()**, and if it returns 0 then push the **jmp_buf** onto the **try**-stack
 - **Note:** Cannot call **setjmp()** in a separate helper function that does these things for us! *Why not?*
 - Left as an exercise for the student...
- When program exits the **try**-block without any exceptions:
 - Need to pop the topmost **jmp_buf** off of the try-stack
 - Can do this in a helper function

NESTED EXCEPTION HANDLERS (3)

- When an exception is thrown:
`void throw_exception(int exception_id)`
 - Helper function that pops the topmost `jmp_buf` off of the try-stack, and then uses it to do `longjmp(exception_id)`
- If an exception isn't handled by a `try/catch` block, or if `catch`-block re-throws the exception:
 - Just invoke `throw_exception()` again with same ID
 - Next enclosing `try`-block's `jmp_buf` is now on top of stack
 - (Do this in `default` branch, or `else`-clause if using `if`.)
- With these tools in place, can easily handle nested exception-passing scenarios
- An example of this is provided in Assignment 4! 😊

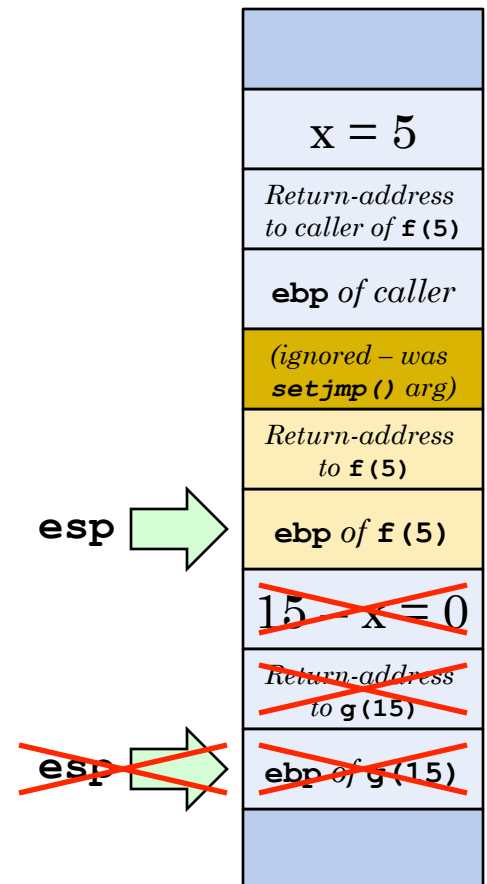
STACK CUTTING

- This kind of exception-handling implementation is called stack cutting
- From previous example:
 - When exception is thrown, the stack is immediately cut down to the handler's frame
- Intervening stack frames are simply eliminated!
- *Very fast* for propagating exceptions...
- Unacceptable if cleanup needs to be done for intervening functions!



STACK CUTTING (2)

- Can perform cleanup for intervening functions, if we keep track of what needs to be done
 - e.g. manage a list of resources that need cleaned up after each function returns
 - When exception is thrown, can use this info to clean up properly
 - Starting to get a bit *too* complicated...
- For languages with GC, don't really have much to clean up from functions
 - Just drop object-references from stack
 - Garbage collector will detect that objects are no longer reachable, and will eventually reclaim the space



STACK UNWINDING

- Another solution to the exception-propagation problem: stack unwinding
 - Solution used by Java Virtual Machine, most C++ implementations, etc.
 - Unlike stack cutting, each stack frame is cleaned up individually. Much better for resource management!
- Remember, the important times in exception handling are:
 - When we are inside of a **try**-block – a thrown exception might be handled by this **try/catch**
 - When an exception is actually thrown
- The compiler generates an exception table for every single function in the program
 - All exception handling is driven from these tables

EXCEPTION TABLES

- Each function has an exception table, containing:
 - A range of addresses [*from_eip*, *to_eip*], specifying the instructions the **try**-block encapsulates
 - An exception that the **try**-block can handle
 - The address of the handler for that exception
- Our example from before:

```
Webpage loadPage(String urlText) {  
    try {  
        Socket s = httpConnect(urlText);  
        ...  
    }  
    catch (MalformedURLException e) {  
        ...  
    }  
    catch (UnknownHostException e) {  
        ...  
    }  
}
```

Exception table for loadPage()

from_eip	to_eip	exception	handler_eip
0x3019	0x315C	malformed_url	0x316B
0x3019	0x315C	unknown_host	0x3188

EXCEPTION TABLES (2)

- When an exception is thrown within a function:
- Two important pieces of information!
 - What is the current program-counter?
 - What is the type of the exception that was thrown?
- Exception table for the current function is searched

Exception table for `loadPage()`

from_eip	to_eip	exception	handler_eip
0x3019	0x315C	malformed_url	0x316B
0x3019	0x315C	unknown_host	0x3188

- Try to find a row where the program-counter is in the specified range, also having the same exception type
 - If found, dispatch to the specified exception handler
- If no matching row is found, the current stack frame is cleaned up, and process repeats in parent frame

NESTED TRY/CATCH EXAMPLE

- Code with nested **try/catch** blocks
- Compiler generates an exception table for each function:

Exception table for f(x)

from_eip	to_eip	exception	handler_eip
0x2005	0x203B	a	0x2043
0x2005	0x203B	b	0x204C

Exception table for g(x)

from_eip	to_eip	exception	handler_eip
0x2116	0x214A	b	0x2159
0x2116	0x214A	c	0x215E

Exception table for h(x)

from_eip	to_eip	exception	handler_eip
----------	--------	-----------	-------------

```
int f(int x) {  
    try {  
        return g(x * 3);  
    } catch (A a) {  
        return -5;  
    } catch (B b) {  
        return -10;  
    }  
}
```

```
int g(int x) {  
    try {  
        return h(8 - x);  
    } catch (B b) {  
        return -15;  
    } catch (C c) {  
        return -20;  
    }  
}
```

```
int h(int x) {  
    if (x > 23)  
        throw new A();  
    else if (x < -15)  
        throw new B();  
  
    return x - 1;  
}
```

NESTED TRY/CATCH (2)

- Call **f(-9)**
- **f(-9)** calls **g(-9 * 3) = g(-27)**
- **g(-27)** calls **h(8 - -27) = h(35)**
- Important point:
 - So far, no overhead for entering **try**-blocks, or any other aspect of exception handling!
- But, we know that **h(35)** is going to throw...

```
int f(int x) {  
    try {  
        return g(x * 3);  
    } catch (A a) {  
        return -5;  
    } catch (B b) {  
        return -10;  
    }  
}
```

```
int g(int x) {  
    try {  
        return h(8 - x);  
    } catch (B b) {  
        return -15;  
    } catch (C c) {  
        return -20;  
    }  
}
```

```
int h(int x) {  
    if (x > 23)  
        throw new A();  
    else if (x < -15)  
        throw new B();  
  
    return x - 1;  
}
```

NESTED TRY/CATCH (3)

- **h(35)** throws exception **A**
- Use our exception tables to direct the exception propagation

- **h(35)** throws **A**. **eip = 0x226C**.

- Check exception table for **h**:

Exception table for h(x)

from_eip	to_eip	exception	handler_eip
----------	--------	-----------	-------------

- Nothing matches. (*duh...*)
- Clean up local stack frame, then return to caller of **h**

```
int f(int x) {  
    try {  
        return g(x * 3);  
    } catch (A a) {  
        return -5;  
    } catch (B b) {  
        return -10;  
    }  
}
```

```
int g(int x) {  
    try {  
        return h(8 - x);  
    } catch (B b) {  
        return -15;  
    } catch (C c) {  
        return -20;  
    }  
}
```

```
int h(int x) {  
    if (x > 23)  
        throw new A();  
    else if (x < -15)  
        throw new B();  
  
    return x - 1;  
}
```

NESTED TRY/CATCH (4)

- Now inside of **g(-27)**
- g(-27)** throws **A**. **eip** = **0x2123**.
 - Check exception table for **g**:

Exception table for g(x)

from_eip	to_eip	exception	handler_eip
0x2116	0x214A	b	0x2159
0x2116	0x214A	c	0x215E

- g** does have entries in its table, but none match combination of exception **A** and **eip** = **0x2123**.
- Again, clean up local stack frame, then return to caller of **g**

```
int f(int x) {  
    try {  
        return g(x * 3);  
    } catch (A a) {  
        return -5;  
    } catch (B b) {  
        return -10;  
    }  
}
```

```
int g(int x) {  
    try {  
        return h(8 - x);  
    } catch (B b) {  
        return -15;  
    } catch (C c) {  
        return -20;  
    }  
}
```

```
int h(int x) {  
    if (x > 23)  
        throw new A();  
    else if (x < -15)  
        throw new B();  
  
    return x - 1;  
}
```

NESTED TRY/CATCH (5)

- Finally, back to **f(-9)**
- **f(-9)** throws **A**. **eip = 0x201B**.

- Check exception table for **f**:

Exception table for f(x)

from_eip	to_eip	exception	handler_eip
0x2005	0x203B	a	0x2043
0x2005	0x203B	b	0x204C

- **f** also has exception table entries, and the first entry matches our combination of exception type and instruction-pointer value!
- Dispatch to specified handler:
 - **return -5;**
- Exception propagation is complete.

```
int f(int x) {  
    try {  
        return g(x * 3);  
    } catch (A a) {  
        return -5;  
    } catch (B b) {  
        return -10;  
    }  
}
```

```
int g(int x) {  
    try {  
        return h(8 - x);  
    } catch (B b) {  
        return -15;  
    } catch (C c) {  
        return -20;  
    }  
}
```

```
int h(int x) {  
    if (x > 23)  
        throw new A();  
    else if (x < -15)  
        throw new B();  
  
    return x - 1;  
}
```

COMPARISON OF METHODOLOGIES

- Stack cutting approach is optimized for the exception-handling phase
 - Transfers control to handler code in one step
 - (Presuming resources don't need to be cleaned up from intervening function calls...)
- Additional costs in the normal execution paths!
 - Need to record execution state every time a **try**-block is entered
 - Need to push and pop these state records, too!
- These costs will *quickly* add up in situations where execution passes through many **try**-blocks

COMPARISON OF METHODOLOGIES (2)

- Stack unwinding approach is optimized for the normal execution phase
 - No exception-handler bookkeeping is needed at run-time...
 - ...because all bookkeeping is done at compile-time!
- Additional costs in the exception-handling paths
 - Each function call on stack is dealt with individually
 - Must search through each function's exception table, performing several comparisons per record
- If a program *frequently* throws exceptions, especially from deep within call-sequences, this will definitely add up

COMPARISON OF METHODOLOGIES (3)

- Typical assumption is that exception handling is a relatively uncommon occurrence
 - *(That's why we call them exceptions!!!)*
 - Additionally, most languages have resources to clean up, within each function's stack frame
 - e.g. even though Java has garbage collection, it also has **synchronized** blocks; monitors need unlocked
- Most common implementation: stack unwinding
- Many other optimizations are applied to exception handling as well!
 - Dramatically reduce or even eliminate overhead of searching for exception handlers within each function