



CS24: INTRODUCTION TO COMPUTING SYSTEMS

Spring 2015

Lecture 12

CS24 MIDTERM

○ Midterm format:

- 6 hour overall time limit, multiple sittings
 - (If you are focused on midterm, clock should be running.)
- Open book/notes/slides/homeworks/solutions
- Open CS24 2014 Moodle, closed people/Internet/etc.
- Can use computer to implement and test your work (but not disassemble or generate IA32 code from C)

○ No collaboration!

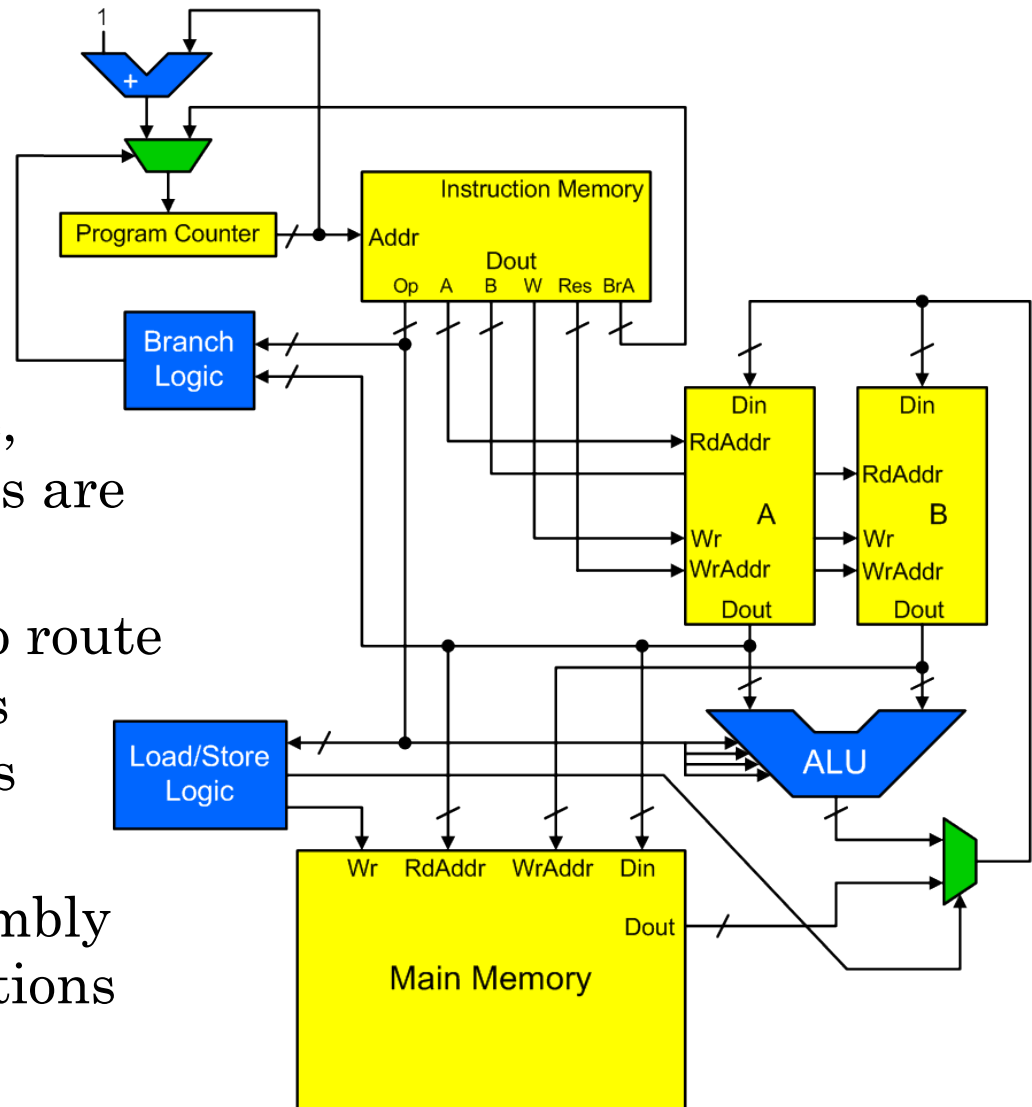
- If you need clarifications, feel free to email Donnie and/or **cs24tas** list (although Donnie has final say)

CS24 MIDTERM (2)

- Potential topics include:
 - **Anything covered in first half of class**
 - Basic processor structure, operation, low-level programming
 - Boolean logic, logical/bitwise operations in C
 - IA32 assembly language programming
 - Flow-control constructs in IA32 – **if, for, while, do**
 - Arrays, structs, heap allocation (**malloc/free**)
 - Explicit heap allocator internals
 - Garbage collection, objects, exceptions

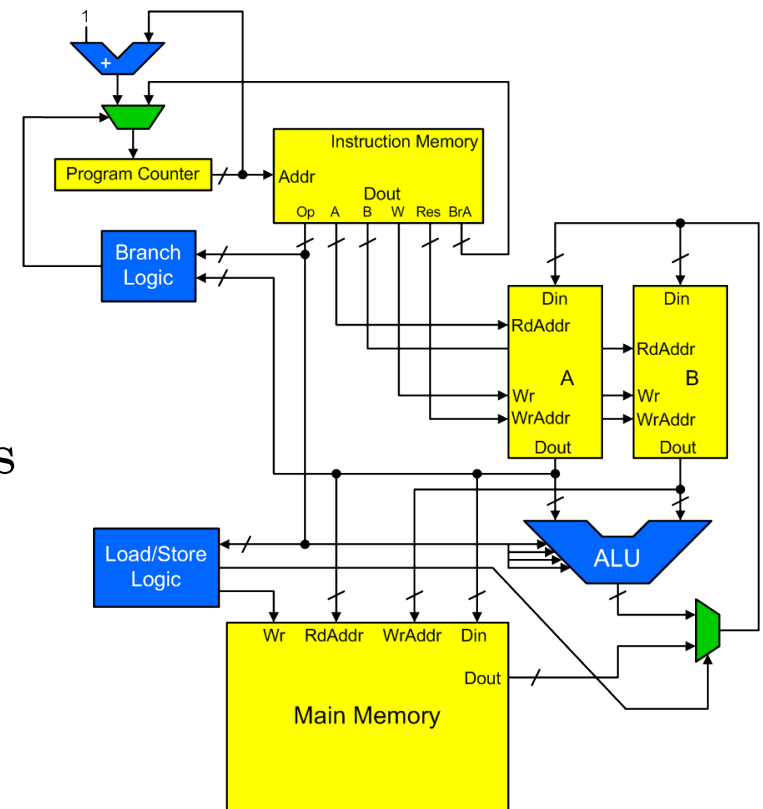
PROCESSOR STRUCTURE

- Covered basic CPU components in lectures 1-4
- General idea:
 - Based on the opcode, different components are enabled or disabled
 - Multiplexers used to route address/data signals between components
 - Feed a sequence of instructions to assembly to perform computations



PROCESSOR STRUCTURE (2)

- Midterm will present a simple variant for you to examine
- Example questions:
 - Write C code to simulate the logic that controls components
 - Write a simple machine-code program to control the CPU
 - Analyze characteristics (i.e. strengths and weaknesses) of the alternate design
- (Similar to problems on HW1)



C LOGICAL AND BITWISE OPERATIONS

- Should be familiar with all C logical and bitwise operations
- C uses integers to represent Boolean values
 - 0 = false; any nonzero value = true
- Logical Boolean operators:
 - Logical AND: **a && b**
 - Logical OR: **a || b**
 - Logical NOT: **!a**
 - Result is 1 if true, 0 if false
- **&&** and **||** are short-circuit operators
 - Evaluated left-to-right
 - For **&&**, if LHS is false then RHS is not evaluated
 - For **||**, if LHS is true then RHS is not evaluated

C LOGICAL AND BITWISE OPERATIONS (2)

- Bitwise operators manipulate individual bits
- Given $a = 00010100_2$ (20_{10}), $b = 00110010_2$ (50_{10})
 - $a \ \& \ b = 00010000$ Bitwise AND
 - $a \ | \ b = 00110110$ Bitwise OR
 - $\sim a = 11101011$ Bitwise negation (invert)
 - $a \ ^ \ b = 00100110$ Bitwise XOR
- Also, bit-shift operators
 - $a \ >> \ 2 = 00000101$ (5_{10})
 - $a \ << \ 1 = 00101000$ (40_{10})
- Shift-right has two versions!
 - Arithmetic shift-right preserves sign (topmost bit)
 - Logical shift-right doesn't preserve sign
 - In C, type of LHS determines if arithmetic or logical

IA32 ASSEMBLY LANGUAGE

- Midterm will include problems writing and/or understanding IA32 assembly programs
- Be familiar with IA32 execution environment
 - Registers: **eax**, **ebx**, **ecx**, **edx**, **esp**, **ebp**, **esi**, **edi**
 - Stack pointer in **esp**; stack grows downward
 - Set up stack frame using **ebp**
- Also, be familiar with cdecl calling convention!
 - Caller-save registers: **eax**, **ecx**, **edx**
 - Callee-save registers: **ebp**, **ebx**, **esi**, **edi**
 - Arguments passed on the stack: **ebp** + *offset*
 - Local variables stored on the stack: **ebp** – *offset*
 - Return-value passed back in **eax**

Lecture 4

Lecture 5

CALLING C FUNCTIONS FROM IA32

- Be familiar with how to call standard C functions from IA32, e.g. **malloc()**
- Important note: these functions also use cdecl ☺
 - Save any caller-save registers you want to preserve!
 - e.g. **malloc()** may change **eax**, **ecx**, **edx**
 - If you need them, save them on the stack!
 - ...*before* you push your arguments to malloc...

Lecture 7

WRITING IA32 ASSEMBLY PROGRAMS

- Very helpful to follow a specific pattern when implementing IA32 functions

```
my_function:
```

```
    # TODO:  Document arguments and offsets here
```

```
    pushl    %ebp          # Save caller ebp
```

```
    movl     %esp, %ebp    # Set up stack frame ptr
```

```
    # TODO:  Save callee-save registers we alter
```

```
    # TODO:  Function body
```

```
    # TODO:  Restore callee-save registers we alter
```

```
    # Clean up stack before returning to caller
```

```
    movl     %ebp, %esp    # Remove local vars
```

```
    popl     %ebp          # Restore caller ebp
```

```
    ret
```

WRITING IA32 ASSEMBLY PROGRAMS (2)

- Usually need to set up and tear down stack frame

`my_function:`

`# TODO: Document arguments and offsets here`

`pushl %ebp # Save caller ebp`

`movl %esp, %ebp # Set up stack frame ptr`

`# TODO: Save callee-save registers we alter`

`# TODO: Function body`

`# TODO: Restore callee-save registers we alter`

`# Clean up stack before returning to caller`

`movl %ebp, %esp # Remove local vars`

`popl %ebp # Restore caller ebp`

`ret`

WRITING IA32 ASSEMBLY PROGRAMS (3)

- Figure out where your arguments are stored
 - cdecl specifies that args are pushed in reverse order!

my_function:

```
# TODO: Document arguments and offsets here
pushl    %ebp          # Save caller ebp
movl     %esp, %ebp    # Set up stack frame ptr
...
```

- Example from lecture 7: vector-add function

vector_add:

```
# a      = 8(%ebp)          (Arg 1, pushed last)
# b      = 12(%ebp)         (Arg 2)
# length = 16(%ebp)         (Arg 3, pushed first)
pushl    %ebp          # Save caller ebp
movl     %esp, %ebp    # Set up stack frame ptr
```

WRITING IA32 ASSEMBLY PROGRAMS (4)

- Easiest to write function body in pseudocode, then translate to IA32

```
my_function:
```

```
...
```

```
# TODO: Save callee-save registers we alter
```

```
# TODO: Function body
```

```
# TODO: Restore callee-save registers we alter
```

```
...
```

- No point in saving/restoring registers until you know what you use in the implementation...
 - Fill in those parts last
 - *(Or, just save them all, but we can do better...)*
- Return-value in **eax**

IA32 AND C FLOW-CONTROL CONSTRUCTS

Lecture 6

- Need to be familiar with how C flow control constructs are implemented in IA32 assembly
 - e.g. given a C program that uses **if**, **for**, **while**, **do**, be able to translate into IA32 assembly with labels and conditional jumps
 - At very least, can take a problem, write it in C-style pseudocode, then translate into IA32
- IA32 conditional jump operations are based on contents of **eflags** register
 - Most relevant flags:
 - CF = carry flag (1 indicates unsigned overflow)
 - SF = sign flag (1 = result is negative)
 - OF = overflow flag (1 indicates signed overflow)
 - ZF = zero flag (1 = result is zero)

IA32 CONDITIONAL JUMPS

- Arithmetic and logical operations update **eflags**
 - Sometimes can use an arithmetic operation to set up for a conditional jump
- Can also update **eflags** with **cmp** and **test**
 - **cmp Src2, Src1**
 - Updates flags as for $\text{Src1} - \text{Src2}$ (i.e. **sub Src2, Src1**)
 - **test Src2, Src1**
 - Updates flags as for $\text{Src1} \& \text{Src2}$ (i.e. **and Src2, Src1**)
 - Src1, Src2 unchanged by comparison/test operation
 - Specify size prefixes, as usual: **cmpl %ecx, \$0**
- **Most important detail:** AT&T syntax imposes confusing order on compare operands!

IA32 CONDITIONAL JUMPS (2)

- A sometimes-useful trick: **test Src, Src**
 - Comparing a value against itself
- Easy way to tell if a value is positive, negative, or zero
 - Causes the Sign Flag and Zero Flag to be set based on Src
 - SF = 1 Value is negative
 - SF = 0, ZF = 1 Value is zero
 - SF = 0, ZF = 0 Value is positive

FOR-LOOP IN IA32

- Remember: **for** loop is a **while** loop with extra features

```
for (i = 0; i < length; i++)  
    result[i] = a[i] + b[i];
```

- Identical to:

```
i = 0;  
while (i < length) {  
    result[i] = a[i] + b[i];  
    i++;  
}
```

- If you start with C code or pseudocode:
 - Can manually perform these transformations to get to IA32 assembly language

FOR-LOOP IN IA32 (2)

- Also turn **while** loops into **do** loops to simplify coding

```
i = 0;
while (i < length) {
    result[i] = a[i] + b[i];
    i++;
}
```

- Identical to:

```
i = 0;
if (!(i < length)) // Factor out first test
    goto end_for;
start_for:
    result[i] = a[i] + b[i];
    i++;
    if (i < length) // Subsequent tests at end
        goto start_for;
end_for:
    ...
```

FOR-LOOP IN IA32 (3)

- ...and then the tricky part: figuring out **cmp**

- C code:

```
if (!(i < length))    /* Or i >= length */  
    goto end_for;
```

- **cmpl Src2, Src1** sets flags as for Src1 – Src2
 - Specify **length** as Src2, **i** as Src1 so that appropriate conditional jump instruction mirrors the C code

- IA32 code: (**i** in **%esi**, **length** at **16(%ebp)**)

```
cmp 16(%ebp), %esi    # if (i >= length)  
jge end_for           #      goto end_for;
```

- First arg in compare is Src1, second arg is Src2
 - Args will appear in reverse order, but the jump will more closely match your C code or pseudocode

SIGNED AND UNSIGNED COMPARISONS

- IA32 includes different instructions for signed and unsigned comparisons!
- Signed: jump if greater/less
 - **jg, jge** – jump if greater, jump if greater or equal
 - **jl, jle** – jump if less, jump if less or equal
 - Instructions examine sign flag, overflow flag, and zero flag to determine signed comparison result
- Unsigned: jump if above/below
 - **ja, jae** – jump if above, jump if above or equal
 - **jb, jbe** – jump if below, jump if below or equal
 - Instructions examine carry flag and zero flag to determine unsigned comparison result

IA32 ARITHMETIC/LOGICAL OPERATIONS

- IA32 provides a variety of arithmetic and logical operations
 - Addition, subtraction, multiplication, division, increment, decrement
 - AND, OR, NOT, XOR, shift left, arithmetic shift right, logical shift right
- Can specify a variety of operand types
 - Source argument can be Immediate, Register, Memory (direct or indirect)
- Destination argument can be:
 - Register, Memory (direct or indirect)
- Both arguments cannot be a memory reference

IA32 MULTIPLY OPERATIONS

- IA32 has a *wide variety* of multiply operations!
 - **imul** = signed multiply
 - **mul** = unsigned multiply
- **mul** only has a 1-operand form
 - Multiply **al**, **ax**, or **eax** by the operand
 - Store result in **ax**, **eax**, or **edx:eax**, respectively
 - **edx** contains high dword, **eax** contains low dword
- **imul** has 1-, 2-, and 3-operand forms!
 - One-operand form is identical to **mul**
 - 2- and 3-operand forms don't implicitly use **al/ax/eax**
 - Result also truncated to same bit-width as arguments

IA32 DIVIDE OPERATIONS

- IA32 has two division instructions
 - **idiv** = signed division
 - **div** = unsigned division
 - Only one form for each: takes one operand
- Instruction divides **ax**, **dx:ax**, or **edx:eax** by the argument
 - Produces both quotient and remainder
- Example: dividing **edx:eax** by some value
 - Quotient is stored in **eax**
 - Remainder is stored in **edx**

C, IA32, AND ARRAYS

Lecture 6

- Need to be familiar with arrays in C and IA32
- For an array declaration: ***T A[N]***
 - ***T*** is the data type
 - ***A*** is the array's variable name
 - ***N*** is the number of elements
- C allocates a contiguous region of memory for array
 - Allocates ***N × sizeof(T)*** bytes for the array
- Variable ***A*** holds a pointer to start of array
 - Can use ***A*** to access various elements of the array
- Can access elements of ***A*** using pointer arithmetic
 - e.g. ***A[i]*** is equivalent to ****(A + i)***
 - ***A + i*** moves pointer forward ***i*** elements, not ***i*** bytes
 - C figures out how large each element is, from type of ***A***

C, IA32, AND ARRAYS (2)

- IA32 memory addressing modes make it *easy* to access arrays from assembly language routines
- Indexed memory access
 - **(RegB, RegI)** accesses $M[\text{RegB} + \text{RegI}]$
 - RegB is the base (i.e. starting) address of a memory array
 - RegI is an index into the memory array
 - **Imm(RegB, RegI)** accesses $M[\text{Imm} + \text{RegB} + \text{RegI}]$
 - Assumes that array elements are bytes
- Scaled indexed memory access
 - With scale factor $s = 1, 2, 4, 8$ (i.e. size of array element):
 - **(, Reg, s)** $M[\text{Reg} \times s]$
 - **Imm(, Reg, s)** $M[\text{Imm} + \text{Reg} \times s]$
 - **(RegB, RegI, s)** $M[\text{RegB} + \text{RegI} \times s]$
 - **Imm(RegB, RegI, s)** $M[\text{Imm} + \text{RegB} + \text{RegI} \times s]$

IA32 MEMORY-REFERENCE SUMMARY

- Summary chart, from IA32 manual:

Base	+	Index	×	Scale	+	Displacement
eax		eax				
ebx		ebx				
ecx		ecx		1		None
edx		edx		2		8-bit
esp	+	<i>(not esp!)</i>	×	4	+	16-bit
ebp		ebp		8		32-bit
esi		esi				
edi		edi				

- Base, Index, Displacement all optional
- Scale only allowed when Index is specified
- Note that **esp** can only be used as a base value

C, IA32, AND ARRAYS (3)

- Again, our vector-add function loop-body:

```
result[i] = a[i] + b[i];
```

- IA32 assembly code:

```
vadd_loop:
```

```
    movl (%ebx, %esi, 4), %edx    # edx = a[i]
    addl (%ecx, %esi, 4), %edx    # edx += b[i]
    movl %edx, (%eax, %esi, 4)    # r[i] = edx
```

- Our arrays contain **int** elements
 - **sizeof(int)** = 4 bytes, so scale = 4
 - **ebx** = starting address of array **a**
 - **ecx** = starting address of array **b**
 - **eax** = starting address of array **result**
 - **esi** = array index **i**

C, IA32, AND STRUCTS


- IA32 addressing modes also make it easy to work with structs
- From lecture 8:

```
struct s1 {  
    int i;  
    char ch;  
    int j;  
};
```

- If **s1 *s** is stored in **%ebx**
 - **s->i** **movl (%ebx), %eax**
 - **s->ch** **movb 4(%ebx), %cl**
 - **s->j** **movl 8(%ebx), %edx**

HEAP ALLOCATORS


- Explicit allocators:
 - Program is responsible for allocating/freeing memory
- Memory fragmentation:
 - Requested amount of memory is available, but no block is actually large enough to satisfy the request
- Example: heap with 6KB total space
 - A = allocate 1KB
 - B = allocate 1KB
 - C = allocate 2KB
 - D = allocate 1KB
 - Free B
- Try to allocate 2KB?
 - 2KB of memory is available, but it's not contiguous!



COALESCING FREE BLOCKS

- Explicit allocators coalesce adjacent free blocks

- A = allocate 1KB
- B = allocate 1KB
- C = allocate 1KB
- D = allocate 2KB
- Free C
- Free B
- Try to allocate 2KB?




- Very important for avoiding false fragmentation!

- In above example, allocator will coalesce B with C when B is freed

PLACEMENT POLICIES

- Allocators employ a placement policy to choose which free block to use for a request
- Example: want to allocate 1KB
 - Two different free blocks to choose from
- First-fit policy: choose the first free block that satisfies the request
- Next-fit policy: like first-fit, but start searching where previous search ended
- Best-fit policy: examine all free blocks; choose the smallest free block that satisfies the request
- For our example:
 - First-fit policy would cause memory fragmentation
 - Best-fit policy avoids fragmentation



IMPLICIT ALLOCATORS

- C language is very powerful, but also allows us to do many dangerous things
 - Buffer overflows, exploits, memory leaks!
- Introduce higher-level abstractions:
 - Eliminate pointers and pointer arithmetic
 - An opaque “reference” to allocated objects
 - Use an implicit allocator to manage memory
- Implicit allocators use garbage collection to find unreachable blocks
 - Mark-and-sweep GC vs. stop-and-copy GC
 - (also reference-counting, but is vulnerable to cycles)

OBJECTS, EXCEPTION HANDLING

- Additionally, can provide more powerful programming abstractions
 - Object-oriented programming
 - Exception handling
- Examined **setjmp()** and **longjmp()**
 - Still works within cdecl, but enables *non-local jumps*
 - Implements an alternate flow-control mechanism by manipulating the stack in well-defined ways
- Midterm may also include topics related to these mechanisms
 - **Familiarity with cdecl, stack frames, etc. will be very important.** (*But, you need that anyway.*)

FINAL MIDTERM NOTES

- Midterm will be available towards end of week
 - Due Thursday, May 7 at 2AM, as usual
- Solution sets for all assignments will be available before midterm is released
 - **Please review the solutions before taking the exam!** I frequently give more background detail in the solution sets.
- Good luck! ☺