



# CS24: INTRODUCTION TO COMPUTING SYSTEMS

Spring 2015

Lecture 20

# LAST TIME: UNIX PROCESS MODEL

- Began covering the UNIX process model and API
- Information associated with each process:
  - A PID (process ID) to identify the process' context
  - Also, a parent process ID and a process group ID
  - Three states: Running, Stopped, Terminated
- Standard API calls to work with processes:
  - **fork()** creates a new process
  - **exit()** terminates a running process
  - **wait()**, **waitpid()** reap terminated (“zombie”) child processes
  - **execve()** loads and runs a program in a process
  - **kill()** sends a signal to another process
- The kernel provides all of these operations

## LAST TIME: THE **init** PROCESS

- Can build powerful facilities using process model
- **init** is the ancestor of all processes
  - Started as the last step of kernel boot sequence
  - PID of **init** is 1
- Purpose of **init** is to manage other processes in the operating system
- Switches between different runlevels
  - Each runlevel has a different (possibly overlapping) set of processes that are running
  - e.g. single-user, multi-user networking, X11
- Also handles some process-termination scenarios
  - (more on this today...)

# LAST TIME: ZOMBIE PROCESSES

- A child process doesn't immediately go away when it terminates
  - Child process terminates with some status value...
  - Parent process may need to find out the child's status
- A terminated child process is called a zombie
  - The process is dead, but it hasn't yet been reaped
- Parent processes reap zombie children by calling:  
`pid_t wait(int *status)`
  - Waits for some child process to terminate  
`pid_t waitpid(pid_t pid, int *status, int options)`
  - Waits for a specific child process to terminate
  - Can also wait on children in a process-group, or all children
  - Both report an error if calling process has no children
  - Helpful functions for extracting details from **status**

# NOTIFICATIONS FROM ZOMBIE CHILDREN

- When a child process terminates, the kernel also sends a **SIGCHLD** signal to its parent
  - Child process calls back to kernel when it terminates, so kernel can inform the process' parent
  - Default behavior is to ignore this signal
- Parent can set up a signal handler for **SIGCHLD** to reap the zombie child process

# PARENT/CHILD PROCESS EXAMPLE

```
#include <unistd.h>
#include <signal.h>
#include <sys/wait.h>
#include <stdio.h>

/* Handle SIGCHLD signals. */
void handle_sigchld(int sig) {
    pid_t pid;
    int status;

    pid = wait(&status);

    /* NOT REENTRANT! */
    /* Avoid in practice! */
    printf("Reaped child %d\n",
           pid);
    sleep(1);
}

int main() {
    int i;
    signal(SIGCHLD, handle_sigchld);

    for (i = 0; i < 3; i++) {
        if (fork() == 0) {
            /* Child-process code. */
            printf("Hello from child %d\n",
                   getpid());
            sleep(5);
            return 0; /* Terminate child */
        }
    }

    /* Parent-process code. */
    while (1) /* Wait for children */
        pause(); /* to terminate. */

    return 0;
}
```

## PARENT/CHILD PROCESS EXAMPLE (2)

- Save, compile, and run from the command line:  
`[user@host:~]> ./reaped`  
`Hello from child 1099!`  
`Hello from child 1101!`  
`Hello from child 1100!`
  - *(5 seconds pass)*  
`Reaped child 1101`
    - *(1 second passes)*  
`Reaped child 1100`
      - *(...and then, nothing else...)*
- Hmm, last child process doesn't get reaped.
  - **ps** reports that process 1099 is a zombie ☹
  - `1099 ttys000 00:00:00 reaped <defunct>`

# OUR EXAMPLE'S SOURCE CODE

```
#include <unistd.h>
#include <signal.h>
#include <sys/wait.h>
#include <stdio.h>

/* Handle SIGCHLD signals. */
void handle_sigchld(int sig) {
    pid_t pid;
    int status;

    pid = wait(&status);

    /* NOT REENTRANT!      *
     * Avoid in practice! */
    printf("Reaped child %d\n",
           pid);
    sleep(1);
}
```

Reaps one zombie  
per **SIGCHLD** received

```
int main() {
    int i;
    signal(SIGCHLD, handle_sigchld);

    for (i = 0; i < 3; i++) {
        if (fork() == 0) {
            /* Child-process code. */
            printf("Hello from child %d\n",
                   getpid());
            sleep(5);
            return 0; /* Terminate child */
        }
    }

    while (1)
        pause();

    return 0;
}
```

- Parent starts 3 child processes.
- Children terminate after five seconds.
- Kernel sends three **SIGCHLD** signals to the parent process.



# BUGGY SIGNAL HANDLER!

- Problem:
  - Parent process is sent three **SIGCHLD** signals in a row
  - Each **SIGCHLD** signal takes one second to handle
    - Due to the **sleep(1)** in the signal handler
  - First **SIGCHLD** received:
    - Parent process handles the **SIGCHLD** signal
    - Kernel blocks other **SIGCHLD** signals while handler is running!
  - Second **SIGCHLD** received:
    - Parent can't receive it yet since it's still in its handler
    - Kernel records the pending **SIGCHLD** signal...
  - Third **SIGCHLD** received:
    - Parent is still in its **SIGCHLD** handler for the first signal
    - Since the process already has a pending **SIGCHLD** signal, the third **SIGCHLD** is discarded
- When parent's **SIGCHLD** handler returns, kernel delivers pending **SIGCHLD**
  - Third, dropped **SIGCHLD** is never delivered

# A BETTER SIGNAL HANDLER

- The kernel does not queue up signals for delivery!
  - Only has a **pending** bit-vector to track next signals to deliver
- Instead, handler should reap as many zombies as it can, each time it's invoked

```
void handle_sigchld(int sig) {
    pid_t pid;
    int status;

    while (1) {
        pid = waitpid(-1, &status, WNOHANG);
        if (pid <= 0)    /* No more zombie children to reap. */
            break;

        /* NOT REENTRANT!  Avoid in practice! */
        printf("Reaped child %d\n", pid);
    }
    sleep(1);
}
```

(CS:APP Figs 8.32-33 don't use **WNOHANG**, but they should...)

# SIGNAL HANDLERS

- Signals cannot be used to count the number of occurrences of an event
- Signals indicate that *at least one* event of that type has occurred
- If expecting multiple signals of a particular type, write the handler to do as much work as possible
  - When your handler is invoked, you know that *at least one* signal of the type was received...

# PROCESS TERMINATION NOTES

- Every process has a parent process
  - The parent must be a currently running process
- If a parent process dies before its children do:
  - The kernel makes **init** (process 1) the parent of the orphaned children
  - When child processes die, **init** receives a **SIGCHLD** and then reaps them
- If a process terminates while it still has zombie children, the **init** process reaps the zombies too
  - The child processes become children of **init**...
  - Then, **init** is informed that it has zombie children
  - **init** reaps the zombie processes

# LOADING AND RUNNING PROGRAMS

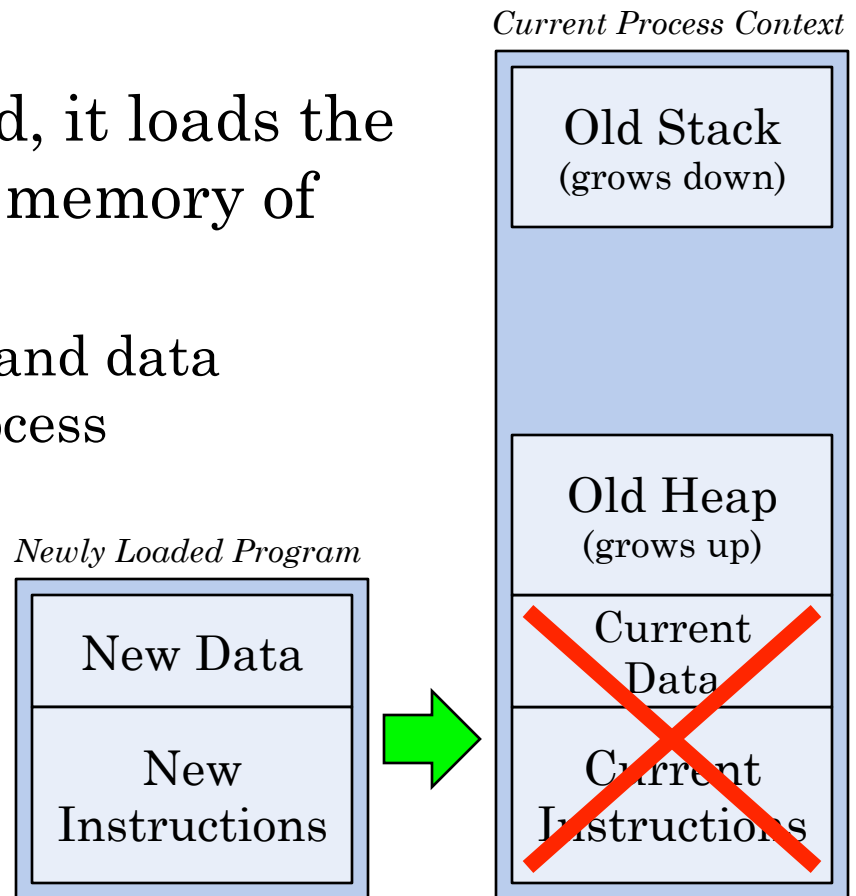
- The **execve()** function is used to load and run a program in the current process context

```
int execve(char *filename,  
           char *argv[], char *envp[])
```

  - **filename** is the binary file (or script) to load and run
    - Scripts must start with “**#! interpreter [args]**” line
  - Program arguments are specified by **argv**
    - These are the command-line arguments
    - Last element in **argv** must be **NULL**
  - The program’s environment is specified in **envp**
    - Each environment variable is a “NAME=VALUE” pair
    - Last element in **envp** must be **NULL**
- On success, loaded program *replaces* the current process’ context and starts running from its start
  - On success, **execve()** does not return!


## LOADING AND RUNNING PROGRAMS (2)

- **execve ()** is an operating system function
  - Kernel has total control over what happens in processes
- When **execve ()** is invoked, it loads the specified program into the memory of the current process
  - Overwrites the instruction and data segments of the current process



# LOADING AND RUNNING PROGRAMS (3)

- Next, **execve()** calls the program's entry point
  - For Linux on IA32, always at address 0x08048000
- For C/C++ programs, the entry point sets up the program's stack and heap
- New stack includes:
  - The new program's environment variables (passed to **execve()**)
  - The program's command-line arguments (also passed to **execve()**)
  - Arguments to the program's **main()** function



# NEW PROGRAM'S STACK

- First, program's environment strings and command-line arguments are copied to top of stack
- Next, program's environment values are pushed onto stack
  - Pointers reference the strings higher up in the stack
  - unistd.h** defines an external **environ** pointer for accessing these values
- Several functions for accessing environment variables

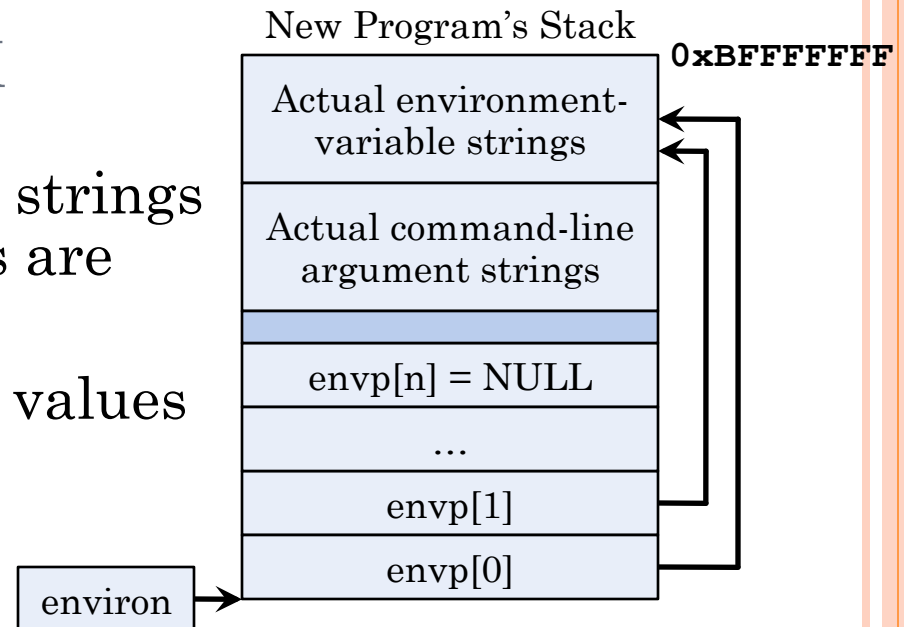
```
char * getenv(char *name)
```

```
int putenv(char *string)
```

```
int clearenv()
```

```
int setenv(char *name, char *value, int overwrite)
```

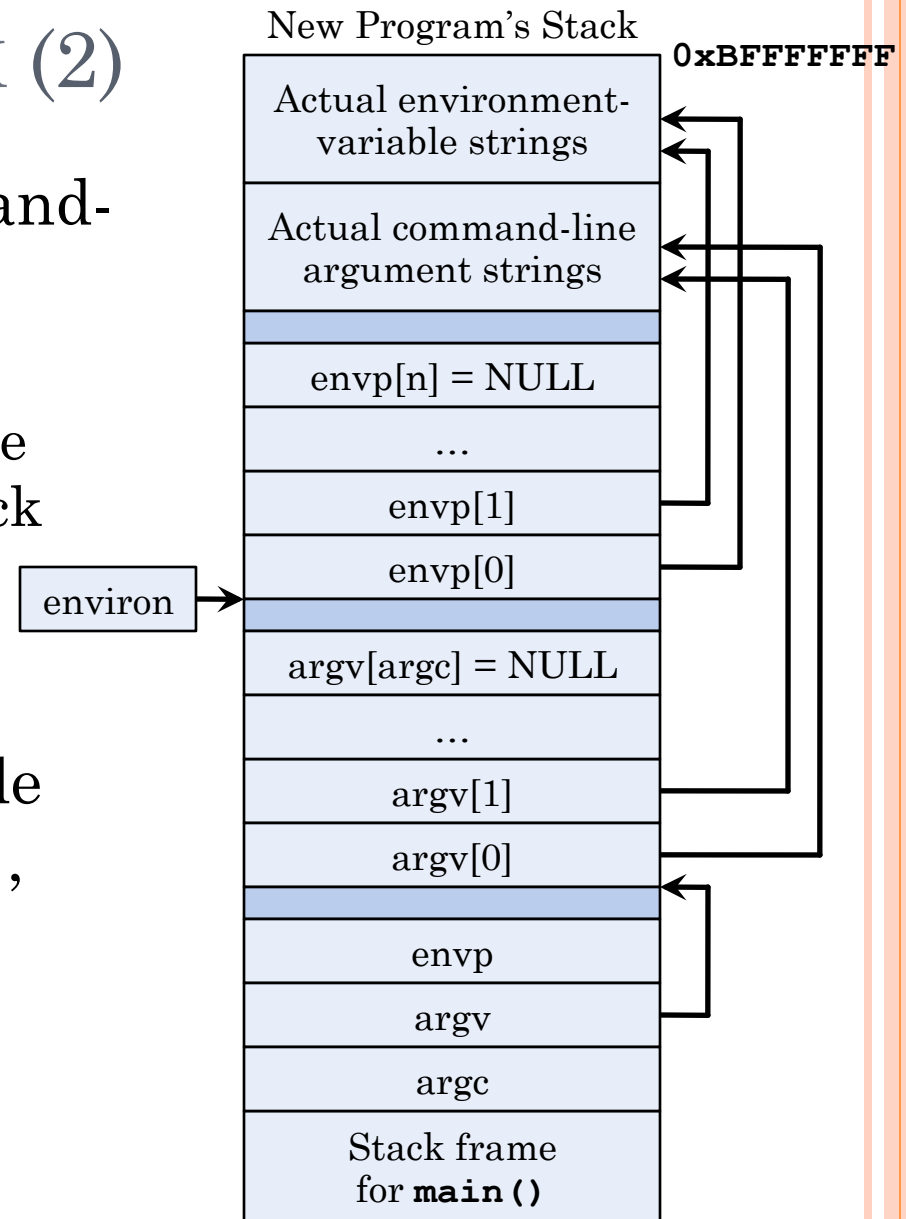
```
int unsetenv(char *name)
```





## NEW PROGRAM'S STACK (2)

- Next, the program's command-line arguments are pushed onto the stack
  - Again, pointers reference the strings higher up in the stack
- Finally, the entry-point code sets up for a call to **main()**, and then calls it!



# fork () AND execve ()

- UNIX shells use a combination of **fork ()** and **execve ()** to run programs

```
void eval(char *cmdline) {
    char *argv[MAXARGS];
    pid_t pid;
    int status;

    parseline(cmdline, argv);          /* Tokenize command */
    pid = fork();                      /* Spawn a process to run command */
    if (pid == 0) {                    /* This is the child process! */
        /* Replace the shell process with the new program */
        if (execve(argv[0], argv, environ) < 0) {
            printf("%s:  command not found!\n", argv[0]);
            exit(1);
        }
    }

    /* Shell process waits for the command to finish. */
    waitpid(pid, &status, 0);
}
```

## `fork()` AND `execve()` (2)

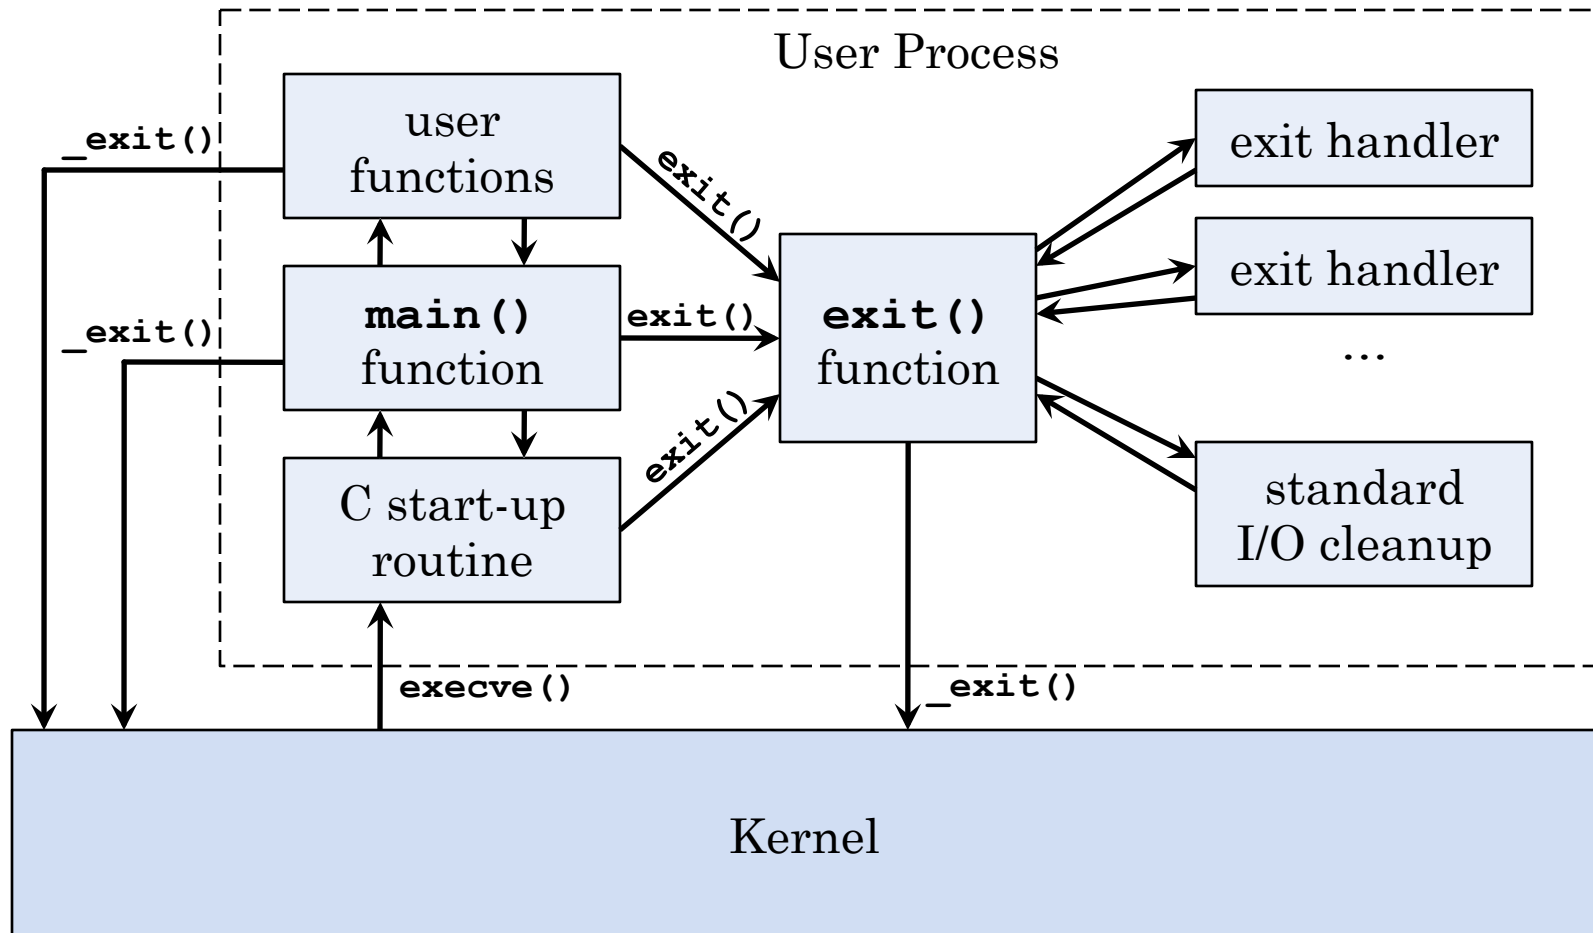
- Basic idea:
  - `fork()` off a child process, then the child `execve()`s the new program
  - New program *completely replaces* the shell state in the child process
- Many useful variations! For example:
  - Background processes
    - If command should be run in background, shell simply doesn't `waitpid()` for the child to terminate
  - Built-in commands
    - Shell can check if specified command is a built-in operation
    - Update shell's internal state instead of forking off a process
  - Input and output redirection
    - Child process can change stdin and stdout file-descriptors before executing the new program

## `fork()` AND `execve()` (3)

- This example code also lacks important features
  - Doesn't include any error checking on UNIX process functions, which is absolutely essential!
  - If we add support for background processes, need to reap them! (Requires a signal handler for this.)
- For a more detailed discussion of shell program structure, see CS:APP §8.4.6, and Problem 8.26
- For a *very* detailed discussion of process control:
  - Advanced Programming in the UNIX Environment (2<sup>nd</sup> edition) by W. Richard Stevens, Stephen Rago

# EXEC/EXIT PROCESS LIFECYCLE

- Simple diagram of process lifecycle, from **execve()** to **exit()**:



# SUMMARY: UNIX PROCESS MODEL

- UNIX process model is simple, but very powerful
  - API is clean and relatively straightforward
  - Can easily build up very sophisticated abstractions and systems using the simple model
- In last two lectures, we have covered:
  - How to start processes, how to handle process termination, how to load and start programs, ...
- This is enough to implement basic interactive-shell support for the computer!
- Apps take these tools provided by OS, and build even more sophisticated services with them

# UNIX PROCESSES: UNDER THE HOOD

- This is how the UNIX process abstraction is exposed to user programs...
- How this model is implemented in the kernel is *much* more involved
  - Process states and state-transitions
  - Details that the kernel must keep track of
  - Process behavior and scheduling
- Many platform-specific design choices to make!
  - This will be a high-level overview of general themes

# PROCESS STATES

- Processes follow a specific series of state transitions
- UNIX process API exposes these states:
  - Running – the process is actively running on a CPU, or is waiting to be executed
  - Stopped – the process will not be scheduled for execution until it is resumed
  - Terminated – the process has permanently stopped executing
- Process model allows multiple running processes...
  - ...but really, the number of CPUs dictates how many processes can actually run at any given time
  - e.g. 4 cores = 4 running processes, and that's it!
- Kernel must distinguish between processes that are running, and processes that are “ready to run”
  - Ready processes are waiting to get a turn on the CPU



## PROCESS STATES (2)


- Running processes can make slow system calls
  - Read or write a block of data to the disk
  - Read or write a data buffer to the network
  - Read user input
  - Kernel mediates these operations, so it can schedule other processes while long-running tasks complete
- Introduce another process state: blocked
  - A blocked process can't continue until some external condition changes
  - e.g. a DMA transfer from disk to memory must complete, and a hardware interrupt must occur
  - While blocked, makes no sense to schedule process!

## PROCESS STATES (3)

- Running processes can be stopped
  - e.g. type Ctrl-Z at terminal, which sends **SIGTSTP** signal to the process
  - Stopped processes are not scheduled until they are resumed, by sending a **SIGCONT** signal to them
- Blocked processes necessarily complicate this...
  - A blocked process can be stopped and resumed
  - While a blocked process is stopped, it can become unblocked
- Kernel needs to keep track of all of these states
- Define a state diagram for processes
  - Specifies the states and transitions for processes managed by the kernel

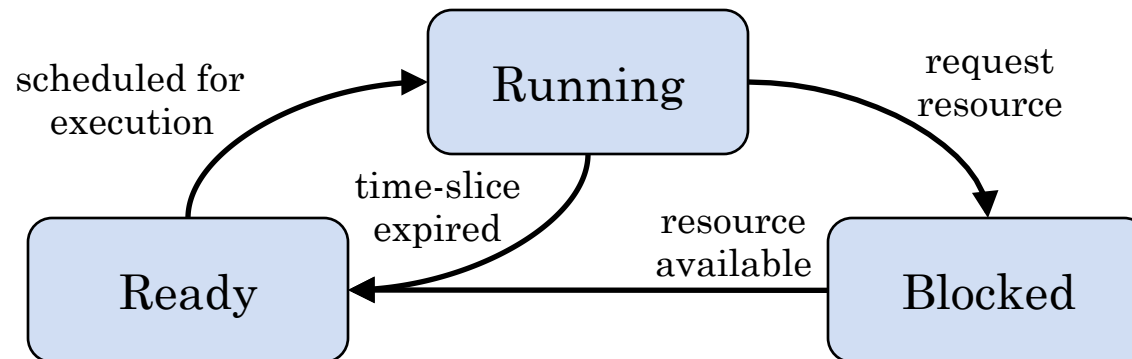
# UNIX PROCESS STATE DIAGRAM (1)

- One process is actually Running, per CPU
  - Ready processes wait for access to the CPU
  - Kernel interrupts the running process to run another



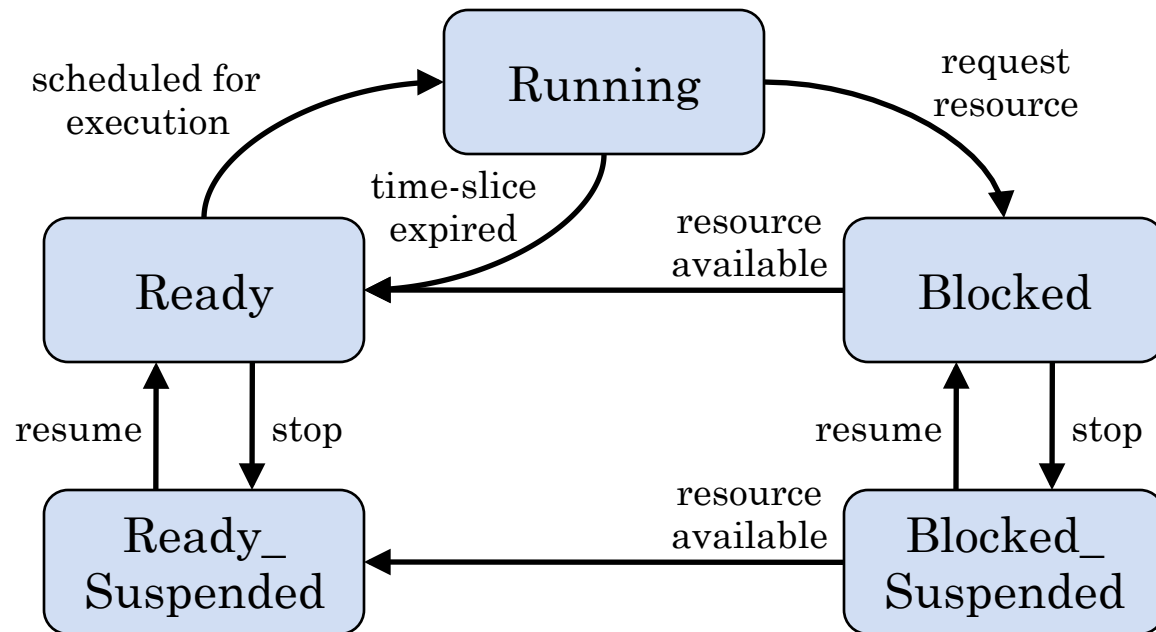
## UNIX PROCESS STATE DIAGRAM (2)

- A Running process can become Blocked
  - Request a resource that won't be available for a while
  - When resource is available, changes back to Ready



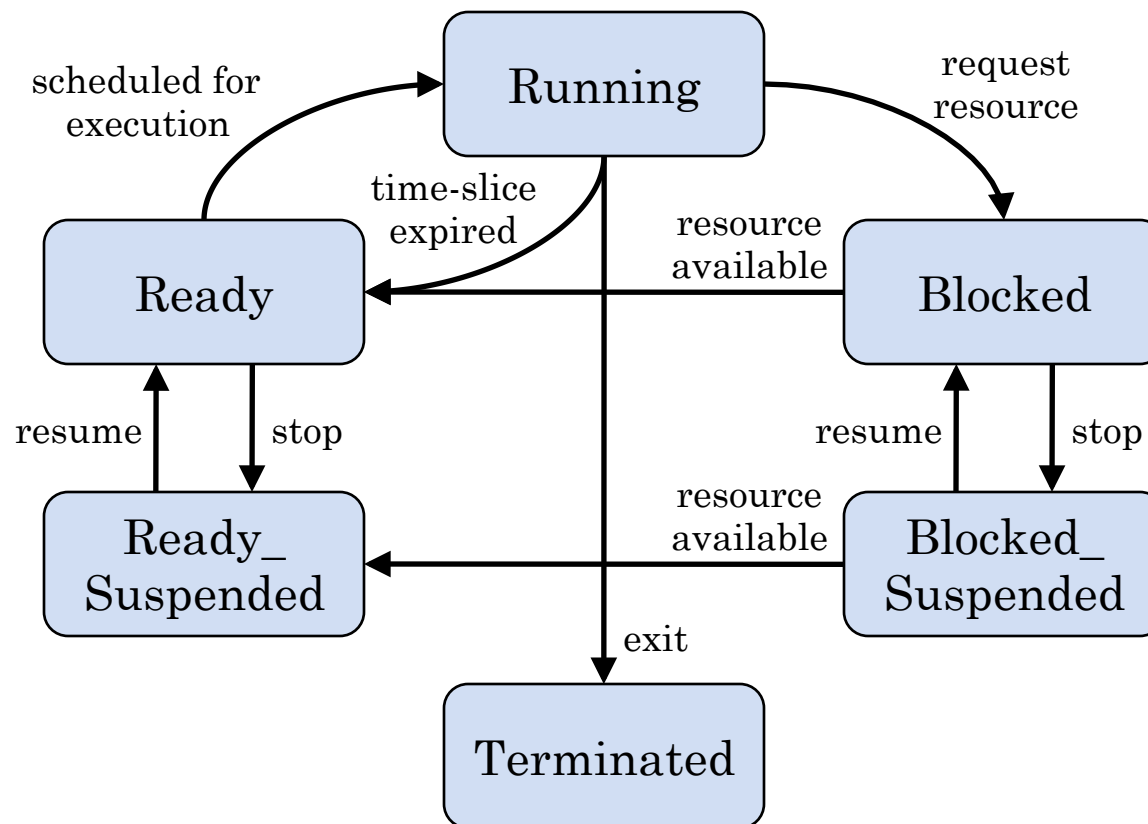
# UNIX PROCESS STATE DIAGRAM (3)

- Stopping a process causes it to become suspended
  - If resource becomes available while suspended, a blocked process changes to ready/suspended



# UNIX PROCESS STATE DIAGRAM (4)

- Running process may also terminate
  - Most resources can be reclaimed by kernel...
  - Needs to preserve status until the zombie is reaped



# STATES AND PROCESSES

- At this level, only one process may be in the Running state for each CPU in the computer
  - e.g. 4 cores = 4 processes in Running state
- Many processes can be in the other states!
- Kernel needs to manage collections of processes in each state
  - Different strategies for managing these processes, so different kinds of collections are employed
- Next Time:
  - Data structures used by kernel to keep track of process context and state information