



CS24: INTRODUCTION TO COMPUTING SYSTEMS

Spring 2015

Lecture 22

LAST TIME: PROCESS SCHEDULING

- Began discussing how to schedule processes
 - When a running process is blocked, suspended or terminated, scheduler chooses another process to run
 - Must be fair: all processes must receive CPU time
 - Must perform its scheduling quickly
- Context-switches are slow
 - Frequently take approx. 5-10 μ s
 - *Multiple tens of thousands of instructions!*

CONTEXT-SWITCH OPERATIONS

- *Lots* of work to do in a context switch:
 - Process is interrupted via exceptional flow control
 - The kernel is invoked...
 - Switch from process code to kernel code
 - Change protection-level (and stack/data areas being used)
 - Kernel must save the process' context into its PCB
 - Kernel picks another process to start running
 - (This is where scheduler does its work. Hopefully quickly.)
 - Kernel must restore the process' context from its PCB
 - Return back from kernel to the next process
 - Again, change protection level (and stack/data areas)
 - Resume executing the next process

LAST TIME: PROCESS BEHAVIORS

- Processes can exhibit a wide variety of behaviors
 - Interactive processes
 - Compute-intensive processes
 - Real-time processes
 - Makes it more challenging to schedule effectively
 - (Plus, programs can change behavior over time...)
- Covered two simple scheduling algorithms
 - Round-robin scheduling – good for compute-intensive processes, bad for interactive ones
 - Shortest jobs first scheduling – good for interactive processes, bad/unfair for compute-intensive ones

SCHEDULING CHALLENGES!

- Generally, if tasks are uniform, it's much easier to solve the scheduling problem
 - If all tasks are interactive, just use Shortest Job First approach
 - If all tasks are compute-intensive, use Round Robin approach with a large time-slice
- The challenges arise when the scheduler must deal with different kinds of processes
 - Must properly balance needs of interactive processes with needs of long-running processes

EARLIEST DEADLINE FIRST SCHEDULING

- Real-time processes are usually deadline-driven
 - *Video player needs to draw 30 frames per second, no matter what!*
- Also tend to be cyclic in their execution
 - Typically don't need the CPU for very long
- Scheduler can estimate several values:
 - When is the deadline for each process? (T_{dl})
 - What is the average run-time for each process? (T_{run})
- Scheduler chooses the process with the smallest $T_{dl} - (T_{now} + T_{run})$
 - (the process most in danger of missing its deadline ☺)
- Variations on this theme
 - Minimize average lateness, minimize max lateness

A GENERAL-PURPOSE SCHEDULER?

○ Problem:

- Interactive processes should run at a high priority, but should have a relatively small time-slice
 - “Well-behaved” interactive processes can be expected to become IO-bound before using up entire time
- Compute-intensive tasks should run at a low priority, but should have a longer time-slice
 - Minimize context-switch overhead for long-running process
- Real-time tasks should run at a very high priority to satisfy timing requirements, but expected to be short

○ Solution:

- As always, build a general-purpose scheduler by combining these strategies

MULTILEVEL FEEDBACK QUEUES

- Most operating systems use a *multilevel feedback queue* strategy for process scheduling
 - Windows NT/XP/Vista, Solaris, BSD variants
 - (Not MacOS X, or Linux 2.5+)
- Idea:
 - Scheduler uses multiple queues to segregate processes of varying priorities and behaviors
 - Each queue has its own maximum time-slice size, and even its own scheduling algorithm if needed
 - Move processes between the queues based on their observed behaviors
 - As the process executes, the scheduler moves it into the best queue for how the process is behaving
 - If process behavior changes, scheduler adapts to this easily

MULTILEVEL FEEDBACK QUEUES (2)

- Example queues:

- Q1: time-slice of 5ms
- Q2: time-slice of 15ms
- Q3: time-slice of 30ms
- Within each queue, it's "first come, first served"

- General rules:

- A new process is put into the highest priority queue
- If a process uses its entire time-slice, it is preempted by the kernel and demoted to the next lower queue
- If a process yields to the kernel before its time-slice is up, it goes to the end of its current queue
- A process can also be promoted if it regularly blocks or yields within the next higher queue's time-slice

MULTILEVEL FEEDBACK QUEUES (3)

- These rules very quickly categorize processes based on their behavior!
 - Different queues contain processes with different behaviors
- Additionally:
 - Want scheduler to give preference to shorter jobs, and to IO-bound processes
 - Want to give longer-running jobs larger time-slices to reduce context-switch overhead
- Scheduler executes Ready processes in Q1 first
 - Then, if Q2 contains Ready processes, these will be executed
 - Finally, if Q3 contains Ready processes, these are executed in round-robin order

MULTILEVEL FEEDBACK QUEUES (4)

- Example queues:
 - Q1: time-slice of 5ms
 - Q2: time-slice of 15ms
 - Q3: time-slice of 30ms
- Q1 ends up with processes that work in <5ms bursts
 - High priority since they will be done fast
 - Are very likely to be interactive processes
- Q2 ends up with processes that work in 5-15ms bursts
 - Lower priority, but longer time-slice
- Q3 contains processes that run in >15ms
 - If a process uses 30ms, it is preempted and sent to back of Q3 (i.e. round robin implementation)
 - If a process starts being regularly IO-bound, it will be promoted back up the queues, based on its run-times

MULTIPLE SCHEDULING ALGORITHMS

- Most operating systems that use multilevel feedback queues also support real-time processes
 - Real-time queue levels are higher priority than standard queue levels
 - The scheduling algorithms are different
 - Real-time processes aren't demoted to below the real-time range of queue levels
- Examples:
 - Windows NT-based systems: levels 0-15 are “normal” processes, 16-31 are soft real-time processes
 - Linux pre-2.5: levels 0-99 are real-time processes, 100-140 are “nice” task levels
 - (Linux versions 2.5+ use various other schedulers)

SUMMARY: SCHEDULING

- UNIX process API is relatively straightforward...
- Implementation is significantly more complex
 - Many details to keep track of for each process
 - Resources that a process is using, pending requests that a process is blocked on, other state information
- Processes can have *very* different performance characteristics!
 - Simple scheduling techniques can handle specific kinds of process behavior
 - Creating a generic process scheduler is more involved
 - Adapt scheduling choices based on past process behaviors
 - Also need to ensure that scheduling is fair! ☺

PROCESS ABSTRACTION: THE SCORECARD



- So far, have covered many aspects of how to implement the process abstraction
- Q: How to enforce differences between kernel and application processes?
 - A: Processor operating modes, hardware protection levels
- Q: What state information to manage for processes? How do we organize and manage it?
 - A: Registers, flags
 - (*What about program memory?*)
 - Use a Process Control Block structure to manage this state
- Q: How to interrupt a running processes?
 - A: Interrupts and exceptions to interrupt logical program flow and let kernel perform various tasks
- Q: How to choose which process should run next?
 - A: Use generalized, fair scheduling algorithms that can manage processes with varying behaviors

PROCESS ABSTRACTION: MISSING PIECES

- Glossed over some pretty important questions!
 - All relating to *how we manage process memory*
- How do we isolate the address spaces of different processes from each other?
- How to provide faster context-switches?
- Some other important questions too:
 - How do we provide access to the kernel and shared libraries in an efficient and uniform way?
 - How do we let processes share memory and coordinate with each other?
 - Very important if our OS is going to support powerful applications and services to be implemented on it

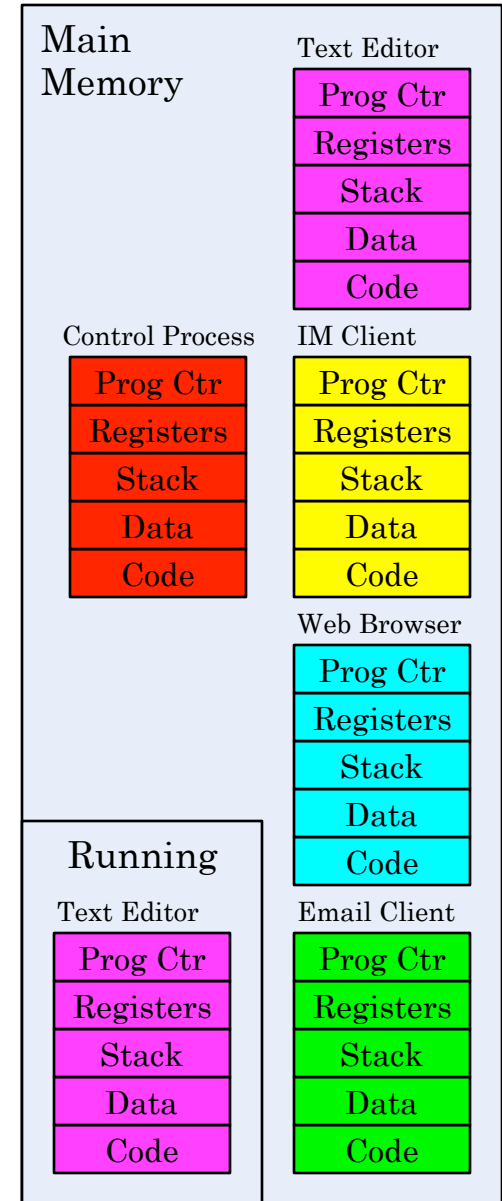
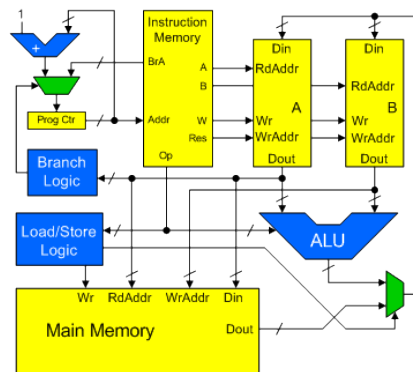
PROCESS ABSTRACTION

- Previous approach:
 - Running processes use a region of memory at bottom of address space
 - Suspended processes occupy other areas
- Context-switch:
 - Kernel interrupts the running program
 - Copies context of running process back to the process' memory area
 - Next, kernel copies context of another process into the running area
 - Finally, starts running the new process for another time-slice




PROCESS ABSTRACTION (2)

- Several big drawbacks with this approach!
- Copying process data will be very slow
 - Accessing DRAM main memory from the CPU takes 50-100ns!
 - Clearly unacceptable for systems running many processes, or ones with large memory footprints
- All processes must fit in main memory
 - System can't handle a large number of processes at same time
 - Severely limits ability of processes to work with large amounts of data



PHYSICAL ADDRESSING

- Main memory is an array of M contiguous bytes
 - Each location has its own physical address
 - Size of main memory is fixed
- When a program accesses memory:
 - Instruction refers to a physical address (direct or indirect)
 - Processor sends this address directly to main memory to retrieve the associated value



VIRTUAL ADDRESSING

- Instead of directly addressing physical memory, introduce a level of indirection
 - Instructions use *virtual* addresses instead of physical addresses
 - Translate virtual addresses into physical addresses as instructions are executed
- Given an address α :
 - Instead of directly accessing $M[\alpha]$, introduce a mapping table T , which maps virtual addresses to physical addresses
 - Use T to translate addresses: $M[T[\alpha]]$
 - Each virtual address α is mapped to a physical address
 - The physical address is used to access physical memory

VIRTUAL ADDRESSING (2)

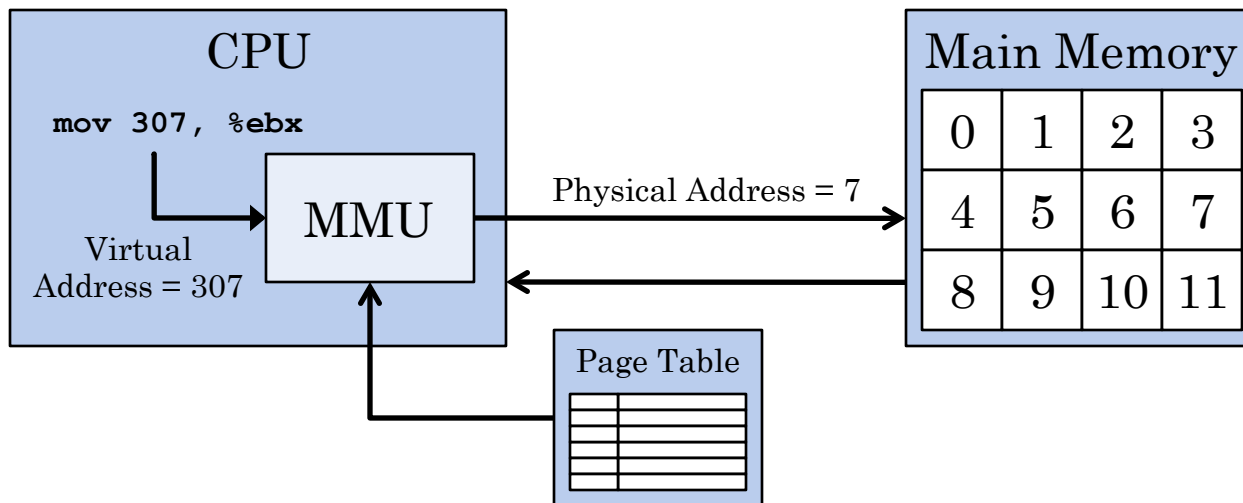
- Clearly prohibitive to map every single virtual address to a distinct physical address...
 - Mapping table would require much more space than the actual memory for the computer
 - Also, programs typically access memory in blocks, exploiting locality
- Divide memory up into pages of size P , $P = 2^p$
 - Choose a page size much larger than 1 byte or 1 word
 - (page-size considerations are discussed shortly...)
- Map each virtual page to a physical page
 - Mapping is specified using a page table
 - Each page table entry maps one virtual page to one physical page

VIRTUAL MEMORY

- Main memory provides a *physical address space*
 - Size of M bytes (frequently, $M = 2^m$, but not required)
 - Computer provides an m -bit address space
- Define a *virtual address space* of size N bytes
 - $N = 2^n$, so this is an n -bit virtual address space
 - Not required that $M = N$, but for now we will assume this is the case
- Virtual memory system must map virtual pages (pages in virtual address space) to physical page frames (pages in physical address space)

VIRTUAL ADDRESSING


- Performing this address translation in software would be horribly slow...
- CPU provides *hardware* support for virtual memory and address translation
 - CPU has a Memory Management Unit (MMU) that performs this address translation on the fly
 - The MMU uses a page table to perform translation



VIRTUAL AND PHYSICAL PAGES


- Virtual addresses specify a virtual page number (VPN) and a virtual page offset (VPO)

- Offset is lowest p bits (pages contain 2^p bytes)
- Virtual page number is upper $n - p$ bits in virtual address



- Physical addresses specify a physical page number (PPN) and a physical page offset (PPO)

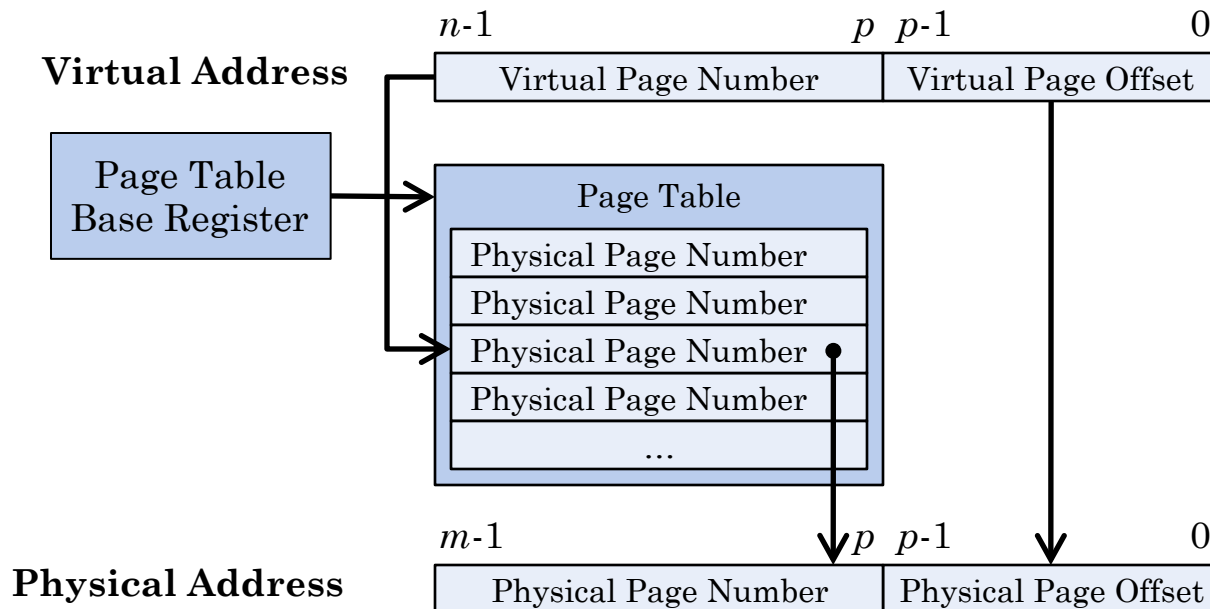
- Offset is lowest p bits, as before
- Similarly, physical page number is upper $m - p$ bits



- Page table maps virtual pages to physical pages

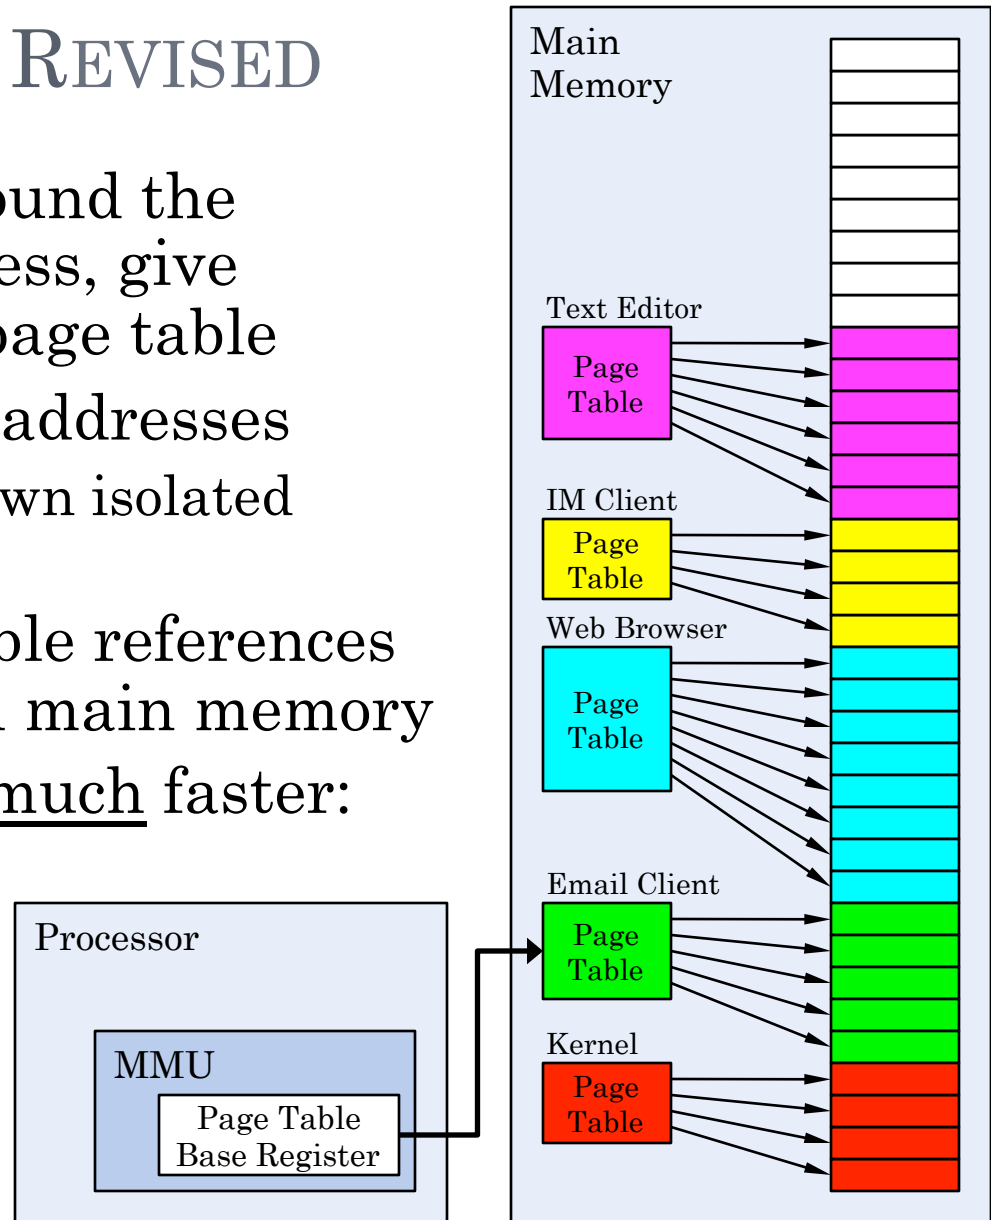
ADDRESS TRANSLATION

- Page table is indexed with virtual page number
 - Page table entry contains the physical page number
 - Combine physical page number with virtual page offset to get physical address
- Start of page table specified in a control register
 - MMU uses this address to look up page table entries



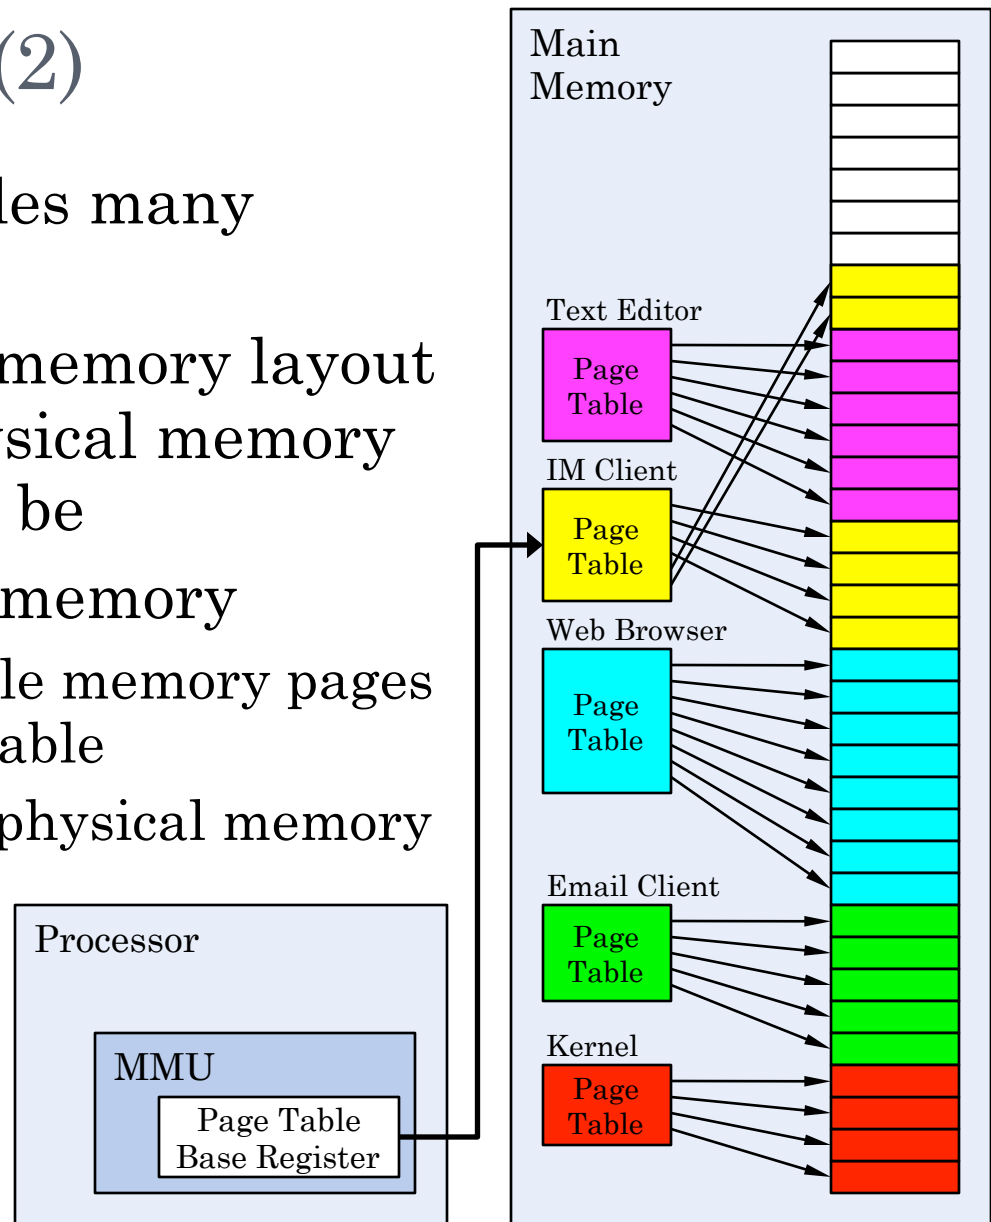
PROCESS MEMORY, REVISED

- Instead of copying around the memory for each process, give each process its own page table
- Programs use virtual addresses
 - Each process has its own isolated address space
- Each process' page table references its own set of pages in main memory
- Context-switching is much faster:
 - Simply change Page Table Base Register to reference the new process' page table!



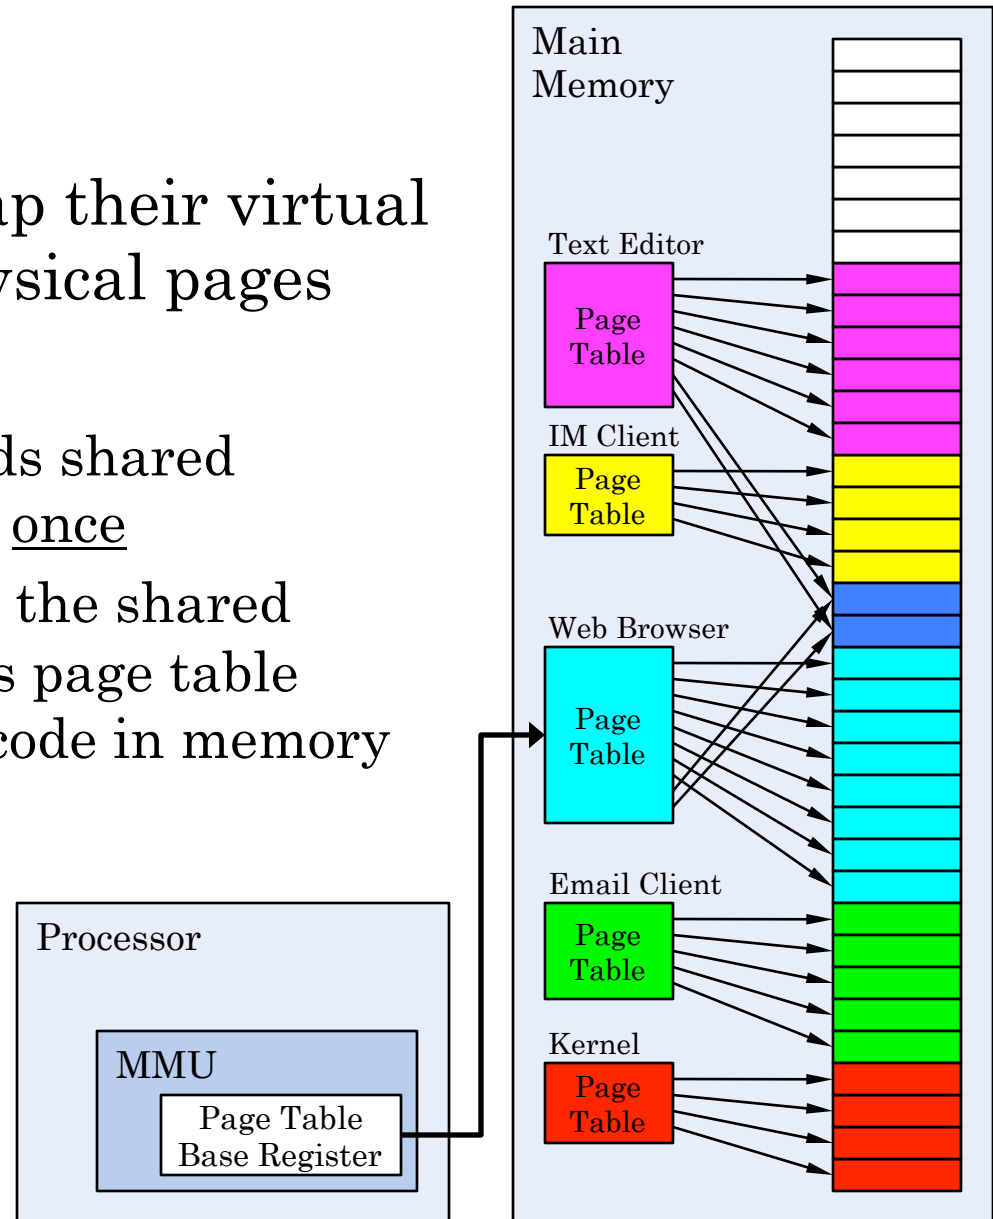
PROCESS MEMORY (2)

- Virtual memory enables many other useful features
- Each process' virtual memory layout is contiguous, but physical memory layout doesn't have to be
- IM client needs more memory
 - Simply assign available memory pages to that process' page table
 - Process doesn't know physical memory isn't contiguous
 - CPU takes care of virtual-to-physical address translation



SHARED PAGES

- Two processes can map their virtual pages to the same physical pages
- Shared libraries:
 - Operating system loads shared libraries into memory once
 - When a process needs the shared library, just update its page table to reference library's code in memory
- Shared memory:
 - Multiple processes collaborate by working in the same memory area



VIRTUAL MEMORY SYSTEM SO FAR...

- Already achieved a lot with this simple idea!
- Much easier to manage memory in processes
 - Context-switches are *much* faster
 - Processes have their own isolated virtual address spaces
 - CPU handles mapping of virtual addresses to physical addresses automatically
 - A process' physical memory layout doesn't have to be contiguous
 - Can map multiple virtual pages to the same physical page
- Problem:
 - Still have to divide up the limited main memory amongst all processes, whether running or stopped
- We know that not all processes are always active...
 - e.g. a process can be stopped and resumed
 - Also, a process might not always be using all of its memory

NEXT TIME

- Incorporate virtual memory into the memory hierarchy:
 - Virtual memory becomes a *cache* for program data stored on disk