



# CS24: INTRODUCTION TO COMPUTING SYSTEMS

Spring 2015

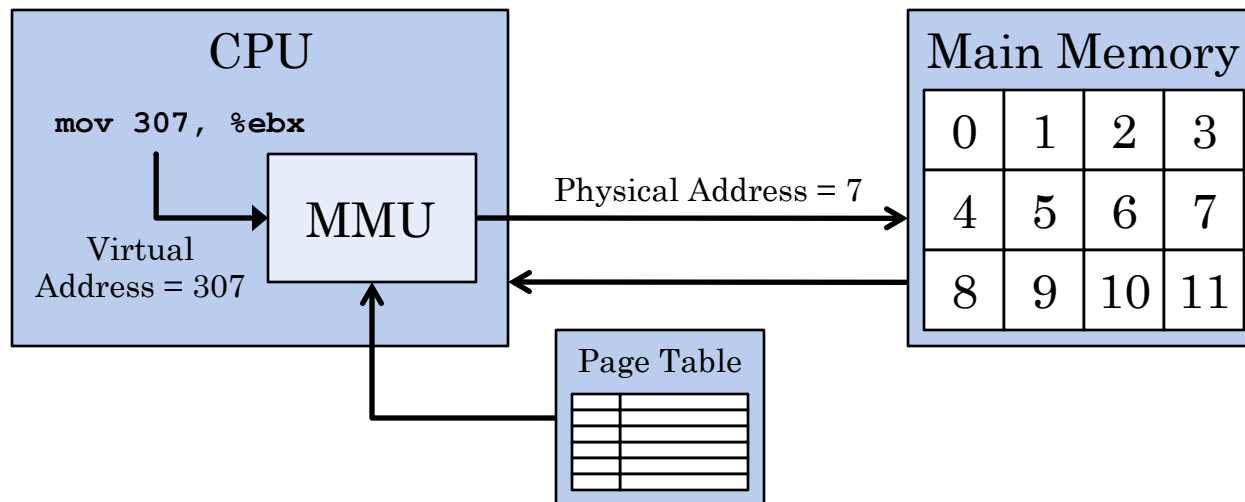
Lecture 23

## LAST TIME: VIRTUAL MEMORY

- Began to focus on how to virtualize *memory*
- Instead of directly addressing physical memory, introduce a level of indirection
  - Programs use *virtual* addresses, not physical addresses
- Divide memory up into pages of size  $P$ ,  $P = 2^p$ 
  - Choose a page size much larger than 1 byte or 1 word
  - (makes mapping table smaller; other reasons too)
- Map each virtual page to a physical page frame
  - Mapping is specified using a page table
  - Each page table entry maps one virtual page to one physical page

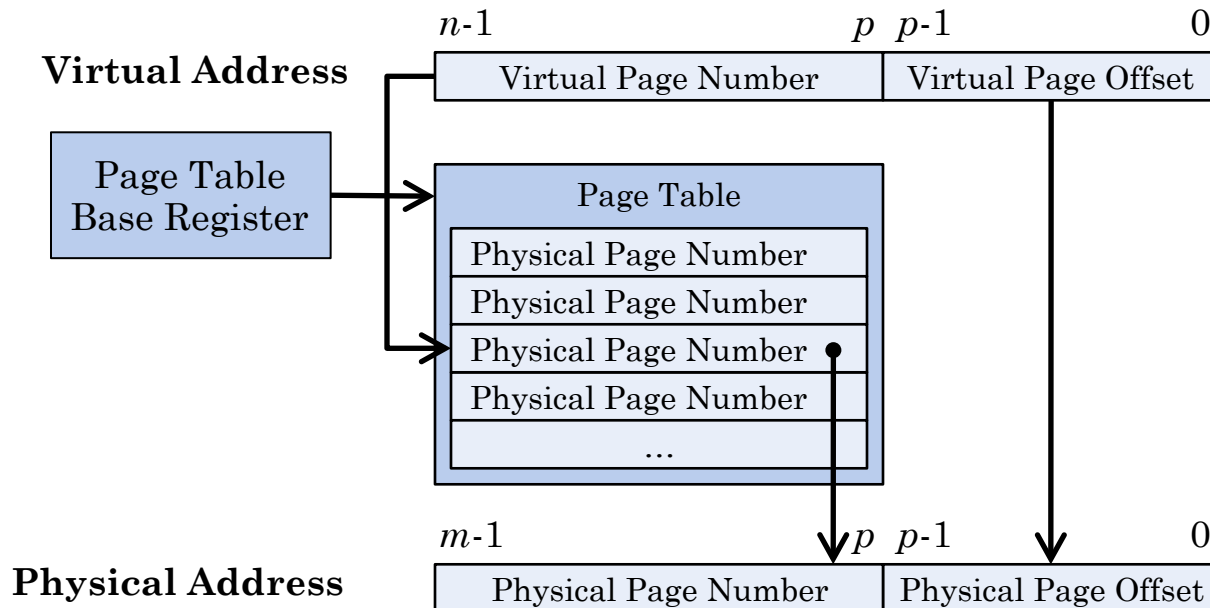
## LAST TIME: VIRTUAL ADDRESSING (2)

- Performing this address translation in software would be horribly slow...
- CPU provides *hardware* support for virtual memory and address translation
  - CPU has a Memory Management Unit (MMU) that performs this address translation on the fly
  - The MMU uses a page table to perform translation



# ADDRESS TRANSLATION

- Page table is indexed with virtual page number
  - Page table entry contains the physical page number
  - Combine physical page number with virtual page offset to get physical address
- Start of page table specified in a control register
  - MMU uses this address to look up page table entries



# VIRTUAL MEMORY: ISSUES

- This approach already simplifies many things
  - e.g. process isolation, context switches, shared memory
- Our virtual memory is still limited:
  - Must share limited main memory amongst all processes, whether running or stopped
- We know not all processes are always active...
  - e.g. a process can be stopped and resumed
  - e.g. a process may only be using a portion of its instructions and data in memory
- Redefine our notion of “virtual memory” to be an array of  $N$  contiguous bytes stored on disk
- Main memory becomes a cache for the local disk
  - Main memory is Level 4 (L4) cache, disk is L5 memory

# VIRTUAL MEMORY AS A CACHE

- Apply same concepts as hardware caches:  
update page table entries to track more details
  - Is a page of memory actually allocated for a specific virtual page number?
  - Is the virtual page cached in DRAM main memory?
  - (similar to extra details stored in cache lines)
- Now, the memory available to all processes can be *much larger* than actual size of main memory
  - Use our caching techniques to get a virtual memory much larger than main memory, but at [close to] DRAM speeds

## VIRTUAL MEMORY AS A CACHE (2)

- DRAM is cache for virtual memory stored on disk
  - DRAM access latency is 50-100ns
  - Disk access latency is 8-30ms (!!!)
  - A virtual memory cache-miss is very expensive
- However, disk throughput is 50+ MB/second
  - Reading the first byte of data from disk is very slow...
  - Reading second and subsequent bytes is *much* faster
- Example: transfer unit of 1 byte
  - Cache-miss penalty is 8-30ms
- How about a transfer unit of *100KB?!?*
  - Cache-miss penalty is 8-30ms + 2ms
  - Not that much more expensive to retrieve!

# VIRTUAL MEMORY: STRATEGIES

- Choose a large page size (e.g. 4KB or 8KB), to compensate for large disk access latency
  - Retrieving additional bytes doesn't add much time to the cost of a cache miss
  - Using a much larger transfer-unit size will reduce the number of cache misses we have in the first place
- Use sophisticated replacement policies in DRAM cache to avoid cache misses as much as possible!
  - Example: DRAM cache is fully associative
- Use write-back policy instead of write-through, when a program writes to a cached DRAM page
  - Don't write a modified page to disk until it is actually evicted from main memory
  - *(One reason why OS needs to be shut down cleanly...)*



# PAGE TABLE ENTRIES

- Page table entries must be enhanced to support this virtual memory model
- Include a *valid* bit, along with each address
- *valid* = 1 means page is cached in DRAM memory
  - Corresponding address is the start of the page in DRAM physical memory
- *valid* = 0 means the page is not cached in DRAM
  - If stored address is 0, the virtual page is not allocated
  - If address is not 0, the virtual page is allocated, and the address is the page's location on disk
- IA32 page-table entries call this the *present* bit, i.e. “Is the page present in main memory?”

# PROCESSOR AND OPERATING SYSTEM

- *A very important detail about virtual memory facility:*
- The processor provides *some* hardware support for virtual memory
  - Address translation is performed in hardware, using the Memory Management Unit
  - MMU must use a page table to perform address translation
- The processor cannot completely support virtual memory in hardware, on its own!
  - What to do when a virtual page is not in main memory?  
CPU has no idea how virtual memory pages are cached.
- The operating system must also provide support for many virtual memory operations
  - CPU does as much as it can on its own...
  - CPU invokes an exception when it needs the kernel's help!

# VIRTUAL MEMORY EXAMPLE

- Example: small virtual memory

- 4 physical page frames (PP)
- 8 virtual pages (VP)

- Some virtual pages are cached in DRAM:

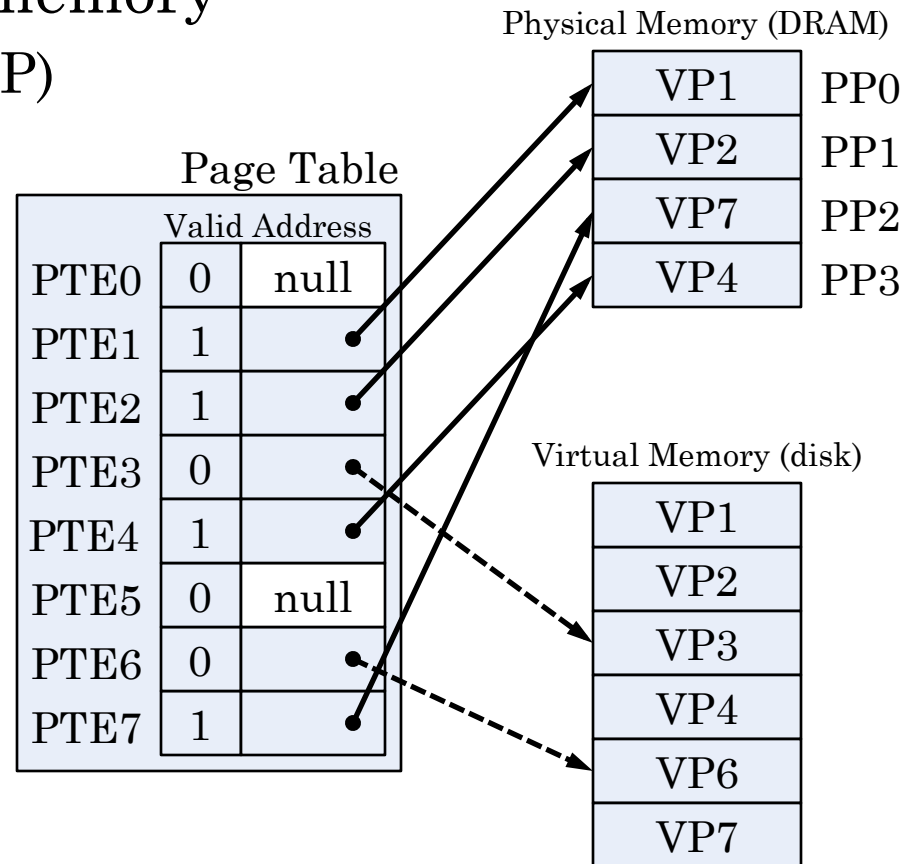
- VP1, VP2, VP4, VP7

- Some virtual pages are only stored on disk

- VP3, VP6

- Two virtual pages have not been allocated

- VP0, VP5




# VIRTUAL MEMORY ACCESSES

- Program accesses a word in Virtual Page 2

- MMU looks in page table for VP2 (PTE2)
- Uses virtual page number as index into Page Table

- Virtual page 2 is stored in physical page 1...

- Valid flag is 1: page is cached in DRAM
- MMU sends physical address to DRAM main memory, and retrieves the data value

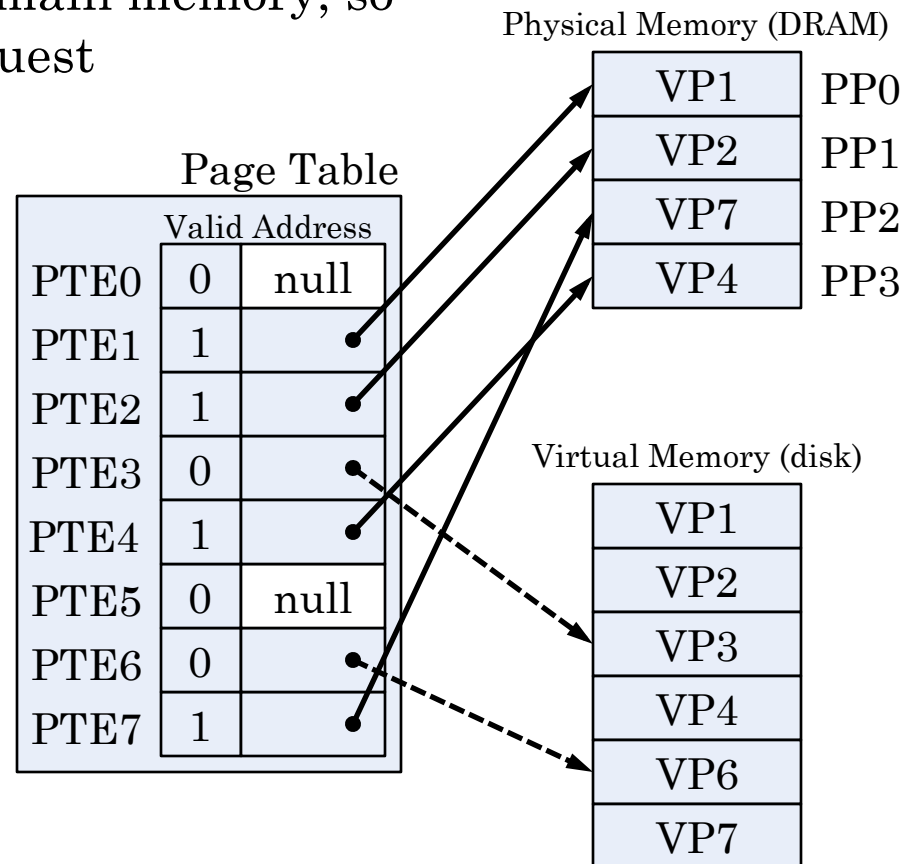


## VIRTUAL MEMORY ACCESSSES (2)

- Now, program accesses a word in Virtual Page 6
  - MMU looks in page table for VP6, but the valid flag is 0!
  - Virtual page isn't cached in main memory, so MMU cannot satisfy the request

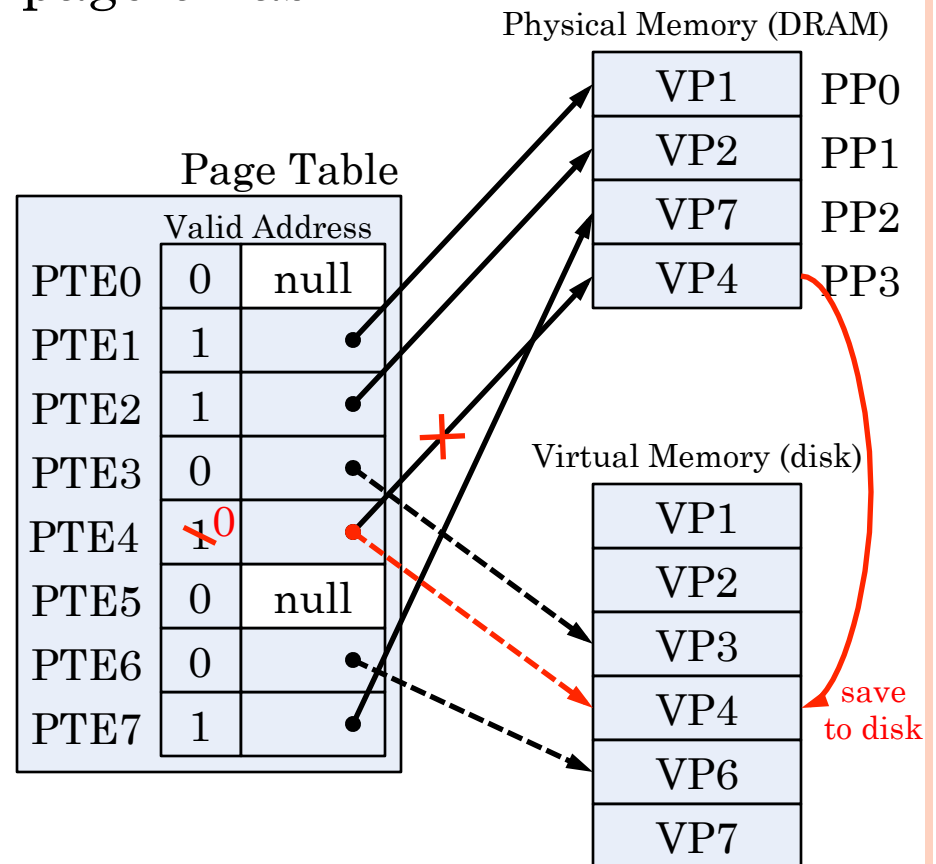
- CPU causes a page fault exception

- It's up to the kernel to resolve the problem!
- Kernel provides a handler for resolving page faults
- The kernel manages the page table, memory pages
  - Processor simply provides support for virtual memory



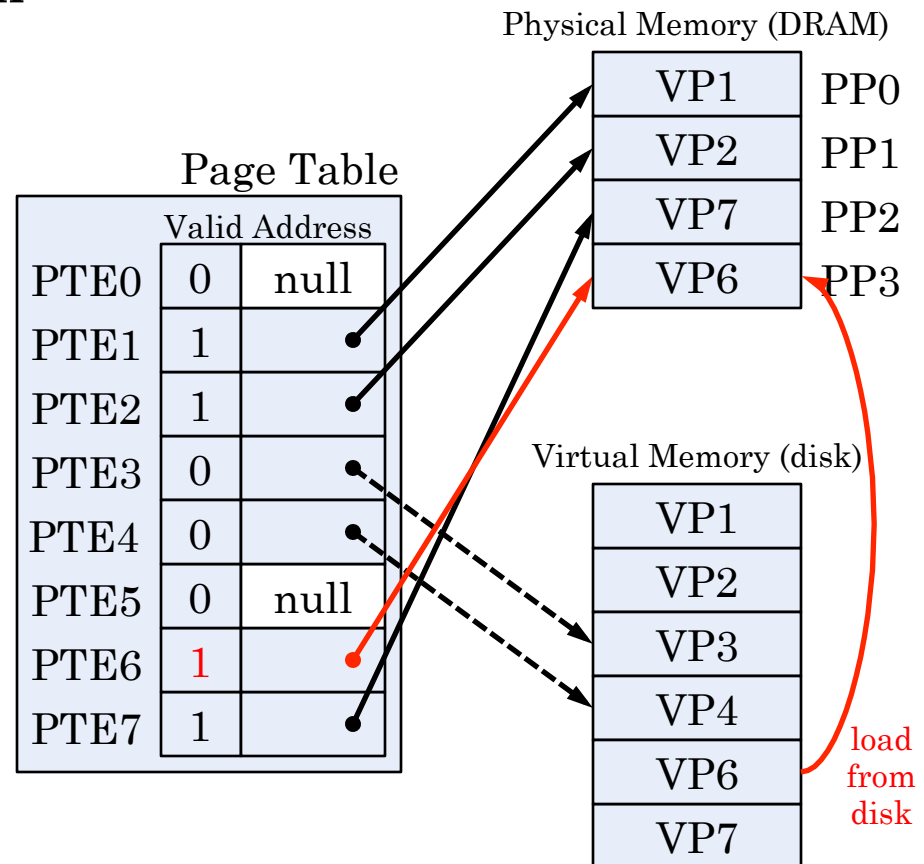
# VIRTUAL MEMORY ACCESSSES (3)

- CPU generates a page fault, and kernel's handler is invoked
- Page-table entry for virtual page 6 has a non-null address...
  - The kernel can load VP6 from disk into memory
- ...but first, some other page must be evicted!
- Kernel selects VP4 as the victim page
  - If VP4 is changed, must also write it back to disk
- Kernel updates PTE4 to show page isn't cached



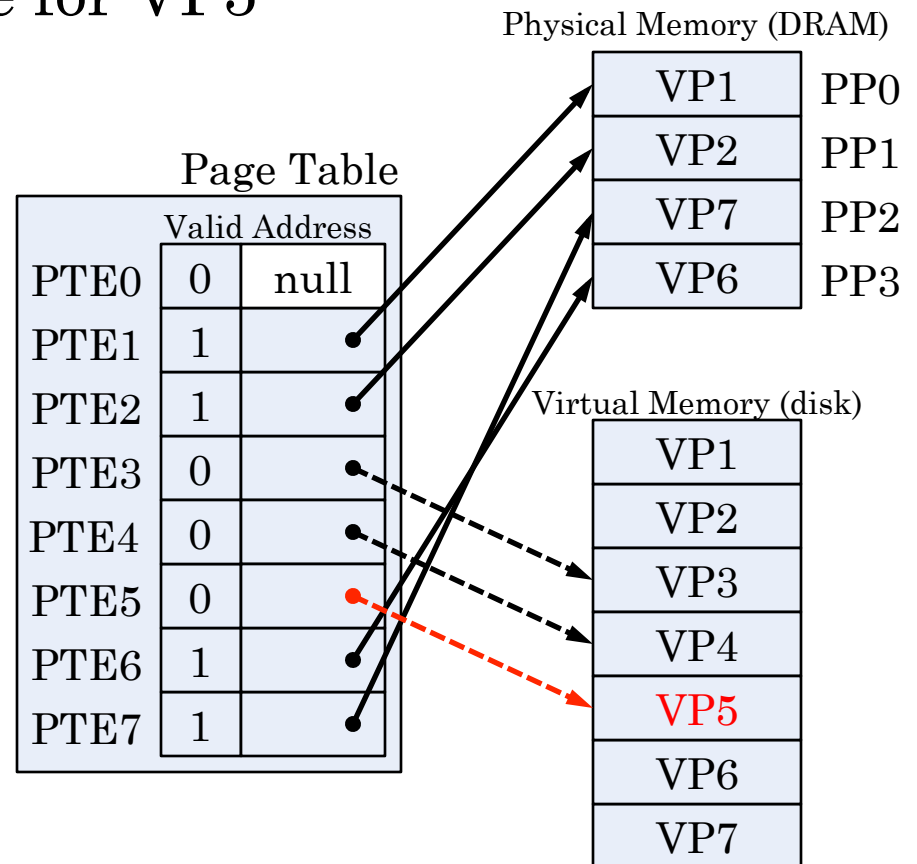
## VIRTUAL MEMORY ACCESSES (4)

- Now kernel loads virtual page 6 into physical page 3
  - Update PTE6 to point to physical page 3 in DRAM memory
- Finally, kernel returns from the page-fault handler
- Since this is a fault:
  - CPU resumes executing the instruction that caused the fault
- Program repeats the access to virtual page 6
  - This time, MMU finds that PTE6 is valid
  - MMU performs address translation, and retrieves value from physical page 3



# ALLOCATING VIRTUAL MEMORY

- If a process needs more memory, kernel can allocate another virtual memory page
- Example: allocate a page for VP5
  - Make room on the disk for virtual page 5
  - Update PTE5 to point to VP5 on disk
  - (**Note:** *no requirement that pages on disk be kept in order...*)
- When program begins using VP5, it can be swapped into physical memory





# VIRTUAL MEMORY AS A CACHE

- DRAM main memory is a cache to virtual memory stored on disk
- Caching terminology is slightly different
  - Virtual memory was invented in 1960s, before SRAM caches were needed between the CPU and DRAM!
  - Caches have *blocks*, but virtual memory has *pages* (virtual blocks) and *frames* (physical blocks)
  - Caches have *cache sets* and *cache lines*, but virtual memory has *page tables* and *page-table entries*
  - Caches have *cache-hits* and *cache-misses*, but virtual memory has *page-hits* and *page-faults*
- Nearly identical concepts, but terminology is slightly different

## VIRTUAL MEMORY AS A CACHE (2)

- Like caches, virtual memory depends on *locality*
  - When programs exhibit good locality, we can get a virtual memory much larger than physical memory, but with same performance as physical memory
- Processes typically reference a specific set of memory pages as they run
  - Called the *working set* of the process, or the *resident set*, since these pages are kept in memory
- If a program has good locality:
  - The working set of pages will usually fit entirely within memory, yielding good performance

## VIRTUAL MEMORY AS A CACHE (3)

- If a program's working set doesn't fit within memory, *thrashing* will occur
  - Pages are constantly swapped in and out of main memory
  - The program runs *very* slowly, since many accesses incur the 8-30ms disk latency penalty
- If a program runs very slowly, and you see the disk is constantly being accessed, it is probably thrashing ☹
- UNIX provides a **getrusage()** function
  - “get resource usage”
  - Reports several details about virtual memory system, such as the number of page faults
  - Also reports number of voluntary and involuntary context-switches, number of signals received, etc.
  - Can monitor how well-behaved your program is! ☺

# VIRTUAL MEMORY AND DISK FILES

- Once virtual memory is disk-based, we can provide some new features
- Loading programs and libraries into memory:
  - Programs and libraries are stored as binary files containing instructions and data
- When we want to start a new program:
  - Kernel allocates a contiguous set of virtual pages large enough to hold the program's code and data
  - Then, the kernel updates the process' page table, pointing these new pages to the binary file stored on disk
  - The new page-table entries are marked as *valid* = 0
    - i.e. the pages are not yet cached in DRAM memory
- As the program runs, virtual memory system loads program's code and data into memory automatically

## VIRTUAL MEMORY AND DISK FILES (2)

- This technique is called memory mapping
- Can map a disk file into memory with the UNIX **mmap()** function

```
void * mmap(void *start, size_t length, int prot,
            int flags, int fd, off_t offset)
```

  - Maps a specific portion of the file **fd** into memory
  - Portion starts at **offset** and contains **length** bytes
- Function works by allocating virtual memory pages, and pointing page table entries to disk file
- Instead of manually performing file I/O, simply perform reads and writes against a memory area
  - Virtual memory system automatically loads and saves data to the file stored on disk

# PAGES AND PERMISSIONS

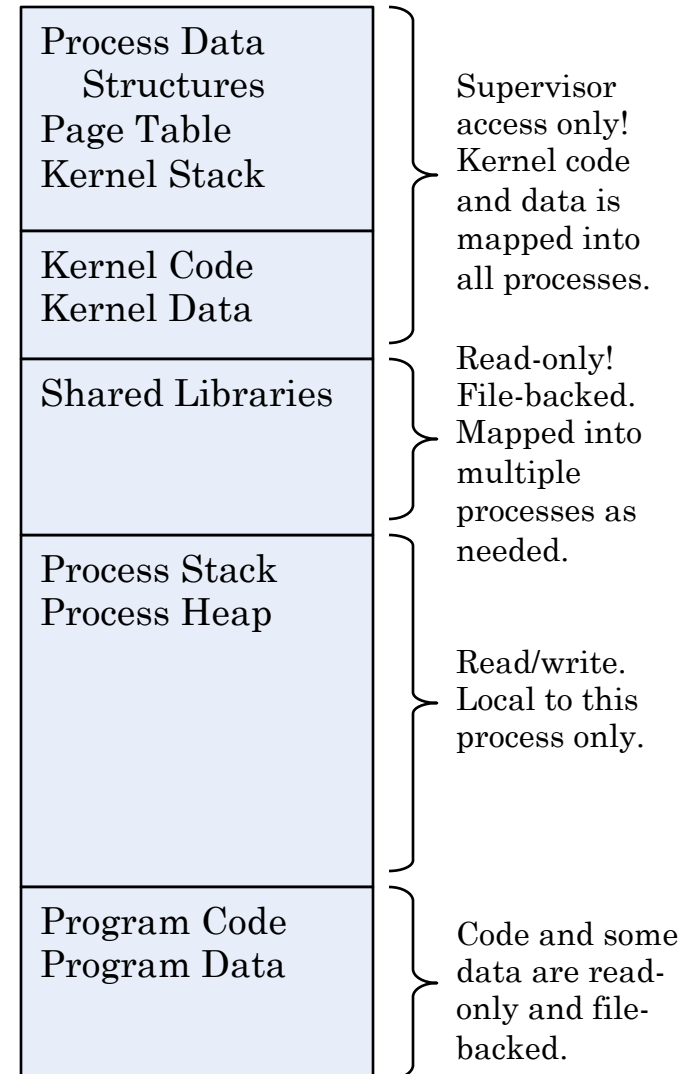
- We can use the virtual memory system to load programs and read-only data into memory...
  - Really don't want to let programs modify their code or read-only data. Or shared libraries, etc.
- Add permission bits to each page table entry, that control read/write access, etc.
- IA32 has two permission bits:
  - Read/Write flag controls whether code can write to the page, or only read from it
  - User/Supervisor flag controls whether code must have a higher privilege level to access the page

Page Table			
<i>valid</i>	R/W	SUP	Physical Page Number
<i>valid</i>	R/W	SUP	Physical Page Number
<i>valid</i>	R/W	SUP	Physical Page Number
<i>valid</i>	R/W	SUP	Physical Page Number
<i>valid</i>	R/W	SUP	...

# PAGES AND PERMISSIONS (2)

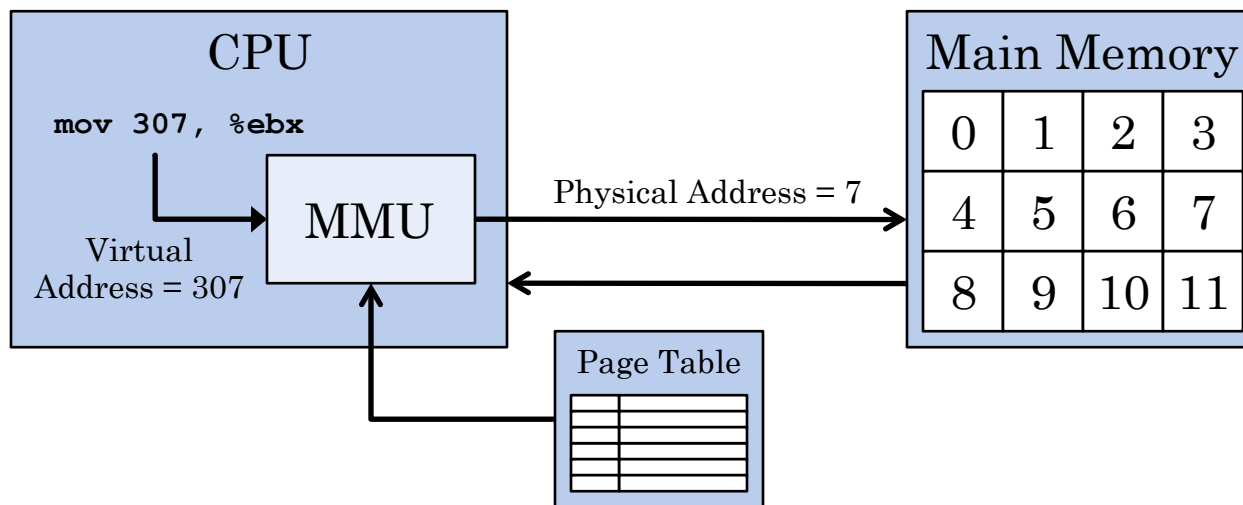
## Process Virtual Memory

- Kernel uses these virtual memory features to load and run programs
  - Load program and shared library code from disk automatically
  - Share kernel and library code across multiple processes
  - Isolate address space of processes
- If code tries to access a page in a way disallowed by permissions, CPU raises a *general protection fault*
- The kernel has a handler for general protection faults
  - Typical response: terminate the process! 😊



# IMPLEMENTATION ISSUES

- Glossed over a couple of implementation issues
- Issue 1: Where do page tables live?
  - Showed page tables as separate from main memory

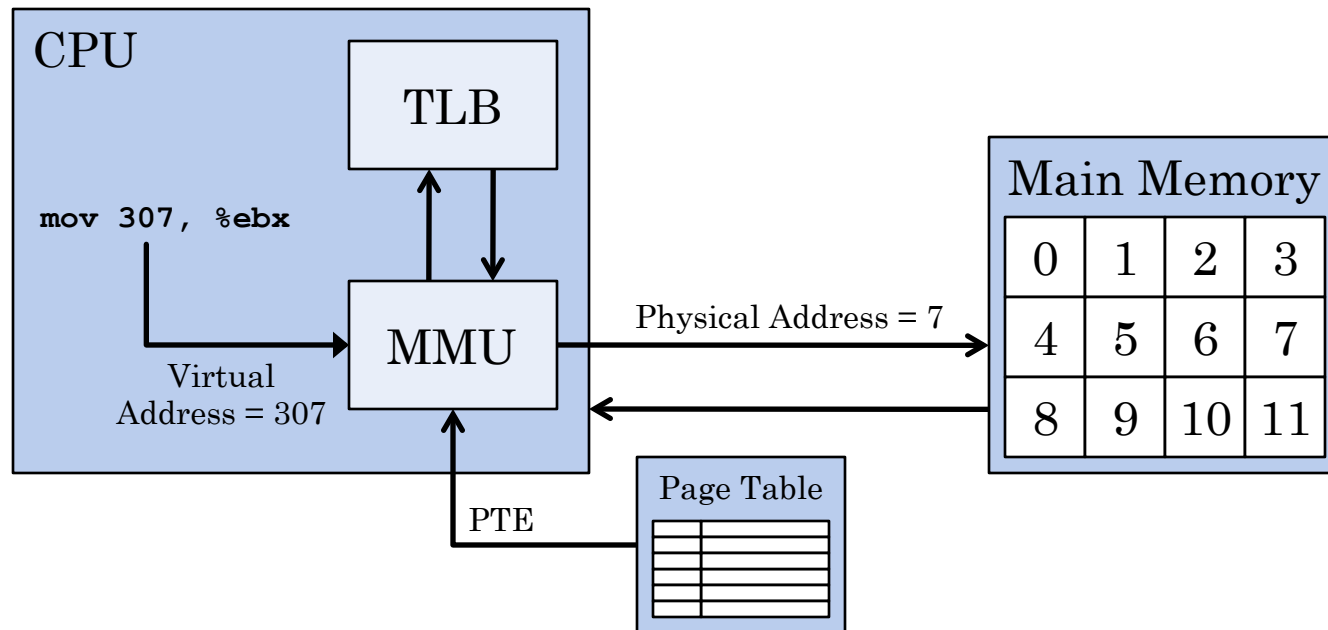


- In reality, page tables also live in main memory
- Why would this be a problem?
  - *Hint: How long does it take to access main memory?*



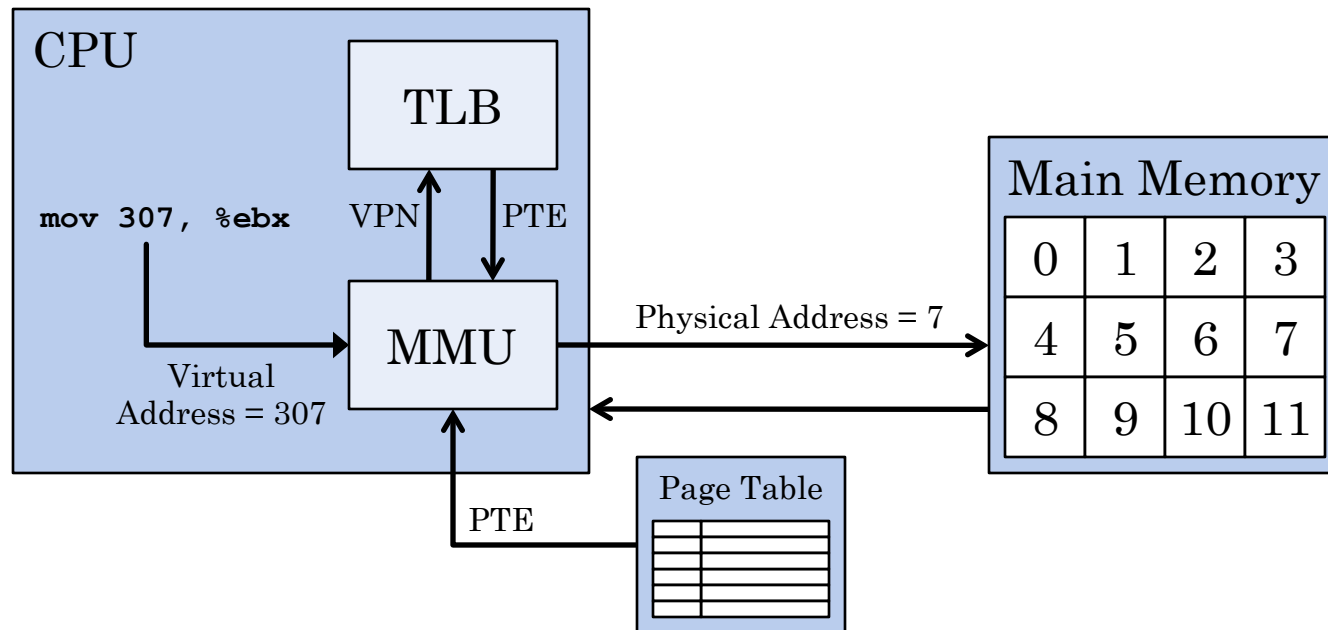
## IMPLEMENTATION ISSUES (2)

- If page tables live in main memory, address translation would frequently incur DRAM access overhead of 50-100ns!
  - If you're lucky, part of page table is in L1 or L2 cache
- CPUs frequently use a Translation Lookaside Buffer (TLB) to cache page table entries
  - Don't even want to incur L1 or L2 penalty, if possible!



# TRANSLATION LOOKASIDE BUFFER

- Translation lookaside buffer stores one page table entry per cache line
  - Input is virtual page number, output is page table entry
- TLB often has a high degree of associativity
  - Maximize chance that TLB contains needed page table entry!

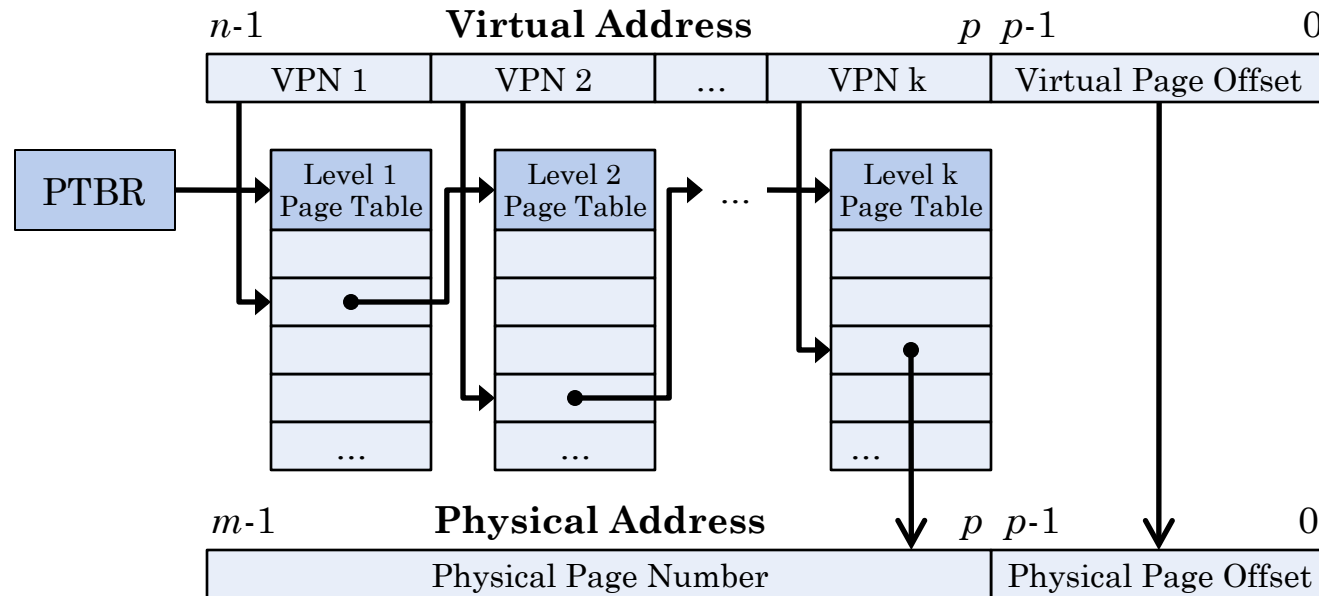


## IMPLEMENTATION ISSUES (3)

- Issue 2: How large are page tables?
  - Showed the processor using a single page table
  - Real-world system doesn't work that way!
- Given a 32-bit address space, and pages of 4KB:
  - Page offsets take 12 bits, page numbers are 20 bits
- To map each virtual page to a physical page, how many page table entries?
  - $2^{20} = 1,048,576$
- If each page table entry is 32 bits, how big is the entire page table?
  - Page table for this address space would be 4MB!
- Every process would require a 4MB page table
  - *Even bigger trouble for larger address spaces!*

# MULTI-LEVEL PAGE TABLES


- Instead of a single page table, introduce a *hierarchy* of page tables



- Idea: typically, most of a process' virtual address-space will be unallocated
  - Hierarchy of page tables will be relatively sparse

## MULTI-LEVEL PAGE TABLES (2)

- Instead of a single page table, introduce a *hierarchy* of page tables



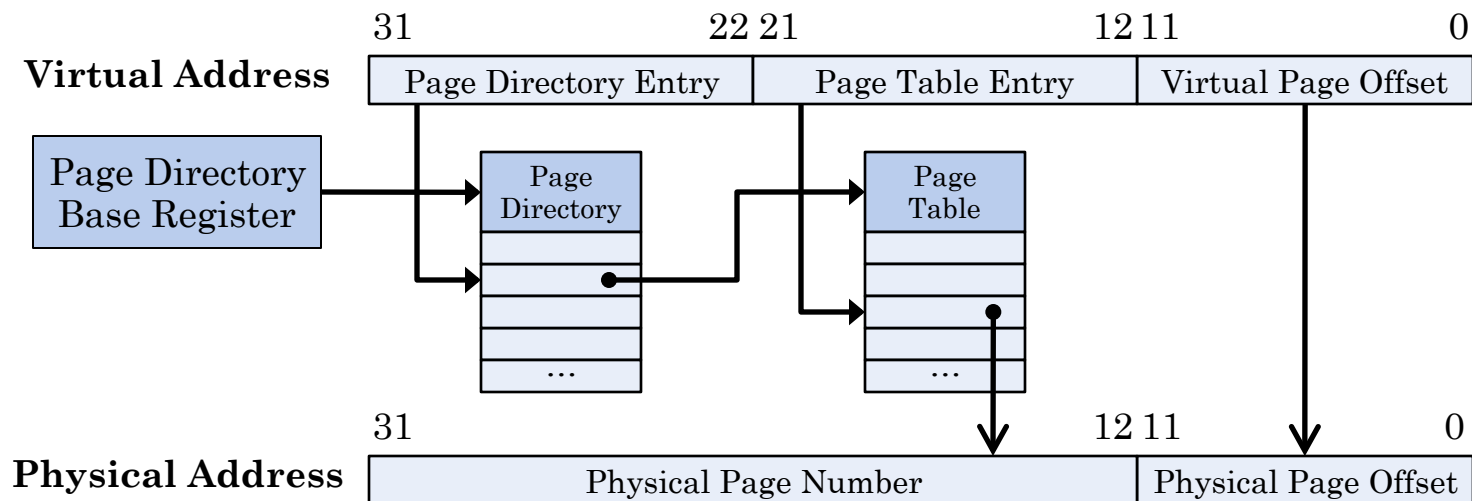
- If an entire page table at level  $i$  is empty:
  - Don't allocate a page for the empty page-table
  - Corresponding table entry at level  $i - 1$  is null

# IA32 VIRTUAL MEMORY SUPPORT

- Intel Pentium-family processors provide hardware support for virtual memory
- Virtual and physical address spaces are 32 bits
- Pages are 4KB in size ( $2^{12} = 4096$ )
  - Pages are identified by topmost 20 bits in address
- Pentium-family processors implement a two-level page table hierarchy
- Level 1 is the *page directory*
  - Entries in the page directory refer to page tables, as long as the page table is not empty
- Level 2 contains *page tables*
  - Entries map virtual pages to physical pages

# IA32 VIRTUAL MEMORY SUPPORT (2)

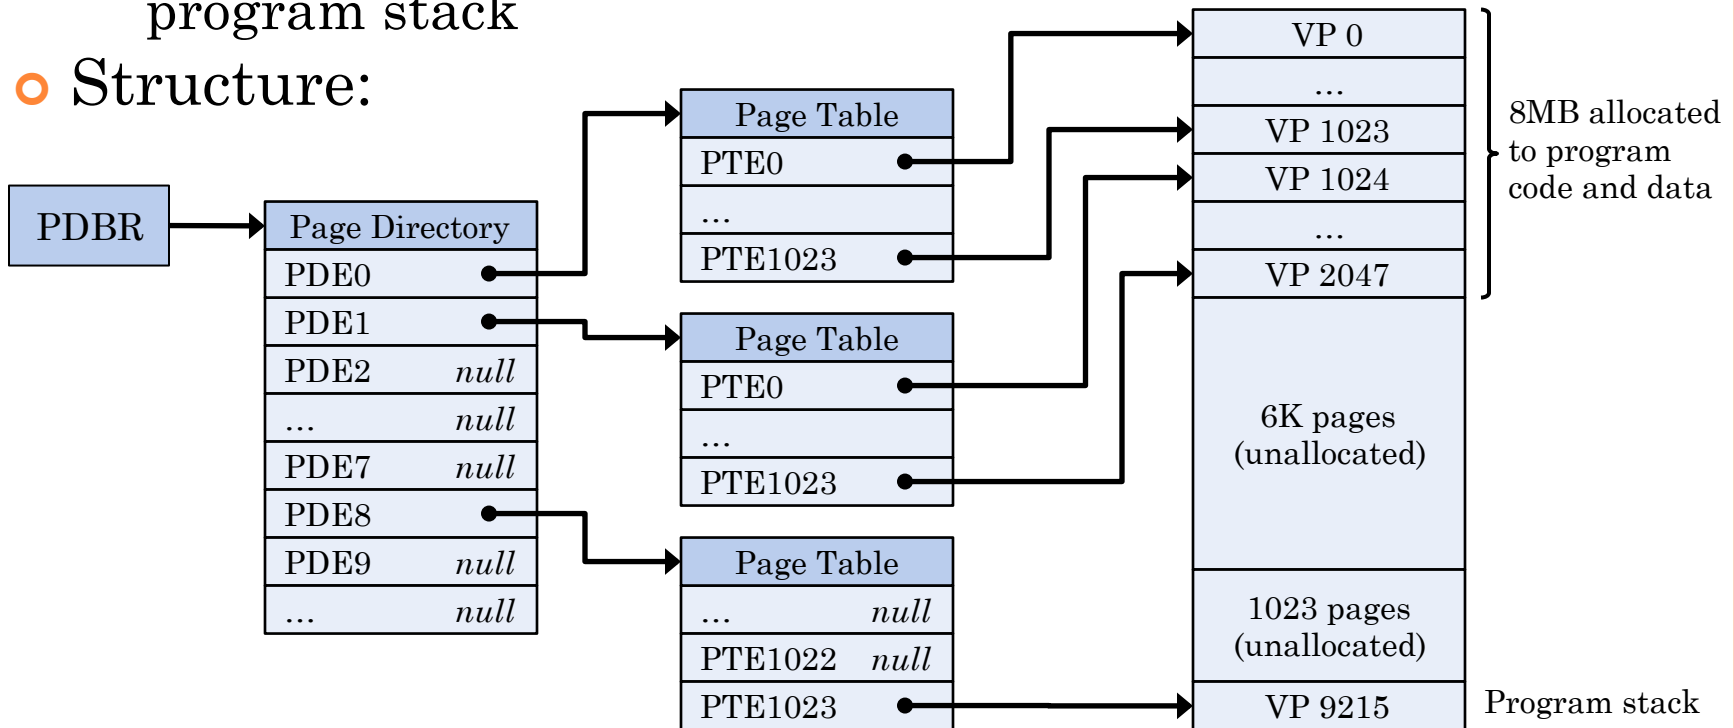
- IA32 page-table hierarchy:
  - Top 10 bits of virtual address is the page directory entry
  - Next 10 bits is used for page table entry
  - Bottom 12 bits used for the virtual page offset



# MULTI-LEVEL PAGE TABLE EXAMPLE

- Program with this memory layout:
  - Bottom 2K pages allocated to program code and data (4KB page size  $\times$  2K = 8MB program and data size)
  - 6K pages left unallocated, for program heap to use
  - 1023 unallocated pages, plus 1 page allocated to program stack

- Structure:





## ASIDE: CORE I7 AND PAGE TABLES

- Intel Core i7 implements a four-level hierarchy
  - 48 bit address space – again, page tables would be too large if not broken into a hierarchy
  - Preserves the 4KB page size of IA32
- Core i7 page-table hierarchy level names:
  - Level 1 = page global directory
  - Level 2 = page upper directory
  - Level 3 = page middle directory
  - Level 4 = page table
- Each level uses 9 bits of virtual address
  - Core i7 page table entries are 64 bits
  - Page tables still work out to be 4KB in size
- Virtual page offset is bottommost 12 bits
- See CS:APP2e §9.7.1 for more details on Core i7

## NEXT TIME

- Continue discussion of IA32 virtual memory
  - More hardware implementation details
  - How Linux uses IA32 virtual memory facilities