



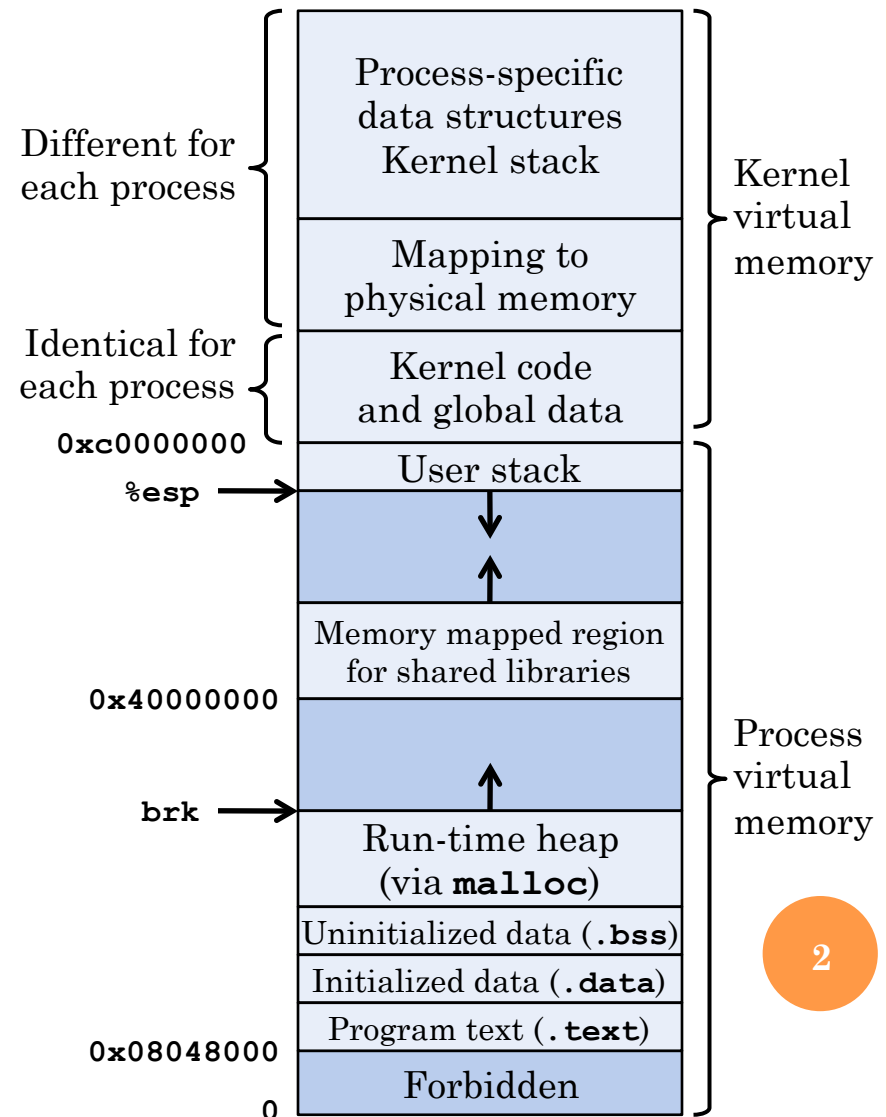
# CS24: INTRODUCTION TO COMPUTING SYSTEMS

Spring 2015

Lecture 25

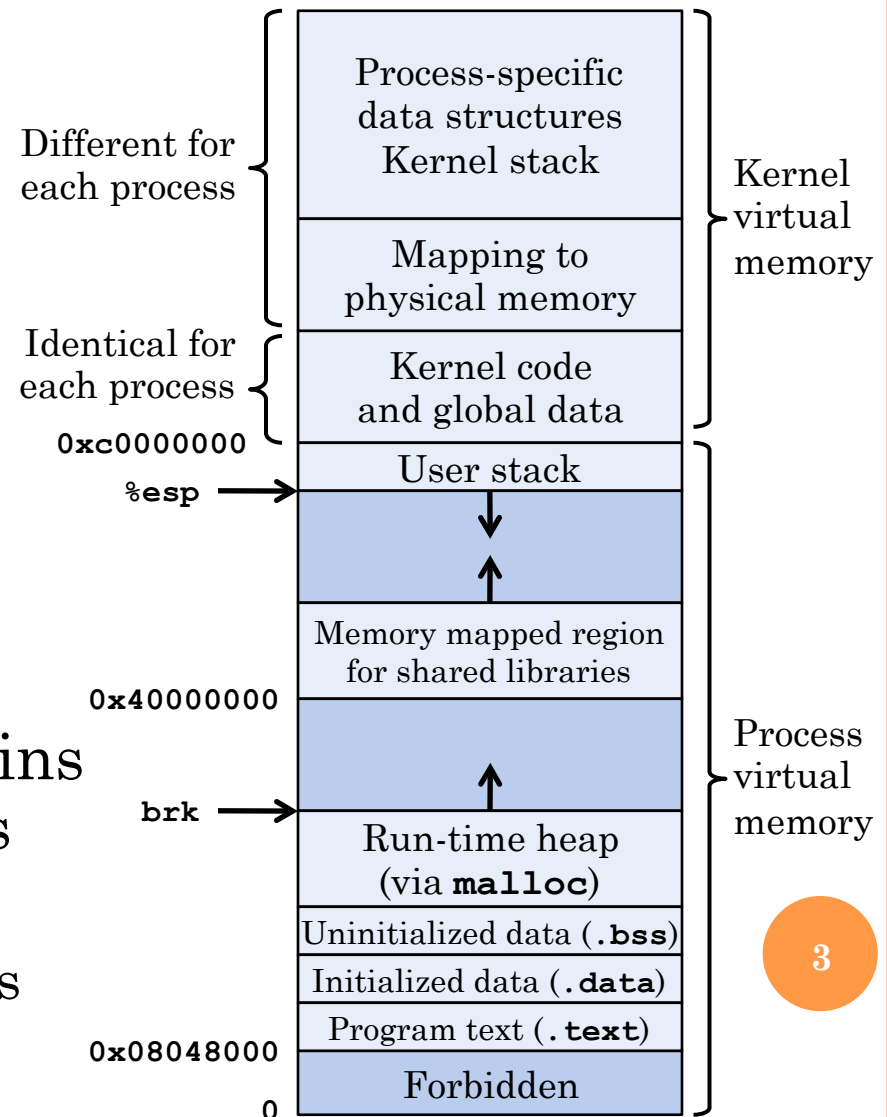
# LAST TIME: PROCESS MEMORY LAYOUT

- Explored how Linux uses IA32 virtual memory
- All processes have a similar memory layout
  - Each process has its own page table structure
  - Processes have isolated address spaces
  - Program entry-point is at **0x08048000**
  - Program's stack grows down from **0xc0000000**



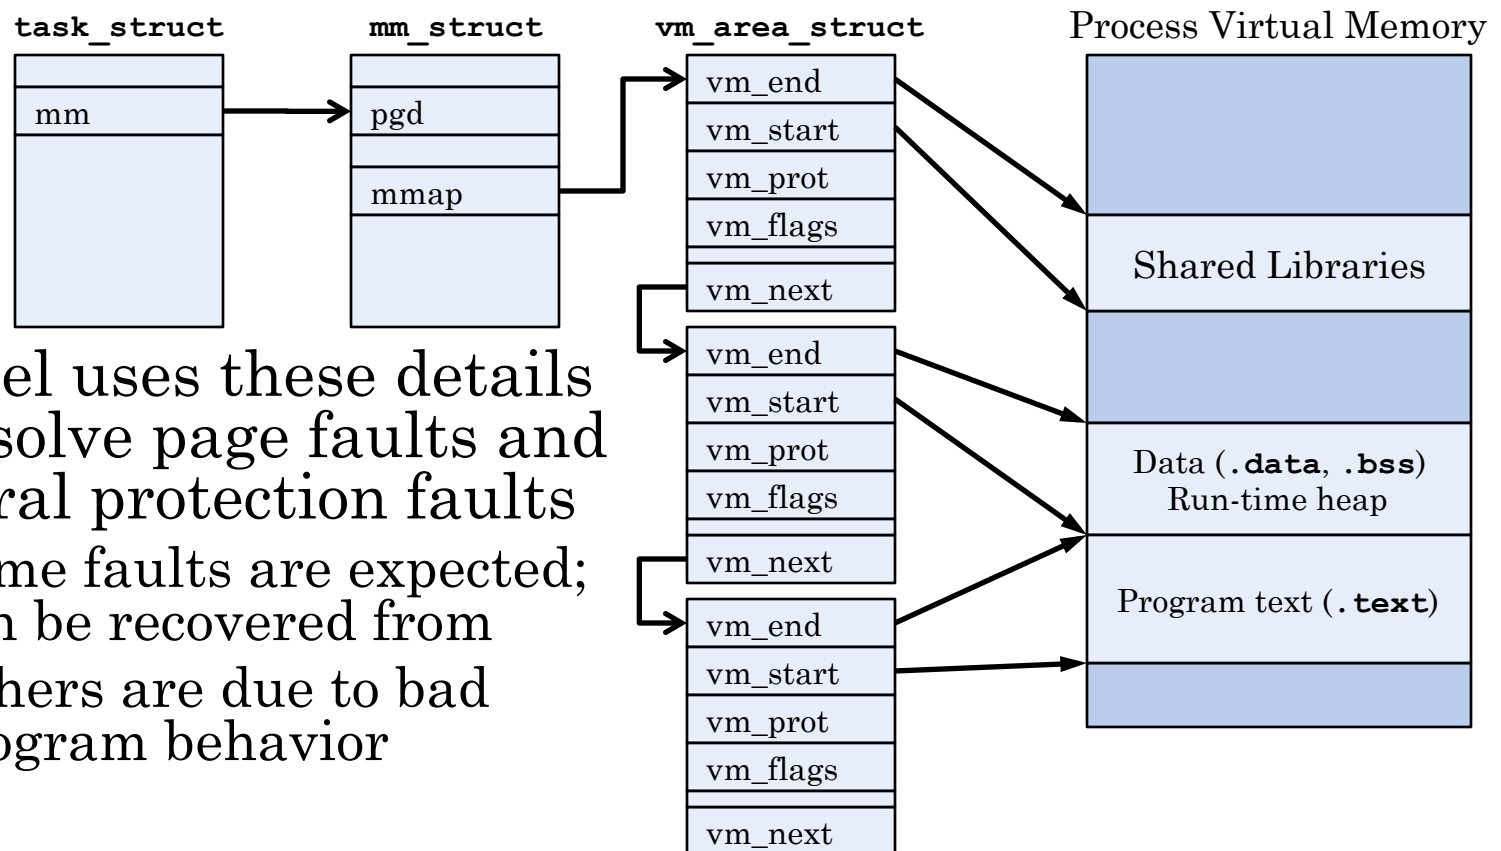
## LAST TIME: PROCESS MEMORY LAYOUT (2)

- Kernel maps some of its own code and data into each process' memory
  - Data structures that all processes need
  - Support for system calls
  - Processes cannot access this memory directly!
    - Must use `int 0x80` trap, or `sysenter/syscall`
- Kernel memory also contains data specific to the process
  - Page table for the process
  - Kernel-stack for the process



# PROCESS VIRTUAL MEMORY AREAS

- Kernel maintains details about virtual memory regions in each process
  - Supplemental details beyond what is recorded in the IA32 page-directory structure



- Kernel uses these details to resolve page faults and general protection faults
  - Some faults are expected; can be recovered from
  - Others are due to bad program behavior

# MEMORY MAPPING OBJECTS

- In Linux, each memory area is associated with an *object* on disk:
  - A *regular file* in the UNIX filesystem, such as a program binary or a shared library
  - An *anonymous file*, presented by the kernel, containing only zero values
    - (Not an actual file, but presented via the file abstraction)
- The anonymous file is used when a process allocates new virtual pages
  - (e.g. when the heap or stack is expanded)
  - A victim page is evicted, and then the page's contents are overwritten with zero values
  - Once a new page is allocated, it is saved to disk in a special swap area

## MEMORY MAPPING OBJECTS (2)

- The anonymous file is used when the kernel allocates new virtual pages to a process
- Reason:
  - A process should never be able to see data from another process, unless it is explicitly shared
- Example: a process that prompts the user for a password
  - Another process allocates a virtual page, which maps to the same physical page in DRAM...
  - ...but the password data wasn't overwritten before the second process gets the page!
- The kernel must ensure that such situations cannot happen

# SHARED AND PRIVATE OBJECTS

- Objects can be mapped into a virtual memory area as either a *shared object* or a *private object*
  - Memory area that the shared object is mapped into is called a *shared area*
  - Similarly, memory area that a private object is mapped into is called a *private area*
- Multiple processes may be running same code (e.g. **/bin/bash**), or using same shared library (**libc.so**)
  - Load the shared object into physical memory once, and then map it into the address space of multiple processes
- Writes to a shared object are visible to all processes accessing the object!
  - Writes will also modify the shared object stored on disk!
- With shared objects, *very* important to enable writes only when appropriate!


## SHARED AND PRIVATE OBJECTS (2)

- When a process writes to a private object, only that process should see the modification
  - Additionally, the change should not modify the private object on disk
- A simple technique:
  - Each time a specific private object is mapped into a process, load another copy into physical memory
- This approach can become *very* expensive
  - Particularly in situations where a specific private object is loaded into multiple processes, but none of the processes are making changes to the object!
- A much better technique: *copy-on-write*




# PRIVATE OBJECTS AND COPY-ON-WRITE

- Like shared objects, a private object is initially loaded into physical memory only once
  - Kernel sets the object's pages to be read-only, and flags the memory area as *private copy-on-write*
- Example:
  - Two processes using a private copy-on-write object
  - Initially, object is loaded into physical memory only once
- Both processes have virtual memory area marked as read-only, private copy-on-write



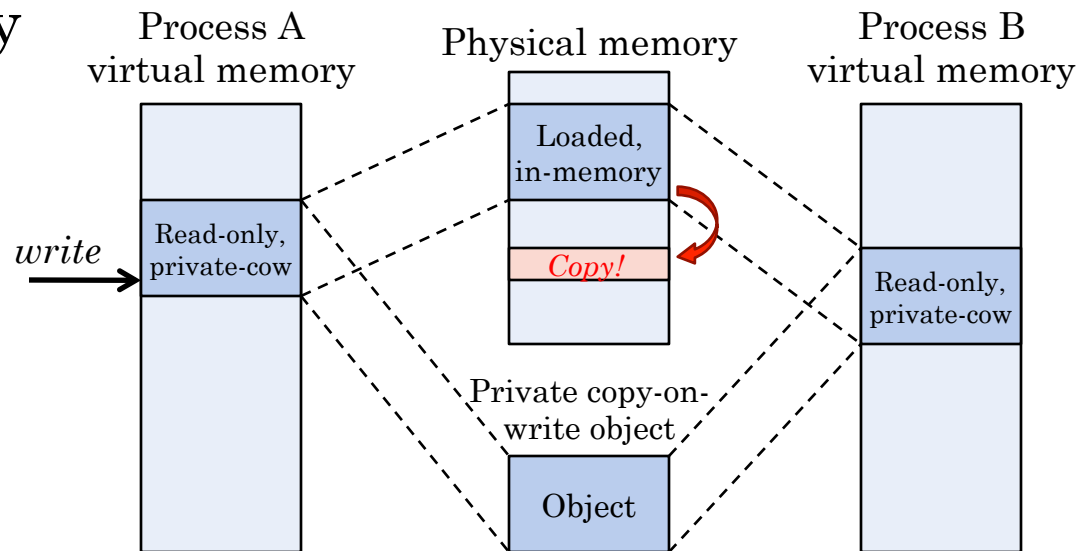
## PRIVATE OBJECTS, COPY-ON-WRITE (2)

- When a process writes to a private copy-on-write page, this generates a general protection fault
  - Process tried to write to a read-only memory page!
- Kernel handles these faults in a special way:
  - If write was to a read-only area that is also copy-on-write, make a copy of the page that was accessed
  - Create a new page
  - Copy data from old page into new page
  - In the process' page table, replace the old page with new page
    - Now change is local to the writing process



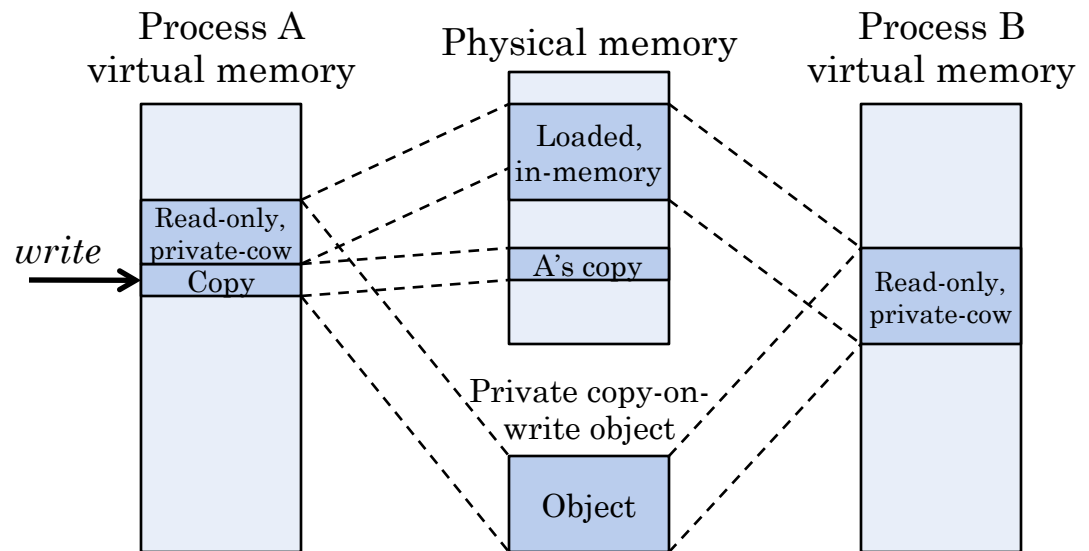
# PRIVATE OBJECTS, COPY-ON-WRITE (3)

- Example: Process A writes to the private object
- Kernel receives a general protection fault
  - Process tried to write to a read-only page...
  - ...but, the page is flagged as “private copy-on-write,” so the kernel performs copy-on-write steps
- Step 1: Allocate a new page, and copy the old data into it



## PRIVATE OBJECTS, COPY-ON-WRITE (3)

- Step 2: Update Process A's virtual memory space to reference the *copied* page, not the original
  - Copied page is also marked read-write, not read-only
- Step 3: Return from the protection-fault handler
  - CPU retries the instruction that caused the fault, and this time it succeeds without any problems

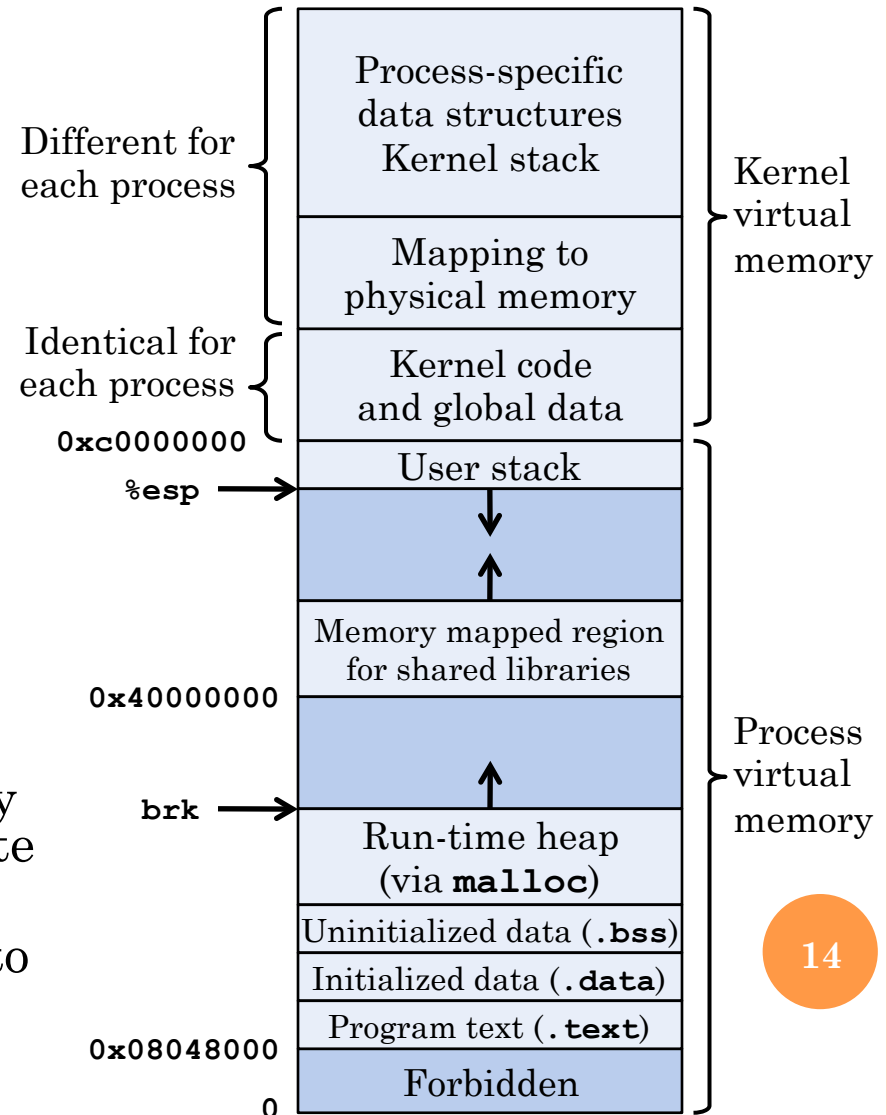


## COPY-ON-WRITE AND **fork()**

- When a process calls **fork()**, it spawns an identical child-process
  - Most significant differences are that the process ID and parent-process ID are different
  - All code, data, and I/O state of parent process is exactly replicated in the child process
- Creating an actual copy of the parent process would be inefficient and slow
  - Parent and child process share the same program text and shared libraries, and these are read-only...
  - Plus, parent and child processes might not actually change all of the data they now share
- Instead, kernel can use copy-on-write technique to create the child process

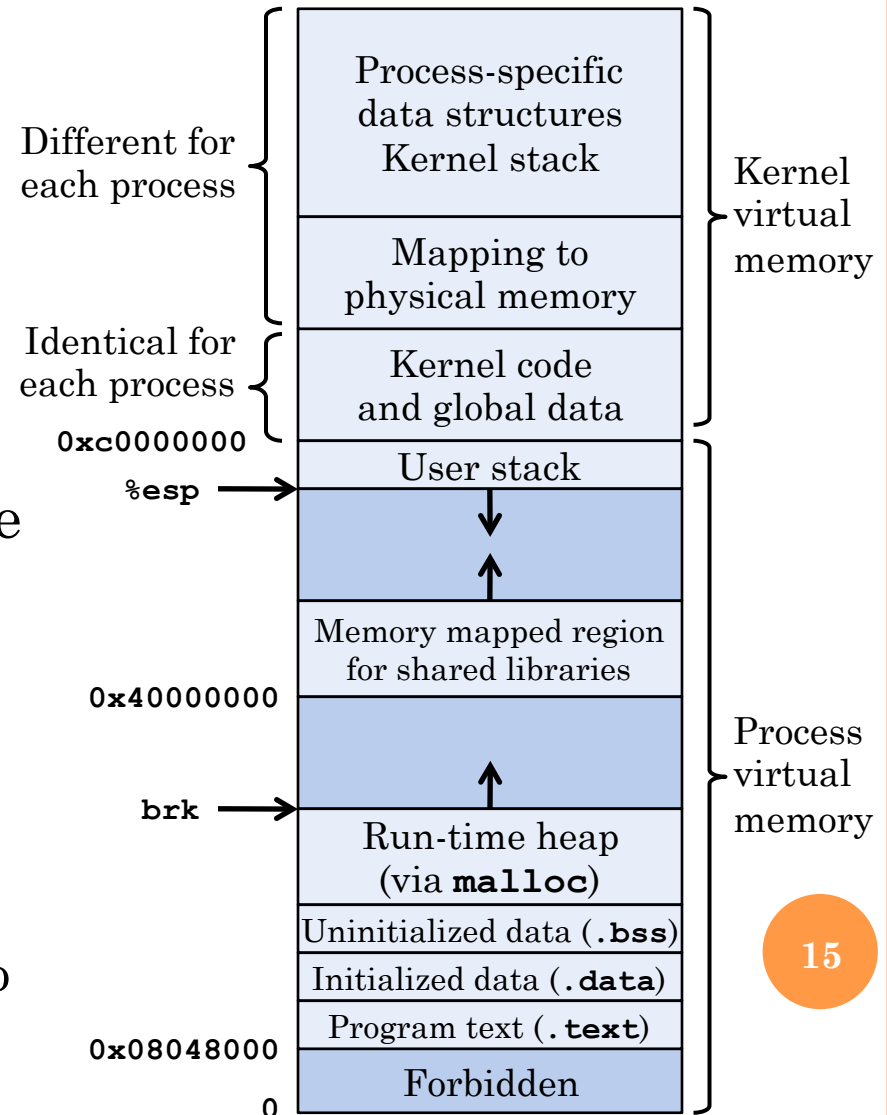
# COPY-ON-WRITE AND `fork()` (2)

- When process calls `fork()`, the kernel can use the copy-on-write technique:
  - Duplicate process-specific data structures, including page tables and `mm_struct` virtual memory details
  - Flag all pages in both processes as read-only
  - Mark all memory areas as private copy-on-write
- When either parent or child process writes to memory:
  - Modified page is automatically duplicated by the copy-on-write mechanism
  - Parent and child still appear to have isolated address spaces



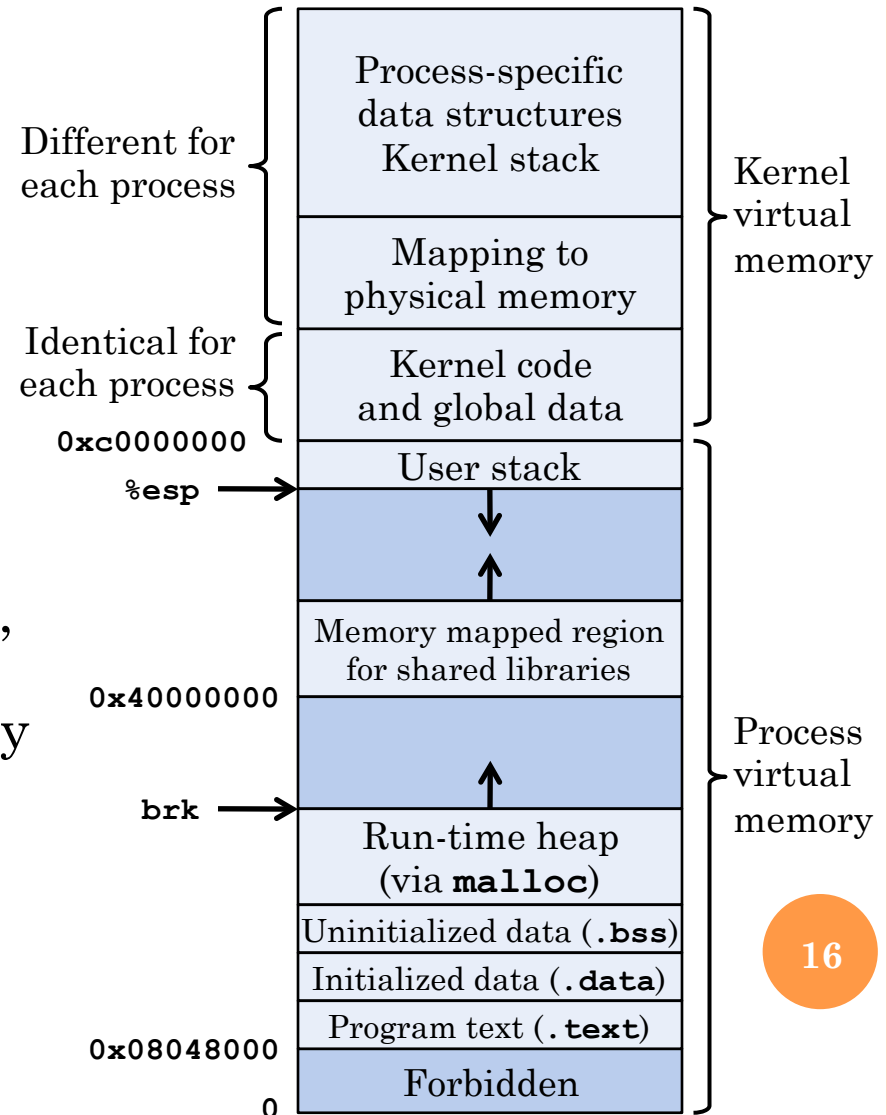
# VIRTUAL MEMORY AND `execve()`

- UNIX `execve()` function also relies heavily on the virtual memory system
  - Loads and runs a new program in the current process context
- Step 1: clean up the current process' virtual memory state
  - Reset the process' virtual address space in preparation for the new program
- Iterate through the `vm_area_struct` list:
  - Unmap virtual memory areas
  - Delete `vm_area_structs`, too



# VIRTUAL MEMORY AND `execve()` (2)

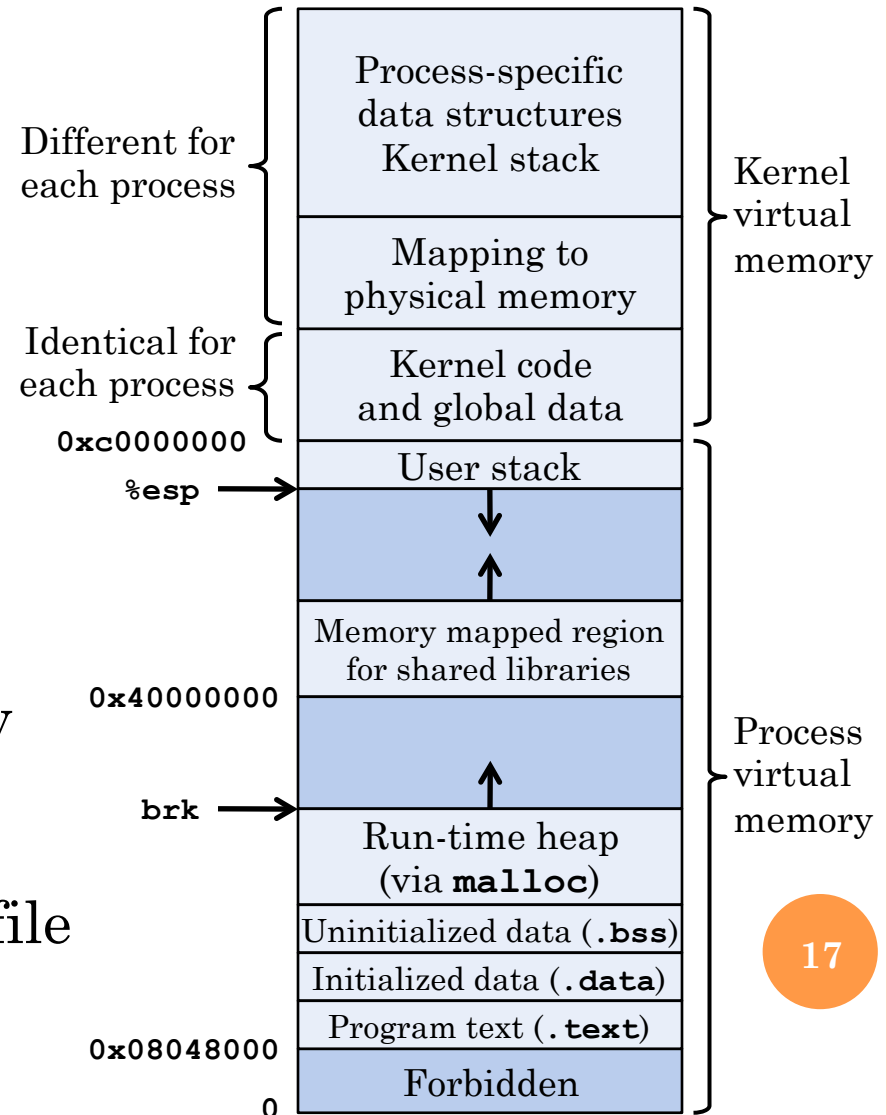
- Step 2: map private memory areas
- Program text (**.text**) and initialized data (**.data**) are specified in binary file
  - Map these virtual pages directly to the appropriate areas of the binary file
  - When pages are referenced, kernel will swap them into main memory automatically
- Both areas are marked private copy-on-write
  - (Or, **.text** may be marked read-only instead...)





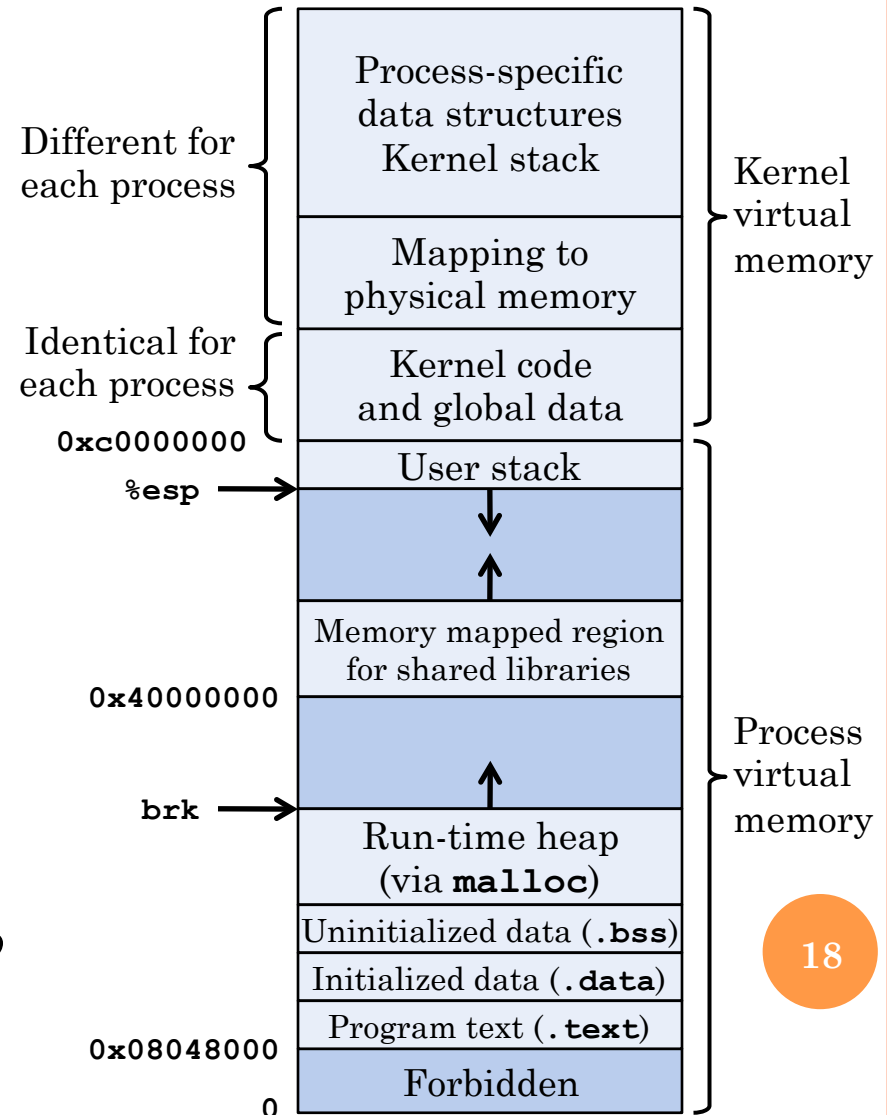
# VIRTUAL MEMORY AND `execve()` (2)

- Step 2: map private memory areas, *cont.*
- Uninitialized data (**.bss**) isn't contained in the binary file
  - Binary simply specifies the size of this region
  - Kernel maps the **.bss** pages to the *anonymous file*, which contains all zero values...
- Uninitialized data is initially set to all zero values
- Kernel also maps user stack and heap to the anonymous file
  - Sizes are increased as needed



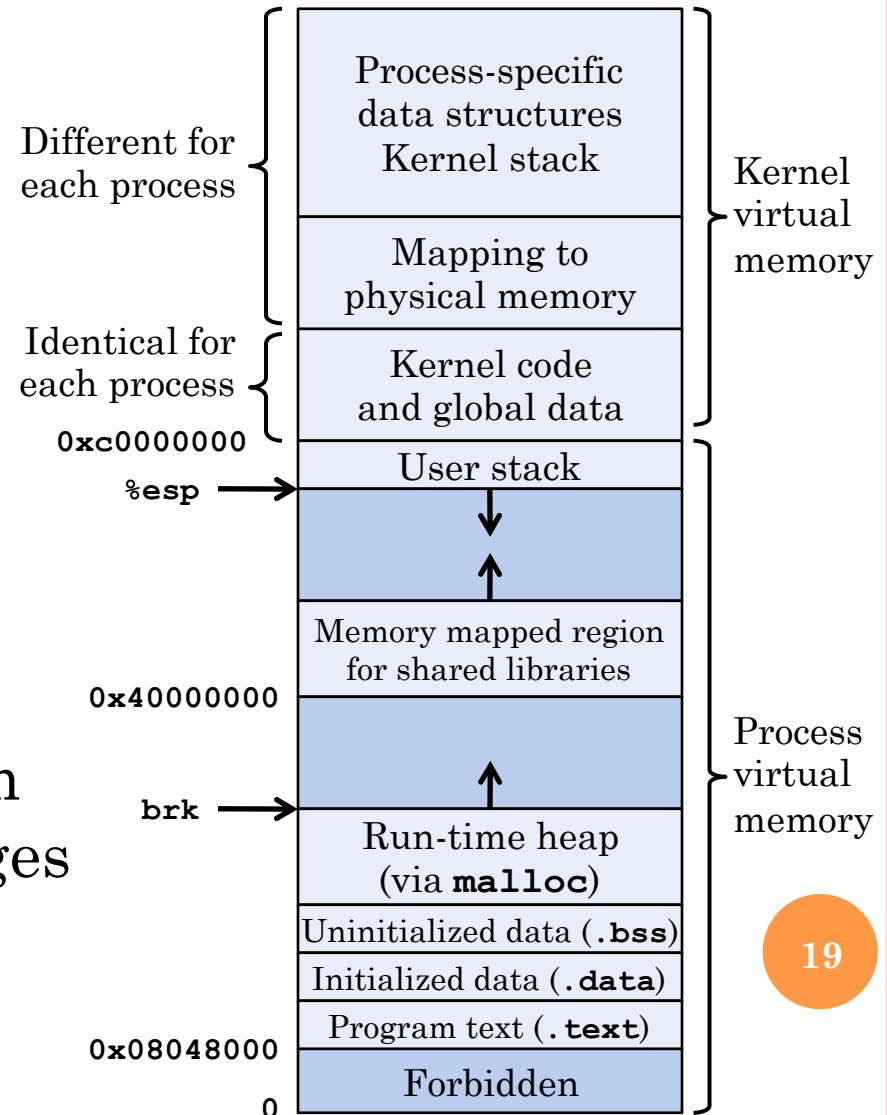
# VIRTUAL MEMORY AND `execve()` (3)

- Step 3: map shared memory areas
- If the program is linked with any shared objects:
  - e.g. `libc.so`, the C standard library
- Libraries are dynamically linked into the program
- Then, the shared objects are mapped into the process' virtual address space
  - e.g. shared read-only
  - *(otherwise, changes are visible to all other processes, and also modify the original file!)*



# VIRTUAL MEMORY AND `execve()` (4)

- Step 4: set the process' Program Counter to the program's entry-point
  - Next time the process is scheduled for execution, it will start running from the entry-point
- As new program executes, the virtual memory system will swap in necessary pages
  - e.g. program text, data, shared library code, etc.



# SUMMARY: VIRTUAL MEMORY

- Virtual memory is an essential component of modern operating systems
- Used extensively to implement many features
  - Process memory and isolation of address-spaces
  - Fast context-switches and process-forking
  - Simplifies loading of programs and libraries into main memory
  - Facilitates efficient memory use by sharing common data across many processes
    - Data is only duplicated when strictly necessary


# SOLID STATE DRIVES AND TRIM

---

Supplemental Material

# File Systems and File Deletion

- File systems generally separate directory information and file contents
  - One or more directories can contain a link to a given file (e.g. hard links, symlinks)
- When a file is deleted, two operations:
  - Remove directory entry to file
  - If no more directory entries reference the file, mark the file's space as available
- Normally, the file's actual data isn't modified by deletion
  - (Can securely delete file data by overwriting it one or more times)
- The operating system simply stops using the disk sectors where the file previously resided...
  - ...at least until the area is used by a new file

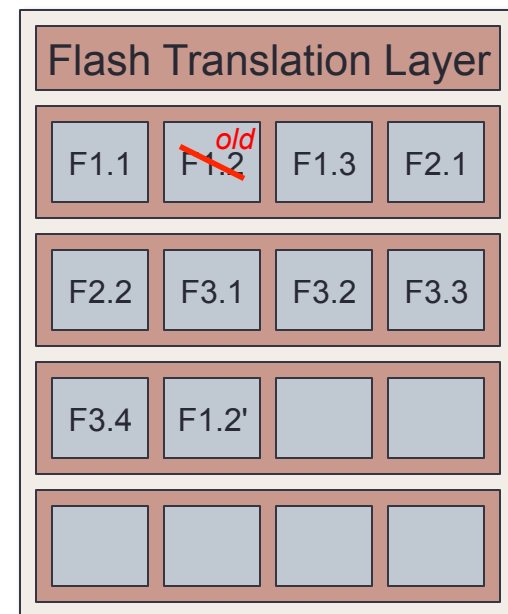


# Free Space and SSDs

- Solid State Drives (SSDs) and other flash-based devices often complicate management of free space
- SSDs are block devices; reads and writes are a fixed size
- Problem: can only write to a block that is currently empty
- Blocks can only be erased in groups, not individually!
  - An **erase block** is a group of blocks that are erased together
- Erase blocks are much larger than read/write blocks
  - A read/write block might be 4KiB or 8KiB...
  - Erase blocks are often 128 or 256 of these blocks (e.g. 2MiB)!
- As long as some blocks on the SSD are empty, writes can be performed immediately
- If the SSD has no more empty blocks, a group of blocks must be erased to provide more empty blocks

# Solid State Drives


- Solid State Drives include a **flash translation layer** that maps logical block addresses to physical memory cells
  - Recall: system uses Logical Block Addressing to access disks
- When files are written to the SSD, data must be stored in empty cells (i.e. old contents can't simply be overwritten)
- If a file is edited, the SSD sees a write issued against the same logical block
  - e.g. block 2 in file F1 is written
- SSD can't just replace block's contents...
- SSD marks the cell as "old," then stores the new block data in another cell, and updates the mapping in the FTL






# Solid State Drives (2)

- Over time, SSD ends up with few or no available cells
  - e.g. a series of writes to our SSD that results in all cells being used
- SSD must erase at least one block of cells to be reused
- Best case is when an entire erase-block can be reclaimed
  - SSD erases the entire block, and then carries on as before




# Solid State Drives (3)

- More complicated when an erase block still holds data
  - e.g. SSD decides it must reclaim the third erase-block
- SSD must relocate the current contents before erasing
- Result: sometimes a write *to* the SSD incurs additional writes *within* the SSD
  - Phenomenon is called **write amplification**



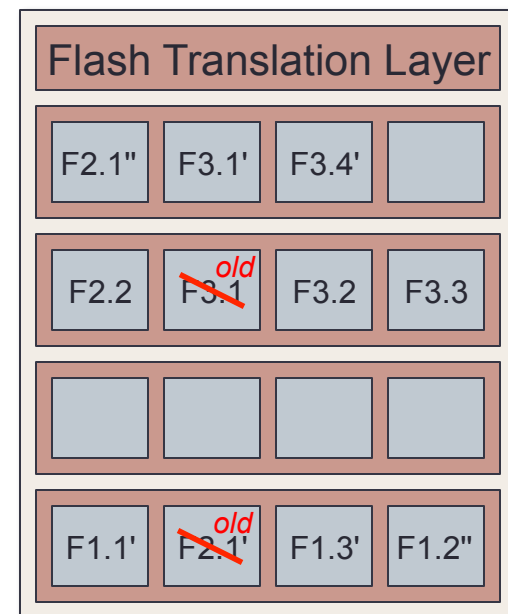
# Solid State Drives (4)

- SSDs must carefully manage this process to avoid uneven wear of its memory cells
  - Cells can only survive so many erase cycles, then become useless
- How does the SSD know when a cell's contents are no longer needed? (i.e. when to mark the cell "old")
- The SSD only knows because it sees multiple writes to the same logical block
  - The new version replaces the old version
  - The SSD knows that the old cell is no longer used for storage



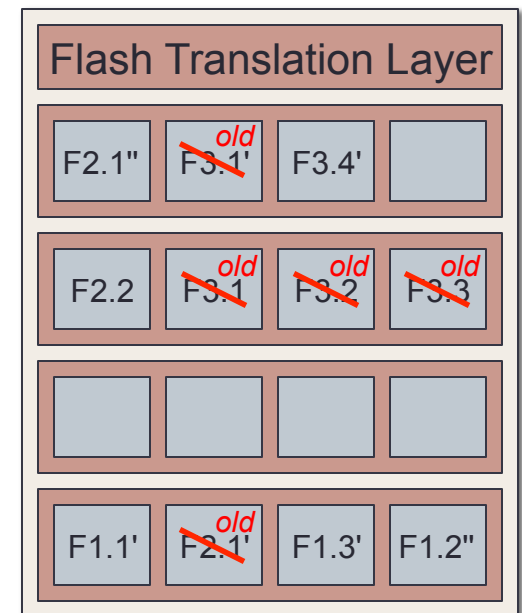
# SSDs and File Deletion

- Problem: for most file system formats, file deletion doesn't actually touch the blocks in the file themselves!
  - File systems try to avoid this anyway, because storage I/O is slow!
  - Want to only update directory entry and other bookkeeping data, and we want this to be as efficient as possible
- Example: File F3 is deleted from the SSD
  - SSD will only see the block with the directory entry change, and maybe a few other blocks
- The SSD has no idea that file F3's data no longer needs to be preserved
  - e.g. if the SSD decides to erase bank 2, it will still move F3.2 and F3.3 to other cells, even though the OS and the users don't care!



# SSDs, File Deletion and TRIM

- To deal with this, SSDs introduced the TRIM command
  - (TRIM is not an acronym)
- When the filesystem is finished with certain logical blocks, it can issue a TRIM command to inform the SSD that the data in those blocks can be discarded
- Previous example: file F3 is deleted
  - The OS can issue a TRIM command to inform SSD that all associated blocks are now unused
- TRIM allows the SSD to manage its cells much more efficiently
  - Greatly reduces write magnification issues
  - Helps reduce wear on SSD memory cells



# SSDs, File Deletion and TRIM (2)

- Still a few issues to resolve with TRIM at this point
- Biggest one is TRIM wasn't initially a queued command
  - Couldn't include TRIM commands in a mix of other read/write commands being sent to the device
  - TRIM must be performed separately, in isolation of other operations
- TRIM must be issued in a batch-mode way, when it won't interrupt other work
  - e.g. can't issue TRIM commands immediately after each delete operation
- This was fixed in SATA 3.1 specification
  - A queued version of TRIM was introduced
- Another issue: not all OSes/filesystems support TRIM (or not enabled by default)

