



How To Debug Programs, Part 1

General Debugging Approaches



CS24 – Spring 2011

CS24 and Debugging

- ▶ Debugging is an essential skill for all programming
- ▶ Particularly important in CS24:
 - ▶ Assignments involve relatively complex systems, frequently with multiple moving parts
- ▶ May also be the first class where you must debug a binary program, not something running in an interpreter
 - ▶ Must use other tools to peer inside this black box as it runs
 - ▶ Tools may be confusing, but you must learn them and use them
- ▶ Knowing how to debug problems effectively pays off:
 - ▶ For CS24, it can easily cut hours off of your assignments
 - ▶ If you end up programming for a living, knowing how to debug well will make you a superstar programmer



CS24 and Debugging (2)

- ▶ These lectures are to help you learn how to debug better
 - ▶ Note: there are many other considerations when you debug in a professional software development environment
 - ▶ This lecture is mainly to help you fix your own bugs
- ▶ Part 1: What are the basic principles and approaches of debugging?
 - ▶ [Mostly] independent of specific language, platform or toolset!
- ▶ Part 2: What tools and approaches can I use to debug my C and IA32 assembly language programs in CS24?
 - ▶ GDB, Valgrind



Bugs

- ▶ What is a “bug” anyway?
 - ▶ Helpful to break into two different components
- ▶ The program contains a defect in the code, possibly due to a design issue or an implementation issue
- ▶ This defect is manifested as a failure of some kind
 - ▶ Program produces an incorrect result, or it crashes, etc.
- ▶ Defects are not always manifested:
 - ▶ May require specific data inputs to produce the failure
 - ▶ May require running the program on a specific platform
- ▶ Defects are also not always manifested immediately!
 - ▶ May require running a program for hours, days, or weeks (!!!) before the failure occurs



Finding Bugs

- ▶ The majority of the work in fixing a bug is finding the actual defect that produces it
 - ▶ Usually, once the defective code is identified, a fix is very easy!
 - ▶ If overall *design* is defective, this can be much more difficult: can require redesigning and reimplementing a large portion of the program
- ▶ The defect is only the *cause* of the failure, but it is not the failure itself!
 - ▶ The defect will immediately begin to affect the program's state, but the effects may not become visible for some time
 - ▶ The greater the separation between defect and failure, the harder it is to diagnose the defect from the failure



Finding Bugs (2)

- ▶ If the defect directly causes the failure:
 - ▶ e.g. your loop's logic is broken and it dereferences a NULL ptr
 - ▶ Happy days! Simply need to identify point that the program fails, and can fix the bug very quickly
- ▶ Often, the defect and the failure are separated by a significant amount of execution time
- ▶ Example:
 - ▶ Function `f1()` mangles a linked list, but completes successfully
 - ▶ Function `f2()` attempts to use the linked list, and fails miserably
- ▶ In this case, the bug itself is not actually in `f2()`!
 - ▶ The defect is in `f1()`; its *manifestation* is in `f2()`
- ▶ How do we determine the actual cause of the failure?



Preemptive Bug Detection

- ▶ One common technique for causing defects to manifest quickly is using assertions
 - ▶ Frequently have conditions that you expect to be true at certain points in your program
 - ▶ Explicitly state these in an assertion, in your code
 - ▶ At runtime, the assertion is checked: if it's false, the program is stopped immediately!
- ▶ In C programs:
 - ▶ `#include <assert.h>`
 - ▶ `assert(condition expected to be true);`
- ▶ Use assertions to check function arguments, return values, and the state managed by your functions
 - ▶ See both the provided code and solution sets for examples!



Detective Work

- ▶ Once you have observed a failure in your program, you have a mystery to solve!
- ▶ One thing that no good detective ever does:
 - ▶ Guess randomly!
- ▶ The programming version of this:
 - ▶ Make random guesses as to the cause, and try various changes
 - ▶ Called “the shotgun approach,” and “monkeys on a typewriter”
 - ▶ (The results are just as good, too.)
- ▶ Considering various clues, you must track down the defect from the indications of failure
- ▶ Once the issue is identified, *then* make your fix
 - ▶ Not as likely to introduce other defects this way, too.
 - ▶ (Hunt bugs with a rifle, not with a shotgun.)



Step 1: Reproduce the Failure

- ▶ Before you can do anything else, you must find a way to reproduce the failure
 - ▶ Recreate the steps that caused the program to fail
 - ▶ Was it specific data inputs? Was it a specific interaction with the user interface? Does the program fail under heavy load?
- ▶ This is very important, for three major reasons:
 1. **So you can watch it fail.** You can see exactly what the program was doing as it crashed and burned.
 2. **So you can zero in on the cause.** If the program fails in some circumstances but not in others, this will give you *hints* as to what part of the program actually contains the defect.
 3. **So you can test if you actually fixed it.** If you can reliably cause the failure, and your fix makes it go away, you win!



Step 1: Reproduce the Failure (2)

- ▶ Which of these reasons actually requires debugging tools?
 1. **So you can watch it fail.** You can see exactly what the program was doing as it crashed and burned.
 2. **So you can zero in on the cause.** If the program fails in some circumstances but not in others, this will give you *hints* as to what part of the program actually contains the defect.
 3. **So you can test if you actually fixed it.** If you can reliably cause the failure, and your fix makes it go away, you win!
- ▶ Debugging tools and techniques allow you to watch your programs fail
 - ▶ ...or to see what they were doing when they failed...
 - ▶ A small but very critical part of the challenge of debugging



Technique: Keep a Debugging Log!

- ▶ When debugging, it's extremely helpful to keep a record of your efforts (hand-written or typed, it doesn't matter)
 - ▶ A general description of the failure
 - ▶ Inputs or circumstances in which the failure occurs
 - ▶ Ideas of potential causes, along with the efforts you made to verify your ideas, and indications for or against each theory
- ▶ Very effective for helping focus your thoughts and ideas
- ▶ Also allows you to set aside an issue, and pick it up later
- ▶ It's generally a very bad idea to debug while exhausted
 - ▶ Thought processes aren't clear; can't reason about the bug
 - ▶ Tend to revert to “monkeys on a typewriter” mode
 - ▶ Put it down, walk away; come back later when you're fresh



Step 2: Isolate the Failure

- ▶ So your program fails. And you know how to make it fail.
- ▶ Problem: the defect is only in a small part of your code
 - ▶ Need to zero in on the part of the code that's actually flawed
- ▶ Goal: try to devise the smallest possible scenario that still causes the failure to occur
 - ▶ If the scenario is small, there won't be very much code involved
 - ▶ The less code that's involved, the easier it is to find the defect
- ▶ Example: a program that processes large log files
 - ▶ A particular log file causes it to crash. The log file is 27MB.
 - ▶ What do you do?



Example: Failure Isolation

- ▶ **Example:** a program that processes large log files
 - ▶ A particular log file causes it to crash. The log file is 27MB.
 - ▶ What do you do?
- ▶ If you have the offending log file, can reproduce the failure
 - ▶ Still, a lot of code is executed before the failure occurs
- ▶ **Want to narrow in on the actual cause of the bug:**
 - ▶ Cut the log file down until you have the *minimal* portion that still causes the crash
 - ▶ Perhaps the program crashes because log contains bad values
 - ▶ If you're lucky, may identify a handful of input lines that causes the problem
 - ▶ Should make it much easier to identify and resolve the defect



Example: Failure Isolation (2)

- ▶ Example: a program that processes large log files
 - ▶ A particular log file causes it to crash. The log file is 27MB.
- ▶ What if you can't get the issue to reproduce by trying smaller parts of the log file in isolation?
- ▶ That would be a much harder bug to track down...
 - ▶ ...but, would indicate that the bug is triggered by some characteristic of the log file *as a whole*
 - ▶ e.g. maybe log's size causes some memory management issue
 - ▶ e.g. maybe the combination of some values read from the log causes the failure
- ▶ You learn more about the nature of the defect by trying to isolate it! (Record these things in your debugging log!)



Step 2: Isolate the Failure (2)

- ▶ Previous example involved trying to isolate the *inputs* that cause the failure
- ▶ Can also try to isolate the general part of your program's code that may cause the failure
- ▶ Example: memory allocator in assignment 3 😊
 - ▶ When you run the allocator testing program, it crashes
 - ▶ Easy to reproduce: every time you run it, it crashes
 - ▶ But, the testing program doesn't have any inputs.
 - ▶ What do you do?



Technique: Understand the System!

- ▶ Knowing where to start with this bug really requires you to understand the system being debugged. For example:
- ▶ What are the major functional components of program?
 - ▶ What portions of the code perform each of these functions?
 - ▶ If failure manifests when a specific feature is used, you know what files to start focusing on.
- ▶ Are there portions of the code that are always executed?
 - ▶ If failure manifests in a range of different scenarios, this would be a likely location for the defect.
- ▶ Are there portions of the code that can easily be disabled or removed, and the test run again?
 - ▶ If part of the code can be removed and the failure still occurs, we know the cause is not in the code that was removed.



Technique: Understand the System! (2)

- ▶ Can ask similar questions for the inputs to your programs
- ▶ Are there portions of the code that only execute with specific data inputs, or sequences of inputs?
 - ▶ e.g. are certain parts of your program executed only when you feed the program specific inputs?
- ▶ Example: Sparse Vector in CSII C++ track
 - ▶ Represented as a linked-list of elements, kept in order of index
 - ▶ By feeding in specific sequences, can target testing more specifically
 - ▶ Add values to the vector with decreasing indexes:
 - ▶ Tests linked-list prepend code
 - ▶ Add values to the vector with increasing indexes:
 - ▶ Tests linked-list append code
 - ▶ Add values to the vector with varying indexes:
 - ▶ Tests linked-list insert code
 - ▶ Each of these is usually implemented in a *different* piece of code



Technique: Understand the System! (3)

- ▶ The whole point of understanding the system:
 - ▶ You are performing experiments with your program, and making observations of its behavior
 - ▶ You need to *correlate* the observed behaviors with various parts of your program's source code
 - ▶ The more effective you are at doing this, the faster you will zero in on the location of bugs
- ▶ (Debugging often involves a lot of thinking...)
- ▶ Both positive *and negative* correlations are helpful!
 - ▶ If you can say, "This behavior is definitely not caused by the code in this part of the program," that also helps narrow down the source of the issue.



Technique: Understand the System! (4)

- ▶ Memory allocators provide two major operations:
 - ▶ Allocate a chunk of memory from the heap. If requested amount of memory isn't available, simply return NULL.
 - ▶ Release a chunk of memory back to the heap.
- ▶ What would happen if I change the “release memory” code to be a no-op?
 - ▶ Allocations will succeed until all memory is consumed, and then subsequent allocations will fail
 - ▶ This is not a problem! Caller expects that allocations may fail!
- ▶ Can disable the “release memory” code, recompile and rerun the test, and see if the crash still occurs



Technique: Disabling Code

- ▶ How to disable a portion of the code depends on the specific language being used, and the program's structure
- ▶ In the C language, can comment out a chunk of code:

```
/* This is my awesome function that doesn't work. */
```

```
void foo() {
```

```
    bar();
```

```
    /* TODO: this may be buggy...
```

```
    abc();    /* Do something amazing! */
```

```
    */
```

```
    return xyz();
```

```
}
```

- ▶ C doesn't support nested block-comments (bit of a pain)



Technique: Disabling Code (2)

- ▶ Can also use *preprocessor directives* to disable a region of code:

```
/* This is my awesome function that doesn't work. */  
void foo() {  
    bar();  
    #if 0        /* TODO: this may be buggy... */  
    abc();      /* Do something amazing! */  
    #endif  
    return xyz();  
}
```

- ▶ The `#if` preprocessor directive evaluates its expression
- ▶ If expression result is nonzero, the block of code is compiled
- ▶ If result is zero, the block of code is excluded from compilation
- ▶ Code within the `#if 0 / #endif` region will be excluded



Example: Memory Allocator

- ▶ First try: comment out the “release memory” code, and try running the allocator testing program again
- ▶ If it still crashes, what do we know?
 - ▶ We know that there’s a crashing bug in the allocation code...
 - ▶ (Can’t conclude that there definitely isn’t a crashing bug in the deallocation code, but first things first...)
 - ▶ Can focus our attention on allocation, looking for defects
- ▶ If it doesn’t crash, what do we know?
 - ▶ Again, cannot conclude that the allocation code is bug-free!
 - ▶ It may simply mangle the heap in such a way that deallocation crashes
 - ▶ Need to use another strategy to isolate the defect
- ▶ It’s an easy check, and will give you more information



Technique: Create Simple Test Cases

- ▶ If you can't chop down the input data, and you can't chop down the program itself:
- ▶ Can create very simple test cases to exercise simple paths through the code
- ▶ Allocator example:
 - ▶ Write a test that allocates one chunk of memory, then frees it
 - ▶ If it crashes, you can better discern what code paths are being followed in your program
 - ▶ If not, create a more complex test case
 - ▶ Write a test that allocates two chunks of memory, then frees the first chunk
 - ▶ Repeat until you have a test case that still causes the failure



Step 3: Identify the Defect Itself

- ▶ If you have successfully completed the first two steps:
 - ▶ You probably have the location of the defect narrowed down to a relatively small portion of your code
- ▶ Now, need to identify possible origins of the failure (may be multiple candidates!), and eliminate them one by one
 - ▶ Requires that you reproduce the failure yet again, and watch what your program does as it goes down in flames
 - ▶ If you can't peer into your program's execution, you cannot identify the exact origins of the bug



Step 3: Identify the Defect Itself (2)

- ▶ **Several different approaches for this step:**
 - ▶ Instrument your code to produce logging/debugging output
 - ▶ Run your program in a debugger to single-step through the failure scenario
- ▶ **Both approaches have the same fundamental goal:**
 - ▶ Examine the state-changes your program is performing, correlated with the lines of code making those state-changes
 - ▶ Determine the exact point in time when your program begins to create invalid state
- ▶ **Using a debugger is a more powerful and less intrusive way of doing this, but either approach will work**
 - ▶ Can introduce other bugs while adding your debug output...



Technique: Printing Out Details

- ▶ A very common approach for debugging C programs:
 - ▶ Add `printf()` statements to the code, then compile and rerun
 - ▶ Then, pore through the debug output to see what happened
- ▶ Beware: the standard output stream (`stdout`) is buffered!
- ▶ Sometimes, program seems to crash in a location not indicated by debug output!
- ▶ Solution: flush unwritten debug output to the console

```
void buggy(int x) {  
    int i;  
    for (i = 0; i < x; i++) {  
        printf("i = %d, a[i] = %s\n", i, a[i]);  
        fflush(stdout);  
        ... /* do buggy stuff with i and a */  
    }  
}
```

- ▶ (Not always necessary, but if you see odd behavior, give it a try.)



Technique: Printing Out Details (2)

- ▶ If you are going to add debug output, might as well print out everything you can think of
- ▶ Common scenario:
 - ▶ The program fails.
 - ▶ Programmer suspects a particular cause of the failure, and adds debug-output to the program to explore that specific cause.
 - ▶ Guess what, it's not that.
 - ▶ Programmer has to go back and print out more details...
- ▶ When adding debug output, print out all details that could be useful to know, so that you have full information
 - ▶ Will allow you to evaluate multiple potential causes in less time



Technique: Printing Out Details (3)

- ▶ Make sure every debug output line is unique in some way!

```
void buggy(int x) {  
    int i;  
    for (i = 0; i < x; i++) {  
        printf("i = %d, a[i] = %s\n", i, a[i]);  
        ... /* do buggy stuff with i and a */  
        printf("i = %d, a[i] = %s\n", i, a[i]);  
        ... /* do more buggy stuff with i and a */  
    }  
}
```

- ▶ In this example, can't easily correlate debug output with the line that produced it!
- ▶ Simplest approach: put numbers or some other unique value at the front of each debug output line

```
printf("buggy 2: i = %d, a[i] = %s\n", i, a[i]);
```



Technique: Printing Out Details (4)

- ▶ Let's say you followed these steps, and had wild success.
 - ▶ All bugs vanquished!
- ▶ And, the program now dumps out a ton of debug info. ☹
 - ▶ You need to get rid of this debug output...
 - ▶ (graders/customers do not want to see it!)
 - ▶ But, you may need this debug output again in the future!
- ▶ Similar to before, you should modify your debug output code so that you can conditionally enable and disable it
- ▶ At the top of your source file, or in a widely-used header file, define a symbol that controls debug output:

```
/* Set to 0 to disable debug output, nonzero to enable. */  
#define DEBUG_INFO 1
```



Technique: Printing Out Details (5)

- ▶ Then, wrap all debug-output lines with an `#if` guard

```
void buggy(int x) {  
    int i;  
    for (i = 0; i < x; i++) {  
        #if DEBUG_INFO  
            printf("i = %d, a[i] = %s\n", i, a[i]);  
            fflush(stdout);  
        #endif  
        ... /* do buggy stuff with i and a */  
    }  
}
```

- ▶ Now, can enable/disable all debug output with one switch!
 - ▶ Make sure debug output is disabled before turning things in...



Step 4: Fix the Defect; Verify the Fix

- ▶ Once actual defect is found, usually straightforward to fix
 - ▶ However, if the program's *design* is defective, may need to rework substantial portions of the code
 - ▶ (This is why it's always good to design up front!)
- ▶ You aren't finished fixing the bug until you verify the fix
- ▶ By this point, you should have some way of reproducing the failure...
 - ▶ Retry your test cases and see if the failure no longer occurs
 - ▶ If no more failures, you're done!
- ▶ Do not assume that if the fix *compiles*, the bug is fixed!
 - ▶ If this is your approach, you actually probably made it worse...



Step 4: Fix the Defect; Verify the Fix (2)

- ▶ Bug fixes can also introduce new defects into the code
 - ▶ Such defects are called *regressions*
- ▶ Usually occurs when:
 - ▶ The actual cause of the original bug is not fully understood
 - ▶ Or, the impact of the bug-fix is not fully understood
- ▶ Good programmers also check for regressions:
 - ▶ Verify that the original bug is fixed
 - ▶ Also run other tests to ensure that no new bugs were introduced
- ▶ Be a good programmer 😊
 - ▶ It will save you tons of time and frustration in the long run



Some Notes about Fixing Bugs

- ▶ Programmers are an imaginative bunch.
- ▶ You must beware of some common pitfalls that are part of *your own nature*!
- ▶ **Believe your observations, not your suspicions.**
- ▶ A common scenario:
 - ▶ The program fails. The programmer *suspects* that the failure is caused by a particular issue, and focuses his attention there.
 - ▶ However, there is no indication that the suspected cause is the *actual* cause of the problem.
 - ▶ Sometimes this can even cause you to miss obvious details that clearly indicate the actual source of the problem.



Some Notes about Fixing Bugs (2)

- ▶ **If you didn't fix it, it isn't fixed!**
- ▶ Another common scenario:
 - ▶ An intermittent failure is occurring.
 - ▶ You have made a few stabs at fixing it, but you still really don't know what causes the problem or how to reproduce it.
 - ▶ You definitely can't correlate any "fixes" with the problem
- ▶ It is very appealing to assume that if a problem hasn't occurred recently, it must be fixed.
 - ▶ You may even avoid focused testing on that issue (mainly because you really don't want to know...)
- ▶ Normally, these problems come back.
 - ▶ Usually during a demo, or when your code is being graded.



Next Time: Debugging with GDB

- ▶ Have covered a relatively large range of general debugging issues and approaches
 - ▶ There is certainly much more where these came from
 - ▶ As always, practice makes perfect
- ▶ Next time, will focus on how to use GDB and Valgrind to debug your C programs

